

IT 494 · RESEARCH AND PROPOSAL · PHASE 1

Persistent Memory for Stateless Models

Proposal, literature review, and project plan

Justin Hoffman
Illinois State University, MS Computer Science
Supervisor: Dr. Xing Fang
August 27, 2026

Literature summaries assembled with AI assistance; all citations verified against source publications.

Index

The proposal, before the research	3
Reading list	7
The field, by problem	12
Where things stand	22
The paper's argument	28
Requirements and testing	37
Design: schema, ingestion, retrieval, delivery	44
The unit contract	57
Fall 2026 plan	60
Spring 2027 plan	67
Schema prior art and references	71
Ways the research could elevate the project	77
Anatomy of the prototype	82
Feasibility, reviewed antagonistically	97
The literature, work by work	104
Appendix: one-page summaries	117

REFERENCE, not superseded. This is the project's intent in Justin's own voice, sealed from the research. The only document that holds it. It is background rather than part of the working set. Index: docs/README.md.

IT 494 project proposal, pre-research draft

What I have running and what it has done

Every conversation I have with an AI is saved as a raw transcript and never edited. At the end of a working session a summary gets written and filed into a tree, organized by area of my life and then by topic. Those summaries roll up, so a higher level can answer a question without reading everything underneath it. A maintenance pass runs on a schedule: it folds in new material, audits a rotating subset of summaries against the original transcripts, and re-summarizes anything whose source changed after it was written. Nothing is deleted. Children fold up into parents and redundant siblings merge, so what has to be read stays bounded while the detail stays recoverable at the bottom. It is wired into my tools, so a new chat starts knowing the map exists and can search it.

This was built inartfully. I had no knowledge of the field when I started, so nothing about it follows a known design. It accreted as I hit problems, and it is a pile of scripts and text files rather than a system anyone would draw on a whiteboard. It is also in daily use and doing real work.

The clearest proof was annual training. I ran separate chats at the same time for the flight schedule, the storyboard, the weather update, and the situation report, all against one shared knowledge base. For the storyboards I pulled images and video out of email and matched them against the mission tracker, AMRS, and the flight schedule by their metadata, then built the boards from a cached template. Each of those threads would have exceeded a single conversation on its own, and they had to agree with each other. The D&D campaign I run is the same shape: long-running, many sessions, dependent on decisions made months ago staying findable and staying correct. When the system gets something wrong, I tell it and it corrects itself; I am entirely hands off, and I never touch the files myself.

The end state I am aiming at

The model is essentially a function. It takes a prompt and predicts the best answer: a decision engine. But an engine does not make an airplane. The memory system is the airframe, the part that carries state from one call to the next, and the airframe is what I am building. The end state in one sentence: a user-level knowledge backend that can be scaled to an organization.

Today the backend serves one user, me. The organizational version I can picture concretely is a chat help desk that learns from its own conversations. An example from an AWS Summit talk: a utility answering billing questions with access to account information. Store the raw conversations, map how conversations actually progress, track what was concluded, digest it into long-term memory. Many users, one organization. My GAO work is why I think this is tractable rather than speculative. Pulling central

ideas out of free human text worked. What failed was asking one tool to discover topics and enforce a rigid classification at the same time; separate those jobs, discovery on one side and the organization's own fixed categories on the other, and it works. The utility already has its categories in its dispute codes, the same way the agency had its survey indexes. I also came out of that work with methods for judging whether this kind of output is any good, and that is what I would bring.

The help desk has something my system does not: an outcome signal. Ticket closed, reopened, escalated, repeat contact. My system cannot learn that it gave a wrong answer. It records filing something in the wrong place, but a wrong answer leaves no trace, so there is nothing to measure improvement against. What I am really after is an efficient framework for ingesting data, organizing it, retrieving it, then injecting it, and maintaining it over time. I think in those five steps, and all five need to be built in phases. The idea is a living tree that ingests data over time, not an index built once.

What I do not know

I have no formal grounding in knowledge graphs or in persistent memory for AI, and I am unlikely to encounter either in a classroom. A model is stateless, so anything that appears to remember is machinery built around it: vector stores, retrieval-augmented generation, caching. I know roughly what these do. I do not know how they compare, where each breaks down, or what else exists. I do not even know that knowledge graphs are the way I want to go.

I do not have real metrics. I set up a fixed set of questions with known answers, and the system answers them, but the set is imperfect and I would not claim it measures much. So I cannot state the system's limitations with any confidence. Problems get patched as they surface, which is fine for my own use and is not evidence of anything.

I do not know whether my structural instincts hold up. I chose written summaries over similarity search on raw text, and the reason is merging. "Jenny is employed by GAO" and "Jennifer works at the Government Accountability Office" are the same fact, so some degree of merging base facts, or inferring until you reach the things that matter, has to happen; a hundred fact tokens might merge into a single attribute in a summary. Two decisions here are already made and I will defend them: written summaries over similarity search, and a record where facts are never overwritten; new facts supersede old ones, and the superseding itself is part of what the system knows. How the merging actually works is a question the reading must settle. The merge cannot be a black box: I would keep a ledger of which tokens, in sequence, merged into an attribute understanding, and a list of the keywords that all map to the same individual, but that is a suspicion, not a design. It would have to understand pronouns too, which is why an LLM has to do the distilling: grabbing facts from documents takes context, and an embedder cannot do that. The bookkeeping around the distilling does not need one. I think the real trick is having the AI organize information on its own, a learned database. The open questions are easiest to see on a concrete corpus. Say the corpus is Harry Potter: which entities get their own pages? What attributes of a character exist on a page, and how do those change over time? How are characters grouped, by house, by school, by allegiance, and which groupings are most relevant for organization? How is all of it structured and searched through, and what data structures are we using? My first approach was chats at the leaves. Then

it became clear the system needed to maintain knowledge about entities that exist outside any topic (me, my wife, my internship, my school), so the organization cannot be one tree; entities have to cut across it. How the rest is organized probably depends on how the text is ingested, and the use case is LLM chat logs and uploaded documents as the corpus. I do not have settled answers past that point, and settling them defensibly is the design work. A drive can be as simple as observe, learn, adapt, improve, but the model does nothing unless directed: my schedulers fire on time, a drive would fire on condition, and my maintenance log already computes conditions it is not allowed to act on. Whether any of this matches how the field thinks about these problems, I cannot say. And one claim under the whole multi-chat use case gets checked before anything else: that concurrent sessions cannot share a context window regardless of its size. The checks that could sink the semester are scheduled below.

How I would spend the semester

Research first, then the proposal: the near work is gathering resources and reading, and I will not commit to a design before that is done. But my own estimate governs the plan: the full five-step pipeline is hundreds of hours, and the semester gives me seventy to a hundred, so it is a year's project. The fall carries reading, design, and measurement; the spring carries the product half, and the spring gets the same treatment the fall is getting here: its own hour budget and its own cut list, before it starts, not after.

Measurement starts now, and it starts on the system I already run, not on the new build. Nothing has to exist first: the logging is a small addition to the mock that is working today, recording for each session what memory was cued, what was read, and what the answer used. The AI adds the logging the same way it built everything else; I direct it and verify the output, so the hands-off property survives, and instrumenting the mock does not break it. I need four weeks of collection running before the numbers mean anything, so this is the one thing that cannot slip. Evidence accrues while I read and design; by November I can put a denominator on how often an answer skipped something the tree already held and how often a stale summary got served. Deciding that an answer skipped something is a judgment I make per logged session, and that judging is hours of my own labor, counted in the budget. Week one also carries three cheap checks that can each move the design: write the undo for a merge on a small hand-built set of facts, measure how often a small local model fails to return valid structured output across a hundred calls, and label thirty answers twice, a week apart, to see whether I even agree with myself. If any of those fails, the design moves now, cheaply. The paper stands on those numbers.

For testing, I think literature works well, because a novel gives me both things my system keeps: the story as it happened, and the facts about people and places that the story changes over time. It is like episodic memory combined with a memory tree: the record of what happened, and the knowledge of what is, kept as two representations of one corpus. A novel supplies both kinds of information and known right answers for each. That is the spring corpus; the fall's only test bed is my own archive. Working out how systems like this are measured, and by what, is what I need the reading for; taking the tree from works-for-me to validated and tested is the year's work, not the fall's.

The fall commitment: the logging and counters on my current system; one comparison on my own archive, with the tree and without it, scored on the facts the record corrected over time. That comparison

is thirty to forty questions, written and scored by me, which costs real hours I am budgeting for, and the metric definitions go to Dr. Fang before I collect anything, so the standard is agreed before the numbers exist. If the week-one check shows I cannot agree with my own labels, the metric is dead and I will say so in the paper. The commitment also carries a synthesis memo from the reading, and a design document that settles the pipeline on paper, including the one step the mock never defined at all: how a long document or a long chat gets cut into pieces before anything else happens. The mock just says “by topic” and I never made it say more. The paper posts by November 15, because the application window follows. Past that, the build proceeds in order as far as the hours allow, ingest and organize first; everything beyond it is stretch, not promise.

One constraint stands regardless of direction: the difference between the mock and the real thing. The mock was built without me examining the execution at all: I told the AI to build it, used the result, and told it to adjust when it was not working the way I wanted. There are about 2,300 lines of code in the current solution and I have not looked at a single line. That was the right way to prove the idea, and it is the wrong way to build an actual solution. For the real system, I can use published tools, but I cannot just vibe code this: the load-bearing parts I build and defend myself, and plumbing can come from libraries. The other part of the reason is that AI-built solutions tend toward complexity: left to direct itself, the tooling grows and never shrinks, and I want this system small enough to defend.

Reading list

The Phase 1 working list, assistant-compiled from its survey with citations verified against source pages, for Justin’s review. Each entry except the nine marked *library* names its reference copy in `papers/` and its one-pager in `summaries/one-pagers/` by slug; *library* entries have neither yet (the digest covers them from abstracts) and carry DOIs in `papers/MANIFEST.md` for an ISU library pull.

The advisor-meeting record of 2026-08-19 keeps the original 44 items with fuller annotations; this working index supersedes it and adds the ten knowledge-graph items requested on 2026-08-25.

Start here

1. Kumaran, Hassabis, McClelland 2016. What Learning Systems do Intelligent Agents Need? Trends in Cognitive Sciences 20(7). *library* · kumaran2016-cls-updated
2. Park et al. 2023. Generative Agents. UIST 2023. arXiv:2304.03442 · park2023-generative-agents
3. Sarthi et al. 2024. RAPTOR. ICLR 2024. arXiv:2401.18059 · sarthi2024-raptor
4. Rasmussen et al. 2025. Zep: A Temporal Knowledge Graph Architecture for Agent Memory. arXiv:2501.13956 · rasmussen2025-zep
5. Edge et al. 2024. From Local to Global: A Graph RAG Approach. arXiv:2404.16130 · edge2024-graphrag

Consolidation and agent memory

1. McClelland, McNaughton, O’Reilly 1995. Psychological Review 102(3). *library* · mcclelland1995-cls
2. Teyler, Rudy 2007. The hippocampal indexing theory. Hippocampus 17(12). *library* · teyler2007-hippocampal-index
3. Tononi, Cirelli 2014. Sleep and the Price of Plasticity. Neuron 81(1). *library* · tononi2014-shy
4. Zacks et al. 2007. Event perception. Psychological Bulletin 133(2). *library* · zacks2007-event-perception
5. Summers et al. 2024. Cognitive Architectures for Language Agents. TMLR. arXiv:2309.02427 · summers2024-coala
6. Rezazadeh et al. 2025. MemTree. ICLR 2025. arXiv:2410.14052 · rezazadeh2025-memtree
7. Talebirad et al. 2026. Toward a Theory of Hierarchical Memory. arXiv:2603.21564 · talebirad2026-hierarchical-theory
8. Wu et al. 2021. Recursively Summarizing Books with Human Feedback. arXiv:2109.10862 · wu2021-recursive-books

Retrieval, vector search, caching

1. Wang et al. 2024. Searching for Best Practices in RAG. EMNLP 2024. arXiv:2407.01219 · wang2024-rag-best-practices
2. Douze et al. 2024. The Faiss library. arXiv:2401.08281 · douze2024-faiss
3. Kuffo, Krippner, Boncz 2025. PDX. SIGMOD 2025. arXiv:2503.04422 · kuffo2025-pdx
4. Malkov, Yashunin 2016. HNSW graphs. TPAMI. arXiv:1603.09320 · malkov2016-hnsw
5. Chan et al. 2025. Don't Do RAG: Cache-Augmented Generation. WWW 2025. arXiv:2412.15605 · chan2025-cag
6. Modarressi et al. 2025. NoLiMa. ICML 2025. arXiv:2502.05167 · nolima2025-long-context
7. Jeong et al. 2024. Adaptive-RAG. NAACL 2024. arXiv:2403.14403 · jeong2024-adaptive-rag

Benchmarks and faithfulness

1. Maharana et al. 2024. LoCoMo. ACL 2024. arXiv:2402.17753 · maharana2024-locomo
2. Hu, Wang, McAuley 2025. MemoryAgentBench. arXiv:2507.05257 · wu2025-memoryagentbench
3. Chang et al. 2024. BoookScore. ICLR 2024. arXiv:2310.00785 · chang2024-boookscore
4. Kim et al. 2024. FABLES. COLM 2024. arXiv:2404.01261 · kim2024-fables

Knowledge graphs: foundations

1. Chen 1976. The Entity-Relationship Model. ACM TODS 1(1). *library* · chen1976-er-model
2. Angles et al. 2017. Foundations of Modern Query Languages for Graph Databases. ACM CSUR 50(5). arXiv:1610.06264 · angles2017-graph-query-foundations
3. Hogan et al. 2021. Knowledge Graphs. ACM CSUR 54(4). arXiv:2003.02320 · hogan2021-knowledge-graphs
4. Snodgrass 1999. Developing Time-Oriented Database Applications in SQL. Morgan Kaufmann, author-released PDF · snodgrass1999-time-oriented-db
5. Weikum et al. 2021. Machine Knowledge. FnT Databases. arXiv:2009.11564 · weikum2021-machine-knowledge
6. Pan et al. 2024. Unifying Large Language Models and Knowledge Graphs. IEEE TKDE. arXiv:2306.08302 · pan2024-unifying-llm-kg
7. Rossi et al. 2021. KG Embedding for Link Prediction. ACM TKDD. arXiv:2002.00819 · rossi2021-kg-embedding

Knowledge graphs: added 2026-08-25

Curated to ten from a verified 28-item survey, weighted toward representation and current practice; the remainder are listed at the bottom. Knowledge graphs are one candidate representation, not a commitment.

1. Angles 2018. The Property Graph Database Model. AMW 2018, CEUR Vol-2100. What a property graph formally is; the model closest to a tree of pages with typed links · angles2018-property-graph-model
2. Hernandez, Hogan, Krotzsch 2015. Reifying RDF: What Works Well With Wikidata? SSWS at ISWC 2015. The four ways to represent qualified, non-binary facts, compared on real data · hernandez2015-reifying-rdf-wikidata
3. Cai et al. 2024. A Survey on Temporal Knowledge Graph. arXiv:2403.04782. Where facts-valid-for-an-interval lives; representation chapters, not the embedding chapters · cai2024-temporal-kg-survey
4. Rost et al. 2021. Bitemporal Property Graphs to Organize Evolving Systems. arXiv:2111.13499. Valid time and transaction time on every node and edge, so corrections never erase history · rost2021-bitemporal-property-graphs
5. Fowler 2021. Bitemporal History. martinowler.com. The practitioner bridge to the same idea. HTML · fowler2021-bitemporal-history
6. Vrandečić, Krotzsch 2014. Wikidata. CACM 57(10). The largest deployed store of time-scoped, source-attributed, supersedable facts · vrandečić2014-wikidata
7. Zhong et al. 2024. A Comprehensive Survey on Automatic Knowledge Graph Construction. ACM CSUR 56(4). arXiv:2302.05019. The construction pipeline end to end · zhong2023-kg-construction-survey
8. Noy et al. 2019. Industry-scale Knowledge Graphs: Lessons and Challenges. CACM 62(8). How Google, Microsoft, Facebook, eBay and IBM actually run them. Free HTML at queue.acm.org/detail.cfm?id=3332266 · noy2019-industry-scale-kgs
9. Xu et al. 2024. RAG with Knowledge Graphs for Customer Service Question Answering. SIGIR 2024 industry track. arXiv:2404.17723. LinkedIn’s deployed graph, directly on the help-desk use case · xu2024-kg-rag-customer-service
10. Balog, Kenter 2019. Personal Knowledge Graphs: A Research Agenda. ICTIR 2019. The academic niche closest to the existing system · balog2019-personal-kg-agenda

Conversation mining and outcomes

1. Kecht et al. 2021. Event Log Construction from Customer Service Conversations. ICPM 2021 · kecht2021-event-log-nli
2. Gung et al. 2023. Intent Induction from Conversations, DSTC 11. SIGDIAL 2023. arXiv:2304.12982 · gung2023-dstc11-intent

3. De Raedt et al. 2023. IDAS: Intent Discovery with Abstractive Summarization. NLP4ConvAI at ACL 2023 · deraedt2023-idas
4. Wegmann et al. 2022. Same Author or Just Same Topic? RepL4NLP 2022. arXiv:2204.04907 · wegmanna2022-style-content
5. Brynjolfsson, Li, Raymond 2025. Generative AI at Work. QJE 140(2). *library* · brynjolfsson2025-genai-at-work
6. Ouyang et al. 2026. ReasoningBank. ICLR 2026. arXiv:2509.25140 · ouyang2026-reasoningbank
7. Liu, Zhang, Choi 2025. User Feedback in Human-LLM Dialogues. EMNLP 2025. arXiv:2507.23158 · liu2025-user-feedback
8. Rezazadeh et al. 2025. Collaborative Memory. arXiv:2505.18279 · rezazadeh2025-collaborative-memory

Shared state across sessions

1. Erman, Hayes-Roth, Lesser, Reddy 1980. The Hearsay-II Speech-Understanding System. ACM CSUR 12(2). *library* · erman1980-hearsay-ii
2. Hayes-Roth 1985. A Blackboard Architecture for Control. Artificial Intelligence 26(3). *library* · hayesroth1985-blackboard-control
3. Salemi et al. 2026. LLM-Based Multi-Agent Blackboard System. arXiv:2510.01285 · salemi2026-blackboard-llm
4. Pollertlam, Kornsuwannawit 2026. Beyond the Context Window. arXiv:2603.04814 · pollertlam2026-beyond-context-window

Build references: added 2026-08-25 for the harness

Twenty-two sources surveyed and verified for the six gaps the harness build opened. These are stage references, consulted when building the stage they govern, not read-through material. Grouping matches the pipeline.

Contracts and cheap models: 54. Willard, Louf 2023. Efficient Guided Generation. arXiv:2307.09702 · willard2023-guided-generation / 55. Geng et al. 2023. Grammar-Constrained Decoding. EMNLP 2023 · geng2023-grammar-constrained / 56. Tam et al. 2024. Let Me Speak Freely? EMNLP 2024 Industry. arXiv:2408.02442 · tam2024-format-restrictions / 57. Chen, Zaharia, Zou 2023. FrugalGPT. arXiv:2305.05176 · chen2023-frugalgpt / 58. Ong et al. 2025. RouteLLM. ICLR 2025. arXiv:2406.18665 · ong2024-routellm / 59. Zhou et al. 2024. UniversalNER. ICLR 2024. arXiv:2308.03279 · zhou2024-universalner

Segmentation and extraction: 60. Hearst 1997. TextTiling. CL 23(1) · hearst1997-texttiling / 61. Koshorek et al. 2018. Text Segmentation as Supervised Learning. NAACL 2018 · koshorek2018-supervised-segmentation / 62. Xing, Carenini 2021. Dialogue Topic Segmentation. SIGDIAL 2021 · xing2021-dialogue-segmentation / 63. Niklaus et al. 2018. OpenIE Survey. COLING 2018 · niklaus2018-openie-

survey / 64. Zhang, Soh 2024. Extract, Define, Canonicalize. EMNLP 2024 · zhang2024-edc-kg-construction / 65. Caufield et al. 2024. SPIRES. Bioinformatics 40(3) · caufield2024-spires

Judging and agreement: 66. Zheng et al. 2023. Judging LLM-as-a-Judge. NeurIPS 2023 · zheng2023-llm-as-judge / 67. Wang et al. 2023. LLMs are not Fair Evaluators. ACL 2024. arXiv:2305.17926 · wang2023-not-fair-evaluators / 68. Panickssery, Bowman, Feng 2024. Self-Preference in LLM Evaluators. NeurIPS 2024 · panickssery2024-self-preference / 69. Kamaloo et al. 2023. Evaluating Open-Domain QA. ACL 2023 · kamaloo2023-open-qa-eval / 70. Artstein, Poesio 2008. Inter-Coder Agreement. CL 34(4) · artstein2008-intercoder-agreement

Narrative ground truth and injection: 71. Kocisky et al. 2018. NarrativeQA. TACL · kocisky2017-narrativeqa / 72. Pang et al. 2022. QuALITY. NAACL 2022 · pang2021-quality / 73. Wang et al. 2024. NovelQA. arXiv:2403.12766 · wang2024-novelqa / 74. Liu et al. 2023. Lost in the Middle. TACL 2024. arXiv:2307.03172 · liu2023-lost-in-the-middle / 75. Jiang et al. 2023. LongLLMLingua. ACL 2024. arXiv:2310.06839 · jiang2023-longllmlingua

Memory foundations: 76a. Tulving 1972. Episodic and semantic memory. In Tulving and Donaldson (eds.), Organization of Memory, Academic Press. *library*. The original episodic-versus-semantic distinction; the classical root of the dual representation (tree as consolidated episodic, facts as semantic) · no slug, book chapter

Coreference: 76. Lee, He, Lewis, Zettlemoyer 2017. End-to-end Neural Coreference Resolution. EMNLP 2017. arXiv:1707.07045. The canonical treatment of the pronoun problem named in the proposal; in the harness this capability is bought (a model call), not built, and this is the reference for what the task is · lee2017-e2e-coreference

Cut from the knowledge-graph survey, available on request

Named Graphs (Carroll et al., WWW 2005); RDF-star foundations (Hartig 2017); OWL 2 Profiles tutorial (Krotzsch 2012); the OneGraph vision paper (Lassila et al. 2023); LLM information-extraction survey; LLM-empowered KG construction survey (Bian 2025); entity-resolution overview (Christophides et al.); Open IE survey; Knowledge Vault (Dong et al. 2014); AutoKnow (Amazon, KDD 2020); GQL and SQL/PGQ pattern matching (SIGMOD 2023); a practitioner piece on why KG projects fail; the PKG ecosystem survey and PKG API tool paper; text-lifelog knowledge-base construction; LifeGraph; PIMO.

REFERENCE, not superseded. This is the field organised by problem. Nothing supersedes it. It is background rather than part of the working set. Index: docs/README.md.

The field, by problem

Assembled with LLM assistance; citations verified against sources. Three works (Brynjolfsson, Li, and Raymond; Erman and colleagues; Hayes-Roth) are characterized from abstracts only.

The literature on persistent memory for language systems is organized here by problem rather than by method or year: eight sections, each tracing one question, and a coda of recurring observations.

Can a Context Window Substitute for Memory?

The simplest persistence is resending everything the model needs, every time. Wang and colleagues (EMNLP 2024) treated the retrieve-and-generate stack as swappable modules and searched greedily for the best method at each position; a small BERT classifier deciding whether to retrieve at all, at 0.95 accuracy, raised the pipeline’s average score from 0.428 to 0.443 while cutting latency from 16.41 to 11.58 seconds per query: the cheapest win was a gate, not a better retriever.

Jeong and colleagues pushed the gate further (NAACL 2024). Adaptive-RAG routes each question to no retrieval, one retrieval step, or an iterative retrieve-and-reason loop, using a 770M-parameter T5 classifier trained on self-generated labels, the cheapest strategy that succeeded on each training query. On GPT-3.5-Turbo-Instruct the router matched the always-iterate baseline (50.91 versus 50.87 F1) at a third of its retrieval steps, spanning 0.35 seconds per no-retrieval query to 27.18 for multi-step. The classifier is only 54.52 percent accurate and an oracle router would reach 62.80 F1, so routing quality, not the strategies, is the current ceiling.

The opposite temptation is to preload everything. Chan, Chen, Cheng, and Huang’s cache-augmented generation (2025) encodes the corpus once into the model’s key-value cache, persists it, and appends only the query at inference. On corpora of 21k to 85k tokens with Llama 3.1 8B, CAG beat sparse and dense RAG in most configurations and answered in 0.85 to 2.26 seconds where re-encoding per query took up to 92.1 seconds; at the largest HotPotQA setting it fell below sparse retrieval, and the authors call the method impractical beyond small collections.

Windows are also far shorter than their specifications. Modarressi and colleagues’ NoLiMa benchmark (ICML 2025) showed needle-in-a-haystack tests had been measuring pattern matching: its questions share almost no words with the hidden fact (ROUGE precision 0.069 against 0.905 for vanilla NIAH). There, 11 of 13 models claiming 128K contexts fell below half their base score by 32K tokens; by the 85-percent criterion most effective lengths are 2K or less, with GPT-4o best at 8K. The same Llama 3.1 70B that reaches 32K effective length on RULER manages 2K here, isolating the lexical crutch. A preloaded 85k-token CAG cache thus sits well inside the region where associative recall has collapsed.

Pollertlam and Kornsuwannawit priced the trade in 2026, comparing full-history long-context inference against a Memo-style fact-extraction store. Long-context GPT-5-mini won factual recall by 33 to 35 points on LoCoMo and LongMemEval, while extraction compressed a 101,600-token history roughly 35:1 and became cheaper after about ten turns even with a 90 percent caching discount granted to the long-context side. What compression loses: temporal markers, coreferences, and updates that never propagate. Accuracy favors resending everything; economics favors extraction; neither wins both.

Unsettled: NoLiMa’s collapse is measured on synthetic needles in book text, not dialogue; the pricing study covers flat extraction, not structured designs; and no work here evaluates a live, incrementally updated store rather than a frozen snapshot.

When Does a Memory Need an Index, and What Does One Cost?

A memory storing more than it can reread must find nearest vectors without scanning everything. Malkov and Yashunin’s Hierarchical Navigable Small World graphs (HNSW, 2016) are the modern base answer: elements insert one at a time, each assigned a maximum layer from an exponentially decaying distribution; search greedily descends from the sparse top layer, a probabilistic skip list with proximity graphs for linked lists; a neighbor-selection heuristic keeping only candidates closer to the new element than to any already-chosen neighbor preserves connectivity across cluster boundaries. HNSW beat Annoy, FLANN, FALCONN, and VP-tree by wide margins, with nearly three orders of magnitude fewer distance computations than the flat NSW graph on one non-metric benchmark. Insertion is genuinely incremental, suiting an append-only history, and the index handles arbitrary metrics; but the paper defers element update and deletion to future work, and its complexity claims were verified only at low dimension, the check at $d=128$ needing infeasibly large datasets.

Douze and colleagues’ 2024 account of Faiss, Meta’s library packaging HNSW alongside a dozen other index types, gives the deployment rules: graph indexes win below roughly one million vectors when memory is unconstrained, inverted-file indexes take over once compression must fit RAM, and graph quality peaks at 64 edges per node. On a 200-million-vector SIFT subset, HNSW beat Faiss product quantization on recall and build time (42 minutes against 11 to 12 hours) while paying in RAM (64 Gb against under 30). Faiss stores no metadata, filters only through id-bit tricks, and handles deletion poorly in its graph indexes, hence the database layer Milvus and Pinecone wrap around it. Its sizing logic puts the threshold where indexing beats brute force near ten thousand vectors, far above any personal corpus.

Kuffo, Krippner, and Boncz sharpened the point at SIGMOD 2025 with PDX, a layout partitioning dimensions across blocks of 64 vectors so that pruning algorithms, which stop reading a vector once a partial distance rules it out, finally beat plain SIMD scans: ADSampling swung from losing to a linear scan to running 4.6x faster, and the exact-search variant PDX-BOND beat Faiss, USearch, and Milvus by 2.5x to 4x with no preprocessing. Retrieval cost is partly a systems problem, and at personal scale exact search stays feasible, sidestepping recall tuning entirely.

When retrieval helps the model is a separate question; Wang and colleagues’ search adds: hybrid sparse-plus-dense retrieval with HyDE gave the best retrieval quality, query rewriting did not help, and placing

the most relevant document nearest the query beat other orderings. Their chunk-size finding rests on sixty pages of a single financial filing, so the recipes are priors, not law.

The admitted limits: HNSW’s missing deletion and update are precisely the operations a maintained memory tree performs, and its logarithmic scaling was never verified at embedding-scale dimensionality; PDX’s results stop at ten million vectors, single-threaded and in memory, graph indexes untested; and the transfer of any recipe to a personal store remains unproven.

What Must an Agent Write Down, and When?

The earliest demonstration that writing can be recursive predates the agents literature. Wu and colleagues at OpenAI showed in 2021 that GPT-3 could summarize entire novels by fixed decomposition: summarize chunks, then concatenations of summaries, to depth 3 for books averaging over 100,000 words. The framing was scalable oversight (locally judged subtasks made each data point over 50 times cheaper than end-to-end collection); the durable finding is that a summary tree compresses an arbitrarily long corpus while keeping facts traceable to source, with depth-1 summaries alone letting a zero-shot reader score competitively on NarrativeQA. The failure modes proved as durable: local context turned a request for a dance in *Pride and Prejudice* into a marriage proposal, book-level output read as event lists, and themes spread thinly across chunks vanished. Any tree distilling leaves independently inherits these risks, so cross-leaf context and audits at composition points are design requirements.

Park and colleagues brought consolidation into agents in 2023. Their generative agents record every experience as a timestamped natural-language object, retrieve by an equally weighted sum of recency, importance, and embedding relevance, and, when summed importance of recent events crosses 150, run reflection: the model poses salient questions, retrieves against each, and writes cited higher-level inferences back into the stream. With 100 raters ranking interview responses, ablations cost believability stepwise, from a TrueSkill μ of 29.89 for the full architecture to 21.21 with no memory, a standardized effect of 8.16 for the fully ablated condition: a flat log plus scored retrieval was not enough. The result holds at 25 agents over two simulated days; planted false memories remain open.

Sumers, Yao, Narasimhan, and Griffiths supplied the organizing frame in 2024. CoALA reads language agents as heirs to production systems and architectures like Soar, describing any agent by memory stores (working, episodic, semantic, procedural), action space, and decision cycle: Park’s reflection becomes a learning action moving content from episodic to semantic memory, and Voyager writes code skills into procedural memory. Its diagnosis, as of a 2022 to 2023 snapshot: almost no agent treated learning as a deliberately chosen action rather than a fixed schedule, deletion and modification were neglected entirely, and writing to procedural memory is the riskiest form of learning.

Rezazadeh and colleagues answered part of that diagnosis in 2025 with MemTree, which makes the write path structural and continuous: new content routes down a tree by embedding similarity against depth-adaptive thresholds, every ancestor’s summary rewritten to fold it in, at $O(\log N)$ per insertion. Being online arguably outweighs the choice of structure: RAPTOR and GraphRAG require full rebuilds to absorb new material, while MemTree pays about 1.4 times their cumulative cost to stay current, roughly matching RAPTOR on QuALITY, beating GraphRAG on MultiHop RAG, and beating full-history

prompting outright on its 70-session conversational extension. The cost of perpetual rewriting is measured drift: ancestor content depends on insertion order and favors recent material.

Talebirad and colleagues generalized in 2026: these systems, from RAPTOR to conversational memories to execution-trace managers, implement one pipeline of extraction, coarsening, and traversal, and each lossy coarsening level strictly decreases entropy, so the atom layer caps what any summary can recover. Their coupling claim, that self-sufficient summaries support collapsed search while referential labels demand top-down refinement, would explain why MemTree’s retrieval works, but lacks a controlled test. The theory also assumes a static hierarchy; insertion, restructuring, and decay break its central argument, their named most pressing open problem, and MemTree’s drift shows the gap is real. Beneath both sits CoALA’s observation that principled forgetting remains nearly unstudied.

What Does a Tower of Summaries Actually Know?

Wu and colleagues’ recursion left two facts every later hierarchy inherits: compression works (the 175B models set state of the art on BookSum), and compression lies in characteristic ways, training on full trees making models worse as lower-level errors compounded upward.

The next wave measured the lies. Chang, Lo, Goyal, and Iyyer built BoookScore in 2024 on a contamination-resistant corpus of 100 recent books, deriving eight coherence error types from 1,193 human annotations and automating them as an LLM-judged metric. Hierarchical merging is almost always more coherent than incremental updating, the rolling-summary pattern most memory systems use, though Claude 2 with an 88K-token window, needing fewer update-and-compress steps, largely erased the gap; annotators still preferred incremental summaries for detail by 83 to 11 percent, so coherence and completeness pull apart. Kim and colleagues pressed harder in 2024 with FABLES, hiring readers who already knew 26 recent novels to verify 3,158 atomic claims from hierarchically merged summaries. Even the best model, Claude-3-Opus, produced only 90.9 percent faithful claims; unfaithful claims clustered around character and relationship states; and verification itself failed, the strongest auto-rater reaching 58.2 F1 at catching fabrications with the whole book in hand, retrieval-based checking far worse. BoookScore finds omission dominating its errors and FABLES finds 33 to 65 percent of summaries omitting key events: the error a summary’s reader cannot see.

Meanwhile the tree was repurposed from output to index. Sarthi and colleagues’ RAPTOR (ICLR 2024) clusters chunk embeddings, summarizes each cluster, re-embeds, recurses, then at query time collapses the tree so nodes from every level compete in one similarity search under a token budget; clustering on embeddings rather than position lets distant passages share a parent, unlike Wu’s positional recursion. The hierarchy did the work: on one QuALITY story, full-tree search beat any single layer by some fifteen points, and the headline QuALITY result with GPT-4 was 82.6 percent against a prior best of 62.3. Edge and colleagues’ GraphRAG (2024) took the complementary path: extract entities and relations into a knowledge graph, partition it with hierarchical Leiden, pre-write summaries for every community at every level. For global sensemaking over million-token corpora, root-level summaries beat vector RAG on comprehensiveness while using 97 percent fewer context tokens than map-reduce summarization, at a

price of 281 minutes of GPT-4-turbo indexing per million tokens. The load-bearing finding is what nested pre-written summaries buy.

Talebirad and colleagues’ theory lands hardest here, because these trees are its main exhibits: self-sufficiency (the fraction of a group’s information its representative preserves) gives the coupling claim its terms, and a derived depth bound predicts when one summary layer stops sufficing as a corpus grows.

Two questions stay open. Every system here builds its hierarchy in one batch, and RAPTOR, GraphRAG, and the theory paper all name incremental insertion and restructuring as unsolved. And if no automatic method reliably catches what a summary fabricates or omits, the upper layers of any memory tree become a record no one can cheaply check.

How Does a Graph Store a Fact That Was True Then, Learned Later, and Corrected Since?

The temporal half was answered before the graph half existed. Snodgrass distilled roughly 1600 research papers into a 1999 book for SQL developers whose central distinction endures (Fowler 2021; Rost and colleagues 2021): valid time is when a fact held in the modeled world, transaction time is when the database learned it, and a table carrying both is bitemporal. His Danish mortgage bank case (Nykredit, nine million loans) shows why both clocks matter: staff repair reported bad data by rolling back in transaction time and reading the valid-time history as it then stood, the repair itself logged because transaction time is append-only. Fowler restated the pattern in 2021 with the terms actual and record: once retroactive change is allowed, actual history stops being append-only but record history never does, preserving the audit trail; bitemporality complicates a system significantly and can be skipped when retroactive change is forbidden.

The graph half took shape independently. Angles and colleagues surveyed graph query languages in 2017, formalizing the property graph (labels and attribute maps on nodes and edges) and cataloging which features are cheap: fixed small patterns over a growing graph sit in logarithmic-space data complexity, while path queries under simple-path semantics turn NP-complete. Angles added a Codd-style formal definition in 2018, no standard specification existing despite Neo4j and Neptune’s commercial dominance. Hogan and sixteen coauthors unified the field in 2021: graphs let schema definition wait while data and scope evolve, and their context machinery (temporal validity, provenance, geographic scope) is the vocabulary for statements true only at a time or per a source.

Attaching that context to a fact is where the engineering fights happened. Vrandecic and Krotzsch described Wikidata in 2014: statements carry qualifiers (“as of 2010”), references, and preferred or deprecated ranks that let contradictory claims coexist. Hernandez, Hogan, and Krotzsch tested in 2015 how qualified statements survive translation into plain RDF, comparing four reification schemes over 57 million Wikidata quads on five engines: singleton properties, the most concise encoding, broke four of the five engines because millions of unique predicates violated indexing assumptions, while named graphs, n-ary relations, and standard reification split the wins with no clear victor. Rost and colleagues joined the halves in 2021 with a bitemporal property graph carrying both Snodgrass intervals on every

vertex, edge, and individual property value, superseding their earlier TPGM, which had to copy an entire element when one property changed; the system remains a prototype with no benchmarks.

Extraction from text is its own literature. Weikum, Dong, Razniewski, and Suchanek set out the four-phase pipeline in 2021 (discovery, canonicalization, augmentation, cleaning): entities must be canonicalized before relations mean anything, and construction is a life-cycle with curators permanently in the loop. Zhong and colleagues surveyed over 300 construction methods in 2023, fixing the vocabulary of entity linking, coreference, and relation extraction, coverage stopping at BERT-era models; Pan and colleagues extended the map to LLM-based construction in 2024. Cai and colleagues surveyed temporal knowledge graph embeddings in 2024, though Rossi and colleagues had shown in 2021 that a cheap rule-mining baseline matches embedding models, which perform worst exactly in the sparse, thin-support regime a personal graph occupies. Rasmussen and colleagues assembled the stack in 2025: Zep builds a bitemporal graph continuously from conversation, four timestamps per fact edge, contradiction handled by invalidation rather than overwrite, beating full-context baselines on 115k-token conversations by 15.2 to 18.5 percent while conceding that below context-window scale the graph buys little.

The named gaps: the personal knowledge base built from one user’s longitudinal footprint (Weikum); entity resolution, with no benchmark yet testing the fusion of conversational and structured data (Rasmussen); temporal-graph methods evaluated only on benchmarks far smaller than real deployments (Cai).

Does the Knowledge Graph Survive Contact with Real Data, and Does It Shrink to One Person?

Wikidata is also the longest-running public deployment of the layered-statement design, presented by Vrandečić and Krotzsch two years after its October 2012 launch as the answer to a duplication problem: the population of Rome lived in hundreds of independently edited Wikipedia articles and disagreed across them. The layered statement carried the load, down to distinguishing “value unknown” and “no value” from mere absence. By August 2014 the system held 15.8 million items and 43.2 million statements, 23.2 million referenced, about 90 percent of its half million daily edits coming from bots under human curation; the numbers documented adoption, not quality, and complex query support did not yet exist.

Noy and colleagues (2019) pooled the experience of five production graphs at Google, Microsoft, Facebook, eBay, and IBM, ranging from Facebook’s 50 million entities to Google’s roughly 1 billion entities and 70 billion assertions. Facebook keeps provenance on every assertion and defines correctness as always being able to explain why an assertion was made; IBM treats source evidence as a primitive and defers entity resolution to query time. The sharpest finding is negative: entity disambiguation remained the top industry challenge despite years of research, Bing’s single entry for the actor Will Smith assembled from 108,000 facts across 41 websites against 200 Will Smiths in Wikipedia alone. The paper is experience: self-reported mechanisms, no benchmarks.

The same year, Balog and Kenter argued that shrinking this machinery to one person is a different research object, not a scaling exercise. Their personal knowledge graph has a spiderweb topology, every node connected directly or indirectly to a central user node, and three properties that invert the industrial assumptions: entities need not be prominent (“my guitar”, “Mom’s dentist”), entity information can be radically sparse, and relations are often short-lived, the opposite of the stability criterion general graphs use for inclusion. Data-hungry neural link prediction is unlikely to transfer, and linking “Jamie” with no external profile entangles linking, population, and NIL-detection at once. There is no system behind the paper, no open PKG dataset existed, and synthetic data may be the only path. Noy and colleagues independently flagged personalized on-device graphs as open: by 2019 industry and academia were pointing at the same gap from opposite ends.

Xu and colleagues supplied deployed evidence in 2024, at organizational scale. Their LinkedIn customer-service system replaced flat-text RAG over Jira tickets with a two-level graph: each ticket parsed into a section tree (summary, root cause, steps to reproduce), trees linked by explicit Jira references and thresholded embedding similarity, retrieval returning a coherent subgraph scored only over nodes matching the query’s intent. With GPT-4 and E5 held fixed in both arms, MRR rose from 0.522 to 0.927 and Recall@3 from 0.640 to 1.000; in a randomized six-month split of the support team, median resolution time fell 28.6 percent. The golden dataset was small and internal, and the graph template was hand-built: every deployed success so far bought its structure with the human schema work Balog and Kenter said a personal graph cannot assume.

Still missing: disambiguation at scale, unsolved for the five largest graphs as of 2019 and strictly harder in the sparse, nickname-heavy personal case; a shared evaluation benchmark, no open personal knowledge graph dataset existing; and the automation Xu and colleagues name as future work, template extraction and dynamic updates, without which their result cannot transfer to a graph nobody curates by hand.

Can a Schema Be Pulled Out of Raw Dialogue, and How Would Success Be Measured?

Hand-built templates are what a conversational archive cannot assume. Whether structure can be induced from the transcripts themselves is the question these four works ask; the collective answer is a qualified yes, the qualifications doing most of the work.

Kecht, Egger, Kratsch, and Roglinger (2021) take the schema as given and apply it cheaply: process-mining event logs built from customer service conversations with no model training, each message the premise of a natural language inference pair, a candidate activity (“This example asks for iOS version”) the hypothesis, a pre-trained BART model returning an entailment probability. The economics are the finding: roughly 100 to 300 hand-labeled messages, used only to calibrate per-label thresholds, structured corpora of 88,000 to 288,000 tweets, 55 of 56 classifications clearing an MCC of 0.72. But hypothesis wording proved brittle (one rephrasing moved MCC from 0.53 to 0.91, and a grammatically broken hypothesis scored 0 where its fix scored 1), and the mined model covered only topics defined in

advance. For open-ended conversational memory that concession is the whole problem: the schema is exactly what is not known ahead of time.

Wegmann, Schraagen, and Nguyen (2022) supply a warning about what conversation-derived representations encode. Models trained on authorship verification can succeed by memorizing what an author writes about rather than how; when distractors were forced to come from the same conversation as the anchor, so topic could no longer separate authors, base RoBERTa picked style over content only 5 percent of the time on their STEL-Or-Content probe, while the conversation-controlled model reached 0.42. Any embedding over a conversation archive silently entangles speaker, topic, and phrasing; disentangling them requires hard negatives mined from conversation structure, and the general discipline is to build a probe where the shortcut and the target disagree rather than trusting the training metric.

Gung and colleagues at AWS confronted the missing-schema problem directly with the DSTC 11 track (2023) on inducing intents from raw call-center transcripts, arguing prior work had evaluated on repurposed classification datasets lacking full dialogues and realistic skew. The three simulated corpora ran 59 to 71 turns per conversation, and the open-induction subtask judged an induced schema not by cluster purity but by the accuracy of a classifier trained on it, a framing that transfers to any induced structure. Of thirty-four teams, the winner reached 69.8 accuracy on clustering against a 55.8 k-means baseline. Two analyses cut at the evaluation itself: propagating labels to one system's unassigned noise points moved it from 33rd place to 9th, and swapping the downstream classifier's encoder shuffled the middle of the leaderboard while the top ten held.

De Raedt and colleagues (2023) attacked the representation problem by another route. IDAS keeps the encoder frozen and moves disambiguation into text: an LLM abtractively summarizes each utterance into a short label, the labels are embedded alongside the utterances, and the combination is clustered. Every labeling strategy tried beat raw utterance encodings by 5 to 19 ARI points, and with zero labels IDAS beat semi-supervised methods given labels for half the intents on CLINC; the summary strips the syntax and noun overlap that otherwise drive similarity. Yet IDAS assumes the true cluster count is known, pays an LLM call per utterance, and found smaller models produced labels too poor to use.

Unfinished: the evaluation instruments are unstable, punishing abstention and wobbling with the downstream encoder; cluster count and per-utterance cost remain unsolved; Kecht's pipeline still needs the event vocabulary in advance. The question a memory system cares about most stands open: discovering a schema when neither the categories nor their number is given.

What Teaches a Memory, and Who Is Allowed to Read What It Learns?

A store of experience becomes a memory only when it changes in response to how things went. The candidate signals split into three lineages: outcome feedback, self-judged experience, and shared structured state.

The cleanest evidence that outcomes are a usable teacher comes from outside the memory literature: Brynjolfsson, Li and Raymond's 2025 field deployment of an AI assistant in customer support, measured

on resolution outcomes. Most conversational systems have no resolution ticket to close, so Liu, Zhang and Choi asked in 2025 whether the implicit feedback buried in chat logs could substitute, densely annotating 109 conversations from LMSYS and WildChat. Feedback appears in more than half of later user turns, but praise is rare, prompts that elicit positive feedback are slightly more toxic than random ones (users praising successful jailbreaks), and training on feedback-conditioned regenerations won on MTBench but lost on the harder WildBench, plain distillation beating both. Implicit feedback marks where intent went unmet; it cannot be consumed raw as a reward.

Ouyang and colleagues took the opposite bet in 2026 with ReasoningBank: drop human feedback entirely, let an LLM judge label each finished trajectory success or failure, and distill both into titled strategy items, successes as validated moves and failures as preventative lessons. On WebArena the loop lifted success from 40.5 to 48.8 percent, and the failure records earned their place: on the WebArena-Shopping ablation, success-only extraction reached 46.5 where adding failures reached 49.7, while the earlier success-only AWM pipeline degraded when fed failures its schema could not hold. The judge was only 72.7 percent accurate, yet varying simulated accuracy from 70 to 90 percent barely moved results, suggesting curation can run without labels at all. Maharana and colleagues’ 2024 LoCoMo finding points the same way, retrieval over distilled per-turn assertions beating both raw transcripts and session summaries: writing memory and reading it converge on small, self-contained, distilled units.

The shared-state lineage is much older: Erman, Hayes-Roth, Lesser and Reddy’s Hearsay-II blackboard (1980), independent knowledge sources coordinating solely through shared structured state, extended to control by Hayes-Roth in 1985. Salemi and colleagues revived the pattern for LLM agents in 2026, for data discovery over large file lakes: the main agent posts what it needs, and partition agents that have pre-digested their holdings volunteer only when capable. Against an otherwise identical master-slave design, self-selection won everywhere (file discovery F1 of 0.561 versus 0.513, and 0.247 for embedding retrieval), the advantage growing with lake size to a 211 percent relative gain on a 1,556-file lake. A coordinator’s index of who knows what goes stale; letting the holders of knowledge nominate themselves does not.

Once memory is shared, someone must be kept out of parts of it. Rezazadeh and colleagues proposed in 2025 stamping every memory fragment with immutable provenance (time, contributing user, agents, resources) and checking it against time-varying permission graphs at read time, so revoking an edge retroactively hides fragments that depended on it; shared memory also paid for itself as a cache, cutting external resource calls by up to 61 percent at high query overlap. This is a formulation with feasibility evidence, not yet a benchmark, and the LLM-based enforcement can still occasionally breach policy.

Three doubts remain. ReasoningBank consolidates by pure addition, no pruning or supersession, and its streams run hundreds of tasks, not years; Hu, Wang and McAuley’s MemoryAgentBench (2025) shows the gap is real, no tested system exceeding 28 percent on its fact-update tasks. Positive feedback encodes bad behavior as readily as good, leaving the polarity of the teaching signal in doubt. And access control enforced by the model being controlled has no hard guarantee.

Coda

Four observations recur. The field builds structures well and maintains them badly: HNSW deferred deletion in 2016, RAPTOR, GraphRAG, and Talebirad and colleagues’ theory name incremental insertion unsolved, MemTree measures its own drift, and no MemoryAgentBench system exceeded 28 percent at superseding stale facts. Small distilled units keep beating raw volume, from Maharana and colleagues’ per-turn assertions to De Raedt and colleagues’ labels to ReasoningBank’s strategy items, with Pollertlam and Kornsuwannawit’s pricing showing why and what the compression costs. Selection beats capacity: the retrieval gate, the router, the self-nominating partition agents, and the NoLiMa collapse all locate the gains in deciding whether and whom to ask rather than in asking harder. And the instruments trail the systems: Kim and colleagues’ best fabrication catcher reached 58.2 F1, Gung and colleagues’ leaderboard moved with the evaluation encoder, and the personal-graph dataset Balog and Kenter found missing in 2019 has not been reported closed by any later paper in this corpus.

Handoff: where things stand

Written 2026-08-28, replacing the 08-27 handoff. Read this first. It records what is settled, what was wrong, and what is still open. The document index is `README.md`.

Read this before opening anything

Check you are on the current clone. This session opened a checkout that was 19 commits behind origin and concluded, reasonably and wrongly, that the handoff documents had never been written. They were on origin the whole time. Two machines are in play and the maintenance script pushes without fetching, so a stale local checkout is the default failure mode here, not an unlikely one.

```
bash git -C ~/it494-memory-project fetch && git -C ~/it494-memory-project status -sb
```

If that shows you behind, pull before reading further. A handoff document is useless if the next session opens the wrong clone.

The plan of attack, locked

A working backend, a published dataset, measured cost and tier numbers, by November 15. The novelty claim this plan used to carry is gone; see the next section. The build below is unchanged by that, because none of it depended on the claim being true.

1. **Clear the administrative gate first.** The approved proposal on file is a survey of LLM and NLP tools on government documents. Everything in this repository is a different project, and nothing on record shows the topic change was ever approved. This blocks everything and is not an intellectual problem.
2. **Preprocess in September**, inside the first open block, and publish the cleaned corpora as a dataset. That publication is the reproducibility story.
3. **Pilot on Oz book 1 before September 27**, across three tiers, for about three dollars.
4. **Build October 19 to November 7.** Batch the cell calls from the start.
5. **Measure and write November 8 to 15.** The free instruments, per-stage tier sensitivity, the two corpus controls, the contamination probe, and LongMemEval as the one paid benchmark. The LitBank head to head is withdrawn.
6. **When a week disappears, cut from the bottom of the order in 05_fall_plan.md**, which is decided in advance precisely so it is not decided under pressure.

The full calendar, hour budget, cost model and cut order are in `05_fall_plan.md`.

The headline: there is no novelty claim left

Nine candidate contributions have now been searched adversarially. **All nine were occupied.** The last one fell on 08-28, after this pass had already rewritten the argument document around it.

The surviving claim was *pre-determined, unit-level entity salience driving extraction, versus per-chunk extraction*. It failed three ways, and the first is not about prior art:

1. **The contrast was factually wrong.** GraphRAG’s own appendix says its extraction prompt “first identifies all entities in the text ... before identifying all relationships.” Zep’s fact-extraction prompt takes an <ENTITIES> block and extracts facts pertaining to them. **Both baselines are already entity-first.** The only difference was the size of the unit the entity pass covers, and that is not a contribution.
2. **The mechanism is published three times over.** iText2KG (2409.03284) accumulates a Global Document Entities set and feeds it to relation extraction. RAKG (2504.09823) does document-wide disambiguation then per-entity relation construction, explicitly against GraphRAG’s ordering. LINK-KG (2510.26486) builds a global alias cache and rewrites every chunk against it.
3. **The ablation has been run, and it came out negative.** iText2KG compared global-cast against local-cast conditioning: global scored about **10 points lower** triplet precision.

And LitBank could not have measured it. Computed, with the parse validated by reproducing LitBank’s published 210,532 tokens and 29,103 mentions exactly: **96 of 100 documents are exactly two chunks** at GraphRAG’s default size. One boundary per document, against an effect that compounds across many. LitBank also has no relation annotation at all. BOOKCOREF (ACL 2025) exists precisely because LitBank is too short for book-scale work.

What this means, and what it does not

It does not mean the project is dead. It means the project is **an engineering project with measurements**, which is what the advisor steered toward twice, and which needs no novelty claim. What survives regardless of framing:

- A working backend, built by hand, defensible line by line.
- Four cleaned public-domain corpora published as a dataset. Nobody occupies that.
- Measured per-stage cost and tier sensitivity across a 146-fold provider spread.
- Instruments that need no gold data: quote gate, rejection rate per stage per tier, duplicate minting, predicate sprawl, plot-versus-cell consistency.

Do not invent a tenth claim. Every one of the nine died the same way, by asserting an absence without searching for it, and three of them were refuted by evidence already sitting in papers/. A fresh formulation with no search behind it would be the same error a tenth time.

One question is genuinely unsearched, and it is written down as a question, not a claim: iText2KG’s negative result was on short semantic blocks and LINK-KG’s gains came from a different mechanism, so whether cast conditioning helps or hurts **at book scale**, across a hundred-plus chunks, is unsettled. It would need its own adversarial search and a book-scale gold resource before anyone builds on it.

Read Narrative World Model (arXiv:2607.05577) in full before deciding anything. It is writer memory for long-form fiction, evaluated on a fiction corpus against Graphiti/Zep and GraphRAG with the reader held constant, and it isolates extraction quality from representation by rebuilding the baseline with its own extractor. That is this project’s niche, baselines and corpus type, published six weeks ago. It is currently an abstract-only read, and it is the single most important unread paper on the list.

What was wrong, and what it propagated into

The point of this list is not the individual errors. It is that each one was load-bearing somewhere else, and fixing the sentence does not fix the argument that rested on it.

Error	What it propagated into
Zep said to lack community summarisation, read from a README	The entire gap argument. If Zep already does hierarchical summarisation and incremental maintenance, “build the composition that covers the requirements” loses most of its force. The argument is now “nothing here is claimed as new,” which is honest and much smaller
The correction to it was itself overstated	“Extends dynamically without full refresh” became “delays but does not eliminate full refresh.” The weaker version is better for us: it means Zep also carries a batch refold
“Best single system covers five or six of nine”	Never computed. Deleted rather than repaired, because a defensible version of that coverage argument already exists across twelve systems (arXiv:2606.24775), and there are seven requirements, not nine
Covering versus partition asserted as unoccupied	Was the surviving secondary claim. RAPTOR soft-clusters by design and this repo’s own one-pager said so; CAM does incremental overlapping clustering. The design stays, the claim goes
MemoryAgentBench “28% on fact update”	Requirement 1’s urgency. It is the multi-hop ceiling only; single-hop reaches 78 to 100. Supersession is less obviously unsolved than the row implied

Error	What it propagated into
“33 to 65% omit key events” attributed to BoookScore	Requirement 4. It is FABLES. Both halves were true, the attribution was not
“RAPTOR, GraphRAG name incremental insertion as unsolved”	Requirement 5. Neither does. MemTree measured it and is the real source, which makes the row stronger than the version that was wrong
Requirement 7 as “no benchmark can express this”	An absence claim, and false. ENPMR-Bench is exactly that benchmark. Same failure mode, third instance
“Both baselines extract chunk-first”	The whole surviving claim, and the committed core measurement. Neither baseline extracts chunk-first: GraphRAG’s prompt does entities then relations, and Zep’s fact extraction is conditioned on a resolved entity list. This was in both papers, on disk, the entire time
Cost model: 70x spread, 1.5M batched output, 3x saving	The tier-sensitivity argument. Real figures are 146x, 1.29M (unchanged by batching), and 2.33x on total input, 5x on the dominant term

All seven of the previous handoff’s “unverified assumptions” are now discharged. Both Graphiti questions came back usable: `EpisodeType.text` ingests plain text, and `add_triplet` accepts pre-extracted entity instances, verified against source at commit 683a853. All four uncertain citations are resolved and citable.

Open, in priority order

1. **What the project is now, given that there is no novelty claim.** This is the decision to walk into the advisor meeting with, and it is not the assistant’s to make. See above.
2. **The administrative gate.** The approved proposal on file is still the government-document tool survey, and nothing on record shows the topic change was approved. Unresolved and blocking. See Phase 0 in `05_fall_plan.md`.
3. **Read Narrative World Model in full.** See above.
4. **data/clean/ does not meet the frozen contract and has no producing script.** It holds 351 records covering 16 of 81 works, committed on 08-27 as collateral inside an unrelated documentation commit. It has no `unit_type` and no `unit_id`, and uses `chapter_ordinal` where `04_unit_contract.md` says `unit_ordinal`. Nothing in the repository or its history produced it, so it is not reproducible. Treat it as a scratch spike and regenerate from a committed splitter.
5. **Two unrevoked tokens in pushed history.** Recorded in the risk register in `reference/00_design_brief.md` and again in `06_spring_plan.md`. Revoking is the fix; deleting the file is not, because the history still carries it. This is unrelated to the project’s argument and should be closed anyway.

6. `papers/MANIFEST.md` is **incomplete**, 72 rows against 89 PDFs, and missing all four sources that `07_references.md` now depends on. Its header is corrected but the table is not rebuilt.
 7. **The Consult logger** is a separate deliverable from switching consult-logging on, and only the latter is scheduled. The rate needs weeks of collection, so if it is going to happen at all it has to start early.
-

What this pass changed

`docs/` went from 45 markdown files to nine live ones, plus `reference/` (still valid, never superseded) and `archive/` (27 files, each carrying a banner naming its successor). Nothing was deleted.

Seven parallel audits re-verified every factual claim against full paper text, source code, or the publisher's page. The corrections are listed above and recorded inside the documents that carried the errors, rather than quietly patched, because the pattern is more useful than the individual fixes.

The build script shipped six superseded documents, missed every current one, and merged the *previous* run's table of contents, so every page number after the fourth entry was wrong. Fixed and verified: 201 pages, contiguous part numbering, and all 16 contents entries checked against real section starts.

`SOURCES.md` still claimed Apollodorus and Water Margin were unavailable when both are on disk. Corrected in place, with the wrong text struck through rather than removed, since they are the same false-absence failure the method rules exist to prevent.

Three of 92 papers are genuinely unreferenced and moved to `papers/_unused/` with their one-pagers. An initial pass said 23; that matched filename slugs only and produced a false positive for every paper cited by project name or arXiv identifier. The corpus is almost entirely in use.

The method rules, which are the actual deliverable

Every substantive error in this project has had the same shape: asserted without checking. These are mechanical and checkable, and they belong in front of the next session.

- **An absence claim requires a documented search.** Never write that nobody has done X without reporting how you looked. Three contributions died to this rule and all three had been asserted without one.
- **Run a control query before trusting an empty result set.**
- **State the source type for every factual claim:** full text, abstract, README, or memory.
- **Never emit a number you did not compute.**

The stronger version is structural rather than dispositional. Generate in one context and refute in another that has not been softened by the conversation. That is what the cold audit does in the prototype,

it is what the seven parallel audits did here, and it is the only thing that has reliably caught this class of error.

The paper’s argument

2026-08-27, revised 2026-08-28. Nine candidate contributions have now been searched adversarially and **all nine were occupied**. Seven fell on 08-26 and 08-27; covering versus partition fell on 08-28; the extraction-ordering claim, which this document was previously built around, fell later the same day.

This document no longer proposes a contribution. It records what was claimed, how each claim died, and what survives that needs no claim at all. Do not reuse an older draft of it: every version before 08-28 asserts things that are false.

There is no surviving novelty claim. Read this before anything else.

The claim this document was built around was **pre-determined, unit-level entity salience driving extraction, versus per-chunk extraction**. It was searched adversarially on 2026-08-28, the same way the other eight were, and it died the same way.

It failed on three independent grounds, and the first one is the serious one because it is not about prior art at all.

1. The contrast is factually wrong. Both baselines are already entity-first. This was checked against the full texts in papers/, not recalled.

- **GraphRAG.** Its element-instance-generation appendix describes “a multipart LLM prompt that first identifies all entities in the text, including their name, type, and description, before identifying all relationships between clearly related entities.” Entity pass, then relations. The paper also pre-empts the duplicate-minting objection directly: GraphRAG “is generally resilient to duplicate entities since duplicates are typically clustered together for summarization in subsequent steps.”
- **Zep.** Its fact-extraction prompt reads `<ENTITIES> {entities} </ENTITIES>`, followed by “Given the above MESSAGES and ENTITIES, extract all facts pertaining to the listed ENTITIES.” Entities are extracted, embedded, hybrid-searched against the existing graph and LLM-resolved **before** edge extraction. Zep is entity-first against a persistent global cast, not a four-message one.

So the difference was never the *order*. It was the *scope of the unit the entity pass covers*, and “run the entity pass over a bigger unit” is not a contribution.

2. The mechanism is published, repeatedly. iText2KG (arXiv:2409.03284) accumulates a “Global Document Entities” set and states that it is “provided as context along with each Semantic Block to the Incremental Relations Matcher,” for the same stated reason: unresolved entities produce redundant relations. RAKG (arXiv:2504.09823) runs per-chunk recognition, then document-wide disambiguation, then per-canonical-entity relation construction, explicitly framed against GraphRAG’s ordering, and

beats it on MINE. LINK-KG (arXiv:2510.26486) builds a global alias-to-canonical cache, rewrites every chunk against it, then runs GraphRAG.

3. The ablation has been run and it came out against us. iText2KG compared global-cast conditioning against local-cast conditioning on two datasets. Global-cast conditioning scored roughly **10 points lower** triplet precision: a richer graph carrying more irrelevant relations. The related line does report gains, but through pre-extraction *coreference resolution* rather than a cast list: LINK-KG cuts node duplication from 27.0% to 10.6% on short documents and 36.0% to 17.8% on long ones, and CORE-KG's ablation shows removing its coreference pass costs 28.25% more node duplication. That is a different mechanism, and it is occupied three times over.

And the measurement instrument does not work either. See the LitBank section below.

What follows from this

Do not repair the claim. Nine consecutive candidates have been searched and nine were occupied, and the failure mode every time was asserting an absence without searching for it. A tenth formulation invented at this point, without a search behind it, would be the same error again.

The decision now is whether there is a project left once the novelty claim is removed, and that decision is not the assistant's to make. It belongs to Justin and his advisor. What can be said is what survives regardless of framing:

- A working backend, built by hand, with a design that is defensible line by line.
- Four cleaned public-domain corpora published as a dataset, which is a real reproducibility contribution and is occupied by nobody.
- Measured per-stage cost and tier-sensitivity numbers across a 146-fold provider spread.
- A set of instruments that need no gold data: the quote gate, rejection rate per stage per tier, duplicate minting per unit, predicate sprawl, plot-versus-cell consistency.

That is an engineering project with measurements, which is what the advisor originally steered toward twice, and it needs no novelty claim to be worth a semester. Whether it also wants a paper, and what that paper claims, is the open question in `HANDOFF.md`.

One genuinely unsearched question survives, and it is recorded here as a question rather than a claim: iText2KG's negative result was measured on short semantic blocks, and LINK-KG's positive results came from a different mechanism. Whether cast conditioning helps or hurts **at book scale**, across a hundred-plus chunks rather than two, is not settled by either. That question would need its own adversarial search before anyone builds on it, and it would need a book-scale gold resource, which LitBank is not.

What the argument is, in four moves

1. The requirements are real, and the field says so, not the author.

Seven requirements, each sourced to a published paper stating the problem is unsolved. They are set out with their citations and their tests in `02_requirements_and_testing.md`.

This is **motivation**, **not contribution**, and it must be labelled that way in the paper. An earlier draft derived the requirements from a two-year personal deployment and validated them against the literature. That is a wish list with a case study attached, and a reviewer answers it with *says who?* The deployment corroborates in one sentence and does not authorise.

Labelling it as motivation is also now forced rather than merely prudent. *Are We Ready For An Agent-Native Memory System?* (arXiv:2606.24775, June 2026) decomposes agent memory into four modules, then evaluates twelve representative systems plus two baselines across five workloads spanning eleven datasets, with ablations on representation fidelity, retrieval precision, update correctness and long-horizon behaviour. That is the requirements-plus-measured-coverage move, executed empirically and at a scale this project will not match. Cite it as motivation. Do not compete with it silently.

2. Nothing in the architecture is claimed as new, and the borrowing is named.

Hierarchical summaries from GraphRAG and RAPTOR. Temporal facts and invalidation from Zep and, underneath it, bitemporal modelling. Community summarisation from GraphRAG and, as corrected below, from Zep as well. Per-document per-entity summarisation from EntSUM. Per-entity chronology from timeline summarisation. Online maintenance from MemTree and Graphiti. Property graph structure and the edge-typing constraint from Angles. Ranks and qualifiers from Wikidata.

Naming the borrowing is the strength. A paper that shows only what was built reads as naive; one that draws the line explicitly reads as someone who did the survey, which is what happened, adversarially, with the results recorded below.

3. Measure what the design decisions cost and buy.

The entity-first versus chunk-first ablation is **withdrawn**, for the three reasons at the top of this document and because LitBank cannot carry it. What is left is free and still worth reporting:

Measurement	Cost	Notes
Tier allocation per stage	free	Falls out of running anything; the cost arithmetic is in <code>05_fall_plan.md</code> . Asymmetric Capacity Allocation (arXiv:2608.21345) covers adjacent ground and must be cited
Duplicate minting per unit, as a descriptive curve	free	Falls out of ingest. Without gold and without a baseline arm it is a description of this pipeline's behaviour, not a comparison, and must be reported as such

Measurement	Cost	Notes
Rejection rate per stage per tier, predicate sprawl, quote gate pass rate	free	Contract instrumentation; no dataset needed
Plot-versus-cell consistency	free	Two independent summaries over identical text, compared as sets. No judge, no ground truth

LitBank cannot measure extraction ordering, and this was computed rather than estimated

The parse was validated by reproducing LitBank’s published totals exactly, 210,532 tokens and 29,103 mentions, before any of the numbers below were taken from it.

Per document (n=100)	mean	median	range
tokens	2,105	2,043	1,999 to 3,419
chunks at GraphRAG’s 1,200-token default	2.0	2	2 to 3
coreference clusters crossing a chunk boundary	6.9	6.5	1 to 16
cross-chunk clusters with more than one proper-name form	3.6	3	0 to 10

96 of 100 LitBank documents are exactly two chunks. There is one boundary per document. Cast drift and duplicate minting compound *across many* chunks, so on this corpus the phenomenon has about three to six chances to appear and exactly one place to appear in. A null result would be uninterpretable rather than informative.

Two further limits: LitBank’s entity layer carries the same 2,000-word ceiling, and LitBank has **no relation or triple annotation at all**, so there is no gold analogue for what a knowledge-graph extractor actually emits. This is known ground: BOOKCOREF (ACL 2025, arXiv:2507.12075) exists because “existing benchmarks, such as LitBank, remain limited in length and do not adequately assess system capabilities at the book scale.”

If a book-scale ordering experiment is ever wanted, BOOKCOREF is the resource (averaging over 200,000 tokens per document, roughly 170 chunks), with the caveat that its annotations are pipeline-produced rather than fully manual. DocRED is **worse** than LitBank here, not better: it has gold relations but its documents are roughly 200-word abstracts.

4. Evaluate on literature, because it has ground truth the assistant case cannot.

And render the output as a wiki, because a page is auditable in a way a chat answer is not.

Arm	What it is	What it measures
Assembled	fact rows and cells, composed mechanically	what the store holds
RAG-generated	a model querying the backend	what an assistant built on it would say
The gap	claims tracing to neither a fact row nor a cell	the unsupported-claim rate of this rendering path on this corpus

The last row is why the wiki belongs in the paper rather than only on a poster. State it narrowly. It is not “the assistant’s fabrication rate”; it is one rendering path on one corpus, and the generalisation from there to assistants at large is exactly the step a reviewer will refuse.

Nor is the assembled arm zero-fabrication. It is **100% traceable by construction**, because a claim with no matching quote is never emitted. Traceable is not the same as true: the edge-typing constraint exists precisely because a fabricated relation can carry a perfectly genuine quote, and `parent_of` between a place and a thing passes every other check in the system. Claim traceability, which is defensible and unusual, and do not claim correctness, which is neither.

Requirement 3 is the strongest position in the paper and it is worth being precise about why. FABLES reports the best automatic faithfulness checker reaching 58.2 F1 with the entire book available. The check here is exact string matching: not a classifier, not a judge, a `find()`. That is not a better classifier, it is a different kind of guarantee, and it should be framed that way rather than as a score comparison.

The framing, and it is load-bearing

“Here is our backend, and we measured it” is a system paper and dies on prior art.

“Here are measured tradeoffs in this design space, demonstrated on a working backend” is an empirical study and survives, because the contribution is the measurement.

The difference appears in the first paragraph, so it has to be decided before one is written.

This framing was checked, and it places. Empirical and design-space studies appear at the venues this work would target: *How Memory Management Impacts LLM Agents: An Empirical Study of Experience-Following Behavior* (ACL 2026), *Searching for Best Practices in Retrieval-Augmented Generation* (EMNLP 2024), *The Power of Noise: Redefining Retrieval for RAG Systems* (SIGIR 2024), *Seven Failure Points When Engineering a Retrieval Augmented Generation System* (CAIN 2024), and the VLDB Experiment, Analysis and Benchmark track as the data-management template. Benchmarks place too: MemoryAgentBench at ICLR 2026.

The flip side of the same finding is the real constraint. This niche is being occupied right now by studies larger than this one. The differentiator cannot be breadth of systems compared. It has to be the instrument: gold-annotation, entity-level scoring of extraction ordering, which none of the large comparisons do.

The novelty record

Source types for everything added on 08-28, per the standing rule that a claim must say how it was read. Full text: RAPTOR (the PDF is in papers/), and the CHI EA 2025 interview study. Machine extracted full text, not verbatim checked line by line: CAM. **Abstract only:** NarrativeXL, StoryBench, Memora, *Are We Ready For An Agent-Native Memory System?*, ENPMR-Bench, CogniFold, ProactAgent, AutoSchemaKG, SCOPE/SCION, EvoTaxo, and the rest of the occupied-contribution citations below. Venues were confirmed through DBLP or the ACL Anthology rather than assumed.

An abstract-only read is enough to establish that a paper *exists and claims a thing*, which is all an occupancy finding needs. It is not enough to quote a number from. Nothing in this document quotes a number from an abstract-only source.

Adversarial full-text queries against the arXiv API, written to find the paper that scoops the project rather than to map the field. The August 19 survey ran seven agents and verified 33 citations, and missed nearly all of the papers below. A survey asking “what is this field” and a search asking “what kills this claim” return different sets, and only the second one protects a contribution. That is a methods finding independent of the result.

Proactive recall. Taken. The original claim was that every published memory benchmark presupposes a question is asked, so a memory that is complete, indexed and never consulted fails no existing benchmark. That sentence is false as of May 2026. CogniFold (arXiv:2605.13438) presents an always-on agent memory that continuously folds fragmented event streams into self-emerging structures, extends Complementary Learning Systems theory, and surfaces intents on concept-cluster density; it ships code and evaluates on LoCoMo and LongMemEval. ProactAgent (arXiv:2604.20572) takes the measurement, learning when to retrieve by comparing paired continuations from identical prefixes with and without retrieval, which is the counterfactual miss rate as a learned policy. ENPMR-Bench (arXiv:2605.27240) is a benchmark for proactive memory retrieval, and its existence alone falsifies the “no benchmark can express this” framing. Also nearby: ProMem (2601.04463), PASK (2604.08000), VitaBench 2.0 (2605.27141), Proactive Memory Agent (2607.08716).

Cold-start structure induction. Taken. AutoSchemaKG (arXiv:2505.23628) is fully autonomous KG construction that induces schemas directly from text, at fifty million documents, reporting 92 percent semantic alignment with human-crafted schemas at zero manual intervention. SCOPE and SCION (arXiv:2607.21610) make schema induction a measured benchmark with released baselines. EvoTaxo (arXiv:2603.19711) builds and evolves taxonomies from temporally ordered streams, which is the living-tree idea applied to taxonomy. From Strings to Things (arXiv:2607.00003) builds personal knowledge

graphs from unstructured conversational data. BoostTaxo (2605.12520) does zero-shot taxonomy induction, and Bian et al. (2510.20345) survey the whole space.

Covering versus partition. Taken, on 08-28. This was carried as the surviving secondary claim on the strength of an assertion that every hierarchical memory system partitions. That assertion is false, and the refutation was on disk in this repository the entire time. RAPTOR: “nodes can belong to multiple clusters without requiring a fixed number of clusters ... thereby warranting their inclusion in multiple summaries” (ICLR 2024, p. 3), and this project’s own one-pager for it already said so. CAM (Li et al., NeurIPS 2025, arXiv:2510.05520) combines incremental overlapping clustering with hierarchical summarisation and online batch integration in agentic memory, evaluated against RAPTOR, GraphRAG, HippoRAG, MemTree, MemGPT and ReadAgent on NovelQA, QMSum, FABLES and MultiHop-RAG. The only daylight was that CAM and RAPTOR derive overlap from embedding-space similarity while this design derives it from a pre-determined entity cast. That is a variant of an occupied idea, and it is the same class of narrow distinction this project already downgraded once. **Keep the covering design, claim nothing for it.**

A literary memory benchmark. Largely taken. NarrativeXL (Findings of EMNLP 2023) is 1,500 hand-curated Gutenberg fiction books with 990,595 questions carrying an explicit retention demand, purpose-built as a long-term-memory dataset. MemoryAgentBench includes temporal recall over novels. CAM scores memory backends on NovelQA and FABLES. StoryBench (2506.13356) builds a long-term-memory benchmark on interactive fiction. The corpus-as-a-resource angle should shrink to one sentence. What is genuinely left is narrow and worth stating precisely: no benchmark found pairs gold structural annotation over literary works with a persistent, incrementally updated backend and supersession scoring. That is a test-instrument gap, not an unused corpus.

Corrections carried into this revision

Four things in the previous version of this document were wrong. They are recorded rather than quietly fixed, because the pattern matters more than the individual errors.

1. **Zep was said to lack community detection and hierarchical summarisation.** False. Its community subgraph nodes “contain high-level summarizations of these clusters,” built with detection that “builds upon the technique described in GraphRAG” using label propagation rather than Leiden. The claim came from reading the Graphiti README instead of the Zep paper. The correction was then itself overstated as “extends dynamically without full refresh”; the paper says communities gradually diverge and “periodic community refreshes remain necessary.”
2. **“The best single system covers five or six of nine” was never computed.** The coverage matrix was drawn with empty cells and a score asserted over them. The sentence is **deleted**, not repaired. Filling forty-nine cells to a standard that survives a reviewer is a research project in itself, and 2606.24775 has already done the defensible version. There are also seven requirements, not nine.

3. **The assembled-versus-RAG gap was called “the assistant’s fabrication rate.”** Narrowed above to one rendering path on one corpus.
4. **The requirement list was said not to exist yet, and to be the author’s to write from his own deployment.** It exists, in `02_requirements_and_testing.md`, sourced to literature.

The standing method rules that came out of these, and out of two occasions where “no records found” was read as absence when the material was on Project Gutenberg the whole time:

- An absence claim requires a documented search. Never assert that nobody has done X without reporting how you looked.
- Run a control query before trusting an empty result set.
- State the source type for every factual claim: full text, abstract, README, or memory. A README is not the paper.
- Never emit a number you did not compute.

Both original contributions were absence claims, and both were false. The third, covering versus partition, was also an absence claim, and it was false too. Absence claims are the weakest thing this project can assert and should be the first thing an outside reader attacks.

What already killed it, and what remains a risk

The first item is no longer a risk. It happened.

- **The extraction-ordering claim is dead.** Both baselines are entity-first, the mechanism is published three times over, the ablation has been run and came out negative, and the instrument cannot measure it. See the top of this document.
- **The surrounding niche is also occupied, and closely.** *Narrative World Model* (arXiv:2607.05577, July 2026) is writer memory for long-form fiction with a typed temporal-state graph, evaluated on a public fiction corpus against **Graphiti/Zep** and **GraphRAG** with a reader held constant, and it isolates extraction quality from representation by rebuilding the baseline with its own extractor. That is this project’s niche, its baselines and its corpus type, published six weeks ago. Read it in full before any re-scoping decision; it is currently an abstract-only read and it is the most important paper on this list.
- **Any novelty check goes stale in about three months.** CogniFold was three months old and invisible both to a seven-agent survey and to an assistant whose training data predates it. Whatever the project ends up claiming, re-check adversarially immediately before posting.
- **One search hole is unfilled.** The Semantic Scholar API returned HTTP 429 on every attempt during the 08-28 pass, so that source contributed nothing. Everything above rests on arXiv, DBLP, the ACL Anthology, web search, and the PDFs on disk.

Standing note on how these were found

Every one of the nine dead claims died to the same rule: **an absence claim requires a documented search.** In three cases the refuting evidence was already on disk in this repository, in papers/ or in the project's own one-pagers. The searches that killed them were run adversarially, in a context that had not been softened by the conversation that produced the claim, which is the only method that has reliably worked here.

Requirements and testing plan

2026-08-27, revised 2026-08-28. Supersedes the earlier evaluation plan (now `archive/20_evaluation.md`), which predates the source flip below and the asymmetric baseline strategy.

The LitBank head-to-head is withdrawn, and with it the GraphRAG comparison. The claim it tested is dead (`01_argument.md`), and LitBank could not have measured it regardless: 96 of its 100 documents are exactly two chunks at GraphRAG’s default size, so a phenomenon that compounds across many chunks has one boundary to appear at. Both changes are marked below. The free instruments and LongMemEval are unaffected.

Every citation in the requirements table was re-verified against full paper text on 08-28, with the single exception marked `[abstract only]`. Three were wrong and are corrected here. The corrections are marked, because the corrected versions are weaker than what they replace and nobody should reintroduce the stronger wording from an old draft.

The source flip, and why it matters

An earlier draft derived requirements from a personal deployment and validated them against the literature. That is a wish list with a case study attached, and a reviewer answers it with *says who?*

Requirements are derived from the literature’s own stated failures. The deployment corroborates; it does not authorise. Every row below cites a published paper saying the problem is unsolved, and no personal detail appears anywhere. The deployment’s entire role is one sentence: *a two-year deployment of a personal memory system encountered each of these failure modes in practice*.

Consequence, stated plainly: this section is **motivation, not contribution**. Its only job is to make a reviewer agree the requirements are real before they are shown numbers. The contribution is the measured comparison in `01_argument.md`.

The seven requirements

#	Requirement	The field saying it is unsolved
1	Supersession and fact update	MemoryAgentBench: on multi-hop fact consolidation, no tested method exceeds 28% . Memora (ACL 2026 Findings) introduces a forgetting-aware metric that penalises reliance on invalidated memory <code>[abstract only]</code>

#	Requirement	The field saying it is unsolved
2	Entity resolution	Noy et al. 2019: disambiguation is named as <i>one of</i> the top challenges, recurring almost across the board in five production knowledge graphs. Bing’s single Will Smith entry is composed from 108,000 facts taken from 41 websites
3	Faithfulness and provenance	FABLES: best model 90.9% faithful; the strongest automatic checker reaches 58.2 F1 with the entire book in hand
4	Omission and coverage	BooookScore: omission errors are the most common error type across both prompting strategies. FABLES separately finds 33.3% to 65.4% of summaries lack mentions of key events
5	Incremental update	MemTree measures the cost directly: incorporating one new observation takes roughly 3,750 LLM calls for RAPTOR and 3,850 for GraphRAG , against 3.27 per insertion for MemTree. Talebirad names evolving hierarchies as “the most pressing open problem” for its own theory
6	Cost accounting	Anatomy of Agentic Memory: “system-level costs are frequently overlooked”
7	Proactive recall	Occupied, not open. See the correction below

Corrections to this table, made 2026-08-28

Requirement 1 was overgeneralised. The 28% is the ceiling on the **multi-hop** FactConsolidation subtask specifically, not on fact-update tasks in general. On single-hop fact consolidation in the same table, GPT-5-mini reaches 78.0 and GPT-4o 60.0; on a 6K-token version, o4-mini reaches 100 on single-hop and 80 on multi-hop. The competency is also named **Selective Forgetting** in the paper, not “fact update.” Quote the multi-hop figure with its qualifier or a reviewer with the paper open will catch it.

Requirement 2 dropped two hedges. Noy et al. write that disambiguation “continues to be *one of* the top challenges in the industry *almost* across the board,” and list it alongside scale, extraction and

knowledge evolution as co-equal recurring challenges. The Will Smith numbers are exact and the company is Bing/Microsoft. Restore the hedges.

Requirement 4 spliced two papers into one citation. “Omission dominates the error types” is BoookScore, correctly. The 33 to 65 percent figure is **FABLES**, not BoookScore. Both halves are true; the attribution was not. The archived topic narrative had it right and the error entered when that sentence was compressed into a table row.

Requirement 5 was two-thirds wrong. RAPTOR does not name incremental insertion as unsolved anywhere in its paper, and neither does GraphRAG; the word appears in GraphRAG only in a bibliography entry. The criticism is real but it comes from **MemTree**, which measured the rebuild cost of both. Talebirad does name it, though as a limitation of its own theory rather than a field-wide verdict. MemTree measures a recency and insertion-order bias in itself, quantified as accuracy rising from 82.1 to 84.2, but never calls it drift. Reattributing to MemTree makes this row both correct and stronger.

Requirement 7 was an absence claim, and it is false. The original wording was that every memory benchmark poses the query, so a store that held the answer and was never consulted fails none of them. ENPMR-Bench (arXiv:2605.27240) is a benchmark for proactive memory retrieval; ProactAgent (arXiv:2604.20572) learns when to retrieve by comparing paired continuations with and without retrieval, which is the counterfactual miss rate as a learned policy; CogniFold (arXiv:2605.13438) is an always-on memory that surfaces intents unprompted. The requirement is still real, and the instrument here is still unresolved, but it must be presented as **hard and occupied**, not as unclaimed ground. See 01_argument.md.

How each is tested

#	Test	Cost	Status
2	Duplicate-minting curve per unit, as a descriptive measurement of this pipeline	free	falls out of ingest. The LitBank head-to-head is withdrawn: see the banner above. Without a baseline arm this describes behaviour, it does not compare
3	Quote gate: every claim must trace to a verbatim string in its unit	free	deterministic
6	Run rows, per stage per tier	free	already in the schema
5	Refold cost and staleness versus full rebuild	cheap	direct comparison; GraphRAG and RAPTOR are batch, Zep is not, and

#	Test	Cost	Status
			MemTree published the per-insertion figures to compare against
1	LongMemEval knowledge-update band; the MemoryAgentBench multi-hop ceiling as cited context; Tip to Ozma as a fixture	moderate	
4	Plot-versus-cell consistency as a cheap proxy; BoookScore only if run	partial	true omission measurement requires knowing what should have been included
7	No clean test. The <code>Consult.used_ids</code> instrument is unresolved	hard	a paired counterfactual doubles every measured query

Say #7 is unsolved rather than fudging it. It is the honest limitation. Note that it is no longer also a limitation of the field, since ProactAgent measures a version of it, so the framing is “we did not solve this,” not “nobody can.”

Requirement 3 is the strongest position

FABLES reports the best automatic fabrication checker reaching 58.2 F1 with the entire book available. The check here is exact string matching: not a classifier, not a judge, a `find()`. On the assembled rendering path it is **100% traceable by construction**, because a claim with no matching quote is never emitted.

That is not a better classifier, it is a different kind of guarantee, and it should be framed that way rather than as a score comparison. It is also traceability and not correctness: the edge-typing constraint exists because a fabricated relation can carry a genuine quote.

Baselines are not symmetric

Baseline	Compete on	Why
Zep	its own metric , LongMemEval	Public, conversational, has knowledge-update bands, and Zep published numbers on it
GraphRAG	nothing, for now	The comparison was going to be LitBank entity F1 and duplicate minting against its “chunk-first

Baseline	Compete on	Why
		ordering.” That premise is false: GraphRAG’s own appendix says its prompt “first identifies all entities in the text ... before identifying all relationships,” so it is already entity-first within a chunk. Its published evaluation is LLM-judged on self-generated questions over podcast and news corpora and cannot be reproduced fairly. There is no honest comparison left to run this semester

Zep’s published LongMemEval numbers, from Table 2 of its paper: 63.8% with gpt-4o-mini against a 55.4% full-context baseline, and 71.2% with gpt-4o against 60.2%. Two qualifiers that must travel with those figures. They are **LongMemEval_s**, the small variant whose conversations average about 115,000 tokens, not the full benchmark. And the stronger half of Zep’s result is not the accuracy: it served those numbers from about 1,600 tokens of context against the baseline’s 115,000, with roughly a tenfold latency reduction. If this project competes on LongMemEval it is competing on that frontier, not on accuracy alone.

DMR: run it because it is cheap and widely cited, and quote Zep’s own caveat. Zep scored 94.8% against a 94.4% full-conversation baseline with gpt-4-turbo. The gpt-4o-mini row is the sharper illustration and is currently unused: 98.2% against 98.0%, a margin of two tenths of a point. Zep’s paper calls the benchmark inadequate in its own words, noting each conversation is only 60 messages and “easily fitting within current LLM context windows.” Cite Table 1 rather than the abstract, because the abstract’s headline compares Zep to MemGPT rather than to full context.

One precision note on GraphRAG. Its evaluation uses four criteria, not two: comprehensiveness, diversity, empowerment, and directness as a deliberate control. The headline wins are in comprehensiveness and diversity, and empowerment showed mixed results, so say “primarily comprehensiveness and diversity” rather than implying there are only two.

Scope, against 70 to 100 hours

Committed core, with no question authoring, no judge calibration and no kappa study:

- The free instruments: quote gate, rejection rate per stage per tier, duplicate minting, predicate sprawl, token cost, plot-versus-cell consistency (requirements 3, 4, 6)
- Refold versus rebuild (requirement 5)
- Per-stage tier sensitivity, which is the one measurement with a real spread behind it (146-fold across providers, costed in 05_fall_plan.md)

Withdrawn 2026-08-28: the LitBank head-to-head, which was the anchor of this list. It tested a claim that turned out to be occupied and mis-stated, on a corpus that could not have measured it. Nothing replaces it, and inventing a replacement without a search behind it would repeat the exact error that killed nine consecutive claims.

If a book-scale version is ever wanted, the resource is BOOKCOREF (ACL 2025), not LitBank, and it needs its own prior-art search first.

The one paid benchmark worth running: LongMemEval. It is where the assistant claim lives and where Zep can be matched or beaten (requirement 1).

Cited, not run: MemoryAgentBench and BoookScore. They supply the multi-hop ceiling and the omission context without costing a run. Note that the 33 to 65 percent figure comes from FABLES, which is also cited and not run.

Not attempted this semester: requirement 7. Consult-logging should still start now, because the rate needs weeks of collection before it means anything, but the measurement is not promised. The logger that emits conforming run rows is a separate deliverable from switching logging on, and only the latter is currently scheduled.

Demonstrations, clearly labelled as such

These produce no comparable numbers. They produce evidence a reader can check by eye, and the paper must not present them as results.

Fixture	Corpus	Shows
Tip becoming Ozma at the end of book 2	Oz	Aliasing, merge, supersession and time-scoped truth in one case, at document 2 rather than document 14
Watson's wound, shoulder in <i>A Study in Scarlet</i> and leg in <i>The Sign of the Four</i>	Holmes	Contradiction with no reconciling reading, inside one author
The deerstalker probe	Holmes	Free generation versus fact-row composition; the second structurally cannot leak, because Doyle never wrote the detail so no fact row exists to compose
Source disagreement rendered inline	Greek	Helen reached Troy per Homer, never per Euripides, both live and both cited
OCR tax and translation tax	Chinese	Same content, one variable. Cite Sun et al. 2026, <i>When Good OCR Is</i>

Fixture	Corpus	Shows
		<i>Not Enough: Benchmarking OCR Robustness for Retrieval-Augmented Generation</i> (arXiv:2605.00911), which is the prior art and is on disk as papers/sun2026-ocr-robustness-rag.pdf
The design session's own transcript	none	Ten interleaved, non-contiguous threads with cross-thread supersession. A figure , not a result

Design: schema, ingestion, retrieval, delivery

2026-08-27, revised 2026-08-28. The design worked out in session on 08-26/27, consolidated from three separate documents (schema, ingestion pipeline, retrieval and delivery). Records the *what* and the *why*; implementation is the student's, per the authorship rule.

Citations resolve against `07_references.md`, where every source is now complete and citable. The unit contract is `04_unit_contract.md` and is frozen.

1. What the store is shaped like

The design records on two axes over the same source text, and neither substitutes for the other.

Plot axis. Summaries per unit, attached to work nodes. Complete by construction: every unit gets one, no threshold, no entity judgment involved.

Entity axis. Per-node narrative cells. Sparse by design: only above-threshold nodes, and only for the units they appear in.

Each covers the other's blind spot. The plot axis loses the individual thread, because a character appearing in twelve chapters is scattered across twelve summaries, none of which is about them. The entity axis loses the causal frame: Ozma's thread says she was Tip and became ruler, but not *why*, because the why is Mombi's enchantment and Jinjur's revolt, which are facts about other things.

Formally the atom is the **cell**, (`unit × node`). Plot summaries are the row marginals; entity narratives are the column marginals. Two complementary views over one matrix.

Threads are a covering, not a partition

GraphRAG partitions by graph topology, and says so outright: each level of its hierarchy “provides a community partition that covers the nodes of the graph in a mutually exclusive, collectively exhaustive way” (Edge et al. 2024, p. 6). TraceMem is equally explicit, its coarse clustering operation “produces a partition into m topic clusters” (Shu et al. 2026, eq. 4), with a KNN step reassigning noise points so nothing is left out and nothing lands in two places. HERCULES recursively applies k-means. Zep uses label propagation. A folder tree assigns by topic path.

Threads do not. A chapter belongs to Tip's thread *and* Jack Pumpkinhead's *and* Mombi's, simultaneously, at full weight in each. The same text is covered N times from N vantage points.

This is a design choice, not a novel one, and the distinction matters. RAPTOR already soft-clusters: “nodes can belong to multiple clusters without requiring a fixed number of clusters ... thereby warranting their inclusion in multiple summaries” (Sarathi et al., ICLR 2024, p. 3). CAM (Li et al., NeurIPS 2025) combines incremental *overlapping* clustering with hierarchical summarisation and online batch integration, evaluated against RAPTOR, GraphRAG, MemTree and MemGPT. Covering is

settled ground in hierarchical memory. Adopt it because it is right for this store, and claim nothing for it. See `01_argument.md`.

The prototype deployment shows the same structure outside literature: leaves mentioning one project sit in six different topic folders, including one in an unrelated personal-health topic. The routing is not wrong, because that leaf genuinely belongs where it was filed. The thread is simply orthogonal to the hierarchy, so no assignment of documents to folders can represent it. That is a structural limit rather than a curation failure.

Consequence: **a thread index is a second index over the same units, not a reorganisation**. The hierarchy stays. Nothing is re-filed.

2. Schema

Eight record types, one controlled vocabulary, and three instrumentation records. Field names are indicative.

Source layer

Document {doc_id, corpus_id, work_id, kind: text|chat, title, source_url, sha256, ingested_at, provenance: {contributed_by, resources, ts}} Raw text sits beside it on disk and is never edited. Ids are content hashes, so re-ingesting the same material is a no-op rather than a duplicate.

provenance is immutable and exists on every derived record below. It carries who contributed the record, what resources it drew on, and when. This is the substrate for read-time access control (Rezazadeh et al. 2025, *Collaborative Memory*).

Unit {unit_id, doc_id, unit_type: chapter|book|play|ode|hui|session, work_ordinal, unit_ordinal, title, span: [start, end], text} The atom. Every derived record points here. The full contract, including the eight per-convention handlers and the three acceptance gates, is `04_unit_contract.md`.

unit_type carries heterogeneity forward as data rather than erasing it. A Euripides play has no siblings to fold with; an Oz chapter does. Normalising to one record shape is correct; flattening the distinction is not.

Node layer

Node {node_id, canonical_name, kind: person|org|place|thing|topic|event, created_from_unit, provenance} **Entities and topics are the same type**. A person is a node, a project is a node, a recurring event such as the Chunin Exams is a node. kind exists because extraction prompts differ per kind, and because it participates in the delta constraint below, not because storage differs.

Works are nodes. “Oz book 2” is a topic node. This removes the need for a separate chapter-summary record type: the plot axis is cells on work nodes, and the book blurb is that node’s abstract.

Type versus instance is not a kind distinction. The Chunin Exams is one node whose sequence holds every exam event; an instance is promoted to its own node only when it accumulates enough to warrant one. Same promotion mechanism as minor entity to major entity. This matches wiki practice, where instances appear as sections until they earn a page.

Alias {alias, node_id, alias_kind: name|nickname|epithet, first_seen_unit, evidence_quote} The surface-form list mapping mentions to one node. Resolution reads it; nothing writes over it.

Fact layer

Fact {fact_id, subject, predicate, object, object_is_node, qualifiers, rank: preferred|normal|deprecated, unit_id, quote, valid_from, valid_to, # [from, to), upper exclusive asserted_at_unit, asserted_at_wallclock, tier, contract_version, provenance}

Append-only. Nothing is overwritten, and there is no stored `superseded_by`: supersession is computed at read time over a functional-predicate list.

rank fills a gap supersession cannot. Computed supersession handles *the world changed*, in which a later assertion collides with an earlier one on a functional predicate. It does not handle *we were wrong*: an extraction error has no later assertion to supersede it, and deleting violates append-only. deprecated means present, preserved, and excluded from reads by default.

Source note, corrected 2026-08-28: the three-value vocabulary (preferred, normal, deprecated) comes from Wikidata's *Help:Ranking*, not from Vrandecic and Krotzsch 2014. That paper describes optional marking as “preferred” or “deprecated” on p. 83 but never uses the word “rank” and never names a “normal” value. Both sources are in `07_references.md`.

An error of exactly this shape occurred in the prototype deployment, where two distinct people were merged under one alias and the fix was to append an is-distinct-from assertion rather than edit the record. Same pattern, done ad hoc.

qualifiers carry n-ary relations, not just timing. Ozma rules Oz *as restored by Glinda*; Zhuge Liang serves Liu Bei *in the capacity of Chancellor*. Without qualifiers, role information becomes either lost detail or invented predicates, and invented predicates are the sprawl that silently breaks supersession. Source: Wikidata qualifiers, Vrandecic and Krotzsch 2014, p. 82.

Interval convention is [from, to), lower bound inclusive and upper exclusive, following SQL. Stating it prevents the off-by-one that appears wherever two intervals abut. Source: Rost et al. 2021, pp. 6 to 7. Cite that work as an arXiv preprint and technical report; it has no conference or workshop venue.

Fact-level validity subsumes property-level. Rost puts validity on individual properties because those vertices carry properties directly. Ours do not, because our facts *are* the properties, so this is already covered. Recorded so it is not re-opened.

Predicate registry (controlled vocabulary, not a record) core: ~15-20 universal predicates, including every functional one corpus: per-corpus vocabulary, namespaced delta: allowed (subject_kind, predicate, object_kind) triples

Namespaced predicates, such as `appearance.height` and `appearance.build`, give attribute grouping for free. That is how a “Physical description” section assembles itself, and it stays human-readable.

Delta is a typing constraint on edges. A table of allowed (`subject_kind`, `predicate`, `object_kind`) triples, with violations logged as rejections rather than stored. Source: Angles 2018, *The Property Graph Database Model*, Definition 2, where delta “defines the edge types allowed between a given pair of node types.”

This catches a class of error **the quote gate structurally cannot see**: a fabricated relation can carry a perfectly genuine quote. `parent_of` between a place and a thing passes every other check in the system. This is also why the assembled path should be described as fully *traceable*, not as zero-fabrication. See `01_argument.md`.

Merge ledger {`attr_id`, `node_id`, `attribute`, `merged_from`: [`fact_ids` in order], `absorbed_node`, `evidence`, `tier`} Which rows, in what order, folded into an attribute. The anti-black-box requirement, stored as data rather than asserted.

Summary layer

Cell {`cell_id`, `node_id`, `unit_id`, `corpus_id`, `work_ordinal`, `unit_ordinal`, `text`, `tier`, `provenance`} One node’s narrative for one unit. The expensive layer, and the thresholded one.

Ordered by (`work_ordinal`, `unit_ordinal`), **never a single integer**, because a node’s cells span multiple works.

Keyed by (`node`, `unit`), which flattens two separate scenes involving the same character in one chapter into a single cell. Accepted deliberately; recorded so it is a decision rather than an omission.

Abstract {`node_id`, `corpus_id`, `text`, `children_hash`, `tier`, `updated_at`, `provenance`} One paragraph per node **per corpus**. Salience is corpus-scoped: Tip is a major entity in Oz and a passing citation in a conversation archive. The node is global; its narrative is local.

`children_hash` is computed over the node’s cells, so the abstract is stale exactly when its cells change. Refold is incremental, never a full pass.

Composition, not storage. A visual description is not a stored artifact, it is a query over `appearance.*` fact rows, rendered. Free to regenerate, automatically current, and every clause traces to a quote. This is also why the Holmes contamination probe passes by construction: there is no deerstalker fact row, because Doyle never wrote one, so nothing composes.

Instrumentation

Rejection {`rej_id`, `ts`, `stage`, `unit_id`, `target`, `contract_version`, `tier`, `error`} Run {`run_id`, `ts`, `arm`, `question_id`, `injected_ids`, `answer_text`, `verdict`, `judge_tier`, `tokens_in`, `tokens_out`, `latency_ms`} Consult {`consult_id`, `ts`, `cued_node_ids`, `retrieved_ids`, `used_ids`} Mention {`mention_id`, `unit_id`, `node_id`, `surface_form`, `quote`, `span`, `tier`}

Four fields here were carried forward from the module specs now in `archive/impl/`, where they were identified as gaps and never absorbed:

- **Rejection.target** makes “what got rejected twice” a query rather than a manual scan. Without it the repeated-failure report cannot be written.
- **Rejection.unit_id** makes a failed run resumable. Without it, a crash midway through a corpus restarts from the beginning.
- **Run.answer_text** lets verdicts be re-judged after a judge-prompt change. Without it, every judge revision invalidates every prior run and the runs cannot be re-scored, only re-executed.
- **Mention** makes the surface-form-to-node binding a first-class record. Quote verification and the hand-check sample both need bindings as data; deriving them after the fact from cells is not the same thing.

Consult is the proactive-recall instrument. It needs to exist from day one because the rate is meaningless without weeks of collection. How `used_ids` gets populated is unresolved and is recorded as such in `02_requirements_and_testing.md`; that requirement is explicitly not promised this semester. Note that the logger that emits conforming rows is a **separate deliverable** from turning consult-logging on, and only the latter is currently scheduled.

The node record also needs `merged_into`, so a node absorbed by a merge resolves to its survivor at read time rather than becoming unreachable.

3. Ingestion

Stage 0. Preprocessing

Raw text to normalised unit records. Per-corpus handlers, one output contract, specified in `04_unit_contract.md`. Deliberately outside this pipeline so the heterogeneity is absorbed once.

Stage 1. High-level pass: identify topics and major entities

One call per unit. Returns the cast: entities and topics present, with a salience signal.

Entity-first, not chunk-first. This is the project’s primary claim. GraphRAG’s pipeline composes text units first and extracts entities from each text unit *separately*; there is no document-level entity pass. Zep extracts per episode with a short message lookback, then resolves against the graph.

That ordering is where duplicate minting comes from. Chunk 3 sees “Tip,” chunk 9 sees “the boy,” chunk 15 sees “Tippetarius”: three independent extractions with no shared context, reconciled afterward if at all. The prototype deployment measured the consequence, a registry that grew to roughly 1,169 names against a 400-name prompt cap, with active duplicate minting.

A unit-level cast list passed into each downstream call fixes it before it happens. The same pass does three jobs at once: de-duplication, cross-chunk coreference resolution, and **determining which cells**

exist, which is what makes the sparse matrix tractable rather than requiring every entity against every unit.

Stage 2. Dedup against the existing registry

Resolve the cast against known nodes and aliases. **This is on the critical path and it is the step the prototype already failed.**

Once the registry exceeds what fits in a prompt, this stops being a prompt and becomes retrieval-then-prompt: embed the mention, pull top-k candidates from the node store, resolve against those few.

Topic dedup is harder than person dedup. A person has a name; a topic has whatever description the speaker reached for. Person aliases converge; topic aliases proliferate, because each mention invents its own label.

Topic granularity has no natural boundary. The prototype's hand-curated boundary rules exist because automatic topic boundaries drifted and had to be adjudicated. Expect the same here: a curated layer over the induced one.

Stage 3. Fact extraction

Dated subject-predicate-object rows with verbatim quotes and unit pointers, validated against the fact contract and the delta table. Rejections logged.

Order matters: **dedup precedes facts**, or facts bind to duplicate nodes.

Stage 4. Narrative cells

One cell per above-threshold node per unit. **The dominant cost, and the batching decision lives here.** Emit all N cells for a unit in one call rather than N calls, or the unit text is re-sent N times. Costed in 05_fall_plan.md.

Stage 5. Maintenance: abstracts refold

An abstract is a fold over its node's cells, so `children_hash` says exactly when it is stale. Refold only nodes whose cells changed: **O(nodes touched)**, not O(all nodes). That is the difference between an affordable nightly pass and an unaffordable one.

Supersession and merges also resolve here, at read time rather than by rewriting rows.

The threshold

All nodes enter the fact layer. Only above-threshold nodes get narrative cells. That places the threshold between a cheap layer and an expensive one rather than between having something and having nothing.

Three tiers, matching wiki practice (confirmed against the Naruto and Harry Potter wikis via their MediaWiki APIs):

Tier	On a wiki page	Scope	Cost
Facts	infobox (born, died, alias)	every node	cheap
Abstract	lead paragraph; Background, Personality, Appearance	every node	O(nodes)
Cells	Biography, per-year, named events; arcs	above threshold	O(nodes x units)

Scabbers before the reveal gets facts and a paragraph (Ron’s pet rat, unusually long-lived) and no biography, because he is not doing anything. After the reveal, Pettigrew has a full segmented one.

A search therefore never returns null. The abstract is the floor. Below-threshold nodes degrade gracefully instead of disappearing. And nothing is lost when a node falls below threshold, because the plot axis records that unit regardless. What is missing is the entity-indexed *access path*, not the content, which makes it measurable rather than silent: ask the same question against each axis and the difference is what the entity layer buys.

What the threshold must not be. Salience-by-action would demote exactly the entities that carry an argument. Tip/Ozma and Scabbers/Pettigrew are the most-referenced concrete examples in this project’s own design discussion and they never act: they are cited, not participants.

Major and minor are scoped per corpus, not properties of the node. Tip is major in Oz and minor in a chat archive. Same node, different salience. This falls out of the keying for free, since cells are keyed by (node, unit) and units belong to corpora.

Promotion candidates, and they are not equivalent: cross-unit recurrence; within-unit salience; retroactive promotion via a `same_as` merge; on demand when a query targets the node.

Record “considered, below threshold” as distinct from “never considered.” The two are indistinguishable later otherwise.

Retroactive promotion is re-narration, not backfill. A wiki’s Pettigrew biography covers 1981 to 1993 as “hiding as a rat,” a framing unavailable to anyone reading those chapters at the time. Cells generated retroactively are written from an epistemic position the source did not have, a second-order supersession where both accounts are valid from different vantage points.

A free consistency gate

The plot-axis summary and the entity cells are **independent summaries over identical text**. If the plot summary carries an event that no cell mentions, either it was invented or stage 1 missed a thread. Set comparison, no ground truth, no judge model, catches both directions.

Model tier per stage

The intuition that higher abstraction deserves a higher tier is **half right**, and the wrong half is expensive.

Stage	Volume	Difficulty	Errors	Tier
Entity spotting	1/unit	easy, NER is solved	local	cheap
Coreference and alias resolution	per ambiguous mention	hard	propagate	high
Fact extraction	1/unit	medium	local, quote-gated	mid
Cell generation	N/unit, dominates cost	medium	recoverable	cheap to mid
Abstract refold	1/node when stale	high	recoverable	high
Judge	low	high	n/a	high, and never the writer's tier

The rule is not “spend more at higher abstraction.” It is: spend where errors propagate, economise where they are recoverable. A mediocre abstract costs one call to regenerate. A corrupted entity registry costs a re-ingest. In the prototype, two entities whose names share a substring cannot currently be told apart, and no amount of frontier-model summarisation fixes that: a top-tier abstract over a merged node just writes a confident paragraph about a hybrid that does not exist.

Both expensive categories are low-volume, so this is affordable. The cost driver is cell generation, which is also the layer where a bad output is one cheap regeneration away.

Risks carried forward

- **Stage 2 is the critical path** and the prototype already failed it at scale.
- **Predicate sprawl** breaks supersession silently, because `is_a`, `species` and `form` never collide. Count distinct predicates per unit; linear growth means the contract needs a controlled vocabulary.
- **Duplicate minting.** Count nodes minted per unit. A rising line is the failure.
- **Contract rejection on cheap tiers** is a format failure, not a comprehension failure, and has a different remedy (grammar-constrained decoding). Measure the two separately.

4. Retrieval

Most retrieval here is a keyed lookup, not a similarity search. Standard RAG embeds the query and cosines against chunks. This store is indexed by node, so the first job is not “find similar text” but “which node is this about.” **Embeddings do entity resolution; they do not fetch content.**

Job	Method
Query to nodes	nearest_nodes over node-name and alias embeddings, then a cheap judgment over the top few
Node to abstract, cells, facts	keyed lookup, ordered
Unit to its text	keyed lookup
“Where do X and Y meet”	set intersection of their cells’ unit ids
Exact name, phrase, number	search_text, BM25
Query resolves to no node	search_text over cells and abstracts, the fallback, and it is not optional

The node-name index is a few thousand short strings, not millions of chunks. This is entity linking, not passage retrieval, and it reuses the stage-2 machinery.

Question shape picks the partition. *Who is X* reads the abstract. *What happened to X* reads the cell sequence. *What happened in chapter 5* reads the unit summary. *Why did X do Y* reads both nodes’ cells, joined on shared units. Facts index into the sequence: a *same_as* firing at unit N tells you which cells to fetch rather than all of them. That join is why both axes are kept.

Two honest weak points. The whole path rests on query-side entity resolution. “The boy who became a princess” names no node, the front end fails, and you are on the similarity fallback, which must therefore be good rather than vestigial. The prototype’s own worst retrieval failure was a *query* failure fixed by multi-query, not a storage failure, which argues the fallback needs query expansion more than a better index. And what gets ranked is **generated text**: cells and abstracts are summaries, and every one points back to its unit, so drilling to source works, but the thing you rank is not the thing that is true.

Raw text is reachable, not ranked. Drill-down is always available. Ranked search over raw is the naive-RAG control arm and stays out of the default path; blend it in and the arms stop being separable.

Drill-down, at execution: cell -> unit_id -> doc_id -> file, sha256 fact -> unit_id, then unit.text.find(quote) -> +/-500 chars Three dictionary lookups and a string find. No search, no ranking, because the pointer was written at ingest time, so all the difficulty is in *storing* it correctly.

Quote-primary, span-as-cache. Span is fast, quote is durable. Any splitter change re-derives units and invalidates every offset, but a quote still finds itself: on mismatch the quote wins and the span is recomputed.

Drill-down and the anti-fabrication gate are the same mechanism. Following a quote to its unit and checking that it lands is a drill-down; doing it for every assertion at once is the gate. A find returning -1 means either the quote was fabricated or preprocessing moved, and it is the same check either way.

Triggers

Three, and the distinction between them is where the design lives.

1. **Boundary hooks.** Session start and prompt submit, firing on new chat, topic shift, idle gap. Unconditional, no question required. Already built and running in the prototype.
2. **Explicit question.** The standard path.
3. **Agent-initiated.** Mid-task lookup.

The hook cues; it does not retrieve. It injects pointers and an instruction. Whether anything is read is a discretionary second step, and that step demonstrably fails: during the 08-26/27 session, memory cues fired on nearly every turn, each naming three leaves, and almost none were followed. The store held the content, the trigger fired correctly, the pointers were in context, and nothing was read. **The trigger is not the broken part. The step after it is.**

The fork worth measuring. Does the hook inject **pointers** or **content**? Pointers are cheap and leave a discretionary step that fails; content is expensive in tokens and removes the discretion. Two arms, no corpus work required, since the hook, the store and a live agent all already exist.

5. Rendering and delivery

A wiki page is **largely assembly of content that already exists**:

Element	Source	Cost
Lead paragraph	stored abstract	zero
Biography sections	stored cells, ordered	zero
Infobox	fact rows	zero
Categories and groupings	query over facts	zero

Only three parts need generation, and all three are low-volume folds: grouping cells into arcs (24 chapter-cells become 5 sections), prose smoothing, and contradiction presentation. **If rendering needs a frontier model every time, the store is not doing its job.** Distillation happens once at write time so reads are cheap; a frontier call per page view inverts that.

Two views of one page. The *provenance view* is appended cells coloured by source work: pure assembly, free, always available. The *article view* is the readable synthesis: one fold per node, cached by `children_hash`, so it regenerates only when its cells change. That split is not new machinery, it is the leaf-and-rollup architecture rendered. It also inherits the prototype's measured finding that defects concentrate in the synthesised layer, where 39 node claims traced back to raw gave 31 confirmed, 4 wrong and 4 unsupported, with every defect in the synthesis and all 31 leaf checks clean. **The provenance view is the audit tool for the article view**, visible to a reader instead of requiring a rotating audit.

What colour-by-source buys. Real wikis partition by source authority because they must, with sections such as “Behind the scenes” and “In Other Media” kept separate from canon. Doing it structurally means the Baum-to-Thompson authorial handover is visible on the entity page itself, and for the Greek corpus source disagreement renders inline: Helen reached Troy per Homer, never reached Troy per Euripides, both live, both cited, neither overwritten. No existing mythology wiki does that, because a human writing one article cannot hold five contradicting sources in the prose.

Two deployments, one logical schema

Separated by a storage port of twelve operations.

A, the research backend. Neo4j, with the vector index and the graph in one store and temporal traversal as a query. Graphiti (Zep’s engine) runs on Neo4j and supplies temporal facts, provenance to raw episodes, incremental writes, and hybrid retrieval, all specified independently in this design.

Zep does perform community detection and hierarchical summarisation. Its community subgraph is the highest level of the graph, and its nodes “contain high-level summarizations of these clusters” (Rasmussen et al. 2025, p. 2). Community detection “builds upon the technique described in GraphRAG” but uses label propagation rather than Leiden (p. 4), chosen for its straightforward dynamic extension. An earlier claim in this repo that Zep lacks this was based on the Graphiti README rather than the paper, and was wrong.

Two precision notes, because the correction was itself overstated once. Label propagation **delays full refresh, it does not eliminate it:** the paper says the communities “gradually diverge from those that would be generated by a complete label propagation run. Therefore, periodic community refreshes remain necessary” (p. 4). That concession is useful rather than damaging, since it means Zep also carries a batch refold. The remaining gap is narrower than first stated: per-entity narrative sequences, unprompted firing, and per-stage cost accounting.

Two Graphiti capabilities were checked directly against source at commit 683a853 (v0.29.3) and both came back usable:

- It ingests plain text, not only chat messages. `EpisodeType.text` is a first-class episode type, so a chapter goes in as one episode. There is no document abstraction, no loader and no splitter on the ingestion path, so unit splitting stays ours, and each episode body must fit the extraction model’s context window. A `content_chunking` module exists but has zero callers in `graphiti_core` at that commit.
- It accepts pre-extracted entities. `add_triplet(source_node, edge, target_node)` takes constructed `EntityNode` and `EntityEdge` objects with no extraction call. Two caveats: it still runs node and edge resolution, so a supplied entity may be merged rather than stored as-is, and it creates no episode node, so triplets added this way carry no episode provenance.

Consequence: the Zep-as-baseline comparison stands on the literary corpora, and the spring product does not need to build its own graph layer.

B, the distributable. JSONL as the store of record, .npy for embeddings, SQLite FTS5 for text search, no server. The canonical data is entirely JSONL: human-readable, git-diffable, one record per line. The two indexes are derived and disposable; delete them and they rebuild. This is justified rather than merely convenient, since brute-force exact cosine is correct below roughly 100K vectors and a personal corpus is nowhere near that.

Reaching a desktop client is MCP: a Python server over stdio plus a config entry, where the tool list is the capability contract. Two such servers already run against the prototype.

One dependency decision outstanding. sentence-transformers pulls torch, roughly 2GB, which is not a folder someone scans a QR code for. Either ship precomputed embeddings with a small ONNX runtime for the query vector, or drop embeddings from the distributable entirely: for a fixed corpus with a built alias table, resolving a name is a dictionary hit plus fuzzy match, and FTS5 covers the rest. The second option gives zero ML dependencies and costs recall only on phrasings that share no vocabulary with any alias.

Not asserted: current OpenAI and Gemini desktop MCP support. Check before promising cross-client anywhere.

Not adopted: LangChain and LlamaIndex. They take ownership of chunking, retrieval strategy and prompt templates, which is exactly the surface that has to be built and defended here.

6. Invariants

1. Raw is immutable. Ids are content hashes.
 2. Nothing is overwritten. Supersession, merges and permissions all resolve at read time.
 3. Every assertion carries a verbatim quote that must appear in its unit. Unverifiable claims are flagged or dropped, never asserted.
 4. Every model call goes through the two interfaces, `embed(texts)` and `generate(prompt, schema, tier)`, and records its tier.
 5. Every contract rejection is logged, so rejection rate per stage per tier is a query.
 6. The abstract is a fold over cells; staleness is a hash comparison, never a wall-clock guess.
 7. Every delta violation is a rejection, not a stored row.
-

7. Open questions

- **Mixed-corpus ordering.** Books order by narrative position, chats by wall clock. A store containing both has no defined sequence, and `(work_ordinal, unit_ordinal)` silently interleaves incomparable things.

- **Nothing demotes or retires.** Promotion is one-way, and the prototype already carries 82 orphaned sheets from exactly this.
- **kind does two jobs**, routing extraction prompts and constraining delta. Those will want to diverge, because event and topic need the same predicates and very different prompts.
- **Conformance testing.** Two adapters, twelve operations, no shared test suite. The realistic failure is that the files adapter is developed against and the Neo4j adapter rots.
- **Embedding-to-node desync has no detector.** The .npy matrix and the node table are written separately, so a partial write leaves row i of the matrix pointing at the wrong node. There is no checksum tying them together, and the failure is silent: retrieval simply returns the wrong entity. Store a row count and a hash of the node-id ordering alongside the matrix and verify on load. Carried from the doc 15 quality pass, item 3, never absorbed.
- **FTS5 is assumed, not checked.** The distributable's text search depends on SQLite being compiled with FTS5, which is not universal. Probe for it at startup and fall back to an in-process BM25 if it is missing, or the distributable fails on a stranger's machine in the one way that cannot be debugged remotely. Carried from the doc 15 quality pass, item 4, never absorbed.

The unit contract

2026-08-27. FROZEN. This is the only schema element that must be settled before preprocessing begins. Everything else in `03_design.md` is provisional and gets settled by contact with real data.

Preprocessing emits records matching this contract and nothing else. Per-corpus handlers absorb the heterogeneity; downstream sees one shape.

The record

```
jsonc { "corpus_id": "oz", // oz | greek | holmes | chinese "work_id": "55", // Gutenberg number
or Archive.org identifier "work_ordinal": 1, // position of the work within the corpus
"work_title": "The Wonderful Wizard of Oz", "unit_type": "chapter", // chapter | book | play | ode
| hui | session "unit_ordinal": 3, // position within the work, 1-based "title": "How Dorothy
Saved the Scarecrow", "text": "...", // verbatim, boilerplate stripped, nothing else altered
"chars": 11482, "source_sha256": "...", // sha256 of the RAW file, from the corpus manifest
"unit_id": "u_a91f0c33" // content hash: cid(corpus_id, work_id, unit_ordinal, text) }
```

Field notes

unit_type carries the heterogeneity forward rather than erasing it. A Euripides play has no siblings to fold with; an Oz chapter does. Normalising to one record shape is correct; flattening that distinction is not, and downstream needs to know which it is holding.

Ordering is (work_ordinal, unit_ordinal), never a single integer. A node's cells span multiple works, and a flat ordinal silently interleaves them.

unit_id is a content hash, so re-running preprocessing on unchanged input produces identical ids and a second run is visible as an empty diff.

text is verbatim after boilerplate stripping. No normalisation, no whitespace collapsing, no smart-quote conversion. The quote gate compares extracted quotes against this text, so anything done here has to be done identically at extraction time or every quote fails.

No span field. Offsets into the raw file are deliberately omitted: any change to a splitter invalidates every offset, and the quote is the durable pointer. Spans are recomputed downstream as a cache, never stored here. See `03_design.md`, `quote-primary` and `span-as-cache`.

Acceptance gates

A work passes preprocessing only if one of these holds. **A failure is a failure**, not a warning, a corpus that is quietly wrong in the middle is worse than one that is visibly short.

Gate	Applies to
Split count equals the table-of-contents entry count	any work carrying a TOC
Markers are monotonic with no gaps	regular-marker works with no TOC (e.g. 第N回)
Count equals 1	single-unit works, plays, odes, standalone stories

Table-of-contents titles do not always match body titles. *The Marvelous Land of Oz* lists “Tip Manufactures Pumpkinhead” and the body reads “Tip Manufactures **a** Pumpkinhead.” Match on count and order, never on string equality.

Convention handlers

Eight known variants, one visible function each, plus a dispatcher that picks by inspection. Ugly on purpose: every line is defensible and a failure points at exactly one handler.

Handler	Form	Seen in
caps_inline	CHAPTER I. TITLE, number and title on one line, all caps	Holmes <i>A Study in Scarlet</i> , Brewitt-Taylor <i>Three Kingdoms</i>
roman_own_line	Chapter I with the title on the following line	Oz book 1
word_own_line	Chapter One with the title on the following line	Oz books 7, 14
contents_indent	I. A Scandal in Bohemia indented under Contents	Holmes <i>Adventures</i>
book_marker	BOOK I.	Iliad, Odyssey
hui_marker	第一回 靈根育孕源流出	Chinese originals, 105 in <i>Journey to the West</i>
bare_title	no marker at all; body titles matched against a LIST OF CHAPTERS block	Oz books 2, 3, 13, 17
single_unit	the work is the unit	Euripides plays, Pindar’s odes

Boilerplate stripping

Gutenberg’s *** START OF *** and *** END OF *** markers bound the text, **but they are not sufficient**. Several files carry producer credits *inside* the START marker, “Produced by Dennis

Amundson”, proofreading team notes, HTML-version pointers.

The rule is explicit and inspectable: after the START marker, drop leading blocks matching known producer patterns until the title line. A heuristic that trims N characters is not defensible and will eat a real first chapter somewhere.

Archive.org scans have no markers at all and open with digitisation boilerplate, the Statius title does not appear until character 39,300. Those need a per-file front-matter rule or manual inspection.

Known quality flags to carry forward

These are properties of specific files, recorded so downstream results can be qualified:

- **OCR, not proofread transcription**, three Chinese files, the Richard *Journey to the West*, both Apollodorus volumes, both Diodorus volumes
 - **Bilingual interleaving**, the Loeb editions mix Greek or Latin into the English: Apollodorus ×2, Ovid’s *Heroides*
 - **Unusable**, the 1767 Statius: 18th-century long-s renders as “f” throughout
 - **Abridged**, Richard’s *Journey to the West* is roughly one sixth; Joly’s *Dream of the Red Chamber* stops at chapter 56 of 120
 - **No whitespace tokenisation**, the Chinese-language originals. Every length-based rule, including chunk size, needs a per-language definition
-

What is deliberately not in this contract

Segments below the unit. If a unit needs subdividing for a context window, that is a downstream decision recorded on the segment, not baked into preprocessing. Preprocessing finds the units the work itself declares; it does not invent boundaries.

Anything derived. No summaries, no entities, no embeddings. This layer is the write-ahead log and everything above it is rebuildable from it.

Fall 2026: backend, evaluation, paper

2026-08-27, revised 2026-08-28. Supersedes `archive/08_plan_of_record.md` for the fall. The spring distributable is a separate document, `06_spring_plan.md`, and a separate project. The cost model, previously a standalone document, is absorbed here and its arithmetic corrected.

The framing that governs every decision here: the output is not a corpus tool. It is **the backend for a desktop RAG assistant**, evaluated on literature because literature supplies ground truth that a personal archive cannot. Every design choice should be made as though a single user will point an AI client at this store, because in the spring they will.

The calendar, which is the binding constraint

Nine usable weeks in two blocks, and they are the only two.

Open: August 25 through September 27, with a homework set taking a bite around the 17th. October 19 through November 15, with a homework set on November 1 and quizzes on the 8th and 15th.

Dead: September 28 to October 4, October 5 to 11, October 12 to 18 (three consecutive weeks of tests, a midterm window, a machine problem and two quizzes), November 16 to 22, and November 30 to December 13, which is the entire PhD application window.

Budget: 70 to 100 hours, concentrated September 1 to November 15, at a ceiling of roughly eight hours in a normal week and near zero in a dead one. This figure is arithmetic over fixed commitments, not a measured study log, and it should be corrected rather than trusted.

Two rules follow, and they are not negotiable.

- **No milestone may be placed outside the two open blocks.** Any plan whose arithmetic exceeds nine weeks at eight hours is rejected on sight.
- **Complete by November 15 or it does not count.** The deliverable must be finished, committed and citable by then so it can appear in applications due December 1 to 15 and in the November letter brag sheets. A plan that targets December targets nothing.

Authorship rule. The student implements. An AI assists in drafting and discussion but is never the unexamined author, and completion is logged. This is the hardest constraint on the build and it is the reason hour estimates cannot be priced at assistant speed.

Positioning: what is occupied, and what this plan claims

As of the 08-27 and 08-28 searches, all of the following are occupied by published work:

Claim	Occupied by
Proactive recall, unprompted retrieval	CogniFold 2605.13438, ProactAgent 2604.20572, ENPMR-Bench 2605.27240
Cold-start schema induction	AutoSchemaKG 2505.23628, SCOPE/SCION 2607.21610, EvoTaxo 2603.19711
Narrative threads from conversation	TraceMem 2602.09712, SOTA on LoCoMo
Entity-centric summarization	EntSUM ACL 2022, EntSUMv2 EMNLP 2023, back to KDD 2012
Thread disentanglement	Active ACL/SIGIR/LREC task
Evaluation gaps, backbone dependence and cost	Anatomy of Agentic Memory 2602.19320
n=1 longitudinal personal corpus	MyLifeBits, and the 08-19 adversarial verdict
Covering versus partition	RAPTOR (ICLR 2024) soft-clusters by design; CAM (NeurIPS 2025) does incremental overlapping clustering. Closed 08-28
A literary memory benchmark	NarrativeXL (Findings of EMNLP 2023), MemoryAgentBench, StoryBench. Closed 08-28
A requirements analysis for assistant memory	Partially: Jones et al. CHI EA 2025, and 2606.24775 across twelve systems. Closed 08-28
Entity-first versus chunk-first extraction	iText2KG 2409.03284 (which ran the ablation and got a NEGATIVE result), RAKG 2504.09823, LINK-KG 2510.26486, CORE-KG 2506.21607. And the premise was wrong: both baselines are already entity-first. Closed 08-28

Every search this plan previously listed as outstanding is now closed, and every one came back occupied. **Nine candidate contributions, nine occupied. There is no surviving novelty claim.** See `01_argument.md` for how each died.

What this plan is now. A working backend, a published dataset, and measured cost and tier numbers. That is an engineering project with measurements and it does not require a novelty claim. Whether it also becomes a paper, and what that paper would claim, is a decision for Justin and his advisor, and it belongs in the Phase 0 conversation below alongside the administrative item.

Phase 0: decisions and freeze (now to September 7)

1. **Resolve the administrative item, and treat it as the gate.** The approved April proposal on file is a survey of LLM and NLP tools on government documents. Everything below is a different

project. A new one-semester proposal on the Fang, then Hasselbring and Tang, routing is still outstanding, and nothing on record shows the topic change was ever approved. This is the single highest-value outcome available and it is administrative rather than intellectual.

2. **Take the positioning question to the advisor, and treat it as a real decision rather than a formality.** Nine candidate contributions were searched and all nine are occupied. There is no novelty claim left. Present that plainly, present what survives without one (a working backend, a published dataset, measured cost and tier numbers), and ask whether the fall deliverable should be an engineering project with measurements or something else. The advisor steered toward a practical project twice before; this finding points the same way.
3. **Freeze the unit contract** (`04_unit_contract.md`, already frozen). It is the only schema element that must be settled before preprocessing.
4. **Decide the backend.** Files (JSONL, npy, SQLite) versus Neo4j. This plan assumes files, because brute-force exact cosine is correct below roughly 100K vectors and a personal corpus is far under that. Neo4j is a spring call.
5. **Turn on consult-logging**, and note that the logger emitting conforming rows is a separate deliverable from switching logging on. The instrument needs about four weeks of collection before its number means anything, so it cannot wait for the pipeline.

Exit: proposal filed, positioning agreed, contract frozen, logging running.

Phase 1: preprocessing (September, inside the first open block)

Raw corpora to normalised units. The student's code, visible, in a notebook.

- **Eight** per-convention handlers, one output contract, per `04_unit_contract.md`
- Table-of-contents count as the acceptance gate where a TOC exists; a monotonic-marker check where it does not; count equal to one for single-unit works such as plays and odes
- Gutenberg producer notes appear *inside* the start marker on some files and must be stripped

Deliverable: `data/clean/<corpus>/chapters.jsonl` for all four corpora, published as a public dataset. That publication is the reproducibility story and is worth more to a reviewer than any backend choice.

Note on the existing partial output. `data/clean/` already holds 351 records covering 16 of 81 works, committed on 08-27 as collateral in an unrelated documentation commit. It **does not meet the frozen contract**: it has no `unit_type` and no `unit_id`, and it uses `chapter_ordinal` where the contract says `unit_ordinal`. No script in this repository produced it, so it is not reproducible. Treat it as a scratch spike, not as progress: Phase 1 regenerates everything from a committed splitter.

Scope valve: Oz alone is sufficient to proceed. Greek and Chinese can follow.

Phase 2: pilot (mid to late September, before the block closes, about \$3)

Full pipeline over **Oz book 1 only**, 24 units and roughly 57,000 tokens, across three model tiers. Produces the sensitivity numbers, the rejection rates and the duplicate-minting curve, and surfaces schema gaps by contact rather than by reasoning. Freeze the remaining schema against what it finds.

This moved. It was previously scheduled for early October, which falls inside the three dead weeks of September 28 to October 18. It must land before September 27 or it slips to October 19 and eats the build block.

Note that book 1 averages about 2,375 tokens per unit, below the 2,800 corpus average, so it is a cheap pilot rather than a representative one. Book 2 is the hard case: it carries the `bare_title` convention and the Tip-to-Ozma fixture, and it should be the second thing run.

Phase 3: build (October 19 to November 7)

Pipeline and storage per `03_design.md`, behind the port described there.

Build order: ingest, then organize, then maintain, with the evaluation spine alongside. **Batch the cell calls from the start.** See the cost model below: batching cuts total input by about 2.3 times and the dominant cell-input term by 5 times.

The assembled renderer is built here, not later. It is a development instrument. A thin assembled page is a visible bug report, and it is the only way to see pipeline defects that a JSONL file hides.

Phase 4: measure and write (November 8 to 15)

The paper posts by November 15. December holds four graded events in another course, an exam window, and eight to ten PhD applications.

Committed measurements, in order of certainty:

1. **The free instruments.** Rejection rate per stage per tier, duplicate minting per unit, predicate sprawl, quote gate pass rate, token cost per arm, plot-versus-cell consistency. These fall out of running the pipeline and need no dataset.
2. **Per-stage tier sensitivity.** Which stages need a frontier model and which run on a cheap one, measured per question band. This is the measurement with a real spread behind it.
3. **The two corpus controls.** OCR tax and translation tax, both carried by Three Kingdoms.
4. **The contamination probe.** Free generation versus fact-row composition, deerstalker incidence.

Withdrawn 2026-08-28: the LitBank head to head, which was item 2 and the anchor of this phase. The claim it tested is occupied and its premise was factually wrong, and LitBank could not have

measured it in any case: 96 of its 100 documents are exactly two chunks at GraphRAG's default, and the effect compounds across many chunks. Nothing replaces it this semester.

The one paid benchmark worth running is LongMemEval, because it is where the assistant claim lives and where Zep published comparable numbers. It was missing from the previous version of this plan, which is a reconciliation error against 02_requirements_and_testing.md.

Cut, not stretched. NarrativeQA across three injection arms and the four-arm ablation are removed. They were listed as stretch goals in the previous version and are not in the committed scope; carrying them as aspirations is how the schedule became a factor of three oversubscribed.

The overlap is real and is stated rather than hidden. Three weeks of build and one week of measurement, with a homework set on November 1 and quizzes on the 8th and 15th, is tight to the point of implausibility. That is what the cut order below is for.

The cut order, decided in advance

The calendar guarantees at least one real week will disappear: four or more drill weekends at unknown positions, another course's midterm and final on unknown dates worth nearly half that grade, an unread family calendar, and a documented pattern of crashes after sustained grind. **A lost week deletes from the bottom.**

1. The signed proposal and the meeting package. Survives anything; nothing may displace it.
2. The instruments on the live prototype: consult-logging, miss rate, stale-serve rate, corrected-facts band. These run in evenings while coursework owns the weeks, and they are the paper.
3. The design document and the segmentation spec. These fit campus gaps and dead stretches by construction.
4. The pipeline slice (interfaces, adapters, splitter, ingest and organize on one corpus), if and only if the October 19 to November 15 block survives intact.
5. Everything else, in any order, because none of it is load-bearing this fall.

When a week vanishes, layer 5 is already gone, layer 4 shrinks stage by stage in reverse build order (organize before ingest), and layers 1 through 3 do not move. Phase 3 as written quietly depends on layer 4 existing; under this order the paper depends only on layers 1 through 3. Anything cut is cut in the log the day it is cut.

Cost model

Rates were fetched 2026-08-27 and are **perishable**. One caveat carried from the source document: the Anthropic rates came from a table cached 2026-06-24, two months older than the stated fetch date, and the Google page rendered a header dated in the future. Re-check every figure before it enters a budget or a paper.

Per unit the pipeline makes one entity pass, one fact pass, and N cell calls where N is the number of above-threshold nodes in that unit. **Each call re-sends the unit text.** Oz sizing: about 600 units, 1.27M words, roughly 1.7M tokens, about 2,800 tokens per unit, about 5 major nodes per unit.

Per unit	Input	Output
entity pass	2,700	300
fact pass	2,700	600
5 separate cell calls	13,500	1,250
total, separate	18,900	2,150
total, batched (one cell call)	8,100	2,150

Over 600 Oz units that is **11.3M input separate against 4.9M batched**, with output unchanged at **1.29M** either way, because batching changes how often the unit text is sent, not how much text comes back. Across all four corpora, roughly 3,200 units, batched: about **26M input and 6.9M output**.

Three numbers in the previous version were wrong and are corrected here. Batched Oz output was given as 1.5M; it is 1.29M, and it is identical to the unbatched figure. The saving was described as “roughly 3 times”; it is 2.33 times on total input and 5 times on the dominant cell-input term. And the provider spread was called “roughly 70 times”; the same table’s endpoints are \$335 and \$2.30, which is **146 times**.

Full four-corpus run, batched where a batch discount exists:

Provider	Model	Full run
Anthropic	Fable 5	\$335
Anthropic	Opus 5	\$168
OpenAI	gpt-5.6-sol	\$134
Google	Gemini 3.1 Pro	\$126
xAI	grok-4.6 (no batch)	\$102
Anthropic	Sonnet 5	\$67
Google	Gemini 3.5 Flash	\$56
Anthropic	Haiku 4.5	\$34
Google	Gemini 2.5 Flash	\$14
OpenAI	gpt-4o-mini	\$4.40
OpenAI	gpt-5-nano	\$2.30

Oz alone, batched: about \$6 on Haiku, \$13 on Sonnet 5, \$32 on Opus 5. **The pilot is about \$3** for all three tiers over Oz book 1.

What the spread changes. A 146-fold gap is not an academic question about which stages need a frontier model, it is the difference between a corpus that can be re-run every time the pipeline changes and one where iteration is rationed. **And the binding constraint remains hours, not dollars,** which the feasibility review concluded from the other side.

Why not run everything on the top tier. The cost is survivable; the methodology is not. The contribution is measurement, and “we ran it on the best model available” reports a ceiling and gives a reader nothing to act on. The spread is the finding. Iterate on the cheapest tier that works, report at minimum two tiers, and reserve the top tier for the judge, which must differ from every writer anyway and is cheap because judging is far fewer tokens than generating.

Two operational constraints. The prototype’s fact layer has failed twice on credit balance, so instrument spend per run and set a hard stop before starting a full pass; the Run record already carries the token counts, so this is a query rather than new logging. And local models will fail on format rather than comprehension: the feasibility review predicts 30 to 40 percent rejection on strict nested JSON from 7B instruct models, and grammar-constrained decoding makes malformed output structurally impossible. Measure format failure and comprehension failure separately, because they have different remedies.

What “ultimately a desktop RAG backend” changes

- **The storage port exists from day one.** The spring swaps the adapter; the fall must not call one directly.
 - **Files, not Neo4j.** The distributable is a folder someone points a client at. A server dependency now becomes a port later.
 - **The MCP shape informs the retrieval API.** The twelve port operations should map cleanly onto tools a desktop client can call.
 - **Single-user, single-workspace assumptions are fine.** Multi-workspace partitioning is filesystem separation in the spring, not a policy layer now.
-

Realistic outcome

One instrument, measured on one corpus, plus a published dataset. That is a workshop paper or an arXiv preprint supporting an application. It is not a top-conference submission, and planning for one produces neither.

The feasibility review priced the honest demand at 225 to 375 hours against 70 to 100 available, and dated the first unrecoverable break at September 15 to 21. Nothing since has changed that arithmetic. What has changed is that the cost side is now known and small: the full four-corpus run is \$34 to \$67 batched, and the pilot is \$3.

Hours are the binding constraint. Cut scope, not measurement.

Spring 2027: the distributable

2026-08-27. A separate project from 05_fall_plan.md, with its own deliverable, its own risks and its own administrative gates. The fall produces a measured backend; the spring turns it into something a stranger can install.

Entry condition: the fall paper exists. This phase never starts at the cost of that.

What it is

A **desktop RAG backend for an AI assistant**, distributed as something a person installs on their own machine and points their client at.

Two access paths over one store on disk:

Path	Needs running	Audience
MCP over stdio	nothing, the desktop client spawns the process	Claude Desktop and equivalents
Local HTTP + browser extension	a process while they browse	people who live in claude.ai in a browser

Same JSONL, .npy and SQLite underneath. “Always-on server” is a requirement of the second path only, not of the product.

The framing, and it is deliberate

Multiple partitioned workspaces on paired devices, for one person.

Not multi-user. One account holder, one set of credentials, several isolated contexts, work separate from personal, a sensitive domain separate from the general store, a project scoped so it does not bleed into unrelated sessions.

This is not a euphemism. The live archive is already ten partitioned userspaces, two of them gitignored as local-only precisely because they should not mix. Context isolation for a single person is the thing that was actually built.

The framing has to be carried through into the implementation, or it does not hold:

Not this	This
username	workspace name
grant access to a spouse	pair another device

Not this	This
partitioned users	isolated contexts
an invite flow	a device-pairing flow

The auth code is for your laptop and your phone, not for another person.

Why this framing and not the household one

It resolves four problems at once: no question about whose API key serves whom, no minors, no IRB exposure from family members, and a stated purpose that matches every provider's terms without argument.

A household deployment remains technically possible and is off-label. Running one on your own machine for your own family is your call; it is not the product story and it does not appear in the documentation or the onboarding.

Architecture

Partition by filesystem. One folder per workspace, each with its own store, its own credentials, its own service links. Nothing shared, so there is no policy engine and nothing to police. This is what closes the multi-tenancy gap from the archived schema quality pass ([archive/15_schema_and_architecture.md](#)) for this deployment.

One process per workspace, not one process routing by token. If a single server picks the folder from a request token, a bug in token validation crosses the partition and the isolation was only ever code discipline. Separate process, separate port, separate working directory makes the boundary real.

Credentials live outside the synced folder, or encrypted at rest. Each workspace holds its own key, so secrets now exist in N places on one machine and every backup contains N sets. The live archive already has form here, an open credential-shape gate and unrevoked tokens in pushed history.

Three ownership layers, so updates never merge:

Layer	Owner	On update
Client code	host	always overwritten
Config, workspace name, persona, preferences, prompts	user	never touched
User extensions, their own additions	user	never touched

And rollback rather than merge. Snapshot, health-check, revert, the mechanism already exists in the Jessica build. Snapshot *per layer* so a bad code push does not cost the user's config and a bad user edit does not revert the host's fix. Keep a bounded number of snapshots.

Know what the guarantee is. Rollback triggers on “won’t load,” not on “loads and behaves wrong.” The user-facing promise is *it always starts*, not *it always works*.

The folder is portable. Copy it to another machine and the memory comes with it. No export, no migration, no account. That is a genuine product property and it falls out of the design.

What is already proven

This is not speculative architecture. Jessica’s Assistant v3 runs it: server on a desktop, Tailscale Funnel for reach, a browser extension as a thin client, an install page gating the download on a pairing code, a PWA for phones, and self-update with device-side snapshot and rollback.

The failures that build hit were **build-and-publish** problems, a zip writer emitting backslash separators, a syntax check running in the wrong parse goal, a client-side scanner pretending to parse JS, three interacting version gates. All fixed, all documented, and **fixing them once fixes them for every N**. They do not recur per user.

The genuinely N-dependent risks are two, and neither is architectural:

- **Auth.** One pairing code for one trusted person is not N identities on a public URL. Tailscale Funnel is publicly reachable; a code checked once at provisioning is not authentication.
 - **Support burden.** N people hitting problems you diagnose remotely without access to their machine.
-

Deliverables

Committed: - The MCP folder, store, server, manifest, install docs. Zero infrastructure, since the client spawns the process - Backend choice exposed through the storage port: files by default, Neo4j as an advanced option - A conformance test suite run against **both** adapters, or the one not developed against rots - Portability demonstrated: the folder moved to a second machine and working

Stretch: - The browser extension and local HTTP path - Multi-workspace pairing across devices - A household deployment as a personal demonstration, clearly off-label

Administrative gates

IRB, if there is a tester study. A department email recruiting testers whose usage informs the work is human-subjects research and needs review *before* recruitment. Weeks, not days. If testers might set it up for family, **scope the study to adults explicitly**, minors in research is a much heavier category requiring parental permission and child assent.

Per-user API keys as the default. Each tester brings their own. That keeps the host out of the accountability path and keeps the design robust to a provider changing its terms.

Chrome Web Store, only if it goes public. For a controlled test group, sideload from your own install page, you already have that flow, and it skips a review cycle that would scrutinise host permissions on claude.ai. If it ever grows past a known group, the store's automatic updating stops being friction and becomes the fix for the exact thing that broke repeatedly.

Symposium

One board is what ISU provides, 36" × 48", one foamboard, clips and an easel, per the 2025 guidelines. **Verify the 2026 guidelines before designing.**

Three columns, left to right: *what the backend holds* → *whether it works* → *what you can take*. Problem, evidence, artifact. The browsable wiki lives on a laptop rather than competing for board space, and the QR codes go to a website (no install), the preprint, and the repo.

Schema prior art, and what to take from it

Originally read against the superseded schema draft, now `archive/15_schema_and_architecture.md`. The changes below are all incorporated in `03_design.md`. Sources are papers already in `papers/`, read directly rather than recalled. Each section states what the source specifies, then what our schema should do about it.

Angles 2018, the property graph model

What it specifies. A property graph is a tuple $G = (N, E, \rho, \lambda, \sigma)$: nodes, edges, an incidence function mapping each edge to a node pair, a labelling function over nodes and edges, and a property function mapping node-or-edge plus property name to values.

Separately, and more usefully for us, a **schema** is $S = (TN, TE, \beta, \delta)$:

- TN , node types
- TE , edge types, disjoint from TN
- β , which properties each node/edge type carries, and their datatypes
- δ , which edge types are allowed between which pairs of node types

What we should take: δ . Our schema has kind on nodes and a free predicate vocabulary, with a note that “some predicates only apply to some kinds.” Angles gives that the formal shape and makes it enforceable. `parent_of` between two person nodes is valid; `parent_of` between a place and a thing is not, and today nothing would catch it.

This is cheap to add and it catches a class of extraction error that no quote gate can see, because a fabricated relation can still carry a genuine quote. Recommend adding a δ table to the predicate registry: `allowed(subject_kind, predicate, object_kind)` triples, with violations logged as rejections rather than silently stored.

Rost et al. 2021, bitemporal property graphs

What it specifies. Two period domains, `VAL_TIME` and `TX_TIME`, each with `FROM` and `TO` bounds. **Lower bound inclusive, upper bound exclusive**, written `[from, to)`, following SQL. Periods are a type definition rather than a data type.

Critically, temporal attributes attach not only to vertices and edges but **to individual properties**: `var1.prop1.TX_FROM` is a valid access path.

What we should take, in two parts.

The interval convention. [from, to) with the upper bound exclusive. This is free correctness and it prevents the off-by-one that appears at every boundary when two intervals abut. Our current schema says “validity intervals” without stating the convention, which is exactly how a project ends up with two conventions.

The property-level question, answered. Rost puts validity on individual properties because his vertices carry properties directly. Ours do not: our facts **are** the properties, so fact-level validity already gives us property-level granularity. Worth recording that we considered it and why we do not need it, rather than leaving it to be re-litigated.

Vrandečić and Krötzsch 2014, Wikidata

What it specifies. Property-value pairs on items can carry subordinate property-value pairs called **qualifiers**. Qualifiers do two jobs: they state contextual information such as validity time, and they **encode ternary relations that the property-value model cannot express**. Their example: “Meryl Streep played Margaret Thatcher in The Iron Lady” becomes `cast_member = Meryl Streep` on the film, plus the qualifier `role = Margaret Thatcher`.

Wikidata also carries **ranks**, preferred, normal, deprecated, which let contradictory statements coexist while marking which is current.

What we should take.

Qualifiers are load-bearing, not decoration. Our schema has a `qualifiers` field documented only as “any timing the source states.” That undersells it. N-ary relations are common in every corpus here: Ozma rules Oz *as restored by Glinda*; Holmes is retained *by a client for a fee*; Zhuge Liang serves Liu Bei *in the capacity of Chancellor*. Without qualifiers those become either lost detail or invented predicates, and invented predicates are the sprawl that breaks supersession.

Ranks fill a real gap. Our design says nothing is overwritten and supersession is computed at read time from a later assertion. That works when the world changed. It does **not** work when a fact is simply wrong, an extraction error has no later assertion to supersede it, and deleting it violates append-only.

This is not hypothetical: the Ash-Catchum-aliased-to-Ursa error in the live archive was resolved by appending an is-distinct-from assertion rather than editing. That is the same pattern, done ad hoc. A rank field makes it explicit and queryable: `deprecated` means present, preserved, and excluded from reads by default.

Recommend adding rank: preferred | normal | deprecated to the fact row, defaulting to normal.

Rezazadeh et al. 2025, Collaborative Memory

What it specifies. Two-tier memory: private fragments visible only to the originating user, and shared fragments. Each fragment carries **immutable provenance attributes**, contributing agents, accessed

resources, timestamps, to support *retrospective* permission checks. Read policies project stored fragments into filtered views at query time; write policies decide retention and sharing. Access is modelled as a **time-varying bipartite graph** over users, agents and resources.

What we should take: this is the answer to quality-pass item 1. Our schema has no multi-tenancy at all, and the product’s stated end state is one user scaling to an organisation.

The important design property is that **permission is enforced at read time against current policy, not stamped at write time**. That means revoking access retroactively hides everything that depended on it, without rewriting stored fragments, which is the same read-time-resolution discipline we already use for supersession and merges. It fits our model rather than fighting it.

Concretely: add immutable provenance to every derived record (who contributed it, which resources it drew on, when), and make the storage port’s read operations take a policy context rather than returning everything. Single-tenant is then the degenerate case where the policy admits all.

Summary of schema changes these sources drove

Change	Source	Why
Add δ , allowed (subject_kind, predicate, object_kind) triples	Angles 2018	Catches fabricated relations that carry genuine quotes, which the quote gate cannot see
State the interval convention as [from, to), upper exclusive	Rost 2021	Free correctness at every boundary
Record that fact-level validity subsumes property-level	Rost 2021	Stops the question being re-opened
Document qualifiers as carrying n-ary relations, not just timing	Wikidata 2014	Prevents predicate sprawl from encoding roles as new predicates
Add rank: preferred/normal/deprecated to facts	Wikidata 2014	Supersession handles “the world changed”; rank handles “we were wrong”
Add immutable provenance + read-time policy context to the port	Collaborative Memory 2025	Multi-tenancy, with single-tenant as the degenerate case

Deliberately not adopted

RDF reification schemes (Hernandez et al. 2015, four schemes benchmarked over 57M Wikidata quads). The question they answer, how to attach metadata to a triple, is native in a property graph, where edges carry properties directly. Their finding that singleton properties broke four of five engines is a warning about RDF stores, not about us.

CoALA’s four-way memory taxonomy (working / episodic / semantic / procedural). A useful frame for describing the system in prose, but it does not map onto storage: our units are episodic, our facts and abstracts are semantic, and we have no procedural memory at all. Adopting it as schema would add a field that never varies usefully.

References

Bibliographic details extracted from the PDFs in papers/, not recalled.

All four items previously marked [verify] were resolved on 2026-08-28. Two were correct and needed only completion; two were materially wrong. Nothing in this file is now uncitable. The resolution method is recorded with each entry.

Angles, Renzo. “The Property Graph Database Model.” In *Proceedings of the 12th Alberto Mendelzon International Workshop on Foundations of Data Management (AMW 2018)*, Cali, Colombia, May 21-25, 2018, eds. Dan Olteanu and Barbara Poblete. CEUR Workshop Proceedings, Vol. 2100, paper 26. <http://ceur-ws.org/Vol-2100/paper26.pdf>

Single author, not “et al.” The paper is unpaginated and has no DOI; those fields are legitimately absent rather than unconfirmed. **Venue confirmed 2026-08-28** three ways: the local PDF is byte-identical to the CEUR-WS copy (both 680,900 bytes, matching SHA1), the CEUR-WS Vol-2100 index lists it, and Rost’s own reference list cites it the same way. Local: papers/angles2018-property-graph-model.pdf.

Do not confuse this with papers/angles2017-graph-query-foundations.pdf, a different six-author ACM Computing Surveys paper that is not cited here.

- *Property graph data structure*, **Definition 1:** $G = (N, E, \rho, \lambda, \sigma)$
- *Schema, including δ* , **Definition 2:** $S = (TN, TE, \beta, \delta)$, where δ is “a partial function that defines the edge types allowed between a given pair of node types”

Rost, Christopher; Fritzsche, Philip; Schons, Lucas; Zimmer, Maximilian; Gawlick, Dieter; Rahm, Erhard. “Bitemporal Property Graphs to Organize Evolving Systems.” Subtitle: “Towards the development of a graph model, database, and query language to represent, store, and query bitemporal graphs.” University of Leipzig / ScaDS.AI Dresden-Leipzig / Oracle. 21 pages. **arXiv preprint arXiv:2111.13499 [cs.DB], 26 November 2021.** DOI: 10.48550/arXiv.2111.13499. <https://arxiv.org/abs/2111.13499>

Resolved 2026-08-28: there is no venue. The document’s LNCS-style layout is misleading. arXiv lists no journal reference, DBLP carries it only under CoRR as an informal publication, and the authors’ own group page labels it a University of Leipzig technical report. Cite it as a preprint and technical report; do not assign it a workshop or conference. Local: papers/rost2021-bitemporal-property-graphs.pdf.

- *Period identifiers and interval convention*, ~p. 6–7: VAL_TIME / TX_TIME domains, bounds VAL_FROM/VAL_TO and TX_FROM/TX_TO, textual form [{from},{to}], lower bound inclusive and upper exclusive, “like in SQL, not a data type but a type definition”
- *Property-level temporal attributes*, same section: var1.prop1.TX_FROM as a valid access path

Vrandečić, Denny; Krötzsch, Markus. “Wikidata: A Free Collaborative Knowledgebase.” *Communications of the ACM* 57(10), October 2014, pp. 78–85. DOI: 10.1145/2629489. Fully verified from the document. Local: papers/vrandecic2014-wikidata.pdf.

- *Qualifiers*, p. 82: the subordinate property-value pair model, including the Meryl Streep, Margaret Thatcher and *The Iron Lady* ternary-relation example
- *Statement marking*, p. 83: contributors may optionally mark statements as “preferred” or “deprecated”

Correction, 2026-08-28: the rank vocabulary is not in this paper. Two independent extractors over all nine pages find zero occurrences of “rank” and no “normal” value; only “preferred” (once) and “deprecated” (twice) appear, both on p. 83. Controls returned results on the same pass (“qualifier” 9 hits, “Meryl Streep” 2), so the absence is real. For the three-value vocabulary, cite the Wikidata help page below instead.

Rezazadeh, Alireza; Li, Zichao; Lou, Ange; Zhao, Yuying; Wei, Wei; Bao, Yujia. “Collaborative Memory: Multi-User Memory Sharing in LLM Agents with Dynamic Access Control.” Center for Advanced AI, Accenture. arXiv:2505.18279v1 [cs.MA], 23 May 2025. Local: papers/rezazadeh2025-collaborative-memory.pdf.

- *Two-tier memory, immutable provenance, read-time policy projection*, abstract and related-work section, ~p. 3: “Each fragment carries immutable provenance attributes (contributing agents, accessed resources, and timestamps) to support retrospective permission checks”
- *Dynamic access graph*, same section: “time-varying bipartite graphs between users, agents, and resources”

Cited in “deliberately not adopted”

Hernández, Daniel; Hogan, Aidan; Krötzsch, Markus. “Reifying RDF: What Works Well With Wikidata?” In *Proceedings of the 11th International Workshop on Scalable Semantic Web Knowledge Base Systems (SSWS 2015)*, co-located with ISWC 2015, Bethlehem, PA, USA, October 11, 2015, eds. Thorsten Liebig and Achille Fokoue. CEUR Workshop Proceedings, Vol. 1457, pp. 32-47. http://ceur-ws.org/Vol-1457/SSWS2015_paper3.pdf

Completed 2026-08-28. Trap to avoid: the PDF’s own /Subject metadata field reads “Scalable Semantic Web Systems 2014”, a stale LaTeX template value. The venue is SSWS 2015. Its substantive claims check out: 57,088,184 Wikidata quads (Table 2), and the Conclusions state that four of five engines “struggled with the high number of unique predicates generated by singleton properties” (Virtuoso was the exception). Local: papers/hernandez2015-reifying-rdf-wikidata.pdf.

Sumers, Theodore R.; Yao, Shunyu; Narasimhan, Karthik; Griffiths, Thomas L. “Cognitive Architectures for Language Agents.” *Transactions on Machine Learning Research*, February 2024. arXiv:2309.02427v3, 15 March 2024. OpenReview forum?id=1i6ZCvf1QJ. Local: papers/sumers2024-coala.pdf.

Added 2026-08-28

Wikidata contributors. “Help:Ranking.” Wikidata, Wikimedia Foundation. Revision of 11 August 2026 (oldid 2530015510). <https://www.wikidata.org/w/index.php?title=Help:Ranking&oldid=2530015510> (accessed 27 August 2026).

Defines all three values: *preferred* (“the most current statement or statements that best represent consensus”), *normal* (“assigned to all statements by default”), and *deprecated* (“statements that are known to include errors”). Use the oldid permalink so the citation does not drift. This is the source for rank in 03_design.md, not the CACM paper.

Standing rule for this file

Every claim attributed to a source above was read out of the PDF or the publisher’s page, not recalled. Anything added later must meet the same bar.

papers/MANIFEST.md does not cover these four sources. It holds 72 rows and has no entry for angles2018, rost2021, hernandez2015 or vrandecic2014, all of which were fetched after the manifest’s own stated fetch time. Its header also claims “34/43 on disk” against 92 PDFs actually present. Repair it before the manifest is used as an inventory.

REFERENCE, not superseded. This is the menu of candidate mechanisms, each traced to a source. Roughly seven of its citations appear nowhere else in the repo. It is background rather than part of the working set. Index: docs/README.md.

Ways the research could elevate the project

Each item is a candidate mechanism for the general design, scoped to its delta beyond the prototype (the evidence base and first tenant), usually the automated, measured, or scaled version of a manual practice. Citations outside the packaged reading list carry arXiv identifiers inline. Effort figures are orders of magnitude; items exceeding the fall budget of 70 to 100 hours are marked spring-scale. Nothing is ranked.

Capture

A revived runtime citation log

A PostToolUse hook appends every tree read to a session log, making a leaf's citations transcription rather than recollection. The prototype built this (`cite_log.py`) but shelved it unrun; OPERATIONS.md argues the case in prose. CoALA (Sumers et al. 2024) treats memory writes as deliberate internal actions; Noy et al. 2019 report provenance-first assertion as the pattern all five industrial graphs converged on. The rebuild feeds the cued-versus-read-versus-load-bearing distinction OPERATIONS names but never mechanized; the counterfactual miss instrument below consumes its output. Effort: days, from git history.

Outcome-signal mining from correction turns

A pass that scans transcripts for implicit negative feedback, the rephrasings and “no, I meant” turns, and converts recurring ones into failure-log candidates. Liu, Zhang and Choi 2025 show later user turns are dense with correction signal; ReasoningBank (Ouyang et al. 2026) shows failure records carry value when the schema stores lessons rather than replays. `_failure_log.jsonl` records only misroutes somebody noticed, not wrong answers; nothing mines the archive for unlogged misses. Effort: a week, including a hand-checked precision sample.

A segmentation quality lane in the stale-leaf slice

An audit of whether multi-topic sessions were cut into leaves at the right boundaries, with permission to re-cut. Event segmentation theory (Zacks et al. 2007) treats the cut as an encoding decision; Reflective Memory Management (2503.08026, ACL 2025) shows rigid granularity fragments conversational structure. The stale-leaf slice re-summarizes in place but never re-segments. The delta is a sampled bad-cut rate plus a re-segmentation lane in the existing burn-down. Effort: days for the sample, a week for the lane.

Structure and representation

A one-era property-graph trial

A representation experiment; knowledge graphs are one candidate, not the destination. Render one era's ledger rows as a typed property graph under the Angles 2018 formalization, whose delta function (an edge-type whitelist per node-type pair) doubles as a vocabulary contract for an extraction pipeline; answer the same golden questions by graph walk versus sheet read, scoring accuracy and hop count. The ledger is already subject-predicate-object rows with qualifiers, so most of a graph exists; the delta is typed edges and multi-hop query. Rossi et al. 2021 show sparse personal graphs are the worst regime for embedding methods, so the trial stays symbolic with lightweight schema (Hogan et al. 2021). Effort: days.

Recorded currency judgments as dated rank rows

When a read-time judgment decides which of two conflicting rows is current, record that judgment as its own dated, sourced row, following Wikidata's preferred and deprecated marks (Vrandečić and Krotzsch 2014) and Fowler's 2021 observation that record history stays append-only even under correction. Nothing is overwritten, so the facts-cannot-be-wrong rule holds; an expensive judgment is made once and cached as data rather than remade at every read. The ledger already stores corrections as new rows and imprecision as date_precision qualifiers; the currency call lives only in the reading agent's head. Effort: days; the stale-serve measurement below can consume these rows.

Entity resolution measured on the existing registry

A labeled set of alias pairs drawn from the corpus (Defy Medical, Defy, the clinic; GAO across four eras) and a measured resolution accuracy over it. Balog and Kenter 2019 call long-tail personal entity linking the defining open problem of personal knowledge graphs; Noy et al. 2019 call disambiguation the top challenge across all five industrial graphs; Zhong et al. 2024 supply the task decomposition. The two entity systems (registry_entities.txt and the ledger's) produced 70 versus 388 entities with only 17 names in common; the corpus demonstrates the problem, nothing measures it. Effort: days to a week.

Currency

A stale-serve rate

How often a recall surface serves a superseded fact. MemStrata (2606.26511) supplies the metric shape and reports plain RAG serving outdated facts 15 to 40 percent of the time; MemoryAgentBench (Hu, Wang and McAuley 2025) finds no tested system above 28 percent on multi-hop fact updates. Drift-debt measures whether a leaf is current against its source; nothing measures whether answers reaching a session were stale. Method: a probe set built from documented supersessions in the corpus (an internship ending, a session renumbering, a topic pivot) run against tree and ledger reads. Effort: days once the probe set exists; the typed gold set item can supply it.

An implicit-conflict probe

Whether the system detects conflicts carrying no explicit negation, where a later observation quietly falsifies an earlier row. STALE (2605.06527) defines the task and reports a best-model score of 55.2

percent; Pollertlam and Kornsuwannawit 2026 name update non-propagation, a “twice a week” fact surviving its revision to three, as a leading failure of flat fact extraction. The tail-scan rule catches corrections inside one conversation, but the validator re-checks only that quotes still appear, never that a claim went false; cross-conversation supersession has no detector. This prices the design choice against Zep-style invalidation-on-write without reversing it. Effort: a week; measurement only, no write-path change.

Retrieval and delivery

A counterfactual miss instrument

How often the store held the answer and the session never consulted it. The project’s strongest publishable claim, so far only a limitation note (access is not recall): every published benchmark poses the query, so this failure cannot occur in their setups, and Remember When It Matters (2607.08716) covers only the intra-task horizon. The failure loop logs misses somebody noticed; missing is the denominator, the misses nobody noticed. Method: seed sessions with tasks whose answers provably sit in the tree, instrument reads through the citation log, count non-consultations. Effort: a week to build, data accruing on normal use; depends on the citation log item.

A retrieval-depth router

A cheap per-query decision among no lookup, one-hop rollup read, and multi-hop descent. Adaptive-RAG (Jeong et al. 2024) matches its always-iterate baseline at a third of the steps and shows usage logs can self-label the router by recording which depth sufficed; Wang et al. 2024 find the decide-whether-to-retrieve gate the single cheapest accuracy and latency win in their module search. Current policy is static: the primer’s standing instruction plus the cue hook’s boundary triggers. The delta buys cheaper recall plus a free labeled dataset of query difficulty over the corpus. Effort: a week.

The endorsed escape-hatch index, built and measured

A brute-force exact-cosine embedding index over leaf text, no vector database, consulted only when tree descent lands on the wrong branch. The design already endorses this role, with the flattened-CBC example as the worked case: at roughly 1,200 files exact search is correct and index machinery is overhead, consistent with Faiss sizing guidance (Douze et al. 2024) placing the brute-force crossover far above this scale. NoLiMa (2025) supplies the motivation: lexical mismatch is the normal case, so retrieval must bridge vocabulary before the model sees anything. The delta is an afternoon of code plus a probe set of vocabulary-mismatch questions measuring wrong-branch recovery. Effort: days.

Self-selection routing across eras (spring-scale)

The blackboard pattern applied to the tree: partition agents that have each pre-digested one era volunteer answers when a query touches their subtree. Salemi et al. 2026 report 13 to 57 percent end-to-end gains over master-slave routing, growing with corpus size. `directory.md` plus the curated boundary lines is exactly the central index whose staleness the blackboard result targets, and the era summaries are already the pre-digestion. It buys an architecture argument for organizational scale-out. Effort: weeks, with payoff scaling on corpus size the project does not yet have; spring-scale on both counts.

Evaluation and measurement

E2, a typed gold set built on published task definitions

An extension of EXPERIMENTS.md: 100 to 150 gold questions typed by the published categories (single-hop, multi-hop, temporal, knowledge-update, abstention) from LoCoMo (Maharana et al. 2024) and LongMemEval; benchmarks supply task definitions and comparability, never a contest to win. `_golden.jsonl` exists as a frozen deterministic set, the adversarial-sampling rule is written, and E1 established the file format; the delta is size, typed coverage, and category-level reporting. Agent fleets make the labeling tractable inside the budget; completion accounting is mandatory, since five of twelve automated audit checks have previously stopped running without raising an error. Effort: a week of fleet runs plus review.

A calibration set for the faithfulness gate

A measurement of the gate itself. FABLES (Kim et al. 2024) shows automatic faithfulness checking reaches 58.2 F1 at best on the indirect multi-hop claims rollups are made of; the cold audit agrees locally: eight defects, all in the rollup layer, all invisible to entailment checking because a hardened sentence is still entailed by its hedged source. The faithfulness metric already runs every pass, and the 39 hand-verified claims from the cold audit are a free labeled seed. The delta scores the gate against those labels, reports precision and recall, and extends the taxonomy to the omission and salience errors BoookScore (Chang et al. 2024) shows dominate hierarchical summarization. Effort: days.

A layer-localized error study of the write path

The 39-claim audit scaled with agent fleets and reported by layer: transcript to leaf, leaf to node, node to era. BoookScore scores only final summaries and cannot localize which level introduced an error; Wu et al. 2021 documented compounding errors at composition points; Talebirad et al. 2026 prove the atom layer sets the information ceiling, making the fold step the place losses concentrate. The rotating source audit is the standing mechanism, and the cold audit ran this design once by hand: 31 confirmed, 4 wrong, 4 unsupported, every defect in the synthesized layer; coverage sits at 2.3 percent. The delta is coverage, statistics, and the by-layer cut: the paper’s motivating study on existing machinery. Effort: a week of fleet time.

The hash-versus-dirty-flag demonstration

The controlled experiment behind the project’s strongest surviving mechanism claim. Hand-edit N leaves out of band, then show that a write-time dirty flag, the scheme MemForest (2605.23986) and every build system use, misses all N while content-hash recomputation catches all N: a flag is only correct if every writer remembers to set it, and a recomputed hash is self-checking. `fold_sig` is implemented, tested, and running; only the comparison arm and write-up are new. Effort: days. It pairs with the layer study as a short paper’s experimental section.

Scale-out

A permission-graph formalization of the existing boundaries (spring-scale)

A recasting of the guard patterns, the holds ledger, and the gitignored eras as the provenance-stamped fragments and retrospective permission checks of Collaborative Memory (Rezazadeh et al. 2025), where revoking an edge in the access graph retroactively hides every fragment derived through it. The single-user version already runs: CUI boundaries, held-local UUIDs, and the guard-your-own-writing rule are hand-enforced policy over the primitives that paper formalizes. Multi-tenant access is the one thing the current design cannot claim. Effort: weeks, spring-scale, though a two-page note mapping current mechanisms onto the formalism would fit a fall afternoon.

A measured cache-tier policy for the primer

Treat the SessionStart primer as the preloaded tier of a two-tier design and set its boundary from data. CAG (Chan et al. 2025) shows that below some corpus size precomputed context beats retrieval outright, and its own large-corpus result shows where that stops; Pollertlam and Kornsuwannawit 2026 supply break-even arithmetic between resending history and extract-then-retrieve, with memory winning after roughly ten turns at 100K tokens. The tiering already exists: the primer preloads the era map, the dossiers load on demand, everything else sits behind search, with no measurement saying what belongs where. The delta varies primer contents and measures recall against token cost per session. Effort: days.

REFERENCE, not superseded. This is the stage-by-stage dissection of the running prototype, with file-and-line receipts. Its by-hand table is the harness component backlog and is the specification baseline the build is measured against. It is background rather than part of the working set. Index: docs/README.md.

Anatomy of the mock: every stage, and what executes it

This document dissects the running claude-archive memory pipeline stage by stage; every claim carries a file-and-line receipt. It is the specification baseline for the engineered harness: corpus in, memory tree plus entity knowledge base out, efficient injection at query time. Of 49 distinct stages, 26 execute end to end as deterministic stdlib Python, but those 26 are plumbing (parsing, hashing, gating, scheduling, serving, alarming). Every stage that creates or changes meaning runs by hand: 23 stages have a human, a scheduled agent, or an in-session LLM as their generative core (1 human-only, 8 with an LLM or agent as sole executor, 14 mixed), all executing prose protocol from OPERATIONS.md rather than code executing a contract. No LLM output in the tree layer is schema-validated at write time; the single write-time contract in the system is the fact ledger's `ledger.append()` validator. Quality is checked only after the fact, by sampling.

The pipeline as it exists

Capture

1. Source acquisition. EXECUTED BY: **human**. Two asymmetric channels. Claude Code sessions self-record: the CLI writes append-only JSONL per session to `C:\Users\jhffm\claude\projects\<project>\<uuid>.jsonl` with no archive code involved. Web-chat corpora (ChatGPT, claude.ai, Gemini) exist only when the operator requests a vendor takeout into `C:\Users\jhffm\claude-archive\raw\`. Judgment: whether and when to request a takeout; the only reminder is weekly-pass step 1 (OPERATIONS.md line 105). Fragility: ChatGPT is frozen at the 2026-07-11 takeout and Gemini at 2026-07-31; nothing measures that staleness, and no freshness alarm watches `raw\`.

2. Nightly trigger. EXECUTED BY: **deterministic-code**. Task Scheduler job `ClaudeArchiveNightlySync` (02:30, 12:30, 19:15, plus wake-catchup) runs `wscript.exe scripts\nightly_hidden.vbs`, which launches a hardcoded python 3.14 path on `scripts\maintain.py` with a hidden window. Heartbeat is the rewritten `_last_maintain.log` (MARKER.write_text, maintain.py line 249). Judgment: none. Fragility: the heartbeat mtime is rewritten even when the push fails, so a “worked but didn’t ship” run reads green to the 2-day staleness rule (OPERATIONS.md line 506); the hardcoded python path dies silently on a Python upgrade (Task Scheduler still reports 0 because wscript succeeded). As of 2026-08-25 both this task and AgentMemoryKit read DISABLED in Get-ScheduledTask; no file records why.

3. Change gating. EXECUTED BY: **deterministic-code**. `raw_changed_since(MARKER)` (maintain.py lines 78-83) compares mtimes of `RAW.glob('*')` against `_last_maintain.log`; the heavy export adapters

run only on change (lines 94-99). Judgment: none. Fragility: the glob is non-recursive (a takeout in a raw\ subfolder never triggers); a zip dropped mid-run can be skipped until touched again; only tail[-1] of each child's output is logged (lines 107-108), so a parser traceback survives as one truncated line while the run continues.

4. Parsing transcripts into corpus files. EXECUTED BY: **deterministic-code**. One parser, four adapters (scripts\parse_sources.py: from_claude_code 177, from_claude_ai 235, from_chatgpt 263, from_gemini 324) feed a shared emit() (113). Identity is <source>:<full id> from the YAML header via conv_key() (65), never the filename; rewrite happens only when the recorded message/event count changes (_recorded_count, 92); drifted filenames are renamed via os.replace; two-plus priors for one id report AMBIGUOUS and touch nothing (conversations\ is immutable by house rule; merging falls to the operator via the printed report). Current corpus: 1060 files (74 ccd, 560 gpt, 219 gemini, ~207 unprefixed claude.ai). Fragility: each export adapter returns after the first archive whose shape matches, newest mtime first, so an older takeout with unique conversations is silently never parsed (bare return at 259/301/369); _load() swallows BadZipFile/JSONDecodeError as “no conversations” (218-231); count-equality cannot see same-count edits; Gemini identity is fabricated as date plus session index (line 365), so a re-export would mint duplicates wholesale; ccd messages starting with [are dropped as noise (195).

5. Code-artifact extraction. EXECUTED BY: **deterministic-code**. Because stage 4 strips tool blocks, session code exists only in the raw jsonl; scripts\extract_code_artifacts.py replays Write/Edit/MultiEdit per session (reconstruct, 76) into gitignored code_artifacts\<id8>\, then --programs rebuilds per-(era, filename) end-states across all chats with a frozen latest-within-85%-of-largest rule, pasted-fence harvesting (_harvest_pasted, 223), and a mechanical audit (_audit_body: ast.parse authoritative) to _programs_AUDIT.md. 2026-08-24 log: 148 snapshots, 1045 end-states, audit OK 659 / SUSPECT 81 / SYNTAX-ERROR 11. Fragility: maintain.py runs the identical snapshot pass twice per nightly (lines 114-120 and 152-158); files created via Bash heredocs, python -c, or generators are invisible; filename attribution for pasted fences has documented misfires; unclassified conversations default to era misc via _conv_era (205).

6. Guard, commit, push. EXECUTED BY: **deterministic-code** (disposition of holds is human; see the inventory). git status --porcelain -uall enumerates changes; allowlisted suffixes are scanned against SECRET_PATTERNS (absolute) and CUI_PATTERNS (diff-aware versus git show HEAD:<path>), with site-specific fragments from gitignored _guard_patterns.txt (missing file: WARNING and reduced coverage). Quarantined paths go to _REVIEW_NEEDED.txt and are unstaged after git add -A; then commit and push, push failure logged, never fatal (maintain.py 160-245). Fragility: as of 2026-08-24 the run ended “PUSH FAILED... fast-forwards”, local commits stranded, heartbeat still green; secrets inside .docx/.png/.pdf are never scanned; novel credential or CUI phrasing auto-pushes (accepted residual risk, docstring 18-21); _REVIEW_NEEDED.txt is only rewritten when something new quarantines, so a drained list reads identically to an ignored one (OPERATIONS.md line 540).

Segment, summarize, file (distill)

7. Rule delivery. EXECUTED BY: **deterministic-code**. The session-close rule is not self-executing; it reaches sessions via the SessionStart hook (mcp_server\session_recall.py, primer strings 160-167) and

the global `~\.claude\CLAUDE.md` standing block, mirrored by hand per `OPERATIONS.md` line 602. Fragility: `OPERATIONS.md` line 31 calls this “the single most load-bearing surface and the easiest to lose”; the hook is fail-silent, an unprimed session simply never writes leaves, and nothing checks the `CLAUDE.md` mirror.

8. Session-close leaf writing. EXECUTED BY: **llm-in-session-per-protocol**. At the end of any substantive session, the in-session LLM writes one leaf per topic segment into `eras\_staging\`: filename `YYYY-MM-DD_<slug>_<convId8>.md`, frontmatter (date, era, topic, people, decisions, code_artifacts, open_threads), body prose, verb-tagged `## Related`, code pointed at rather than inlined (`OPERATIONS.md` lines 40-48, 44, 75-99). Judgment: substantive or not; segment boundaries (the whole specification is “segment by topic”); era/topic routing; what survives summarization; citation from recollection; whether a mid-conversation figure was later retracted (tail-scan rule). Fragility: nothing enforces this at write time; skips surface only as coverage debt (v2 baseline 67.4%, 244 un-leafed, line 294); the `_live-session` misname happened three times, two leaves now unkeyable (line 44); the runtime citation layer (`cite_log.py`) was never wired and was deleted 2026-08-21 (line 450); no validator ever runs on live-session leaf frontmatter.

9. Import digestion (weekly steps 1-2). EXECUTED BY: **scheduled-agent**. The weekly agent diffs `conversations\` against `eras\MANIFEST.md` (PowerShell enumeration, never Glob, per the ripgrep false-negative gotcha), then per unclassified transcript: segment, write staged leaves, add `MANIFEST` rows (`| File | Proposed topic folder | Gist | Source |`; held-local rows carry no gist, lines 574-578); leaf-unworthy conversations go to `eras\_no_leaf.md` as reasoned decisions the coverage gate reads. This is the one place raw is legitimately re-read (line 16). Fragility: on 2026-07-30 a pass announced its worklist and died writing nothing, invisible for five days (leaf `eras\claude-archive-meta\maintenance-and-review\2026-07-30_weekly-run-died-mid-pass_cbfff1ba.md`); nothing catches a bad split; `MANIFEST` is a large hand-edited table with documented drift (495 rows once named dead folders).

10. Promotion and routing. EXECUTED BY: **mixed**. Clean staged leaves move into era/topic folders; routing is LLM judgment against `eras\_boundaries.md` discriminator lines merged into the regenerated `eras\directory.md` (`rollup_tools.py` `directory`, `OPERATIONS.md` line 150); cwd never decides era (line 109); a misroute earns a failure-log row and a new boundary line. Fragility: “reviewed/clean” has no checklist; `_staging` is an unbounded-latency buffer; the boundary parser silently clobbers duplicate keys (its own header warning); routing is only ever corrected reactively.

11. Leaf baselining (stamp-leaves). EXECUTED BY: **deterministic-code**. After promotion, `rollup_tools.py` `stamp-leaves` writes `<!-- rollup: {"src_sha":..., "src_stamped":...} -->` into the leaf (write at ~1324-1330); scan later reads `stale/current/UNMEASURED` (never “clean”) from it; existing baselines are protected and `--force` is banned as a drift launderer. Fragility: unstamped leaves stay `UNMEASURED` indefinitely; doc-import leaves keyed `doc-<hash8>` sit outside the id-keyed drift machinery; the stamp records currency, not correctness (line 296).

12. Document-import variant (one-off, 2026-08-23). EXECUTED BY: **mixed**. ~1,042 filesystem documents became 73 leaves through a 14-step pipeline (`scripts\imports\2026-08-23-doc-import\README.md`): deterministic CUI scan, extraction, backup, clustering, validation (`validate_leaves.py`: REQUIRED frontmatter list, hardcoded KNOWN node map, `source_doc`

resolution), em-dash sweep, promotion, additive fold; LLM cores at triage (13 agents), distillation (13 agents), and adversarial verification (9 agents, 9 cross-doc-bleed findings fixed). Fragility: this is the only mechanical frontmatter/routing validation anywhere, and it is a one-off with an already-stale node map that never runs on live leaves; triage manifests were not committed, so a re-run reconstructs judgment from scratch.

Fold

13. Scan. EXECUTED BY: **deterministic-code**. `cmd_scan` (rollup_tools.py 742-978) rehashes every conversation (sha, 88; multi-file ids get a combined hash, 768-773) and every tree file with markers stripped (sha_nomarker, 96-101), computes recursive per-node `fold_sig` over (relpath, content-hash) pairs (789-805, duplicated in `_compute_fold_sig` 1124-1137), classifies node and leaf staleness from in-file markers, and computes coverage, refcheck, compression, golden, boundaries, fact-layer metrics, and MCP health into `_ledger.json` and `_metrics.json`. Fragility: two `fold_sig` implementations could drift apart; coverage percent covers the measurable only.

14. Node staleness detection. EXECUTED BY: **deterministic-code**. A node is stale iff its marker's `fold_sig` differs from the recomputed one; no marker means stale; `cmd_idempotency` (1051-1062) always rescans first (the 2026-07-26 stale-ledger bug) and prints the refold worklist. Because markers are stripped from hashes, stamping never cascades; a real child edit does. Fragility: `fold_sig` proves the children changed, not that the summary is good.

15. Refold (node summary prose). EXECUTED BY: **scheduled-agent**. The weekly agent rewrites each stale node's ERA_SUMMARY/NODE_SUMMARY from child summaries only, never raw (OPERATIONS.md steps 4 and 7), curating promoted links under the altitude test. No function in `rollup_tools.py` writes summary prose (verified across all 2906 lines). Judgment: everything (compression, spine, link promotion, split/merge). Fragility: quality checks are after the fact and narrow; faithfulness and retrieval are self-graded by the same agent; the pre-remediation mutation study found 5 of 10 real breakages produced no signal.

16. Stamp. EXECUTED BY: **mixed**. `cmd_stamp` (1140-1167) recomputes `fold_sig` and writes {`fold_sig`, `n_children`, `last_folded`} into the node marker, trusting the agent's claim that a refold happened. `stamp -a11` is banned outside a logged re-baseline: it once marked ~90 rollups current without rewriting any (OPERATIONS.md step 6b). Fragility: stamping is a trust declaration; a lazy agent turns every mechanical check green.

17. Leaf drift and the stale slice. EXECUTED BY: **mixed**. Detection is deterministic (`cmd_scan` 821-845, `stale_leaf_ages` 1438-1490, escalation at 21 days, bulk clears logged to `_stale_cleared.jsonl` as "the launderer's signature"); the re-summarise itself is agent work: a dated addendum in place, then clear exactly that leaf's `src_sha` for re-baselining. Fragility: the clock records first observed stale, understating true age; currency is not correctness.

18. Directory generation and boundary curation. EXECUTED BY: **mixed**. `cmd_directory` (2439-2495) deterministically merges disk structure with hand-kept rules from `eras\boundaries.md`, rendering stale rules in a labeled section (the 2026-07-31 cutover hid 11 dead rules for two weeks under the old

logic). The rules themselves are operator-adjudicated, grown one line per logged misroute. Fragility: duplicate-key clobber documented but unenforced; rule quality depends on failures actually being logged.

19. Cross-context links. EXECUTED BY: **mixed**. Agents propose links as JSON; `cmd_links` (1195-1293) validates through a rejection ladder including verbatim-evidence-in-source (“EVIDENCE NOT IN SOURCE (possible fabrication)”, the stated anti-fabrication guarantee) and writes `## Related` bullets only with `--apply`; `links promote` nominates, the fold agent decides via the altitude test. Fragility: leaf-birth citations depend on recollection because `cite_log.py` never ran; link writes churn hashes and cascade refolds.

Entities (retired layer)

20. Entity-page build. EXECUTED BY: **deterministic-code**, RETIRED 2026-08-21. Nightly pages per recurring name in `eras/entities/` from leaf prose plus the curated alias map `eras/entities.md`, threshold-gated, index-not-copy (recoverable at git 5c84ec8; now a no-op shim at `rollup_tools.py` 1341-1352 because the governance-gated nightly still calls it). The failure that killed it: rebuilt nightly, consumed by no recall surface (`archive_mcp` excluded it; `recall_cue`’s dated-filename filter could never match it). Fragility for a rebuilder: `OPERATIONS.md` lines 275-289 still describe the layer as staying, contradicting the retirement record at lines 50-73; the current text wins.

21. Quote-index invariant. EXECUTED BY: **deterministic-code**, RETIRED with stage 20. Every quoted evidence sentence had to still appear verbatim in its leaf (git 5c84ec8 lines ~830-842); the invariant was dropped and nothing replaced it for the ledger: fact rows carry sources, not quotes, so “does the source support this row” is checked only by the monthly human review queue. The pattern survives in `cmd_links`’ evidence check (1175-1199).

Facts (the ledger)

22. Source registration. EXECUTED BY: **mixed**. `ledger/registry_sources.txt` maps portable prefixes to roots with (excluded) markers (`ledger.py source_roots/resolve_source/is_excluded`, 161-204); inclusion policy is human (health was un-excluded 2026-08-23; army-reserve stays out). Fragility: the registry is the distiller’s only privacy boundary; an exclusion typo silently opens a corpus.

23. Queue planning. EXECUTED BY: **deterministic-code**. `extract.undistilled()` (102-155) orders changed-fingerprint docs first (`sha256[:16]` vs `_distilled.json`), then never-read docs ranked by mentions of the top-150 known entities. Fragility: `.md` files only; the rich-get-richer ranking is by design.

24. Prompt build. EXECUTED BY: **deterministic-code**. `extract.brief()` (185-196) assembles RULES contract, predicate vocabulary, entity digest capped at 400 names, source path, and body capped at 24,000 chars. Fragility: the registry is ~1,169 names against the 400 cap, so the in-prompt duplicate-minting warning path is active (“Build candidate retrieval before trusting this batch”, 179-182); long-document tails are silently dropped.

25. LLM extraction. EXECUTED BY: **scheduled-llm-call** (or a live session via `extract_queue\` plus `ingest`). A model emits bare pipe rows `asserted_on | subject | predicate | object | qualifiers | source`; `run_api()` (295-364, `claude-opus-5`, effort low) cleans and hands rows to the validator. This is the only place non-deterministic choice enters the ledger: durability, dating, entity binding versus minting, predicate selection. Fragility: a live double stall; the User-scope `ANTHROPIC_API_KEY` gap left days of “daily left DUE”, and the 2026-08-24 heartbeat records a `BadRequestError` for exhausted API credits with ~1,397 documents undistilled.

26. Validation and append. EXECUTED BY: **deterministic-code**. `ledger.append()` (433-511) with `validate()` (229-296): six fields, full ISO date, predicate registered, entities resolve or mint typed unknown, key:value qualifiers, source resolves and is openable right now, no excluded prefixes, subject `!= object`, byte-duplicate suppression; rejects logged with reasons to `_dropped_rows.txt`, barren docs to `_barren.txt`. “Nothing here trusts the model’s output” (`extract.py` docstring). Fragility: minted unknown-typed entities await a review nothing schedules; `OPERATIONS.md` line 261’s “no row is appended unattended” is false of the running system (`run_api` appends directly).

27. Rendering. EXECUTED BY: **deterministic-code**. `ledger.render()` (388-396) rebuilds ~1,120 sheets and project handoffs, each headed “Nothing here is computed. Read the dates to work out what is still true”, with a content signature enabling content-keyed staleness (`stale_projects`, 399-417). Conflict handling is deliberate absence: append-only, no supersession engine (an earlier one “produced four separate classes of confidently-wrong answer”, `ledger.py` 7-10); later-wins is a read-time judgment.

28. Tick scheduler. EXECUTED BY: **deterministic-code**. `AgentMemoryKit` runs `ledger\tick.cmd` every 20 minutes into the kit’s own `maintain.py tick` (243-290): elapsed-time dispatch so every run retries every missed run, lock recovery, validate-before-touch, split `last_tick/last_success` heartbeat, `LEDGER_AUTOCOMMIT=0` so commits route through the archive’s guarded nightly (the 08-20 to 08-21 guard-bypass hole). Fragility: productivity hinges on one machine-local credential; the task currently reads `DISABLED`.

29. Repair, propose, review. EXECUTED BY: **mixed**. Tier 1 deterministic repairs apply and roll back on regression (`repair.py run`, 118-161); tier 2 drafts split-identity “same as” proposals into `proposals.txt` (“Nothing here appends itself”); tier 3 judgment questions land in audit reports; `monthly_review.write_queue()` asks per assertion whether the source supports the claim, verdicts appended, never edited. Fragility: 131 proposals pending with no forcing function beyond the weekly habit, several obvious similarity false positives; review coverage rides on the monthly habit.

30. Fact serving (MCP). EXECUTED BY: **deterministic-code**. `memory_mcp.py` (38-166) serves `find_entity`, `about`, `assertions_about`, `resume_project`, `vocabulary`, and `remember` (one row through the same validator, `mint=False`). Fragility: no mint tool, so a session cannot record a fact about a new entity without hand-editing the registry; proposals and the review queue are not exposed.

31. Health integration into the archive. EXECUTED BY: **deterministic-code**. The archive watches the ledger at exactly one read-only seam: `fact_layer_metrics()` (566-647), the User-scope key probe (2120-2146), and `append_only_verdict()` (1493-1578), a monotonic byte-plus-sha watermark

chosen because both git-HEAD comparison and silent re-baselining were demonstrated laundering paths on 2026-08-21. Fragility: the watermark file is an untracked trust anchor beside the data it guards; 36h OVERDUE means a two-day stall is the earliest loud signal.

Retrieve and inject

32. SessionStart priming. EXECUTED BY: **deterministic-code**. `session_recall.py` prints, in order: the `gate_alert()` health block (76-134), the “LONG-TERM MEMORY IS LIVE” standing orders, the fact-layer counts, the dossier line, and one line per era (`leaf_count` plus a 150-char first-prose gist). Measured live 2026-08-25: 4,860 bytes, roughly 1,100-1,250 tokens, a map plus orders, never leaf content. Fragility: fail-silent by design; registration is machine-local; this primer is also the system’s only alerting channel, so losing the hook loses all alarms.

33. Cued recall. EXECUTED BY: **deterministic-code**. `recall_cue.py` fires at boundaries only (new session, 30-min idle, or 2 new salient terms, 120s cooldown) and injects at most a reason tag plus 3 leaf paths from a substring search over dated leaf filenames (constants 38-41, `_search` 98-136). Pointers only. Fragility: “indistinguishable from working” (any error prints nothing, exits 0; the hooks are not gated, only the MCP servers are); a subject cues once per session ever; bare substring matching means vocabulary mismatch yields silent zero.

34. Ranked search and reads (MCP). EXECUTED BY: **mixed**. The engine is deterministic (`archive_mcp.py`: `_terms` 118-132, `_score` 135-146 doing lowercase substring counts, `_importance` hand-weighted 189-209, `_recency` half-life on `_access_log.json` 212-227, relevance-guard warning 336-338); when to call it and how to word queries is in-session LLM judgment steered by docstrings. Fragility: search covers leaf summaries only, never raw, so wrongly summarized or un-leafed material is unfindable; the access log is machine-local; the importance weights are hand-set and unvalidated.

35. Drill-down and altitude navigation. EXECUTED BY: **llm-in-session-per-protocol**. No code decides drill-down: the primer’s CONSULT-IT order, tool docstrings, the altitude ladder, and “Following links is ADVISORY, not mandatory” (`OPERATIONS.md` 485-491) are all prose. Fragility: nothing verifies compliance or logs whether a cued leaf was ever opened (the layer that would have, `cite_log.py`, was deleted); the founding failure (2026-07-26, twenty turns argued without consulting the tree) recurs whenever the model ignores the prose.

36. Browser-Claude surface. EXECUTED BY: **llm-in-session-per-protocol**. `eras\BROWSER_RECALL.md` is paste-in prose for a `claude.ai` Project with a GitHub connector: recall before responding, era to node to leaf, cite paths, no ranked search. Fragility: “drift, with nothing to detect it”, and the drift is live: `OPERATIONS.md` line 35 still describes the file as saying health is off GitHub while the file was rewritten 2026-08-23 to say health syncs.

37. Remote HTTP surface. EXECUTED BY: **deterministic-code**, RETIRED, never verified deployed. `remote_server.py` (surviving only at `code_artifacts\40bacd93\remote_server.py`) wrapped the MCP app in bearer auth that `claude.ai`’s connector flow could not use; carried for weeks as “PRESENT, DEPLOYMENT UNVERIFIED” pointing at a `DEPLOY.md` that never existed. The kept lesson: the local/cloud serving split is a real security boundary; rebuild deliberately.

38. Registration and bootstrap. EXECUTED BY: **mixed**. `scripts\install_surfaces.py --apply/--status` registers both hooks and both MCP servers (verified 2026-08-25: all four at USER scope, redundant project approvals in `dnd-campaign`); OPERATIONS mandates `verify-by-observation` after three wrong prose claims about registration. Fragility: on a fresh clone the tree arrives complete and every surface is off, silently; only the servers are probed, the hooks are not.

Maintain and verify

39. Refcheck. EXECUTED BY: **deterministic-code**. `_refcheck` (981-998) reports dangling `[[wikilinks]]` and missing backticked `eras/` citations on every scan. Fragility: `basename-only` slug resolution; prose paths are deliberately skipped.

40. Golden regression. EXECUTED BY: **deterministic-code**. `_golden_check` (1011-1036): 13 frozen rows in `eras\_golden.jsonl`, each passing iff all `expect_keywords` appear as substrings of the first file matching the source `relpath` substring; run in every scan and report; never edit an answer to pass. Fragility: 13 substring probes over ~1000 leaves; `hits[0]` on an ambiguous source name could test the wrong file; growth depends on the `add-when-durable` habit.

41. Content gates. EXECUTED BY: **deterministic-code**. `hedge_stripping` (217-309) flags a node hardening a date its children only hedge; `coverage_regression` (312-380) flags summaries that stop naming their topics or reaching their leaves. Added 2026-08-21 after the mutation study. Fragility: deliberately narrow; they catch hardened dates and dropped topics, not wrong facts or causality, and faithfulness scoring structurally cannot see hedge-stripping.

42. Invariant gates and the failure log. EXECUTED BY: **deterministic-code**. `_run_invariants` (1901-2227): 24 gates spanning privacy (`gitignore`, `held-local`, `held-gist`), guard invocation, tracked-credential sweep, drift measured, truncated-code pointers, count and topic recomputation, push age under 26h, content gates, pip check, the 21-day stale-leaf clock, alarm-wired, coverage-counts-everything, real MCP handshake-plus-tool-call probes (v5 of a probe wrong four ways), distiller key at User scope, autocommit off, the append-only watermark, freshness, and the `GATE_PROOF` meta-gate requiring every gate to declare its mutation-test proof (`test_gates.py`). Non-PASS auto-opens `_failure_log.jsonl` rows; `reconcile_failure_log` (1647) auto-closes them; a crash writes a FAIL gate (`run_invariants`, 1876). UNKNOWN is never PASS. Fragility: `maintain.py` never reads report's exit code (comment 2394-2397); the alerting chain hangs on the `SessionStart` hook, which is checked by one of these very gates, a partially circular guarantee.

43. Alarm surface and response. EXECUTED BY: **mixed**. `write_gate_state` (1618) persists failures; `session_recall.py` renders them first in every session, including "ARCHIVE HEALTH: STALE" past 36h. A human is the consumer; nothing emails, blocks, or retries. Fragility: the human-fix loop has no deadline (the credential-sweep gate has been failing since 2026-08-22, row still open); the alarm depends on the very nightly it reports on.

44. Weekly maintenance pass. EXECUTED BY: **scheduled-agent**. The named task 'weekly-memory-tree-maintenance' fires an agent whose entire prompt is "run the weekly pass per OPERATIONS.md": pass open, scan, import digestion, promotion, folds and stamps, one rebalance, one

backlog item, the stale-leaf slice, rollups, the rotating audit, one log line, quarantine drain, pass close. Fragility: NOT FOUND in Windows Task Scheduler on this machine; the last two log entries (08-21, 08-23) were interactive sessions, not the agent; the entire quality layer rides on an LLM following a 621-line protocol whose only in-repo backstop is the 7-day log-staleness habit.

45. Rotating source-verification audit. EXECUTED BY: **llm-in-session-per-protocol**. `audit-pick` (1079) suggests the least-recently-audited subtree; the agent verifies its leaves against raw transcripts, actively refuting dates, figures, quotes, and causality, subject to the mandatory self-audit exclusion rule; `audit-record` (1090) writes `acc/fid/cmp` into `_audit_state.json` and the node marker; `audit-claims` (2498) mechanizes the node-versus-raw packet where 8 of 8 measured defects lived. Fragility: this is the only mechanism that tests whether summaries are TRUE, and coverage is 5.7% (5 subtrees); scores are unverified self-reports; the exclusion rule lives only in protocol.

46. Metric gate and the failure loop. EXECUTED BY: **mixed**. Before a weekly diff stands: rerun scan plus golden; bars are golden 100% and retrieval at least 10/12; on regression, revert the pass and log a failure. Golden is mechanical; faithfulness and retrieval are LLM-judged strings the agent records into `_metrics.json`, which scan merely carries forward (934-936). Safety-critical process fixes are queued for the operator, never self-applied. Fragility: nothing mechanically blocks a bad commit, the nightly commits whatever is staged; a pass that skips the sampled scoring leaves stale scores that look current.

47. Pass open/close marker. EXECUTED BY: **mixed**. `cmd_pass` (2257-2283) records open time in gitignored `_pass_in_progress.json`; scan warns past 6 hours, converting died-mid-run from invisible to visible, but only when the agent remembered to call open. Fragility: nothing forces the call; the marker protects one machine.

48. Selftest. EXECUTED BY: **deterministic-code**. `cmd_selftest` (2604-2859) rebinds state globals to a temp fixture and asserts ~40 named checks pinning every historical defect; mandated after any edit to `rollup_tools.py`. Fragility: convention, not a pre-commit hook; nothing blocks committing with a red selftest.

49. Staleness watch. EXECUTED BY: **mixed**. Three clocks converge on the operator: the 7-day maintenance-log rule (an LLM-in-session obligation from `~\.claude\CLAUDE.md`), the 2-day nightly heartbeat mtime with its read-the-log-body caveat, and the deterministic gate alarm. Fragility: every surface fails silently by design; the ultimate backstop is an LLM following prose, the least mechanical link in the chain.

Direct answers

Is there semantic search anywhere? No. Every retrieval path is lexical substring matching: `archive_mcp.py` `_score` does `low.count(t) / low.find(t)` over lowercased full text with hand-set importance and recency tiebreaks; `recall_cue.py` does `t in text`; `search_manifest` reuses `_score`. A `grep` for `embed|vector|cosine|semantic|faiss|sentence-transformer|tfidf` across `mcp_server\*.py`, `scripts\*.py`, and `agent-memory-kit\prototype\*.py` returns only incidental prose, and `OPERATIONS.md` names embeddings only as hypothetical: “embeddings (if ever added) are an index, never the memory.”

How are documents ingested? Three ways, none unified. (1) Claude Code sessions: continuous free capture to `~\.claude\projects\*\*.jsonl`, swept three times daily by `parse_sources.py` into `conversations\*.md` keyed by conversation id. (2) Web-chat corpora: manual vendor takeouts into `raw\`, parsed only when the mtime gate notices, only the newest archive per format ever read, currently frozen at 2026-07-11 (ChatGPT) and 2026-07-31 (Gemini). (3) Filesystem documents: the one-off 2026-08-23 import pipeline (`scripts\imports\2026-08-23-doc-import\`), 1,042 docs to 73 leaves through scripted scaffolding around three multi-agent LLM workflows, not repeatable without reconstructing judgment.

How are fact rows generated? `registry_sources.txt` declares the corpus; `extract.undistilled()` picks documents; `extract.brief()` builds the prompt; a model (claude-opus-5 via `run_api()` unattended, or a live session via `extract_queue\`) emits pipe-delimited rows; `ledger.append()` validates every field against the registries and the live filesystem, appends clean rows to `facts.txt`, mints unknown-typed entities, and logs rejects with reasons; sheets re-render deterministically. The LLM decides only row content; code enforces everything checkable. Conflicts are never resolved: append-only by design, later-wins at read time, corrections as new dated rows.

Where do the tree layer and the fact layer connect? Almost nowhere, deliberately. They share no code: the retired entity layer lived in `claude-archive\scripts\rollup_tools.py`, the ledger lives in `C:\Users\jhffm\agent-memory-kit\prototype\` parameterized by `LEDGER_HOME`; they disagreed on what an entity is (threshold-gated 70 versus mint-on-sight 388, 17 names in common, `OPERATIONS.md` 281-285). The ledger became system of record 2026-08-21 and the entity layer was retired outright. Remaining touchpoints are one-directional and read-only: `fact_layer_metrics()` and `append_only_verdict()` read ledger files for health gating, and `registry_sources.txt` excluded `eras\entities\` so the ledger could never cite a derived view of itself. The written rebuild instruction: if a per-entity index is wanted again, build it FROM the ledger and wire it into a recall surface before trusting it (`OPERATIONS.md` 70-73).

The by-hand inventory

One row per stage whose generative core is a human or an LLM improvising against protocol today. This table is the build backlog for the harness.

#	Stage	What the hand does today	What the harness must do instead	Hardest part of automating
1	Source acquisition (S1)	The operator requests takeouts by habit and drops them in <code>raw\</code>	Scheduled export pulls per vendor API where possible, plus a freshness alarm on every source channel	Vendors without export APIs; detecting silent corpus staleness
2	Ambiguous-duplicate merge (S4)	The operator merges	Deterministic merge policy	Choosing a merge rule safe for an

#	Stage	What the hand does today	What the harness must do instead	Hardest part of automating
	residue)	AMBIGUOUS same-id transcripts from the printed report	(superset wins, else union with provenance) executed by code	immutable corpus
3	Quarantine drain (S6 residue)	Weekly human disposition: commit deliberately or hold local	Ticketed queue with deadlines, and machine-checkable disposition records replacing the hand-cleared txt file	Judging CUI/sensitivity is inherently human; the harness can only enforce the deadline
4	Session-close leaf writing (S8)	In-session LLM segments, summarizes, routes, cites, entirely from a prose rule it must remember	A SessionEnd-hooked extraction job: contract-schema leaf output, validated frontmatter, machine-supplied conversation id, retraction-aware tail scan	Segmentation has no specification anywhere; defining a testable boundary function is new research, not porting
5	Import digestion (S9)	Scheduled agent leafs the backlog and hand-edits MANIFEST rows	Same extraction contract as row 4 run in batch, MANIFEST as generated artifact not hand-edited table	Same segmentation problem at scale, plus leaf-worthiness (the <code>_no_leaf_bar</code>) as a classifier
6	Promotion and routing (S10)	LLM routes each leaf against prose boundary rules; “clean” is undefined	Routing as a constrained classification call validated against the machine-readable boundary set, with a defined promotion gate	Boundary rules are failure-grown prose; making them a testable ruleset without losing nuance
7	Document-import triage/distill/verify (S12)	One-off multi-agent workflows with uncommitted judgment manifests	A repeatable cold-ingest pipeline: any foreign corpus in, triaged and distilled	Cross-doc-bleed prevention (the verify pass caught 9); triage verdicts are judgment

#	Stage	What the hand does today	What the harness must do instead	Hardest part of automating
			under the same contracts	
8	Refold prose (S15)	Weekly agent rewrites node summaries from child summaries	Contracted summarize-of-summaries calls with entailment checks against children before acceptance	Verifying faithfulness mechanically; today it is self-graded by the writer
9	Stamp discipline (S16)	Agent honestly stamps only what it refolded	Stamping performed by the refold job itself, atomically with the accepted rewrite	Trivial once refold is a job; the hard part is row 8
10	Stale-leaf re-summarise slice (S17)	Agent picks and rewrites stale leaves with dated addenda	Drift-triggered re-summarization jobs, ordered by a scoring rule, re-baselined atomically	Deciding what changed in the source and whether the leaf is actually misleading
11	Boundary-rule curation (S18)	The operator adjudicates new discriminator lines after misroutes	Misroute detection feeding proposed rule diffs for one-click approval	The adjudication itself stays human; the harness closes the loop around it
12	Link proposal and promotion (S19)	LLM proposes links; fold agent applies the altitude test	Proposal generation as a batch job through the existing validator; promotion as a scored, thresholded decision	The altitude test (“would a reader stopping here be misled?”) resists mechanization
13	Fact extraction (S25)	Scheduled LLM call, currently stalled on a machine-local credential	Same call under harness-managed credentials, retry budgets, and a candidate-retrieval step replacing the 400-name digest cap	Entity binding versus minting at scale; the digest-cap duplicate-minting problem needs retrieval, not a bigger prompt cap

#	Stage	What the hand does today	What the harness must do instead	Hardest part of automating
14	Proposal approval and review verdicts (S29)	Human/session reviews proposals.txt (131 pending) and monthly review queue	Batched review sessions with escalation on backlog age; auto-reject for scored-obvious false positives	“Does the source support this claim” is the same truth problem as row 20
15	Registry curation (S22, S26 residue)	Human edits source inclusion and entity types; minted unknowns await a review nothing schedules	Typed-entity resolution jobs and a scheduled unknown-type triage queue	Identity resolution (same-as) without silent identity blending
16	Query planning (S34)	In-session LLM decides when to search and how to word queries	Automatic retrieval at query time (the RAG arm): the harness retrieves, the model consumes	Query formulation from conversational context; today mitigated only by a warning string
17	Drill-down navigation (S35)	LLM follows the altitude ladder by prose discipline	Harness-side context assembly to a budget: pre-fetch node plus leaf content instead of hoping the model drills	Choosing the stopping altitude per query; nothing today measures whether pointers were followed
18	Browser surface (S36)	Pasted prose instructions, hand-updated, drifting	Generated surface config from one source of truth, or retire the surface into the served API	Keeping any un-hooked surface in sync is unsolvable; generation is the fix
19	Bootstrap and registration (S38)	Human runs install_surfaces.py and verifies by observation	Idempotent provisioning in setup, with the registration gates as post-conditions	Nothing hard; it is discipline made into code
20	Weekly pass orchestration (S44)	A scheduled agent interprets a 621-line prose protocol	A real orchestrator: each step a typed job with inputs,	Decomposing OPERATIONS.md into jobs without

#	Stage	What the hand does today	What the harness must do instead	Hardest part of automating
			outputs, and completion records; the protocol becomes code	losing its accumulated failure wisdom
21	Source-verification audit (S45)	LLM refutes leaves against raw, self-scores acc/fid/cmp	Independent-model verification jobs with claim extraction (audit-claims already mechanizes the packet) and adversarial scoring	Ground truth: verifying summaries are TRUE is the system's only truth check and the least automatable judgment in it
22	Metric gate scoring and revert (S46)	Agent self-grades faithfulness and retrieval, decides revert	Held-out scored evaluation by a non-author model; revert as an automatic transaction on failing bars	Building an honest retrieval benchmark; today the bars have no mechanical enforcement at all
23	Alarm response (S43)	A human reads the primer and fixes what it names	Escalating notification with deadlines per failure class; auto-remediation for the mechanical subset	Repairs like credential revocation and history purges stay human; the harness enforces the clock
24	Staleness habit (S49)	An LLM obeys a 7-day prose rule in CLAUDE.md	Watchdog jobs supervising every scheduler, including the schedulers themselves (both tasks are DISABLED today and nothing said so)	Supervising the supervisor on a single machine with no external monitor
25	Golden-set growth (S40 residue)	Agent/human adds a question when a durable fact lands	Auto-proposed golden rows at leaf acceptance, human-approved	Selecting durable facts worth freezing; 13 probes over ~1000 leaves is the current total

What the mock never had

Stages a rigorous pipeline needs that exist today in no form. A semantic retrieval arm: every path is substring matching, so vocabulary mismatch fails silently. Outcome and citation capture: nothing records whether an injected cue or search hit was ever opened or helped; the one attempt (`cite_log.py`) was never wired and was deleted 2026-08-21, so leaf citations are end-of-session recollection. Injection budgeting: the primer is a fixed ~1,200-token block and the cue a fixed 3 paths; nothing measures a context budget, ranks what deserves injection under it, or adapts payload to query. Write-time contracts on LLM output: the fact validator aside, no leaf, MANIFEST row, node summary, or audit score is schema-checked when written; the harness's central move is making every LLM output pass a contract the way fact rows already do. Enforcement of the session-close rule: no hook fires when a session ends without a leaf; omissions surface only as later coverage debt. Repeatable cold ingest of a foreign corpus: takeouts parse only the newest archive per format, Gemini has no stable identity so a re-export mints mass duplicates, and the document import was a one-off whose judgment manifests were not kept. Capture freshness measurement: nothing alarms on a corpus frozen at its last takeout. An honest retrieval benchmark: golden is 13 substring probes, and the two LLM-judged bars are self-graded by the summary author with stale scores carried forward indefinitely. Independent truth verification at speed: the rotating audit covers 5.7% of the tree at one subtree per week; between it and two narrow content heuristics, no fast check exists that any node is true. Supervision of the schedulers themselves: both Windows tasks read DISABLED today, the reason is recorded nowhere, and the only remaining detector is a session-start alarm whose own registration is machine-local and unprovisioned on a fresh clone.

REFERENCE, not superseded. This is the antagonistic feasibility review. Its conclusions held up and its cut order is carried into 05_fall_plan.md. It is background rather than part of the working set. Index: docs/README.md.

Feasibility, reviewed antagonistically

This document was written by the assistant at the student's request, under instructions to prosecute the plan of record rather than defend it. Three hostile reviews were run independently, on time, on technical risk, and on scope, each grounded in the design brief (00), the proposal (02), and the planning documents (05 through 12). This is their merge. Verdicts are stated as found, not softened. The short version: the fall as planned is oversubscribed by roughly a factor of three, its Phase 3 violates the governing brief's own November 15 rule, its hand-labeling program is priced at assistant speed for work the authorship rule makes non-delegable, and the one deliverable the plan declares safe, the miss-rate paper, depends on an instrumentation step no phase builds. The plan survives, but only after the cuts named below are made now, in writing, before the August 26 meeting.

The time verdict

The phase headers sum to roughly 83 hours: Phase 0 at 4, Phase 1 at 18, Phase 1.5 at 8, the dead-window reading at about 6 (2 hours a week for three weeks), Phase 2 at 35, Phase 3 at 12. That fits the 70 to 100 hour budget only if every box is exact, and the plan itself confesses the boxes are false: Phase 2's preamble calls them "task-decomposition floors, the sums of the happy path" and cites the student's own estimate that the full build runs to hundreds of hours, then leaves every box standing, relabeled "sequence, not schedule." That is not absorption of the correction; it is a false balance sheet with a footnote saying the numbers are false. Every downstream document inherits it: Phase 3's entry condition is "harness end-to-end," 06 lists the harness with tenant zero as a fall deliverable, and 05's step 9 runs the certified set across three arms and a sensitivity table. Apply even a conservative 2x and Phase 2 alone demands 70 hours inside a four-week block whose realistic ceiling, at 6 to 10 hours a week minus HW3 (Nov 1), quizzes (Nov 8 and 15), and one probable drill weekend, is 24 to 32 hours. At the student's own hundreds-of-hours velocity, the seven-stage Phase 2 pipeline is a 10 to 17 week job compressed into 4. The budget carries zero contingency for the brief's named unknowns: four-plus multi-day drill holes at unknown positions, the IT 483 midterm and final (48 percent of a grade, dates unknown), and a family calendar that could not be read.

The first break is datable. The week of September 15 to 21 requires, per 08: reading three literature areas, the metric definition sheet v1 (the artifact Fang must agree to before any result exists), the contracts v0 with named rejection criteria, and the five to eight page synthesis memo that Phase 1 exits on. The same week holds CS 425 HW1 on September 17 and the quiz on September 20. That is 15 to 20 real hours of reading and advisor-grade writing, under a voice rule that forbids ghost-written prose and historically costs correction rounds, inside a project allowance of 4 to 6 hours. Nothing downstream recovers the slip, because the following week (September 22 to 27, quiz on the 27th) boxes the full pipeline design

document, the segmentation specification that 07 explicitly calls new research rather than porting, dry adapters passing tests on real files, and four frozen JSON-schema contracts with validators, at 8 hours: a 20 to 30 hour week priced at 8. By September 27 the plan is two to three weeks behind, the first open block is over, and the dead window forbids catching up by its own rule.

The hand-work is the hidden second budget. Priced at human speed, non-delegable under the authorship and voice rules: the 20-document segmentation gold set, 10 to 15 hours, boxed inside a 5-hour INGEST block that also builds both segmenter arms; the 150 distiller hand-checks (fifty bindings, fifty fact rows, fifty topic paths) plus five merge-ledger walkbacks, 8 to 10 hours, inside the 13-hour ORGANIZE block; the 100 judge-calibration labels with kappa, 5 to 8 hours, inside a 5-hour evaluation-spine box; question authoring at 2 to 3 candidates per certified survivor, so 150 to 250 authored and verified questions to field 60 to 100, 15 to 30 hours; and the prerequisite no document prices anywhere, reading the several-hundred-thousand-word corpus that 06 calls “ground truth obtainable by reading,” 20 to 25 hours appearing in no phase, no box, no valve. Residue (composition reads, smoke tests, arm-disagreement sampling, wiki adjudication by the text): 5 to 10 hours. Hand-work subtotal: 60 to 95 hours, the entire fall budget consumed before a single build stage, reading synthesis, or paper paragraph. Full honest demand: reading 35 to 50, design and advisor-facing writing 25 to 40, fall’s build slice under the authorship gate 80 to 150, hand-labeling 60 to 95, runs and paper 25 to 40; roughly 225 to 375 hours against 70 to 100 available. Oversubscribed by a factor of three.

Two structural violations close the case. First, Phase 3 runs November 16 to December 7 and targets arXiv by December 5, inside weeks the design brief declares dead (Test 2 on 11/18, Thanksgiving, HW4 on 12/3, the application window 12/1 to 15) and against Rule 2’s own sentence: “a plan that targets December targets nothing.” Certification runs across three tiers, wiki-structure scoring adjudicated by rereading a several-hundred-thousand-word serial, the sensitivity table, and the entire paper are not “runs and prose, no construction.” The real deadline is November 15 and nothing in the phase plan honors it. Second, the paper’s own anchor is unscheduled: the counterfactual miss rate “runs on the live mock with consult-logging added,” and no phase in any document contains the step that adds consult-logging to the mock. Proposal 3 in 05, the instrument’s origin, requires roughly four weeks of live collection before the rate means anything. For data to exist by November 15, logging must be running by mid-October, which means it must be built now, this week, not in October. The one deliverable the plan says cannot be held hostage by the build is currently hostage to a line item that does not exist.

The technical risks, ranked

Each risk carries a week-one de-risking probe, none of which is in the plan, and a fallback. The three most likely to eat the semester come first.

1. **Entity identity as a whole: duplicate minting plus an un-undoable merge.** The mentions call carries the current alias ledger in the prompt, and the mock already ran this experiment: its registry grew to roughly 1,169 names against a 400-name prompt cap, with active duplicate minting as the documented consequence and the mock’s own comment saying build candidate retrieval before trusting the batch. Harry Potter alone blows past any cap by mid-book,

sooner on the local tier. Duplicates then corrupt fact binding, supersession (rows on harry_e01 and boy-who-lived_e14 never collide), and the wiki coverage number. On top of that, the same_as merge is the plan's showpiece and its least-designed operation: "superseded where it conflicts" is a semantic judgment, not the deterministic subject-plus-predicate collision the maintain stage defines; after Scabbers re-attaches to Pettigrew, (pettigrew, is_a, rat) versus (pettigrew, is_a, wizard) collide and later-assertion-wins declares the rat row false when the design's own gloss says it stays true of the disguise period, and the schema has no validity interval to express that. Worse, re-attachment as described rewrites fact rows, violating the append-only rule, has no undo, and cascades: a wrong merge changes the alias ledger, which changes every subsequent binding, which poisons downstream extraction, surfacing last, at the wiki render or the coverage scoring. The plan reproduces a failure the mock already documented. Probe: hand-write twenty fact rows including one correct and one wrong same_as into a fixture, implement merge as a pure function, then try to write the un-merge; if the undo cannot be written, the design is wrong now, cheaply. Separately, seed a ledger and run the mini and local tiers over one chapter, counting mint versus bind; the duplicate rate is measurable in an afternoon. Fallback: never rewrite rows; keep same_as as an equivalence relation resolved at read time (union-find at query), so a merge is one retractable row and identity is time-scopable later; candidate retrieval via the embed interface that already exists (top-k candidates per mention, not the whole ledger); bind only named mentions and leave pronouns at leaf level, out of the fact layer.

2. **Local-tier contract rejection is a schedule bomb with a ten-call fuse.** Every segment costs three-plus schema-validated calls, each with one retry, and 7B instruct models are weak at strict nested JSON. If the local tier rejects 30 to 40 percent, the sensitivity table's local column is holes and wall-clock per segment doubles across thousands of segments. The build plan's smoke test is ten calls per tier, which cannot measure a rejection rate, so this detonates in the ORGANIZE block, October 19 onward, exactly where the 35-hour box is already admitted to be a floor against a hundreds-of-hours estimate and there is zero slack before November 15. Probe: a hundred calls per tier against the real contracts drafted early, rejection rate and seconds per segment, multiplied out to corpus size; one evening of API spend. Fallback: Ollama's grammar-constrained JSON decoding; flatten the contracts to one fact per call; or demote the local tier to segmentation only and report the sensitivity axis as mini versus frontier, which an existing valve already permits.
3. **The aggregate hand-labeling budget.** Twenty boundary-labeled documents, one hundred judge-calibration answers, fifty bindings plus fifty fact rows plus fifty topic paths, five merge walkbacks, per-question closed-book certification across every tier, miss-rate labeling: 300-plus personal judgments, tens of hours of exactly the drudgery the record says he abandons, assigned to the one person whose hours are the binding constraint, non-delegable under the QA rule, and scheduled largely into the window that also holds Test 2, Thanksgiving, and the paper, against a documented pattern of multi-day productivity crashes after sustained grind. Kappa on 100 items with a skewed verdict distribution is unstable, and the judge-never-wrote rule collides with the tier-swap table, in practice forcing the frontier model to judge everything against a live record of credit exhaustion. Probe: label thirty answers in week one against the intended judge tier; below

roughly 0.7 kappa the verdict column is noise and the design must move. Fallback: the question record already carries fixed_answer; make the leaf and rollup bands extractive and string-scored, reserving the judge for the synthesis band only; thirty labels now, full kappa study in spring.

4. **Segmentation has no ground truth, and the plan's fix does not produce one.** The mock's entire spec is "segment by topic," and the promised Pk and WindowDiff numbers rest on "twenty documents the student marks himself." A single annotator produces no agreement reference; topic-boundary labeling is low-agreement even between trained annotators, and there is no second annotator anywhere in the plan. If his own boundaries are not reproducible, Pk cannot discriminate the model arm from TextTiling and the arm comparison is noise with a decimal point. Probe: label three documents twice, a week apart, compute self-agreement Pk; if he cannot agree with himself the metric is dead on arrival. Fallback: books carry author-given boundaries for free (chapters, scene breaks) and chats carry session and turn-block structure; fix segmentation as policy and drop it as a measured claim.
5. **The corrected-facts band is a conjunction of every upstream risk.** It requires that the fact pass extracted both the early and the late assertion, bound both to the same entity through the alias ledger, normalized both to the same predicate so they collide (is_a versus species versus form never fires supersession), and that the merge worked. Predicate drift alone can leave the band silently empty, discovered in Phase 3 when the question set is built, after the build window has closed. Probe: run the intended fact contract on the frontier tier over the two Scabbers passages and check whether the emitted rows actually collide under the schema as written. Fallback: extend the planted Scabbers fixture to a seeded set of hand-authored supersession pairs, so the instrument reports planted plus organic cases with denominators stated.
6. **The wiki comparison contains an unmentioned matching problem.** Entity coverage and relation recall against a fan wiki presuppose a mapping from harness entity ids to wiki pages: redirects, disambiguation, granularity mismatches (family pages, event pages, object pages the harness will never mint), and relational content living in infoboxes and categories rather than SPO rows, which makes comparability a second extraction task the plan never budgets. The alignment judgment can dominate the reported numbers. Mostly spring-gated, so it degrades the paper rather than the semester, unless Phase 3 step 1 attempts it cold and eats the writing week. Probe: hand-map twenty wiki pages and extract twenty infobox relations into the fact schema, timed. Fallback: align by exact and alias name match only, report coverage on the matched subset with the unmatched rate stated, and move relation recall to spring.

Committed core versus aspirational shell

The two-question test: does the paper need it, does the August 26 meeting need it. The paper, per the plan's own anchor, needs the counterfactual miss rate on the instrumented mock, the corrected-facts band, the metric sheet agreed with Fang, and a design narrative. The meeting needs the corrected proposal, the written topic-change agreement, the new one-semester form on the Hasselbring and Tang routing, and the one-paper-or-two ruling. Everything else is held to those two lights. One further caution

from the brief itself, Risk 9: docs 09 through 12 are a finished architecture, schema, build plan, and demo scope, prepared by the assistant, for an advisor whose actual assignment was to write down what was discussed and read some articles; under the three-bucket voice rule none of it can yet be shown to him as the student's thinking.

The committed core, defensible to a committee:

- The August 26 meeting package and the signed one-semester proposal: the administrative gate everything else waits behind, and it survives anything.
- Consult-logging on the live mock, implemented this week: mock-side, so the authorship rule does not bind it, and the miss-rate clock needs roughly four weeks of collection, so September, not October.
- The instrument suite on the mock, logged as E2 in the existing EXPERIMENTS.md: counterfactual miss rate, stale-serve rate, and the hash-staleness demonstration (hand-edit N leaves, show dirty-flag schemes miss them), the brief's own strongest defensible claim.
- The metric definition sheet to Fang early: the agreement artifact, and it is cheap.
- One two-arm comparison on tenant zero: a corrected-facts band on the archive only, roughly 30 to 40 certified questions, control versus tree, one corpus, no third arm, no sensitivity table.
- One synthesis memo replacing the four Phase 1 artifacts, covering judge validity, narrative QA, and segmentation: the three areas the instruments actually cite.
- The design document with the segmentation spec, defended as design on paper, not shipped as code: the intellectually best parts of 10 belong here.
- The paper, complete and posted by November 15, framed as measurement instruments for deployed memory systems with a two-year production system as subject, reporting whatever pipeline slice exists as secondary.

The aspirational shell, to be labeled so in writing before Fang sees any of it:

- The entire cold literature corpus: 20 to 25 hours of reading, 150 to 250 authored questions, per-question certification; 60 to 100 hours, all spring by the plan's own portability logic, and 06 already admits it can be stated as future work.
- The wiki demo and generated portraits: 12 says itself it renders whatever exists; it can render in spring, and its one defense (Risk 10, the fall must not be boring) makes it a morale item, which belongs after the committed core.
- The corpus ladder past its first rung: Lovecraft's cross-document resolution and Holmes's adaptation-contamination story are paper ideas for a lab with three students.
- The three-tier sensitivity axis and the Ollama install: the table needs a working multi-stage pipeline to swap tiers in, and the pipeline will not exist this fall.
- The three injection arms and the entity and fact layers as running code past extraction: they answer the proposal's real open questions on paper, in the design document, this semester.

- The public benchmark suite (“scoped as a week”, 10 percent of the fall budget) and the replication corpus: a generality checkmark the December paper does not claim.
- Hand-implementing TextTiling from the 1997 paper: an authorship-rule tax of roughly a day purchased to avoid one NLTK dependency, inside an arm inside a segmenter inside a phase that will not finish; the rule is for the load-bearing core, not plumbing.
- The full hand-labeling program at depth: twenty gold documents (sample five), one hundred judge labels (thirty now, kappa in spring), certification across every tier.
- Phase 3 as scheduled and everything in Phase 4: all fall measurement moves inside November 1 to 15; December contains nothing.

The graceful-degradation order

The named valves in the plan total about six hours of cuts against a build the student prices in the hundreds: valves, not a ladder. The calendar guarantees at least one real week will disappear (four-plus drill weekends at unknown positions, an IT 483 midterm and final on unknown dates worth 48 percent of that grade, an unreadable family calendar, a documented pattern of crashes after sustained grind, two children). The cut order is therefore decided here, in writing, and a lost week deletes from the bottom:

1. The meeting package and the signed proposal. Survives anything; nothing may displace it.
2. The metric sheet and the instruments on the live mock: consult-logging, miss rate, stale-serve rate, corrected-facts band, the hash-staleness study. These run in evenings while coursework owns the weeks, and they are the paper.
3. The synthesis memo and the design document with the segmentation spec. These fit campus gaps and dead stretches by construction.
4. The pipeline slice (interfaces, adapters, segmenter, ingest and organize on one corpus), if and only if the October 19 to November 15 block survives intact.
5. Everything else: the shell list above, in any order, because none of it is load-bearing this fall.

When a week vanishes, layer 5 is already gone, layer 4 shrinks stage by stage in reverse build order (organize before ingest), and layers 1 through 3 do not move. Phase 3 as written quietly depends on layer 4 existing; under this order the paper depends only on layers 1 through 3. Anything cut is cut in the log the day it is cut, per the plan’s own standing rule.

Summary for an outside reader

The plan of record budgets 83 hours for work honestly priced at 225 to 375, against 70 to 100 available; it first breaks the week of September 15 to 21 and is unrecoverable by September 27. Its Phase 3 schedules the paper into December weeks its own governing brief declares dead; the real deadline is November 15. Its paper’s anchor instrument requires consult-logging on the existing mock that no phase builds and roughly four weeks of collection, so instrumentation must land this week. Its three largest technical risks (entity-identity merge without undo, local-tier JSON rejection, a 300-judgment hand-labeling program)

each have a one-evening week-one probe, none currently planned. Re-scoped to the committed core (instruments on the mock, one two-arm archive comparison, one synthesis memo, the design document, a paper frozen November 15), the semester is finishable by one person and produces the application artifact on time.

REFERENCE, not superseded. This is the working bibliography, one paragraph per work. It is background rather than part of the working set. Index: docs/README.md.

The literature, work by work

A per-work reference: one paragraph per work, for looking a paper up; the narrative document serves the read-through.

Assembled with LLM assistance; entries marked as summarized from abstracts are paywalled works whose identifiers appear in the reading list.

Start here

Kumaran, Hassabis, and McClelland 2016. Complementary learning systems updated for AI. A Trends in Cognitive Sciences review revisiting complementary learning systems theory for artificial agents. Two stores: a fast episodic store recording individual experiences as they happen, and a slow structured store holding generalized knowledge, with replay carrying material from the first into the second so the slow store integrates gradually rather than being disrupted. The biological warrant for raw logs under a distilled summary layer with a scheduled process between them. (Summarized from the published abstract.)

Park et al. 2023. Generative Agents. A UIST paper populating a Sims-like town with twenty-five gpt-3.5-turbo agents to test believable long-horizon behavior. Each agent keeps a memory stream of timestamped natural language records; retrieval sums recency (exponential decay at 0.995 per game hour), LLM-rated importance (1 to 10), and embedding relevance; reflection fires when recent importance passes 150, writing cited higher-level inferences back into the stream. In a 100-rater evaluation the full architecture scored TrueSkill μ 29.89 against 21.21 with no memory, each ablated component costing believability; over two simulated days, party knowledge spread from one agent to thirteen with no hallucinated claims of knowing. A flat log with scored retrieval is not enough, and the authors' memory-hacking worry is at bottom an access-control question.

Sarathi et al. 2024. RAPTOR. An ICLR paper replacing the flat chunk store of standard retrieval augmentation with a bottom-up summary tree: 100-token chunks embedded, soft-clustered with Gaussian mixtures over UMAP-reduced embeddings, summarized by gpt-3.5-turbo, recursing until clustering fails; clustering on embeddings rather than position lets distant passages share a parent. The preferred query strategy collapses the tree so every node at every layer competes in one cosine search under a 2,000-token budget. The tree improved SBERT, BM25, and DPR on all three benchmarks tested; with GPT-4 it reached 82.6 percent on QuALITY against a prior best of 62.3, about 4 percent of summaries containing minor hallucinations. Recall should span abstraction levels so each query picks its granularity; the index is batch-built, an incremental variant open.

Rasmussen et al. 2025. Zep. A preprint on a commercial memory service and its engine Graphiti, a temporally aware knowledge graph built continuously from conversation. Messages persist as non-lossy

episode nodes; an LLM pipeline extracts entities and facts, resolves duplicates against the existing graph, and maintains community summaries incrementally by label propagation. Each fact edge carries four timestamps for when the system learned it and when it held true; contradicting edges invalidate rather than overwrite. On LongMemEval, conversations averaging 115k tokens, Zep beats full-context prompting by 18.5 percent with gpt-4o (71.2 versus 60.2) while cutting latency about 90 percent; on the small DMR benchmark full-context nearly ties it, so the graph earns its cost mainly at scale. Close to a reference design: episodic raw storage under a derived semantic layer with explicit fact validity intervals.

Edge et al. 2024. GraphRAG. A preprint targeting questions needing global understanding of a whole corpus, beyond what fetching similar chunks can answer. An LLM extracts entities, relationships, and claims from 600-token chunks into a weighted knowledge graph; hierarchical Leiden partitions it, every community at every level gets a pre-written summary, and queries map the summaries to partial answers in parallel before reducing to one. On two roughly million-token corpora with GPT-4 as judge, every global condition beat vector RAG on comprehensiveness (72 to 83 percent win rates) and diversity; root-level answers used 97 percent fewer context tokens than direct map-reduce summarization (26,657 versus about a million); indexing cost 281 minutes of gpt-4-turbo per million tokens, incremental update open. Summarize once, query summaries forever; the persona-driven evaluation transfers to any memory system lacking gold answers.

Consolidation and agent memory

McClelland, McNaughton, and O'Reilly 1995. Why there are complementary learning systems. The Psychological Review paper founding complementary learning systems theory, the argument behind nearly every two-store memory design since: a single learning system cannot both acquire new material quickly and hold structured knowledge stably, so labor divides between one system capturing experience fast and another integrating it slowly. The original justification for splitting raw capture from consolidated knowledge. (Summarized from the published abstract.)

Teyler and Rudy 2007. The hippocampal index revisited. A Hippocampus paper updating hippocampal indexing theory: the hippocampus stores an index pointing back to the cortical activity patterns of an experience, not a copy; recall reactivates the original pattern through the pointer, so the index stays compact while content remains where it first formed. Translation: build summaries and graph nodes as pointers back to source transcripts rather than duplicating content upward, keeping every derived record traceable and re-expandable. (Summarized from the published abstract.)

Tononi and Cirelli 2014. Sleep and synaptic homeostasis. A Neuron perspective arguing that a core function of sleep is synaptic downscaling: connections strengthened during waking are broadly weakened during sleep, and what survives the subtraction is what mattered; forgetting is part of the mechanism, not a failure of it. Translation: a store that only appends degrades under its own accumulation, so a periodic offline pass that prunes, downweights, and compacts keeps the distilled layer trustworthy as an archive grows. (Summarized from the published abstract.)

Zacks et al. 2007. Event segmentation theory. A Psychological Bulletin paper on event perception: people segment continuous experience into discrete events as it unfolds, and the boundaries are where

memory gets organized, the segments becoming the units later remembered. For conversational history this nominates the event, not the message or the fixed-size chunk, as the storage unit: segment sessions at topic boundaries and let those segments become the leaves the summary layers stand on. (Summarized from the published abstract.)

Sumers et al. 2024. Cognitive Architectures for Language Agents. A TMLR conceptual paper reading LLM agents as heirs to the production-system lineage through Soar. CoALA describes any agent along three dimensions: memory (working plus episodic, semantic, and procedural long-term stores), an action space split between external grounding and internal actions (retrieval, reasoning, learning), and a decision cycle of propose, evaluate, select. Table 2 casts SayCan, ReAct, Voyager, Generative Agents, and Tree of Thoughts into the frame; learned retrieval and deliberate unlearning are named as nearly unstudied. The export is the memory typology plus the reflection step turning accumulated episodes into reusable semantic facts.

Rezazadeh et al. 2025. MemTree. An ICLR paper storing an agent’s accumulating history as a dynamically grown tree, abstraction increasing toward the root. New content descends by cosine similarity against a depth-adaptive threshold, attaches as a new leaf where no child matches, and every ancestor summary is LLM-rewritten to fold it in; insertion is $O(\log N)$, retrieval one collapsed top-k search over all nodes. On the authors’ 70-session MSC-E extension MemTree (82.5) beats full-history prompting (78.0), and it roughly matches the offline RAPTOR on QuALITY (59.8 versus 59.0) despite building incrementally, at about 1.4x cumulative cost. Ancestor summaries drift with insertion order, and tested corpora top out at hundreds of documents. Online versus offline matters more than the specific structure: a modest ongoing tax keeps the index current.

Talebirad et al. 2026. A theory of hierarchical memory. An ICLR theory paper formalizing what RAPTOR, GraphRAG, and a dozen conversational and agent memory systems share: extraction into atomic units, iterated coarsening into a multi-level hierarchy, budgeted traversal at query time. Each lossy coarsening level strictly decreases entropy, so the atom layer caps what any query can recover. The central contribution is self-sufficiency, $SS = I(G; \rho(G)) / H(G)$, and its coupling to retrieval: self-sufficient summaries support collapsed search across all levels, referential ones demand top-down refinement, and mismatching the two wastes budget, though the coupling awaits a controlled test. Appendices bound branching factor and required depth from corpus size and context window. Choose summary style and search strategy jointly; the depth bound predicts when one summary layer stops sufficing.

Wu et al. 2021. Recursively summarizing books with human feedback. An OpenAI preprint training GPT-3 (6B and 175B) to summarize whole novels by fixed recursive decomposition: chunks, then summaries of summaries, to depth 3, one model at every level, trained by behavioral cloning then RL against human comparisons; locally judged subtasks make labels over 50x cheaper than end-to-end collection, the scalable-oversight framing. On 40 unseen 2020 books the best summaries averaged well below human-written ones, but the 175B models set state of the art on BookSum, and depth-1 summaries let a 3B reader score competitively on NarrativeQA. The failures instruct: local context turned a dance request into a marriage proposal, and themes spread thinly across chunks vanished. Hierarchical

summarization compresses arbitrarily long corpora; its errors make cross-leaf context and audits at composition points design requirements.

Retrieval, vector search, caching

Wang et al. 2024. Searching for best practices in RAG. An EMNLP paper treating the RAG pipeline as swappable modules, greedily optimizing one at a time with Llama2-7B over 10 million Wikipedia passages plus 4 million medical texts. The cheapest win is a BERT query classifier deciding whether to retrieve at all: the average score rises from 0.428 to 0.443 while latency falls from 16.41 to 11.58 seconds per query. Best retrieval combines HyDE with hybrid sparse-dense search (BM25 at alpha 0.3); monoT5 balances reranking quality and cost while TILDEv2 runs 225 times faster with some loss; reverse repacking, best evidence nearest the query, beats other orderings. Two recipes result, best-performance and a balanced one at roughly 1.5 seconds per query. The gate, hybrid retrieval under a light reranker, and evidence ordering are priors, the numbers coming from public QA corpora and a 7B generator.

Douze et al. 2024. The Faiss library. The design document for Meta’s open-source approximate nearest neighbor library, organizing a dozen index types around vector compression and non-exhaustive search, under one contract: approximate exact search under the metric the embedding makes meaningful. Selection rules: graph indexes (HNSW, NSG) win below roughly 1M vectors when memory allows; inverted-file indexes take over once compression must fit RAM; search time grows as N to the 0.29 at 50 percent recall and N to the 0.45 at 99. A production deployment indexes 1.5 trillion vectors in an 83 TiB memory map at about a second per query. Faiss stores no metadata and handles deletion poorly, so a database layer belongs on top. A personal corpus sits below the threshold where indexing beats brute force; an organizational one crosses into graph-index territory.

Kuffo, Krippner, and Boncz 2025. PDX. A SIGMOD paper on a storage layout grouping vectors into blocks with each dimension’s values contiguous within a block, the way Parquet holds columns, so pruning algorithms, which stop reading a vector once its partial distance disqualifies it, finally beat plain SIMD scans. Auto-vectorized PDX kernels averaged 2.0x over hand-written SIMD on the standard layout, 5.5x when eight or fewer dimensions are read, the case pruning creates; ADSampling on PDX ran 4.6x its own horizontal version and up to 7.2x over FAISS. The companion PDX-BOND is exact, needs no preprocessing of the vectors, and beats FAISS, USearch, and Milvus on exact search by 2.5x to 4.0x on average. A store appending embeddings daily cannot afford retraining on ingest; exact search stays feasible at personal scale.

Malkov and Yashunin 2016. HNSW. The TPAMI paper behind the graph index inside most current vector stores. Elements insert one at a time into layered proximity graphs, each element’s top layer drawn from an exponentially decaying distribution; search greedily descends from the sparse top layer like a probabilistic skip list; a neighbor-selection heuristic keeps only candidates closer to the new element than to any already-selected neighbor, preserving connectivity across cluster boundaries. The authors argue $O(\log N)$ search; on ann-benchmarks HNSW beats Annoy, FLANN, FALCONN, and VP-tree by wide margins, and against Faiss product quantization on 200M SIFT vectors it reaches higher recall at faster query times, building in 42 minutes versus 11 hours at more than double the RAM.

Insertion is truly incremental with no reshuffling; the deferred problems, element update and deletion plus awkward distributed search, are exactly what surfaces as a store grows.

Chan et al. 2025. Cache-augmented generation. A WWW Companion paper proposing to skip retrieval when the knowledge base fits a long-context window: preprocess the collection once into the model’s key-value cache, persist it, and append only the query at inference. With Llama 3.1 8B on SQuAD and HotPotQA corpora of 21k to 85k tokens, CAG beats the best RAG configuration in most settings (0.7951 versus 0.7676 on HotPotQA-small), answering in 0.85 to 2.26 seconds against up to 92 seconds for re-encoding per query. The margin inverts at the largest corpus, where sparse RAG top-5 wins, tracking known long-context degradation. Below some corpus size retrieval infrastructure is pure overhead; the durable idea is a tiered design, hot stable knowledge preloaded in cache, the long tail behind a retriever.

Modarressi et al. 2025. NoLiMa. An ICML benchmark of needle-in-a-haystack questions whose needles share almost no words with the question (ROUGE-1 overlap of 0.069 versus 0.905 for vanilla NIAH), so finding the fact requires a latent association, like knowing the Semper Opera House is in Dresden. Across 13 models claiming at least 128K context, short-context scores mostly exceed 84 percent, yet at 32K tokens 11 of 13 fall below half their base score; by an 85-percent-retention criterion most models sit at 2K or below, GPT-4o reaching 8K of a claimed 128K, and a single lexically similar distractor cuts that to 1K. Dumping transcripts into context fails well before the advertised window whenever the query is phrased differently than the conversation; retrieval must do the semantic bridging first.

Jeong et al. 2024. Adaptive-RAG. A NAACL paper routing each question, before any answering begins, to no retrieval, a single retrieve-then-read step, or an iterative loop interleaving reasoning with repeated retrieval. The router is a T5-Large classifier trained on automatic labels, the cheapest strategy that answered each sampled query correctly, backfilled from dataset provenance. Averaged over six QA datasets with GPT-3.5 it reaches 50.91 F1 at 1.03 retrieval steps per query, matching the always-iterate baseline (50.87) at about a third of its 2.81 steps; per-query cost spans 0.35 seconds to 27.18. The classifier is only 54.52 percent accurate and an oracle router reaches 62.80 F1, so routing quality is the ceiling. Retrieval depth becomes a cheap per-query decision, self-labeling letting usage logs generate their own routing data.

Benchmarks and faithfulness

Maharana et al. 2024. LoCoMo. An ACL paper of 50 long synthetic two-person conversations, averaging about 300 turns and 9,200 tokens over as many as 35 sessions, generated by persona-seeded gpt-3.5-turbo agents over temporal event graphs and then human-edited. Tasks cover QA across five reasoning types, event summarization, and multimodal dialogue generation. Every tested configuration sits far below the human QA benchmark of 87.9 F1; GPT-4-turbo manages 32.1. Longer context helps ordinary recall but collapses on adversarial questions (gpt-3.5-turbo-16k falls to 2.1 F1), hallucinating answers and misattributing speakers. The finding that travels furthest: RAG over distilled per-turn observations beats RAG over both raw logs and session summaries (top-5 observations reach 41.4 F1), and accuracy degrades as more are retrieved: precision beats volume. The adversarial collapse warns that context scale is not memory.

Hu, Wang, and McAuley 2025. MemoryAgentBench. An ICLR benchmark feeding context to memory agents incrementally, chunk by chunk with memorize-this instructions, so storage mechanisms actually fire before any question arrives. Four competencies, accurate retrieval, test-time learning, long-range understanding, and selective forgetting, tested through 2,071 questions over contexts of 103K to 1.44M tokens; no method masters all four. RAG wins retrieval (HippoRAG-v2 at 65.1 against the 49.2 backbone), long-context models win the holistic tasks (GraphRAG scores 0.4 on novel summarization), and forgetting breaks everyone: no system exceeds 28 percent on multi-hop fact updates. Commercial memory products trail a plain long context badly (Memo at 21.1 and Zep at 24.0, versus 60.6 for GPT-5-mini). Explicit update and consolidation machinery is required equipment.

Chang et al. 2024. BoookScore. An ICLR study comparing the two workflows of any long-corpus compressor: hierarchical merging of chunk summaries versus incremental updating of a running summary. On a contamination-resistant corpus of 100 recent books averaging 190K tokens, 1,193 human span annotations yield an eight-type coherence error taxonomy, automated by a GPT-4 prompt whose 78.2 percent precision essentially matches the annotators' own 79.7. Hierarchical merging is almost always more coherent, but Claude 2 with an 88K chunk lifts incremental updating from 78.6 to 90.9, fewer compression steps largely closing the gap; annotators still preferred incremental summaries for detail, 83 to 11 percent, and omission errors dominate. The comparison maps onto rolling-summary versus tree designs, argues for large update windows, and supplies a reference-free LLM-judge protocol for auditing summaries without gold answers.

Kim et al. 2024. FABLES. A COLM paper reporting the first large-scale human evaluation of faithfulness in book-length summarization: 3,158 atomic claims from 130 summaries of 26 novels published in 2023 or 2024, labeled by annotators who had already read the books, collected for \$5.2K. Claude-3-Opus leads at 90.9 percent faithful claims, yet only five of 130 summaries were entirely accurate. Unfaithful claims cluster on character and relationship states (38.6 percent) and events (31.5 percent), half requiring multi-hop reasoning to refute, and automatic verification largely fails: the best auto-rater, Claude-3-Opus reading the entire book, reaches 58.2 F1 at catching unfaithful claims, far ahead of BM25 retrieval at 24.9. Models also omit key events and over-weight book endings. Retrieval-then-verify collapses on claims distributed across a long record, a summarizing archive inherits silent omission and recency bias, and the reader-annotator protocol is the workable audit template.

Knowledge graphs: foundations

Chen 1976. The entity-relationship model. Published in ACM Transactions on Database Systems 1(1), this paper introduces the entity-relationship model: data as entities and the relationships among them. It stands at the head of the lineage every graph data model in this digest descends from. (Summarized from the published abstract.)

Angles et al. 2017. Foundations of graph query languages. A survey in ACM Computing Surveys 50(5) organizing graph querying by feature rather than by language, bridging database theory and SPARQL, Cypher, and Gremlin. It fixes two data models, edge-labelled and property graphs, and builds a feature hierarchy from basic graph patterns (equivalent to conjunctive queries) through complex patterns and regular path queries, with worked queries in all three languages. The complexity catalog

carries the weight: basic patterns with projection are NP-complete in combined complexity but logarithmic space in data complexity; optional and difference lift complex patterns to PSPACE-complete; RPQs run in linear time $O(|G| \text{ times } |L|)$ under arbitrary-path semantics yet turn NP-complete under simple-path semantics even in data complexity. Fixed small query templates over a growing graph sit in the tractable regime, and multi-hop recall stays linear so long as queries return endpoints rather than enumerating paths.

Hogan et al. 2021. Knowledge graphs. A tutorial survey by seventeen authors in ACM Computing Surveys 54(4), the field's first unified introduction. It defines a knowledge graph inclusively as a graph of data accumulating real-world knowledge, then walks the stack: data models (RDF-style edge-labelled graphs, property graphs, named-graph datasets), query languages, schema, identity, context, reasoning, and lifecycle. It separates three schema kinds (semantic, validating, emergent) and fixes reference numbers: Freebase held over three billion edges at its 2015 shutdown, Cyc represents 900 person-years yielding 2.2 million facts, NELL extracted 120 million edges, and enterprise graphs share six traits including billion-edge scale and deliberately lightweight taxonomies. Graphs postpone schema commitment while scope evolves, identity machinery decides when two mentions are one entity, and the lightweight-taxonomy finding cautions against over-engineering the ontology.

Snodgrass 1999. Developing time-oriented database applications in SQL. A Morgan Kaufmann book translating roughly 1600 temporal-database research papers into guidance for working SQL developers, organized around three orthogonal triples: temporal types (instant, interval, period), kinds of time (user-defined, valid, transaction), and kinds of statement (current, sequenced, nonsequenced). SQL-92 supports none of the temporal semantics, so the discipline falls on the application, illustrated through case studies including Nykredit, a Danish mortgage bank tracking nine million loans with both valid and transaction time on each row. The payoff is a repair procedure: roll back to the transaction time of the error, read the valid-time history as it then stood, and fix it, the fix itself logged because transaction time is append-only. Exactly what a conversational memory backend needs when a stored belief turns out to be wrong.

Weikum et al. 2021. Machine knowledge. A 289-page survey in Foundations and Trends in Databases of how large knowledge bases are built and maintained automatically, organized as a four-phase pipeline (discovery, canonicalization, augmentation, cleaning) with case studies of YAGO, DBpedia, NELL, Wikidata, and industrial graphs. Engineering conclusions: error rates below 5 percent (ideally under 1), the correctness-coverage tension resolved by seeding from premium sources before riskier open extraction, the primacy of canonicalized entities and types, and construction as a life-cycle with humans permanently in the loop. Public KBs hold millions of entities and up to billions of statements; proprietary ones run one to two orders larger. The wrap-up names the personal KB, an augmented memory built from a longitudinal digital footprint, as an open research problem, nearly a specification for a conversation-derived store.

Pan et al. 2024. Unifying LLMs and knowledge graphs. A roadmap survey in IEEE TKDE organizing roughly a hundred systems from 2018 to 2023 into three frameworks: KG-enhanced LLMs, LLM-augmented KGs, and synergized combinations. Documented trade-offs: pre-training injection

performs best but forbids updates without retraining, while inference-time retrieval handles changing facts at some cost; encoder-style completion cannot handle unseen entities while generative methods may hallucinate nonexistent ones. The probing evidence is sobering: scaling fails to appreciably improve memorization of tail facts, and one causal analysis found models complete facts from positionally close words rather than knowledge-bearing ones. The LLM-driven construction pipeline (entity discovery, coreference, relation extraction, iterative correction as in PiVE) is the machinery for distilling structured memory from transcripts; the graph-traversal line (StructGPT, Think-on-Graph) shows path-citing answers, the auditability a memory backend requires.

Rossi et al. 2021. KG embedding comparative analysis. An ACM TKDD study retraining 16 embedding-based link prediction models plus the rule-mining baseline AnyBURL on five standard benchmarks, correlating per-fact accuracy with structural features of the training graph. AnyBURL ranks at or near the top nearly everywhere while training in 100 to 1,000 seconds against 1 to 300 hours for embeddings; only ComplEx with N3 regularization is consistently comparable. Graph structure largely determines difficulty, and benchmark leakage is severe: a two-hop path-support score alone reaches 93.55 percent Hits@1 on WN18. Unreported score-tie policies produce incomparable numbers, CapsE swinging from 73.01 to 1.93 Hits@1 on FB15k. A graph distilled from one person’s conversations will be sparse and skewed, the regime where every tested model does worst, so cheap interpretable rules are the sensible start; the tie-policy result is a standing caution for retrieval evaluation.

Knowledge graphs: representation, time, practice

Angles 2018. The property graph database model. A ten-page AMW paper giving the first formal definition of the model behind Neo4j, Neptune, and Cosmos DB, supplying Codd’s three components: data structure, integrity constraints, query language. The structure is a tuple $(N, E, \rho, \lambda, \sigma)$ with partial labeling and property functions permitting multiple labels and multivalued properties; a schema tuple names node and edge types, and a delta function whitelists which edge types may connect which node-type pairs, with four validity conditions for schema-instance consistency. Query semantics is given by translation to non-recursive safe Datalog with negation, recursion (hence path queries) out of scope. The delta function is the shape of a vocabulary contract for an extraction pipeline, and the all-partial functions match a store growing before its ontology settles.

Hernandez et al. 2015. Reifying RDF for Wikidata. An ISWC workshop paper asking how Wikidata’s qualified statements survive translation into plain RDF, reduced to encoding quads (s, p, o, i) where the statement identifier i carries qualifiers. Four encodings (standard reification, n -ary relations, singleton properties, named graphs) are built from a February 2015 dump of 57.1 million quads and tested with 14 queries on five SPARQL engines. Singleton properties, the most concise scheme, broke four of the five engines because millions of unique predicates violated indexing assumptions; named graphs was fastest in 17 of 70 engine-query cases with reification and n -ary at 16 each, no clear winner; only Virtuoso handled all four. The reification choice determines which stores and query features remain usable; quads plus a statement identifier are the clean conceptual core.

Cai et al. 2024. Temporal KG survey. An arXiv survey (2403.04782) of temporal knowledge graph representation learning, where facts are quadruples (head, relation, tail, timestamp). It catalogs roughly

forty models from 2017 to 2024 in a ten-category taxonomy spanning transformation, tensor decomposition, graph neural networks, autoregression, temporal point processes, language-model methods, and few-shot learning, and documents the benchmarks: ICEWS14 with 7,128 entities, 230 relations, and 90,730 facts; GDELT with 2,278,405 facts over 2,751 timestamps. The interpolation-versus-extrapolation split, completing missing past facts versus forecasting future ones, maps onto recording when something was true and predicting what still is; the few-shot chapter addresses the sparse new entities a personal graph accumulates. Scaling to real-world graph sizes remains open.

Rost et al. 2021. Bitemporal property graphs. A report (arXiv 2111.13499) on a one-year Oracle and University of Leipzig cooperation presenting four artifacts: TPGM+, a bitemporal property graph model; T-PGQL, a temporal extension of PGQL; CGN, continuous event detection; BiTeGra, a relational-backed prototype. The core is bitemporality carried to the property level: every vertex, edge, and property value gets separate valid-time and transaction-time intervals, avoiding whole-element copies when one property changes. T-PGQL exposes the intervals through SQL:2011 period predicates and a FOR TX_TIME clause with AS OF, FROM TO, and BETWEEN variants for time travel. No benchmarks are reported. The two-clock design distinguishes “he changed jobs in May” from “I learned this in July,” and FOR TX_TIME AS OF is ready-made vocabulary for auditing what a system believed at any past moment.

Fowler 2021. Bitemporal history. A pattern essay on martinowler.com working one payroll example: Sally is paid on her recorded \$6000 salary, HR later reports a back-dated raise to \$6500, then corrects it to \$6400. Time has two dimensions, actual history (what happened in the world) and record history (what the system believed when), Fowler’s terms for Snodgrass’s valid and transaction time. The machinery is small but exact: a query takes two time parameters, salaryAt(‘2021-02-25’, ‘2021-03-25’), and while actual history stops being append-only once retroactive change is allowed, record history stays append-only, a reliable audit trail of the corrections themselves. Bitemporality complicates systems and is avoidable when actions capture their inputs. The pattern for correcting extracted facts without destroying the audit trail; generalizing record time into perspectives sketches multi-user ingestion.

Vrandecic and Krotzsch 2014. Wikidata. A CACM systems article by the project director and data model lead on the collaboratively edited knowledgebase Wikimedia launched in October 2012 to end fact duplication across 287 language Wikipedias. The representational contribution is a layered statement model rather than plain triples: statements carry qualifiers (“as of 2010”, “method: estimation”), references, and preferred or deprecated ranks so conflicting claims coexist, “value unknown” distinguished from mere incompleteness, properties themselves community-governed first-class pages. By August 2014 it held 15.8 million items and 43.2 million statements (23.2 million referenced), about 90 percent of the half million daily edits coming from bots under human curation. A memory backend can adopt the qualifier-and-reference structure, the rank mechanism for coexisting contradictions, and stable never-reused IDs; the volunteer schema governance must be replaced by curation elsewhere.

Zhong et al. 2024. KG construction survey. An ACM Computing Surveys 56(4) survey of more than 300 methods for building knowledge graphs automatically, organized into acquisition (entity discovery, coreference resolution, relation extraction), refinement, and evolution. It defines a knowledge

graph as $G = \{E, R, T, F_k\}$, the background-knowledge component F_k , a rule set or schema, separating a knowledge graph from a mere data graph. The construction-from-text chapters fix the pipeline vocabulary, from named entity recognition and linking through open, schema-bound, and distantly supervised relation extraction, and catalog resources with sizes (Wikidata at 96M+ items, YAGO at 2B+ facts); method coverage stops at BERT-era models. A memory system doing entity linking and coreference over chat logs runs exactly these tasks; the flagged open problems, long intricate contexts, knowledge dynamics, and federated privacy, are the organizational-scale versions of the same work.

Noy et al. 2019. Industry-scale knowledge graphs. An experience paper in ACM Queue and CACM in which practitioners from Google, Microsoft, Facebook, eBay, and IBM describe their production graphs. The scale numbers frame everything: Google roughly 1 billion entities and 70 billion assertions, Bing about 2 billion and 55 billion, Bing’s single Will Smith entry assembled from 108,000 facts across 41 websites against 200 Will Smiths in Wikipedia alone, which is why disambiguation ranks as the top industry challenge. The mechanisms are architectural: Facebook keeps provenance on every assertion and rebuilds the graph daily through an idempotent build, eBay funnels writes through a replicated log feeding specialized stores, IBM defers entity resolution to query time so context can settle names. Provenance-first assertions, late-binding disambiguation, and the log-plus-stores pattern all transfer to conversational facts; personalized on-device graphs are named an open direction.

Xu et al. 2024. KG-RAG for customer service. A five-page SIGIR industry-track paper from LinkedIn on a deployed question-answering system over historical Jira tickets. Each ticket is parsed into a section tree (summary, steps to reproduce, root cause, comments) by rules plus a template-guided LLM; trees are linked by explicit Jira references and thresholded title-embedding similarity, stored in Neo4j with embeddings in Qdrant; query-time entity and intent parsing scopes retrieval to matching section nodes, then a generated Cypher query walks the winning ticket’s tree. With GPT-4 and E5 held fixed in both arms, MRR rises from 0.522 to 0.927 and Recall@3 from 0.640 to 1.000; in a randomized six-month team split, median resolution time falls 28.6 percent. The dual-level pattern, parse each unit into a tree then connect units, gives coherent subgraph retrieval instead of arbitrary chunks; the hand-built template marks what must become automatic at scale.

Balog and Kenter 2019. Personal KG agenda. A four-page ICTIR position paper from two Google researchers defining the personal knowledge graph as a distinct research object: structured knowledge about entities of personal rather than general importance, with a mandatory spiderweb topology in which every node connects to one central user node. Three properties separate PKGs from general KGs: entities need not be prominent (“Mom’s dentist”), entity information can be radically sparse, and relations are often short-lived. Four research questions cover sparse ephemeral representation, long-tail entity linking where linking, population, and NIL-detection intertwine, automatic population without curators, and continuous two-way integration with external sources; no open PKG dataset exists. A checklist of the hard parts; the spiderweb constraint is a concrete schema decision, though the multi-user case, where the single central node no longer holds, goes unaddressed.

Conversation mining and outcomes

Kecht et al. 2021. Event logs from conversations via NLI. An ICPM paper building process-mining event logs from raw customer service conversations by zero-shot natural language inference: each message is the premise, a candidate topic or activity (“This example asks for iOS version”) the hypothesis, pre-trained BART returning an entailment probability, per-label thresholds tuned by Matthews correlation on 100 to 300 hand-labeled messages; no model training occurs. On Twitter support corpora (288,828 AmazonHelp tweets, 231,683 AppleSupport, 88,774 SpotifyCares), 55 of 56 classifications exceed 0.72 MCC, and three discovery algorithms recover the same coherent process model from the exported XES log. Hypothesis wording is brittle: one rephrasing moved MCC from 0.53 to 0.91. The labeling economics transfer, roughly 300 examples structuring 200,000-plus messages into a typed event timeline, with the caveat that the event schema must be known in advance.

Gung et al. 2023. DSTC 11 intent induction. A SIGDIAL paper from AWS AI Labs on a shared-task benchmark for inducing customer intents from raw transcripts, releasing three simulated call-center corpora (948, 1,000, and 2,000 conversations with 22, 29, and 39 intents, averaging 59 to 71 turns). Task 1 is intent clustering with the reference count withheld; Task 2 is open induction, judged by training a downstream classifier on the induced schema and scoring it on an independent test set. The winner reached 69.8 and 76.3 accuracy against baselines of 55.8 and 63.6, deep embedded clustering appearing in nine of the top ten systems; reassigning one system’s HDBSCAN noise points moved it from rank 33 to 9, exposing the metrics’ own weakness. Judge an induced structure by the system it trains, not cluster purity, and expect long-tailed request distributions to concentrate any induced taxonomy on a few head topics.

De Raedt et al. 2023. IDAS. An NLP4ConvAI paper on intent discovery that keeps the sentence encoder frozen and disambiguates in text: GPT-3 abtractively summarizes each utterance into a short label (“when is next payday”), bootstrapping demonstrations by labeling K-means prototypes first, then labeling remaining utterances with their 8 nearest already-labeled neighbors as in-context examples; utterance and label encodings are averaged, smoothed, and clustered at the true K. Against MTP-CLNN, the unsupervised state of the art, IDAS gains an average of 3.94 ARI across Banking, StackOverflow, and a private transport dataset; on CLINC it reaches 79.02 ARI with zero labels, beating two semi-supervised methods given labels for half the intents, and any use of generated labels beats raw encodings by 5 to 19 ARI points. Summarize-then-embed organizes an archive by meaning rather than surface form; the known-K assumption and per-utterance LLM cost bite as an archive grows.

Wegmann et al. 2022. Style versus content. A RepL4NLP paper attacking a confound in writing-style embeddings: authorship-verification models can score well by memorizing what an author writes about rather than how. The fix is a contrastive setup sampling the different-author candidate from the same conversation, so topic cannot separate authors, trained on 210k tasks from 100 subreddits. The decisive probe, STEL-Or-Content, pits a same-style paraphrase against a same-content different-style sentence: base RoBERTa picks style 5 percent of the time, the conversation-controlled model reaches .42 on the formal/informal dimension, and clustering its embeddings surfaces signatures like terminal punctuation and lowercase i. Same-source signals silently entangle speaker with topic; conversation structure supplies hard negatives free at archive scale, and the discipline of probing where shortcut and target disagree is one persistent-memory evaluation mostly lacks.

Liu, Zhang, Choi 2025. User feedback in human-LLM dialogues. An EMNLP paper asking whether implicit feedback in chat logs (a user saying “could you label the y-axis?”) can be detected and used for training. The authors densely annotate 109 conversations from LMSYS-chat-1M and WildChat with a six-way ontology (Cohen’s kappa 0.70 binary), finding feedback in more than half of later user turns, skewed heavily negative, refusals explaining under 3 percent; prompts drawing praise are slightly more toxic, traced to users celebrating jailbreaks. The detector roughly doubles prior binary accuracy (41.6 versus 31.5 percent). Training on feedback-conditioned regenerations is mixed: it wins on MTBench (6.86 versus 6.38) but loses on WildBench, plain distillation doing best. Later turns are dense with correction signal marking unmet intent, but the signal cannot be consumed raw.

Brynjolfsson et al. 2025. Generative AI at work. A field deployment of AI assistance in customer support, published in the Quarterly Journal of Economics 140(2), measured on resolution outcomes. It stands in this digest as the outcomes-side evidence that conversational AI assistance can be evaluated on what issues actually get resolved, the deployment-grade standard an organizational memory backend would be held to. (Summarized from the published abstract.)

Shared state across sessions

Erman et al. 1980. Hearsay-II. The Hearsay-II speech understanding system, described in ACM Computing Surveys 1980, introduced the blackboard architecture: independent knowledge sources coordinating through shared structured state rather than direct calls. The origin of the shared-state pattern the modern multi-agent work revives, and the ancestral form of any memory layer multiple agents read and write. (Summarized from the published abstract.)

Hayes-Roth 1985. A blackboard architecture for control. Published in Artificial Intelligence 1985, this work extends the blackboard from the domain problem to control: the choice of what to do next is also mediated through shared structured state, the step from a common store of facts to one that coordinates which contributor acts when. (Summarized from the published abstract.)

Ouyang et al. 2026. ReasoningBank. An ICLR 2026 memory framework for LLM agents facing task streams without ground-truth feedback. Each finished episode is distilled into a structured item (title, one-sentence description, content holding the strategy or pitfall); an LLM judge labels trajectories success or failure, and both feed extraction, successes as validated strategies and failures as preventative lessons; retrieved items are injected into the system instruction, and memory-aware test-time scaling spends extra rollouts per task to curate better items. On WebArena with Gemini-2.5-flash, success rises from 40.5 to 48.8, and to 51.8 with scaling at $k=5$; SWE-Bench-Verified resolve rate rises 34.2 to 38.8, steps per task dropping up to 26.9 percent on successful runs. The judge is only 72.7 percent accurate, yet results barely move across simulated accuracies of 70 to 90 percent. Distilled, self-contained lessons beat raw transcripts, and curation runs on noisy self-judgment; consolidation here has no pruning.

Rezazadeh et al. 2025. Collaborative Memory. An Accenture framework (arXiv 2505.18279) for sharing memory across users and agents without leaking beyond entitlements. Permissions are two time-varying bipartite graphs (user-to-agent, agent-to-resource); memory splits into private and shared tiers; every fragment carries immutable provenance (time, contributing user and agents, resources touched),

and a fragment is readable only if its provenance falls within current permissions, so revoking an edge retroactively hides dependent fragments. On MultiHop-RAG with five users and six domain agents, fully shared memory holds accuracy above 0.90 while cutting external resource calls by up to 61 percent at 50 percent query overlap; a science-QA scenario shows accuracy tracking grants and revocations with no logged fragment crossing a boundary. The provenance-stamped fragment plus retrospective permission check is the primitive a single-user memory tree needs to grow organizational tenants.

Salemi et al. 2026. Blackboard LLM system. An arXiv paper (2510.01285) reviving the Hearsay-II blackboard for LLM multi-agent data discovery: answering data science questions when the relevant files sit unidentified in a large heterogeneous data lake. Instead of a master assigning subtasks, the main agent posts requests to a shared board; helper agents, each pre-assigned one cluster of the lake and having studied it offline, volunteer only when they judge themselves capable. Across KramaBench and discovery-augmented DS Bench and DA-Code, the blackboard beats RAG and an otherwise-identical master-slave variant by 13 to 57 percent relative, with file-discovery F1 of 0.561 versus 0.247 for embedding retrieval, the advantage growing with lake size (211 percent relative on a 1,556-file lake) at about 2.3 times RAG’s cost. Self-selection beats a central index of who knows what because the coordinator’s model goes stale; partition, pre-digest, and broadcast scales in the direction an organization’s archive grows.

Pollertlam and Kornsuwannawit 2026. Beyond the context window. A cost-performance study (arXiv 2603.04814) comparing fact-based memory (the Memo pipeline: extracted atomic facts, pgvector retrieval, top-20 into a GPT-5-mini reader) against resending full history to long-context models, on LongMemEval, LoCoMo, and PersonaMem v2. Long-context wins factual recall decisively (92.85 versus 57.68 on LoCoMo, 82.40 versus 49.00 on LongMemEval), but the gap shrinks to 7.3 points on stable persona attributes. The economics invert with use: extraction compresses a 101,600-token history roughly 35:1 for a one-time \$0.0430 write, then about \$0.0013 per query, undercutting cached long-context after about ten turns at 100k tokens (nine at 500k). The error analysis names what compression loses: precise temporal markers, implicit coreferences, and updates that never propagate, a concrete checklist for any fact schema.

One-page summaries of the sources

91 summaries, alphabetical by source slug. Each states at its foot how much of the source it was written from.

Foundations of Modern Query Languages for Graph Databases

Angles et al. 2017, ACM Computing Surveys 50(5), arXiv:1610.06264

This survey organizes graph querying by feature rather than by language, aiming to bridge database theory and the practical languages SPARQL, Cypher, and Gremlin. It fixes two data models: edge-labelled graphs (the RDF style, edges as label triples) and property graphs, for which the paper gives a new formalization as a tuple $(V, E, \rho, \lambda, \sigma)$ with labels and attribute maps on both nodes and edges. On top of these it builds a small feature hierarchy: basic graph patterns (equivalent to conjunctive queries), complex graph patterns adding projection, join, union, difference, optional, and filter, then regular path queries (RPQs) for arbitrary-length navigation, and their combinations into navigational graph patterns. Each feature is illustrated with worked queries in all three languages.

The load-bearing content is the semantics catalog and the complexity that follows from it. Pattern matching can run under homomorphism semantics (SPARQL, Gremlin), isomorphism variants (Cypher uses no-repeated-edge), or cheaper simulation semantics; path queries can return arbitrary, shortest, simple, or edge-distinct paths, each as bags or sets. Evaluating basic patterns with projection is NP-complete in combined complexity under both homomorphism and isomorphism semantics, but polynomial under simulation, and only logarithmic space in data complexity, where the query is fixed and only the graph grows. Adding optional or difference lifts complex patterns to PSPACE-complete. RPQs returning node pairs under arbitrary-path semantics run in linear time $O(|G| \text{ times } |L|)$, yet the same query under simple-path semantics is NP-complete even in data complexity. The authors also flag what is unknown: Gremlin is Turing-complete in full, and Cypher's no-repeated-edge semantics and path unwinding (already NP-hard) had, as of 2017, essentially no formal analysis.

For a personal knowledge backend over conversational history, the practical lesson is that the store will be a property graph in all but name (entities, typed relations, provenance attributes on edges), and that the expensive parts are choices, not fate. Fixed, small query templates over a growing graph sit in the logspace data-complexity regime, which is the scaling case an org-wide memory would actually hit; tree-shaped patterns stay tractable where cyclic ones do not (treewidth); and multi-hop recall queries (who told whom, transitively) are the linear-time RPQ case so long as the system returns endpoints rather than enumerating simple paths. The semantics table also explains why memory layers built on Neo4j versus RDF stores can return different answers to the same conceptual query.

Read from: pages 1-25 and 37-41 of 59

The Property Graph Database Model

Renzo Angles, AMW 2018, CEUR Vol-2100 paper 26

This short paper (ten pages) gives a formal definition of the property graph database model, the abstraction behind Neo4j, Amazon Neptune, Azure Cosmos DB, and Titan. Angles observes that despite the commercial dominance of property graphs there is no standard specification of the model, which fragments the market and blocks research that would apply across systems. Following Codd's framing of what a database model must contain, the paper supplies the three components in turn: a data structure (Section 2), integrity constraints (Section 3), and a query language with a Datalog-based semantics (Section 4). It is a definitional paper, not an implementation or evaluation; nothing is benchmarked and no system is built.

The data structure is a tuple $G = (N, E, \rho, \lambda, \sigma)$: finite node and edge sets, an incidence function mapping each edge to an ordered node pair, a partial labeling function assigning sets of labels to nodes and edges, and a partial property function assigning sets of atomic values to (object, property name) pairs. The definition deliberately permits multiple labels per object and multiple values per property. A schema is a second tuple (TN, TE, β, δ) naming node types, edge types, typed properties, and which edge types are allowed between which node-type pairs; schema-instance consistency is then four validity conditions checked per node, edge, property, and endpoint pair. Mandatory versus optional properties, unique attributes, and cardinality constraints are sketched as extensions rather than fully developed. The query language combines SQL, SPARQL, and Cypher syntax (SELECT, FROM, MATCH, WHERE, plus UNMATCH for safe negation and UNION), and its semantics is given by translation to non-recursive safe Datalog with negation: a graph becomes Node, Edge, Label, and Prop facts (property identifiers exist specifically to support multivalued properties), and each query becomes a set of rules. The translation is shown by worked example only, not proved in general, and the language omits recursion, so path queries, the signature graph operation, are out of scope.

For a personal knowledge backend over conversational history, the schema machinery is the useful part: the delta function, which whitelists edge types between node-type pairs, is exactly the shape of a vocabulary contract for an extraction pipeline writing entities and relations into a store, and the optional-schema stance (all functions partial) matches how such a store grows before its ontology settles. The Datalog encoding also shows that a property graph reduces cleanly to four relational predicates, a useful mental model for judging whether a graph database is buying anything over a relational one at small scale. As a reading of how knowledge graphs are represented, this is the cited formalization of the property graph side, though its query fragment is weaker than what Cypher or the paper's own G-CORE successor provide.

Read from: full text

Survey Article: Inter-Coder Agreement for Computational Linguistics

Ron Artstein and Massimo Poesio, Computational Linguistics 34(4), 2008; DOI 10.1162/coli.07-034-R2 (ACL Anthology J08-4004)

This survey is the standard treatment of how to measure whether human annotators agree with each other enough for their labels to count as reliable data. Working from a decade of computational linguistics practice after Carletta's 1996 kappa proposal, the authors expose the mathematics and assumptions behind the main chance-corrected coefficients: Bennett's S, Scott's pi, Cohen's kappa, weighted kappa, and Krippendorff's alpha. All follow one formula, observed agreement minus expected agreement over one minus expected agreement; they differ only in how chance is modeled. S assumes a uniform category distribution (and can be gamed by adding categories nobody uses), pi estimates one shared distribution from the pooled data, and kappa gives each coder an individual distribution, which folds systematic coder bias into the chance term. Alpha generalizes pi through distance functions, handling multiple coders, missing values, and graded disagreement, and for nominal data it converges to multi-pi as items or coders grow.

The concrete guidance matters as much as the algebra. Raw percent agreement misleads: with a 95/5 category split, two coders agree 90.5 percent of the time by pure chance, so an observed 90 percent is worse than chance. On skewed data, reliability reduces to agreement on the rare category, roughly δ over $\delta + \epsilon$ in their two-parameter analysis. The pi versus kappa debate is settled by purpose, not by marginal divergence: single-distribution coefficients (pi, alpha) measure reproducibility of the annotation procedure and are right for reliability claims, while kappa speaks to trustworthiness of the specific coders; the numerical gap shrinks as coders are added, so they recommend more than two coders. On interpretation they are candid about limits: their experience says values above 0.8 mark reasonable quality, some useful discourse corpora sit near 0.7, and no single cutoff fits all purposes, so report the methodology and the confusion matrix, not just the score.

For the memory harness this paper governs the judge-calibration step inside LLM-judge scoring, after the three injection arms are run and answers need grading. Before trusting a schema-validated judge call to score arm outputs at scale, the judge must be checked against Justin's hand labels on the small calibration set, and that check should be reported as chance-corrected agreement, not raw match rate, since correct labels will be common and inflate percent agreement. Krippendorff's alpha is the right coefficient there (it tolerates a small item count via its unbiased estimator and admits partial credit for near-miss verdicts), with 0.8 as the working reliability bar and anything near 0.7 flagged as provisional.

Read from: pages 1-20 and 37-42 of the 42-page PDF (full coefficient development, bias and prevalence analysis, interpretation guidance, and references; the annotation-task case studies on pages 20-36 were skipped).

AutoSchemaKG: Autonomous Knowledge Graph Construction through Dynamic Schema Induction from Web-Scale Corpora

Jiaxin Bai et al., arXiv:2505.23628, 2025-05

Abstract (verbatim)

We present AutoSchemaKG, a framework for fully autonomous knowledge graph construction that eliminates the need for predefined schemas. Our system leverages large language models to simultaneously extract knowledge triples and induce comprehensive schemas directly from text, modeling both entities and events while employing conceptualization to organize instances into semantic categories. Processing over 50 million documents, we construct ATLAS (Automated Triple Linking And Schema induction), a family of knowledge graphs with 900+ million nodes and 5.9 billion edges. This approach outperforms state-of-the-art baselines on multi-hop QA tasks and enhances LLM factuality. Notably, our schema induction achieves 92\% semantic alignment with human-crafted schemas with zero manual intervention, demonstrating that billion-scale knowledge graphs with dynamically induced schemas can effectively complement parametric knowledge in large language models.

Bearing on this project

The schema-induction scoop. Fully autonomous knowledge graph construction with no predefined schema, inducing schemas directly from text while extracting triples. 50M documents, 900M nodes, and 92 percent semantic alignment with human-crafted schemas at zero manual intervention. That is cold-start structure induction, at web scale, with a number.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Personal Knowledge Graphs: A Research Agenda

Balog and Kenter, ICTIR 2019, DOI 10.1145/3341981.3344241

This is a four-page position paper from two Google researchers that defines the personal knowledge graph (PKG) as a distinct research object and lays out an agenda for studying it. The paper contains no system, dataset, or experiment; its method is conceptual. It defines the PKG, contrasts it with prior work (Montoya et al.'s PIM knowledge base, Gyrard et al.'s personalized health graph, Yen et al.'s life-event extraction from tweets, and personalized views over general KGs), then works through four problem themes, each anchored by an explicit research question, and closes with notes on evaluation, implementation, and utilization.

The definition it establishes: a PKG is structured knowledge about entities and relations of personal rather than general importance, with a mandatory “spiderweb” topology in which every node connects, directly or indirectly, to one central user node. Three properties separate PKGs from general KGs: entities need not be globally prominent (a friend, “my guitar”, “Mom’s dentist”); entity information can be radically sparse, sometimes just a type plus the relation to the user; and relations are often short-lived, the opposite of the stability criterion general KGs use for inclusion. The four research questions cover representation under sparse, ephemeral predicates (RQ1); entity linking when the target entity has little or no structured information and may not be a proper noun, which the authors frame as long-tail linking where linking, population, and NIL-detection intertwine (RQ2a, plus RQ2b on when to link against the PKG versus a general KG); automatic population and maintenance without human curators, where they note that data-hungry neural link prediction is unlikely to transfer (RQ3); and continuous, one-to-many, potentially two-way integration with external sources, with the user in the loop for conflicts (RQ4). On evaluation they are blunt: no open PKG dataset exists, logging real user interactions is costly and privacy-fraught, and synthetic data may be the only path.

For a personal knowledge backend over conversational history, the paper is a checklist of the hard parts rather than a solution: first-mention population, resolving “Jamie” or “my guitar” without any external profile, deciding what to store (only user-relevant attributes, not full entity records), and expiring ephemeral relations all map directly onto conversation-derived memory. The user-centered spiderweb constraint is a concrete schema decision a backend can adopt, though the paper says nothing about multi-user or organizational graphs, where the single central node no longer holds. As a representation reference it is a compact framing of how KGs are used (entity linking, population, verification, completeness prediction) with pointers into that literature, not a treatment of any mechanism in depth.

Read from: full text

LLM-empowered knowledge graph construction: A survey

Haonan Bian et al., arXiv:2510.20345, 2025-10

Abstract (verbatim)

Knowledge Graphs (KGs) have long served as a fundamental infrastructure for structured knowledge representation and reasoning. With the advent of Large Language Models (LLMs), the construction of KGs has entered a new paradigm-shifting from rule-based and statistical pipelines to language-driven and generative frameworks. This survey provides a comprehensive overview of recent progress in LLM-empowered knowledge graph construction, systematically analyzing how LLMs reshape the classical three-layered pipeline of ontology engineering, knowledge extraction, and knowledge fusion. We first revisit traditional KG methodologies to establish conceptual foundations, and then review emerging LLM-driven approaches from two complementary perspectives: schema-based paradigms, which emphasize structure, normalization, and consistency; and schema-free paradigms, which highlight flexibility, adaptability, and open discovery. Across each stage, we synthesize representative frameworks, analyze their technical mechanisms, and identify their limitations. Finally, the survey outlines key trends and future research directions, including KG-based reasoning for LLMs, dynamic knowledge memory for agentic systems, and multimodal KG construction. Through this systematic review, we aim to clarify the evolving interplay between LLMs and knowledge graphs, bridging symbolic knowledge engineering and neural semantic understanding toward the development of adaptive, explainable, and intelligent knowledge systems.

Bearing on this project

A survey of LLM-empowered knowledge graph construction. Its existence is itself evidence the space is mapped; useful as a related-work anchor.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

A Survey on Temporal Knowledge Graph: Representation Learning and Applications

Cai, Mao, Zhou, Long, Wu, and Lan; arXiv preprint 2403.04782, 2024.

This is a survey of temporal knowledge graph representation learning (TKGRL): methods that learn low-dimensional embeddings for entities, relations, and timestamps in graphs whose facts are quadruples (head, relation, tail, timestamp) rather than static triples. The authors position it as the first broad survey of the area; prior surveys covered static KGs with a short temporal section, or only the completion subtask. The paper proceeds from problem formulation, datasets, and metrics, to a ten-category taxonomy of methods, to three application areas, and closes with future directions. The taxonomy chapters catalog roughly forty models from 2017 to 2024 at the level of representation space, encoder, and scoring function, without new experiments or head-to-head benchmark numbers, so the survey establishes structure rather than empirical rankings.

The standing infrastructure it documents is concrete. Four benchmark families dominate: ICEWS event data (ICEWS14 has 7,128 entities, 230 relations, 365 timestamps, 90,730 facts), GDELT (2,278,405 facts over 2,751 timestamps), Wikidata, and YAGO; evaluation is by mean reciprocal rank and Hit@k. The ten method categories are transformation (translation and rotation, extending TransE and RotatE with time hyperplanes or temporal rotations), tensor decomposition (order-4 CP and Tucker variants such as TComplEx and TuckERT), graph neural networks, capsule networks, autoregression (RE-NET, RE-GCN: snapshot sequences encoded by relational GCNs plus GRUs), temporal point processes for irregular intervals (Know-Evolve, Graph Hawkes), interpretability via subgraph reasoning or reinforcement learning, language-model methods (in-context forecasting, retrieval plus fine-tuning), few-shot learning for rare entities and relations, and a residual class including copy-generation and neural ODEs. A useful task split runs throughout: interpolation (completing missing past facts) versus extrapolation (forecasting future ones), which demand different machinery.

For a personal knowledge backend built over conversational history, the relevant contribution is the framing itself: facts that hold only during an interval, the quadruple representation, and the interpolation versus extrapolation distinction map directly onto memory that must record when something was true and predict what still is. The autoregressive snapshot view suggests treating accumulated conversation state as a timestamped graph sequence, and the few-shot chapter addresses exactly the sparse new entities a personal graph accumulates. The scalability discussion is candid that current benchmarks are far smaller than real-world graphs and that most methods ignore scaling, so organizational use remains an open problem; the LLM-integration section (using language models for relation descriptions and semantic enrichment) is directional rather than settled. As a map of how temporal knowledge is represented and scored, the survey is a serviceable entry point and reference index rather than a source of design-ready results.

Read from: full text, pages 1 through 25 of 36 (everything before the references), with the per-model embedding sections in chapter 3 skimmed rather than studied.

Ask Only When Needed: Proactive Retrieval from Memory and Skills for Experience-Driven Lifelong Agents

Yuxuan Cai et al., arXiv:2604.20572, 2026-04

Abstract (verbatim)

Online lifelong learning agents must decide not only how to act but also when to consult prior experience to continually improve on long-horizon tasks. Existing methods typically retrieve memories passively, such as at task initialization or after each step, and therefore miss knowledge gaps that arise during interaction. We propose ProactAgent, an experience-driven lifelong learning framework for proactive retrieval over a structured Experience Base. ProactAgent continually improves through ExpOnEvo, which jointly updates policies and refines memory, organizing past interactions into factual, episodic, and skill repositories. It further introduces ProactRL, which treats retrieval as an explicit policy action and learns when and what to retrieve. By comparing paired continuations from identical interaction prefixes with and without retrieval, ProactRL provides step-level process rewards that encourage retrieval only when it improves task outcomes or efficiency. Experiments on SciWorld, AlfWorld, and StuLife show that ProactAgent consistently outperforms all baselines, achieving up to 32% relative improvement in success rate and over 33% reduction in interaction rounds. Our code will be publicly available at GitHub.

Bearing on this project

Takes the measurement. Observes that existing methods retrieve passively and therefore miss knowledge gaps arising during interaction, and learns when to retrieve by comparing paired continuations from identical prefixes with and without retrieval. That paired counterfactual is the core of the counterfactual miss rate, expressed as a learned policy rather than a reported rate.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Structured Prompt Interrogation and Recursive Extraction of Semantics (SPIRES): a method for populating knowledge bases using zero-shot learning

J. Harry Caufield, Harshad Hegde, et al. *Bioinformatics* 40(3), 2024. DOI 10.1093/bioinformatics/btae104 (arXiv:2304.02711)

SPIRES is a zero-shot method for populating knowledge bases from raw text, built on one commitment: every LLM output must conform to a user-defined schema, and every named entity must ground to an identifier in an external ontology or vocabulary. The user writes the schema in LinkML, a formal language whose classes carry typed attributes, cardinality constraints, per-attribute prompts, and value sets defined by ontology queries. Given a schema, an entry-point class, and a text, SPIRES generates a pseudo-YAML template prompt, sends it to the model, parses the completion heuristically (line splitting on colons, since responses are not guaranteed valid YAML), and recurses: any attribute whose range is a nested class triggers a fresh SPIRES call on that value alone. Leaf entities are then grounded through the Ontology Access Kit against resources like FOODON, MeSH, and MONDO, and validated against identifier-prefix and value-set constraints. The implementation ships as the open-source OntoGPT Python package.

The grounding numbers make the argument. On 100 random Gene Ontology terms, SPIRES with GPT-3.5-turbo returned 98 correct identifiers; the same model prompted directly returned 3, hallucinating plausible-looking IDs for nearly all 100. GPT-4-turbo often refused to ground at all, and on MONDO disease terms SPIRES with GPT-4-turbo managed only 18 of 100 against 97 for GPT-3.5-turbo, so a newer model tier is not uniformly better at this stage. On the BC5CDR chemical-disease relation benchmark, with no fine-tuning, SPIRES scored F 41.16 (chunked, GPT-3.5) and 43.80 with GPT-4 unchunked, just below the average of the original challenge's 18 trained systems and near BioGPT's 44.98; the trained state of the art reached 0.57, so zero-shot extraction trades several F points for zero training data. Hallucinated relations were rare but real, including a technically true statement (lisinopril induces angio-oedema) that the source text never made, and the authors insist LLM output be validated before entering a KB.

For the memory harness, SPIRES is the published precedent for the fact-extraction stage, where the corpus becomes schema-shaped entity and dated-fact rows. It shows that schema validation plus external grounding converts silent hallucination into countable rejections, exactly the measured-rejection design of the harness's schema-validated calls; that model tiers are swappable behind one prompt-completion interface but must be measured per stage rather than assumed; and that recursive per-class prompting handles nesting that flat triple extraction cannot. Its scope is single-document extraction only: it says nothing about tree consolidation, injection arms, or judge scoring.

Read from: pages 1 through 20 of 21 (the full body and conclusion end on page 11; the rest is references), extracted with PyMuPDF.

Don't Do RAG: When Cache-Augmented Generation is All You Need for Knowledge Tasks

Chan, Chen, Cheng, and Huang; WWW Companion 2025; arXiv:2412.15605

The paper proposes cache-augmented generation (CAG) as a replacement for retrieval-augmented generation when the knowledge base is small enough to fit in a long-context model's window. Instead of retrieving passages per query, the whole document collection is preprocessed once into the model's key-value cache, which is stored on disk or in memory; at inference time only the query is appended, and the model answers with the entire corpus already resident in its inference state. The method has three phases: offline preloading (a one-time encoding cost regardless of query count), inference over the cached context, and a cache reset that truncates the appended query tokens rather than reloading from disk. The claimed gains are zero retrieval latency, no retrieval errors, and a simpler architecture with no retriever to tune.

Experiments use Llama 3.1 8B on subsets of SQuAD and HotPotQA sized from 21k to 85k tokens (3 to 64 documents), against BM25 sparse retrieval and OpenAI dense embeddings via LlamaIndex at top-1 through top-10, scored with BERTScore. CAG wins in most configurations: on HotPotQA-small it scores 0.7951 against 0.7676 for the best sparse RAG setting, and on SQuAD-large 0.7734 against 0.7658. The margin narrows as the corpus grows; at HotPotQA-large CAG's 0.7407 falls below sparse RAG top-5 at 0.7535, which the authors connect to known long-context degradation. Timing on HotPotQA shows CAG generation at 0.85 to 2.26 seconds across sizes, while plain in-context learning (encoding the reference text per query) takes 9.3 to 92.1 seconds, a roughly 40x gap at the large size. Sparse RAG beating dense RAG throughout suggests the benchmarks were not retrieval-hard. All results are for one 8B model on corpora under 100k tokens; the authors state plainly that the method is impractical for significantly larger datasets, naming small-company knowledge bases, FAQs, and call centers as the intended fit.

For a personal knowledge backend over conversational history, CAG marks one endpoint of a design spectrum: below some corpus size, retrieval infrastructure is pure overhead and precomputing the context is both faster and more accurate. A personal archive might start under that threshold; an organizational one will not, and the paper's own large-corpus result shows the accuracy cost of stuffing everything into context. Its concluding suggestion, a preloaded foundation context with retrieval reserved for edge cases, is the more durable idea for memory systems: a tiered design where hot, stable knowledge lives in cache and the long tail stays behind a retriever. The precomputed-KV technique itself is memory-adjacent engineering worth tracking, since it treats an LLM's inference state as a persistable artifact.

Read from: full text

BooookScore: A Systematic Exploration of Book-Length Summarization in the Era of LLMs

Chang, Lo, Goyal, and Iyyer, ICLR 2024, arXiv:2310.00785

The paper studies how well LLMs summarize books longer than 100K tokens, a task that forces the document to be chunked and the chunk summaries combined. The authors compare two prompting workflows: hierarchical merging, where chunk summaries are recursively merged until one remains, and incremental updating, where a running global summary is updated and compressed as the model steps through the book. Because existing benchmarks like BookSum sit in LLM pretraining data, they build a contamination-resistant corpus of 100 recently published books (average 190K tokens, no public summaries found for any of them) and collect 1,193 span-level human annotations on GPT-4 summaries, at a cost of about \$3K and 100 annotator hours. From these they derive a taxonomy of eight coherence error types, including two not seen in earlier taxonomies for fine-tuned summarizers: causal omissions and salience errors. They then automate the annotation as BooookScore, a GPT-4 prompt that checks each summary sentence against the taxonomy; the metric is the fraction of error-free sentences.

Validation shows the automatic annotations have 78.2 percent precision against human judgment versus 79.7 percent for the human annotations themselves, and system-level scores computed from human and LLM annotations nearly coincide. Applying the metric across models (a \$10K API study), Claude 2 and GPT-4 lead (hierarchical scores of 91.1 and 89.1 at chunk size 2048), Mixtral-8x7B roughly matches GPT-3.5-Turbo, and LLaMA2-7B-Instruct trails badly at 72.4 with 36 percent repeated trigrams and cannot perform incremental updating at all. Hierarchical merging is almost always more coherent than incremental updating, but the exception is instructive: Claude 2 with an 88K chunk jumps from 78.6 to 90.9 on incremental updating, meaning fewer update-and-compress steps largely erase the gap. Annotators still preferred incremental summaries for detail (83 to 11 percent) while preferring hierarchical ones overall (54 to 44), and overall preference did not correlate with the coherence score. The taxonomy and metric are both derived solely from GPT-4 outputs on English trade books, a scope limit the authors flag themselves.

For a personal knowledge backend built over conversational history, this is essentially a study of the two compression strategies such a system must choose between. Incremental updating is exactly the rolling-summary pattern most memory systems use, and the paper shows it degrades coherence through repeated compression unless each update window is large; hierarchical merging is more coherent but loses long-range dependencies and detail. The finding that omission errors (entity, event, causal) dominate the error distribution tells a memory designer what to audit for, and the reference-free, taxonomy-grounded LLM-judge protocol is a workable template for evaluating memory summaries where no gold answer exists.

Read from: pages 1-12 of 35 (full main text; pages 13-35 are references and appendices)

FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance

Lingjiao Chen, Matei Zaharia, James Zou; arXiv (cs.LG), 2023 (later TMLR); arXiv:2305.05176

Chen, Zaharia, and Zou survey the pricing of twelve commercial LLM APIs from five providers and find fees that differ by two orders of magnitude: 10M input tokens cost \$0.20 on Textsynth's GPT-J but \$30 on GPT-4. From this they lay out three families of cost-reduction strategy. Prompt adaptation shrinks what gets sent (prompt selection, query concatenation); LLM approximation substitutes cheaper infrastructure for the expensive model (completion caches, fine-tuning a small model on the big model's outputs); and LLM cascade sends each query through an ordered list of models, stopping at the first answer a learned scorer deems reliable. They implement the third strategy as FrugalGPT: a DistilBERT regression scorer rates each generation, and a specialized optimizer (pruning candidate lists by answer disagreement, interpolating the objective from samples) picks the three-model chain and per-model thresholds under a user-set budget.

The empirical results are specific. On HEADLINES, a financial news classification task, the learned cascade (GPT-J, then J1-L, then GPT-4, with score thresholds 0.96 and 0.37) beats GPT-4's accuracy by 1.5 points at one fifth its cost; matching GPT-4's accuracy costs \$0.60 against GPT-4's \$33.10, a 98.3 percent saving. Savings on OVERRULING and COQA are 73.3 and 59.2 percent, with accuracy gains up to about 4 to 5 percent at matched cost. The mechanism is generation diversity: for roughly 6 percent of HEADLINES queries GPT-4 errs where GPT-J succeeds. The results come with real caveats stated by the authors themselves: the cascade needs labeled training examples drawn from the same distribution as the test queries, learning it is an upfront cost that only pays off on large query volumes, and the evaluation covers short-answer classification-style tasks, not long-form generation.

For the memory harness this is the prior art governing the model-tier assignment across pipeline stages, the stage where schema-validated LLM calls (tree building from the corpus, entity extraction, dated-fact extraction, the three injection arms, LLM-judge scoring) get routed through swappable model tiers. FrugalGPT frames the exact question the stage-wise sensitivity experiment answers: which stages a mini-class model can own outright and where escalating to a strong model buys accuracy per dollar. Its distribution-match caveat carries over directly, since a routing rule tuned on one corpus may not transfer, and its labeled-data requirement suggests the harness's judge scores could double as the reliability signal a cascade would need.

Read from: full text of all 13 pages, extracted with pypdf.

VitaBench 2.0: Evaluating Personalized and Proactive Agents in Long-Term User Interactions

Yuxin Chen et al., arXiv:2605.27141, 2026-05

Abstract (verbatim)

Large language models (LLMs) have evolved into interactive agents that collaborate with users in real-world tasks. Effective collaboration in such settings increasingly depends on understanding the user beyond what is explicitly stated, as user intent is often reflected in fragmented daily interactions and requires both personalized modeling and proactive interaction. However, existing agent benchmarks primarily evaluate reasoning and tool use, largely overlooking the challenges of inferring and leveraging user preferences in realistic scenarios. To address this gap, we introduce VitaBench 2.0, a benchmark for evaluating personalized and proactive agent behavior in long-term user interactions. In VitaBench 2.0, tasks are organized as temporally ordered sequences for individual users, where preferences are embedded in fragmented and heterogeneous interactions. Successful completion of tasks requires the agent to continuously extract, utilize, and update user preferences from these interactions. We further evaluate proactiveness through tasks that require agents to recognize missing information and actively acquire it from users or environments before making decisions. To support systematic analysis, we provide an extensible memory interface that enables controlled comparison across different memory architectures. We benchmark a diverse set of frontier proprietary and open-source LLMs. Results show that real-world personalization remains highly challenging even for state-of-the-art models, revealing a substantial gap between current capabilities and practical requirements. Extensive analysis further reveals the failure modes and capability bottlenecks of current agents in real-world personalized decision-making, providing insights for future model improvements.

Bearing on this project

Evaluating personalized and proactive agents in long-term user interaction. Benchmark, part of the same cluster.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

From “Strings” to “Things” for Personal Knowledge Graphs: Evaluating LLM Triple Extraction for Recommendation Systems

Abhirup Dasgupta et al., arXiv:2607.00003, 2026-04

Abstract (verbatim)

Personal Knowledge Graphs (PKGs) offer a privacy-preserving framework for modeling user preferences, yet constructing them from unstructured, decentralized conversational data remains a challenge. This paper bridges the gap between conversational “strings” and semantic “things” by presenting a reproducible pipeline for extracting structured user-preference triples using lightweight Large Language Models. We evaluate Qwen- and Gemma-based models on their ability to extract RDF-compliant triples linked to Wikidata identifiers from conversational data for PKG construction. Our evaluation assesses both the semantic extraction fidelity and the utility of the resulting graphs in a downstream recommendation task. We found that certain models performed well and had proportionally high downstream performance relative to their triple extraction performance.

Bearing on this project

Personal knowledge graph construction from unstructured conversational data, privacy-framed, with evaluation of extraction fidelity and downstream utility. This project’s corpus type and privacy angle.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

IDAS: Intent Discovery with Abstractive Summarization

De Raedt et al. 2023, IDAS: Intent Discovery with Abstractive Summarization, NLP4ConvAI at ACL 2023

The paper tackles intent discovery: partitioning a pile of unlabeled utterances into clusters that each share a conversational goal, the usual first step in building a chatbot without paying for annotation. Prior methods train an unsupervised sentence encoder on the domain data, then cluster; the authors argue that what makes two encodings similar is opaque and often driven by syntax or noun overlap rather than intent. IDAS instead keeps the encoder frozen and does the disambiguation in text: an LLM (GPT-3 text-davinci-003, temperature 0) abtractively summarizes each utterance into a short label such as “when is next payday”. Because the task has no labeled examples for in-context learning, IDAS bootstraps them. It runs an initial K-means, takes the utterance nearest each centroid as a prototype, has the LLM label the prototypes with a five-word-max instruction, then labels every remaining utterance using its 8 nearest already-labeled neighbors as ICL demonstrations, growing the demonstration pool as it goes. Utterance and label encodings are averaged, smoothed over n' nearest neighbors (n' picked by silhouette score), and clustered with K-means at the true K.

Against MTP-CLNN, the unsupervised state of the art, IDAS gains an average of 3.94 ARI, 2.86 NMI, and 3.34 accuracy points across Banking (77 intents), StackOverflow (20 topics), and a private 1,257-utterance transport dataset (42 intents), with the largest single gain +7.42 ARI on Transport; on Banking the two methods are roughly tied when compared in matched settings. On CLINC (150 intents) IDAS reaches 79.02 ARI and 85.48 accuracy with zero labels, beating two semi-supervised methods that were given labels for 50% of the intents, and trailing only at the 75% ratio. Ablations show every strategy using generated labels beats raw utterance encodings (by 5 to 19 ARI points), nearest-neighbor demonstration selection clearly beats random, and doubling K in the prototype step barely hurts. The results all assume the true intent count is known, and the test sets are small (1,000 to 3,080 utterances); smaller models (curie, babbage, Flan-T5-XL) produced labels too poor to use.

For a personal knowledge backend over conversational history, the transferable idea is that summarize-then-embed beats embed-raw-text when you need clusters organized by meaning rather than surface form, and that a seed-and-grow labeling loop turns an unlabeled corpus into its own demonstration set. The K-must-be-known and per-utterance-LLM-call costs are exactly the constraints that bite at organizational scale, and the paper flags both as open problems.

Read from: full text

The Faiss Library

Douze et al., arXiv preprint, 2024, arXiv:2401.08281

This paper is the design document for Faiss, Meta’s open-source C++ library (with Python bindings) for approximate nearest neighbor search over embedding vectors. Rather than proposing a single new algorithm, it explains how the library organizes a dozen index types around two tools, vector compression and non-exhaustive search, and how a user should trade off speed, memory, and accuracy under a fixed budget. The authors frame the whole problem as an “embedding contract”: the embedding model makes distances meaningful, and the index’s only job is to approximate exact search under the agreed metric. They are equally explicit about what Faiss is not: it extracts no features, runs as no service, and provides no database machinery such as sharding, transactions, or concurrent writes.

The benchmarks establish practical selection rules. On datasets from SIFT features (BIGANN, 128 dimensions) up to 768-dimensional Contriever text embeddings of English Wikipedia, additive quantizers win at small code sizes (around 8 bytes per vector), product-additive hybrids take over at larger budgets, and plain scalar quantizers suffice for long codes. For non-exhaustive search, graph indexes (HNSW, NSG) are recommended below roughly 1M vectors when memory is unconstrained, while inverted-file (IVF) indexes are the only option once compression is needed to fit in RAM; graph quality peaks at 64 edges per node and degrades beyond. Fitting a scaling model on BigANN at 1B vectors, search time grows as N to the 0.29 at 50 percent recall@1 and N to the 0.45 at 99 percent, so cost stays below the square root of N but climbs steeply with the accuracy target. On Deep10M under ANN-benchmarks, Faiss beats SCANN and roughly matches DiskANN’s Vamana. A production case study indexes 1.5 trillion 144-dimensional vectors at 54 bytes each (PCA to 72 dimensions plus a 6-bit scalar quantizer), sharded across hundreds of servers into an 83 TiB memory map, reaching about 1 second per query.

For a personal memory backend over conversational history, the field guide here is the sizing logic: a personal corpus sits far below the 10k-vector threshold where indexing beats brute force, while an organizational one crosses into graph-index territory, so the same Faiss code covers both ends without redesign. The gaps matter too. Faiss stores no metadata, filters only through id-bit tricks, and its graph indexes handle deletion poorly; a memory system needs a database layer on top, which is exactly the Faiss-plus-DBMS split that Milvus and Pinecone embody. The retrieval-augmented language model work it cites (Atlas, kNN-LM, RA-DIT) is the direct lineage of persistent AI memory.

Read from: full text

From Local to Global: A GraphRAG Approach to Query-Focused Summarization

Edge et al. 2024, From Local to Global: A Graph RAG Approach to Query-Focused Summarization, arXiv:2404.16130

The paper targets a failure mode of conventional vector RAG: questions that require global understanding of a whole corpus (“what are the main themes here?”) are query-focused summarization tasks, not retrieval tasks, and retrieving a handful of similar chunks cannot answer them. GraphRAG’s answer is to spend indexing-time compute building structure. An LLM extracts entities, relationships, and factual claims from 600-token chunks, aggregates the duplicates into a weighted knowledge graph, then partitions that graph with hierarchical Leiden community detection. Each community, at every level of the hierarchy, gets a pre-written LLM summary, with higher levels built from the summaries below them. At query time the community summaries are shuffled, chunked, mapped in parallel to partial answers with self-rated helpfulness scores, and reduced into one global answer.

The evaluation compared four community levels (Co root through C3 leaf) against direct map-reduce summarization of source text (TS) and vector RAG (SS), on two roughly one-million-token corpora: podcast transcripts (1,669 chunks, graph of 8,564 nodes and 20,691 edges) and news articles (3,197 chunks, 15,754 nodes). With 125 persona-generated sensemaking questions per dataset and GPT-4 as judge, every global condition beat vector RAG on comprehensiveness (72 to 83 percent win rates, $p < .001$) and diversity (62 to 82 percent), while vector RAG won the directness control criterion, as expected of shorter answers. A second experiment counting and clustering extracted factual claims (47,075 claims via Claimify) broadly confirmed the LLM-judge results, agreeing on 78 percent of non-tied comprehensiveness comparisons. The efficiency finding matters most: root-level Co answers used 97 percent fewer context tokens than TS (26,657 versus about a million for podcasts) while still beating vector RAG at 72 and 62 percent. The authors state plainly that all of this is established only for these two corpora, this token scale, and GPT-4.

For a personal knowledge backend over conversational history, the transferable idea is the shape of the index: entities and relationships distilled once, then summarized at nested levels of abstraction, so that broad questions read a few cheap high-level summaries instead of re-digesting raw text. That is essentially a mechanized version of a summarize-once, query-summaries-forever memory tree, and the Co token numbers quantify what the hierarchy buys. The paper also contributes evaluation machinery (persona-driven question generation, comprehensiveness and diversity criteria, claim-based validation) that applies directly to measuring any persistent memory system where no gold answers exist. Indexing cost is the caveat: 281 minutes of gpt-4-turbo calls for one million tokens, so incremental update strategy is left open.

Read from: pages 1-16 of 26 (full main text and references; appendices skipped)

Bitemporal History

Martin Fowler, [martinfowler.com](https://martinfowler.com/articles/bitemporal-history.html), April 2021, <https://martinfowler.com/articles/bitemporal-history.html>

This is a pattern essay, not an empirical study. Fowler works a single payroll example through the whole piece: on February 25 the company pays Sally her recorded \$6000 monthly salary; on March 15 HR reports that she was actually raised to \$6500 back on February 15; on April 5 HR corrects the raise to \$6400. He uses the example to motivate treating time as two dimensions. Actual history is what happened in the world given perfect information; record history is what the system believed at each moment. He prefers the terms actual and record over the Snodgrass valid and transaction terminology used in SQL:2011, on the grounds that workshop audiences found the latter confusing. The essay proceeds through tables and plots (actual time on one axis, record time on the other), then generalizes and closes with implementation options.

The concrete machinery is small but exact. A bitemporal query takes two time parameters, in his notation `sally.salaryAt('2021-02-25', '2021-03-25')`, read as what we believed on March 25 her February 25 salary was. Horizontal slices of the two-axis plot give the actual history as known at some record time; vertical slices give how belief about one date evolved. The key structural claim is that while actual history stops being append-only once retroactive changes are allowed, record history stays append-only: corrections are layered on rather than overwritten, which yields a reliable audit trail of the modifications themselves. He is explicit about scope limits: bitemporality complicates a system significantly, is avoidable when retroactive change is forbidden or when actions capture their inputs (the check records the salary it used), and by itself cannot compute the downstream consequences of a correction, such as back pay or tax effects. For storage he sketches two schemes, a relational table with paired record and actual date ranges, and event sourcing where each event carries both an actual and a record timestamp and simpler read models are projected from the log.

For a knowledge backend built over conversational history, the bearing is direct. Facts extracted from conversations get corrected later, and the record dimension is what lets the system answer both what is true and what was believed when an action was taken, without destroying the audit trail. Fowler's generalization of record time into perspectives, where HR and Payroll hold different record histories of the same fact and a perspective can even be a counterfactual scenario, maps onto multi-agent or multi-user ingestion at organizational scale, though he cautions that stacking many perspective dimensions is rarely worth the complexity. This is design guidance from one practitioner's example, with no benchmarks or formal treatment, so it constrains schema choices rather than settling them.

Read from: full text

ENPMR-Bench: Benchmarking Proactive Memory Retrieval for Emotional Support Agents

Xing Fu et al., arXiv:2605.27240, 2026-05

Abstract (verbatim)

Memory-augmented language agents are increasingly deployed in affective applications such as emotional support, where understanding and responding to users' latent emotional needs is critical. However, existing research often treats memory as a tool for factual retrieval, overlooking its role in shaping users' emotional experiences. In this work, we introduce ENPMR- Bench, a benchmark for evaluating Emotional Need-aware Proactive Memory Retrieval (ENPMR), a core capability that enables agents to infer users' latent emotional needs and proactively retrieve appropriate memories to support empathetic interaction. Grounded in Maslow's hierarchy of needs, ENPMR-Bench includes over 1,800 memory-augmented dialogues and defines structured mappings between emotional needs and supportive memory types. Experimental results demonstrate that current retrieval paradigms, including both embedding-based and LLM-driven approaches, exhibit substantial deficiencies, with empathy scores significantly lagging behind golden memory conditions. While chain-of-thought prompting improves the alignment between inferred emotional needs and retrieved memories to some extent, a notable performance gap remains. Together, these findings reveal critical limitations in current agents and outline directions for advancing personalized emotional support through need-sensitive memory retrieval.

Bearing on this project

A benchmark for proactive memory retrieval. Narrower domain (emotional support), but its existence alone falsifies the claim that no benchmark can express the proactive-recall failure - a sentence that appears in IT494_PLAN.md.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Grammar-Constrained Decoding for Structured NLP Tasks without Finetuning

Saibo Geng, Martin Josifoski, Maxime Peyrard, Robert West; EMNLP 2023; DOI: 10.18653/v1/2023.emnlp-main.674

Geng and colleagues argue that grammar-constrained decoding (GCD) is not a niche trick for parsers and code generators but a unified framework for structured NLP. The idea is to describe a task's entire output space as a formal grammar, then run an incremental parser (Grammatical Framework's, in their implementation) as a completion engine during decoding: at each step the parser reports which tokens can legally extend the partial output, and the model's probability distribution is pruned to that set. They write character-level grammars, compile them into tokenizer-specific token-level grammars, and show fourteen tasks expressible this way, from information extraction to coreference to extractive summarization. Their main extension is the input-dependent grammar (IDG), where the legal output space is rebuilt per input; thirteen of the fourteen tasks need one. The method works with any autoregressive model that exposes its vocabulary distribution, which excludes logit-hiding APIs.

The experiments use few-shot LLaMA and Vicuna models (7B to 33B) with no finetuning. On closed information extraction over SynthIE-text, constrained to 2.7 million Wikidata entities and 888 relations, unconstrained LLaMA-33B scores 17.5 F1 while the constrained version reaches 36.0, beating the finetuned GenIE T5-base at 34.8. On entity disambiguation across six benchmarks, constrained LLaMA-33B averages 80.3 accuracy against 54.1 unconstrained, and the input-dependent grammar beats the input-independent one by twelve points; the finetuned state of the art still leads at 89.4. Constituency parsing shows the limit: IDG guarantees 100 percent valid trees but LLaMA-33B tops out at 54.6 bracketing F1 against 95-plus for supervised parsers, since the task needs syntactic understanding no constraint can supply. GCD adds negligible per-token latency for two tasks and roughly 69 ms on cIE, and it induces a likelihood misalignment where the empty string becomes the top beam, fixable by length normalization.

For the memory harness, this is direct prior art for the extraction stages: pulling entity records and dated fact rows out of a corpus is exactly the extraction-shaped, restricted-vocabulary work the paper tests, and it grounds the claim that a cheap constrained model can approach a stronger or finetuned one there. It governs the schema-validated LLM call layer beneath the swappable model tiers: when a local model serves tree-node, entity, or fact extraction, decode-time constraints make schema-invalid output structurally impossible rather than merely counted, which cleans up the stage-wise sensitivity comparison across tiers. Its scope limits transfer too: constraints repair format and vocabulary, not comprehension, so stages needing genuine reasoning, and the LLM-judge scoring of the three injection arms, should not expect constraint-driven gains, and API-tier models cannot use the mechanism at all.

Read from: pages 1-9 of the 21-page PDF (full main text, introduction through conclusion and best practices), plus citation metadata from gap_survey_entries.json.

Examining the State-of-the-Art in News Timeline Summarization

Demian Gholipour Ghalandari et al., arXiv:2005.10107, 2020-05

Abstract (verbatim)

Previous work on automatic news timeline summarization (TLS) leaves an unclear picture about how this task can generally be approached and how well it is currently solved. This is mostly due to the focus on individual subtasks, such as date selection and date summarization, and to the previous lack of appropriate evaluation metrics for the full TLS task. In this paper, we compare different TLS strategies using appropriate evaluation frameworks, and propose a simple and effective combination of methods that improves over the state-of-the-art on all tested benchmarks. For a more robust evaluation, we also present a new TLS dataset, which is larger and spans longer time periods than previous datasets.

Bearing on this project

The entity axis has a field and this is its survey. Timeline summarization is the task of producing a dated chronological summary for an entity from a corpus - which is what this project calls an entity's cell sequence, in the news domain. Read this before describing the entity axis as new.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Intent Induction from Conversations for Task-Oriented Dialogue Track at DSTC 11

Gung et al. 2023, Intent Induction from Conversations (DSTC 11), SIGDIAL 2023, arXiv:2304.12982

This paper, from AWS AI Labs, describes a shared-task benchmark rather than a single method: the DSTC 11 challenge track on inducing customer intents from raw conversation transcripts. The authors argue that prior intent-mining work evaluated on repurposed classification datasets (BANK77, CLINC150, SNIPS) that lack full dialogues and realistic intent skew. They release three simulated call-center corpora: Insurance (948 conversations, 22 intents, used for development) and Banking (1,000 conversations, 29 intents) and Finance (2,000 conversations, 39 intents) for evaluation, averaging 59 to 71 turns per conversation, three to five times longer than MultiWOZ or SGD. Two subtasks are defined. Task 1 is intent clustering: assign a cluster label to each pre-tagged intentful turn, with the number of reference intents withheld (only a 5 to 50 range given). Task 2 is open intent induction: no turn labels, only noisy automatic InformIntent predictions; systems must produce a labeled training set for an intent classifier, and evaluation trains a logistic regression classifier on frozen all-mpnet-base-v2 embeddings and scores its predictions on an independent test set. Clustering accuracy via Hungarian matching is the primary metric.

Thirty-four teams entered Task 1 and 19 entered Task 2, all using clustering-based pipelines. The winner, T23, reached 69.8 ACC on Task 1 and 76.3 on Task 2, against baselines (k-means over sentence embeddings, with silhouette-based selection of k) at 55.8 and 63.6. Deep embedded clustering variants, particularly SCCL, dominated the top ten; nine of ten top systems used them, and DSE-RoBERTa-large plus supervised pre-training on public intent data marked the winner. Two analyses qualify the numbers. HDBSCAN systems that left noise points unassigned were punished by clustering metrics: propagating labels to T11's noise instances moved it from 33rd to 9th, exposing a weakness of clustering-style evaluation itself. And rerunning Task 2 scoring under nine different classifier encoders left the top ten ranks stable but shuffled ranks 11 through 15, so prediction-based evaluation carries measurable noise. All results are scoped to simulated (not genuinely live) English conversations in three financial-adjacent domains.

For a personal knowledge backend over conversational history, the relevant lesson is methodological. Task 2's framing, judge an induced schema by the downstream system it trains rather than by cluster purity on the source turns, is directly transferable to evaluating induced topic or era structures over an archive: a smaller, cleaner distillation can beat exhaustive coverage. The corpus statistics also matter at organizational scale: real request distributions are long-tailed, so any memory taxonomy induced from usage will concentrate on a few head topics, and the granularity trade-off the metrics here are built around (too many fine categories versus too few coarse ones) is the same one a schema over conversation history must manage.

Read from: pages 1-9 of 18 (full main text and results; remainder is references and appendices)

TextTiling: Segmenting Text into Multi-paragraph Subtopic Passages

Marti A. Hearst, Computational Linguistics 23(1), 1997, ACL Anthology J97-1003

TextTiling divides expository text into contiguous, nonoverlapping multi-paragraph segments, each covering one subtopic under the document's main topic, using nothing but patterns of word repetition. The algorithm runs in three stages. Tokenization lowercases the text, strips 898 stop words, reduces tokens to morphological roots, and regroups them into fixed-size pseudosentences of w tokens (default 20), because real sentences vary too much in length for fair comparison. Scoring assigns a lexical score to every pseudosentence gap; the strongest variant, block comparison, computes normalized inner products of raw term-frequency vectors over adjacent k -unit blocks (default $k = 10$) slid one unit at a time. Boundary identification smooths the score curve, computes a depth score at each valley (the summed drop from the flanking peaks), and marks boundaries where the depth exceeds a cutoff derived from the mean and standard deviation of the document's own depth scores, snapped to the nearest paragraph break.

The evaluation is small but careful. Seven readers segmented 12 magazine articles of 1,800 to 2,500 words; interjudge kappa averaged .647, at the low end of acceptability, which caps how precisely any algorithm can be scored. Against a three-of-seven majority gold standard, the blocks version reached precision .66 with recall .78 at the liberal cutoff, beating a 39 percent random baseline (.44/.42) but trailing the judges themselves (.81/.71). An earlier tf.idf weighting hurt rather than helped, as did thesaural relations; plain within-block frequency proved more dependable. On 44 concatenated Wall Street Journal articles the blocks version hit a .67 precision/recall break-even for article boundaries, though Hearst notes this task violates the algorithm's single-document assumption. Known failure modes include summary paragraphs that echo earlier vocabulary, quoted speech, and topic shifts carried by stop-listed pronouns.

For the memory harness, TextTiling governs the corpus-to-tree segmentation stage, the first cut before entity extraction, dated-fact capture, and the schema-validated LLM calls that build the tree. It is the natural deterministic arm to run beside the model-tier segmenters: fully reproducible, free per document, and parameterized by exactly four knobs, so any LLM-based chunker (and the downstream three-arm injection comparison scored by an LLM judge) must beat this baseline to justify its cost. Its named weakness, cohesion troughs missing topic shifts that share vocabulary, predicts where an LLM arm should win: narrative or dialogue-heavy sources rather than the expository prose the algorithm was designed and validated on.

Read from: full text of all 32 pages of the Computational Linguistics article (introduction, applications, discourse model, related methods, algorithm, evaluation, and the summary and future work section).

Reifying RDF: What Works Well With Wikidata?

Hernandez, Hogan, Krotzsch; SSWS workshop at ISWC 2015

The paper asks how best to represent Wikidata, whose statements carry qualifiers and references, in plain RDF so that off-the-shelf SPARQL engines can query it. Since a Wikidata statement annotates an entire subject-predicate-object relation, the authors first reduce the problem to encoding quads (s, p, o, i), where i is a statement identifier that functionally determines the triple; qualifiers then attach to i as ordinary triples. They compare four established encodings: standard reification, which spends three triples per quad on r:subject, r:predicate, and r:object; n-ary relations, the Exleben et al. scheme Wikidata itself was using, which needs $2(n + m)$ triples for n quads over m distinct properties; singleton properties, which mint a fresh predicate per statement and need only $2n$ triples; and named graphs, which store the quad directly. They build all four datasets from a February 2015 Wikidata dump (57.1 million quads plus 237.6 million triples common to every format, so between 57 and 171 million model-specific tuples) and run 14 benchmark queries, expanded into model-specific versions, against 4store, BlazeGraph, GraphDB, Jena TDB, and Virtuoso.

The empirical picture is uneven in an instructive way. Singleton properties, the most concise encoding, broke four of the five engines: 4store and GraphDB could not index the millions of unique predicates, and Jena timed out on every query. 4store also failed to load the named graphs dataset. Only n-ary relations and named graphs could express the SPARQL property paths two queries required. Virtuoso handled all four models reliably; across all engines, named graphs was fastest in 17 of 70 engine-query cases, with standard reification and n-ary relations at 16 each, so no clear winner emerged among those three. The authors note the caveats themselves: cold-cache runs only, a 60-second timeout, one hardware setup, and a preliminary scope. They also observe that 108 two-triple encodings are possible in principle; only the well-known four were tested.

For a personal knowledge backend over conversational history, the load-bearing problem here is the same one: facts need provenance, temporal scope, and confidence attached to them, and the choice of how to reify those annotations determines which stores and query features remain usable. The concrete lessons transfer directly. Elegant-on-paper schemes can violate engine assumptions (predicate cardinality being the killer for singleton properties), quads with a statement identifier are a clean conceptual core, and the identifier join is highly selective, so the extra joins that reification adds cost less than expected but do complicate query planning. The paper is also a compact tour of how a real large knowledge graph, Wikidata at 65 million statements in 2015, actually structures qualified claims.

Read from: full text

Knowledge Graphs

Hogan et al., ACM Computing Surveys 54(4), 2021, arXiv:2003.02320

This is a tutorial survey by seventeen authors, grown out of a Dagstuhl seminar, that gives the field of knowledge graphs its first unified introduction. The authors adopt a deliberately inclusive definition: a knowledge graph is a graph of data meant to accumulate and convey knowledge of the real world, with nodes as entities and edges as relations. From there the paper walks through the full stack: graph data models (directed edge-labelled graphs such as RDF, property graphs as in Neo4j, heterogeneous graphs, and named-graph datasets), query languages (SPARQL, Cypher, Gremlin, G-CORE) and their shared primitives from basic graph patterns up through regular path queries, then schema, identity, and context, deductive reasoning with ontologies and rules, inductive methods, and the lifecycle of creation, quality assessment, refinement, and publication. A running example, a Chilean tourism knowledge graph, carries the formalism.

The paper's substance is organizational rather than experimental, but it fixes useful distinctions and facts. It separates three kinds of schema: semantic (RDFS and OWL, which license inferences under the open-world assumption), validating (ShEx and SHACL shapes, which check completeness of specified portions of the graph), and emergent (quotient graphs computed from latent structure). On identity it treats IRIs, owl:sameAs links, datatypes, and blank nodes as the machinery for saying when two nodes are the same thing. Its practice section quantifies the landscape as of 2021: Freebase held over three billion edges when it went read-only in March 2015, largely migrated to Wikidata; Cyc represents over 900 person-years of manual effort yielding 2.2 million facts and rules; NELL extracted 120 million edges from web text. Drawing on Noy et al., it lists six traits enterprise graphs (Google, Amazon, Facebook, LinkedIn, Bloomberg) share, including billion-edge scale, continuous refinement, and deliberately lightweight taxonomic ontologies rather than heavy formal ones.

For a personal knowledge backend over conversational history, the load-bearing ideas are the ones the paper opens with: graphs let you postpone schema definition while data and scope evolve, which is exactly the position of a system ingesting unpredictable conversations; the identity machinery answers when two mentions are one entity; and the context section (temporal validity, provenance, geographic scope) is the vocabulary for statements true only at a time or per a source. The framing of a knowledge graph as an ever-evolving shared substrate of knowledge within an organisation or community, plus the enterprise-scale trait list, is a direct sketch of what a personal store scaled to an organization would have to become, and the note that industry favors lightweight taxonomies over expressive ontologies is a useful caution against over-engineering the schema.

Read from: pages 1-20 and 74-78 of 135

SCOPE and SCION: A Benchmark and an Auditable Reference Pipeline for Schema Induction and Fusion from Text

Miaobo Hu et al., arXiv:2607.21610, 2026-05

Abstract (verbatim)

Schema graphs are an upstream bottleneck of schema-grounded information extraction and knowledge graph construction, yet most extraction systems assume the schema is already available. We introduce SCOPE (Schema Construction and Ontology-induction Pipeline Evaluation), a train-text- only benchmark for corpus-to-schema induction and optional schema fusion from raw text, built from 24 public information extraction sources (15 RE and 9 EE) normalized into evaluation-only gold schema graphs; its core event-extraction target covers event types and within-event argument roles, with inter-event links reported separately. We present SCION (Schema Construction and Induction with Ontology Normalization), an auditable reference pipeline rather than a new extraction architecture; it constructs candidate spaces from train text and restricts naming, merging, filtering, validation, and conservative fusion to candidate-linked evidence under strict JSON contracts. On the SCOPE core suite, SCION-lite attains the highest F1 among released source-schema references, Text2Onto-style, LLM-only, and matched extract-then-aggregate baselines under Literal, Fuzzy, Continuous, and Graph schema-graph metrics, while the compact open-model SCION-RL variant reduces reliance on proprietary LLM schema engineers. These results are reported against normalized typed-edge targets rather than as claims that induced schemas surpass human ontology design; the release includes evidence-linked outputs, parse/fallback logs, candidate retention/merging logs, run manifests, code, and benchmark packages at https://github.com/wandugu/paper_scion.

Bearing on this project

A benchmark AND an auditable reference pipeline for corpus-to-schema induction. So schema induction is not only solved but measured, with released baselines to beat.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Adaptive-RAG: Learning to Adapt Retrieval-Augmented Large Language Models through Question Complexity

Jeong et al., NAACL 2024, arXiv:2403.14403

Adaptive-RAG routes each incoming question to one of three answering strategies based on a predicted complexity level: no retrieval at all (the LLM answers from parametric knowledge), a single retrieve-then-read step, or an iterative multi-step loop that interleaves chain-of-thought reasoning with repeated retrieval. The router is a small T5-Large classifier (770M parameters) that maps a query to label A, B, or C before any answering begins, so the LLM and retriever themselves never change. Since no query-complexity dataset exists, the authors build training labels automatically two ways: run all three strategies on sampled queries and assign the label of the cheapest strategy that got the answer right, then fill remaining gaps using dataset provenance (single-hop benchmarks get B, multi-hop benchmarks get C). They train on about 400 queries per dataset and evaluate on a mixed pool of three single-hop sets (SQuAD, Natural Questions, TriviaQA) and three multi-hop sets (MuSiQue, HotpotQA, 2WikiMultiHopQA), 500 test queries each, with BM25 as the retriever throughout.

The headline result, on GPT-3.5-Turbo-Instruct averaged over all six datasets, is 50.91 F1 at 1.03 retrieval-and-generate steps per query, essentially matching the always-iterate multi-step baseline (50.87 F1) at roughly a third of its 2.81 steps, and beating binary adaptive baselines like Mallen et al.'s Adaptive Retrieval (48.20) and Self-RAG (22.98, though trained on a different base model). The pattern holds for FLAN-T5-XL (3B) and XXL (11B). The efficiency stakes are concrete: 0.35 seconds per no-retrieval query versus 27.18 for multi-step. Two caveats travel with these numbers. The classifier is only 54.52 percent accurate over the three labels, and an oracle router reaches 62.80 F1 at half a step per query, so routing quality, not the strategies, is the current ceiling. Classifier size barely matters; a 60M model performs within about one F1 point of the 770M one.

For a personal knowledge backend over conversational history, the transferable idea is that retrieval depth should be a per-query decision made before retrieval starts, by a component far cheaper than the answering model. Most lookups against a memory tree are single-hop ("what did we decide about X"), some need iterative traversal across sessions, and many need no lookup at all; paying multi-hop cost uniformly is exactly the waste this paper quantifies. The self-labeling trick also matters: usage logs can label their own routing data by recording which strategy sufficed, no annotation budget required. The scope limit is real, though: results cover factoid QA over Wikipedia-style corpora with BM25, not dense retrieval over informal conversational text.

Read from: pages 1-11 of 15

LongLLMLingua: Accelerating and Enhancing LLMs in Long Context Scenarios via Prompt Compression

Huiqiang Jiang et al., ACL 2024 (arXiv Oct 2023), arXiv:2310.06839

LongLLMLingua is a question-aware prompt compression method for long-context use of black-box LLMs. It starts from the observation that long prompts create three coupled problems: cost, degradation from irrelevant content, and position bias (the lost-in-the-middle effect). Building on LLMingua, which prunes low-perplexity tokens using a small language model, the authors add question awareness at two granularities. Coarse-grained compression scores each retrieved document by the perplexity of the question conditioned on that document, with a restrictive statement appended to curb small-model hallucination, and keeps only high-scoring documents. Fine-grained compression then scores individual tokens by contrastive perplexity, the shift in a token's perplexity when the question is added to its conditioning context, which the appendix shows is equivalent to conditional pointwise mutual information. Two further pieces complete the pipeline: the surviving documents are reordered by their coarse importance scores so the most relevant land at the edges of the prompt, and a subsequence recovery step restores entity spans (names, dates, places) that token-level compression mangled, by matching response substrings back to the original prompt.

The numbers are strong within a specific regime. On NaturalQuestions with 20 documents and GPT-3.5-Turbo, LongLLMLingua beats the original uncompressed prompt by up to 21.4 points at roughly 4x fewer tokens, and its coarse scoring metric alone out-recalls BM25, OpenAI embeddings, SBERT, and reranker baselines. On LongBench at a 2,000-token budget it scores 48.3 versus 44.0 for the full 10k-token prompt, and end-to-end latency improves 1.4x to 2.6x at 2x to 6x compression. The ablations show every component earns its place; removing question awareness collapses performance to near-LLMLingua levels. The limitations are stated plainly: compression is question-specific, so the context cannot be cached across queries, the method costs about twice the compute of LLMingua, and the coarse question-aware scoring may miss subtler context-question relationships than the tested QA benchmarks exercise.

For the harness, this paper governs the budgeted injection stage, the point where scored memory rows from the tree (entities, dated facts) are packed into a fixed token window before the schema-validated call goes out through a model tier. It supplies two cited design choices the injection arms can adopt or test: select what enters the budget by question-conditioned relevance rather than raw entropy, and order the survivors by score with the best at the edges. It is also the natural compression baseline for the model-sensitivity comparison across tiers, with the caching caveat carried alongside: per-question compression trades reuse for density, which matters if the same tree serves many questions.

Read from: full paper text extracted via pdfplumber, pages 1 through 11 of 20 (complete main text, conclusion, and limitations; remaining pages are references and appendices).

Anatomy of Agentic Memory: Taxonomy and Empirical Analysis of Evaluation and System Limitations

Dongming Jiang et al., arXiv:2602.19320, 2026-02

Abstract (verbatim)

Agentic memory systems enable large language model (LLM) agents to maintain state across long interactions, supporting long-horizon reasoning and personalization beyond fixed context windows. Despite rapid architectural development, the empirical foundations of these systems remain fragile: existing benchmarks are often underscaled, evaluation metrics are misaligned with semantic utility, performance varies significantly across backbone models, and system-level costs are frequently overlooked. This survey presents a structured analysis of agentic memory from both architectural and system perspectives. We first introduce a concise taxonomy of MAG systems based on four memory structures. Then, we analyze key pain points limiting current systems, including benchmark saturation effects, metric validity and judge sensitivity, backbone-dependent accuracy, and the latency and throughput overhead introduced by memory maintenance. By connecting the memory structure to empirical limitations, this survey clarifies why current agentic memory systems often underperform their theoretical promise and outlines directions for more reliable evaluation and scalable system design.

Bearing on this project

A survey that names, as known limitations, two things this project was going to contribute: backbone-dependent accuracy (the stage-wise tier sensitivity) and overlooked system-level cost. Naming is not measuring, so per-stage measurement may still be open, but the gap can no longer be presented as unnoticed. Also useful in the other direction: it says the field's empirical foundations are fragile - underscaled benchmarks, metrics misaligned with semantic utility - which is the argument for careful measurement being worth more than another architecture.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Evaluating Open-Domain Question Answering in the Era of Large Language Models

Ehsan Kamalloo, Nouha Dziri, Charles L. A. Clarke, Davood Rafiei; ACL 2023 (long paper, 2023.acl-long.307); arXiv:2305.06984

This paper asks whether lexical matching, still the default scoring method for open-domain QA, tells the truth about model quality once answers come from generative systems. The authors take 12 open-domain QA models, spanning retriever-reader pipelines (DPR, several FiD variants), end-to-end systems (EMDR2, EviGen), and closed-book InstructGPT in zero-shot and few-shot modes, and have their answers to a random 301-question subset of NQ-OPEN (1490 question/answer pairs) judged four ways: exact match, the supervised semantic matcher BEM, an InstructGPT-based judge prompted with question, gold answer, and candidate, and human annotation with a third-rater tiebreak (Fleiss kappa 72.8 percent). A parallel study on CuratedTREC 2002 adds regex-specified gold answers and 20-year-old TREC systems as baselines.

The headline finding is severe underestimation. InstructGPT zero-shot scores 12.6 percent under exact match but 71.4 percent under human judgment, a gap of nearly 60 points, and the few-shot variant reaches 75.8 percent, a new state of the art on the subset. Rankings shift, not just magnitudes: EMDR2 loses its top spot only under human evaluation. Semantic equivalence accounts for 50.3 percent of exact-match failures, with list-style questions (20.6 percent) and granularity mismatches (15.0 percent) next. Model judges correlate far better with humans than lexical metrics (Kendall tau 0.75 for the LLM judge and 0.70 for BEM, against 0.23 for EM), yet they fail exactly where it matters most for generative output: of 47 hallucinated long-form InstructGPT answers, the LLM judge wrongly accepted 30, BEM 18, and GPT-4 as judge still 9. Scope is explicitly limited to factoid questions with short gold answers; multi-hop, discrete, and causal reasoning are untested, and the human annotators were the authors themselves.

For the memory harness this paper governs the LLM-judge scoring stage that scores the three injection arms against the fixed gold answers of the certified-unknown novel's question set. It justifies the design directly: put the gold answer in the judge prompt rather than asking for open-ended grading, expect judge scores to run a few points below human truth, and treat any arm that emits long, elaborated answers as a known blind spot, since attribution failures slip past every automated judge tested. It also argues for regex or alias expansion of gold answers as a cheap first-pass filter before the model judge runs, since most mismatches are shallow syntactic variation that patterns catch.

Read from: full extracted text of the 16-page ACL PDF, pages 1 through 11 (body, conclusion, limitations); remaining pages are references and appendix.

Event Log Construction from Customer Service Conversations Using Natural Language Inference

Kecht, Egger, Kratsch, and Roglinger; ICPM 2021; doi:10.1109/ICPM53251.2021

The paper builds standardized process-mining event logs out of raw customer service conversations, using natural language inference instead of a trained classifier. Each customer inquiry or agent reply becomes the premise of an NLI pair; the hypothesis is a short sentence naming a candidate topic or activity, such as “This example asks for iOS version.” Facebook’s pre-trained BART model, called through the transformers library, returns an entailment probability for every message-hypothesis pair. The workflow has five steps: define domain topics and activities, hand-label a small sample (100 to 300 messages) with binary tags, run five-fold stratified cross-validation to pick a per-label decision threshold from candidates between 0.70 and 0.999 by maximizing the Matthews correlation coefficient, optionally rewrite hypotheses that score poorly, then classify the full corpus and export it as an IEEE XES log via PM4Py. No model training happens anywhere; this is zero-shot classification in the sense of Yin et al. 2019.

Evaluation uses English Twitter support conversations from AmazonHelp (288,828 tweets), AppleSupport (231,683), and SpotifyCares (88,774) after filtering, with a keyword matcher as baseline. With tuned hypotheses, 55 of 56 topic and activity classifications reached an MCC above 0.72 on the labeled samples; on 100 AppleSupport outbound tweets, four of seven activities scored a perfect 1.00 across all metrics. Hypothesis wording mattered a great deal: “This example is about prime” hit 0.91 MCC where the default phrasing managed 0.53, and one grammatically broken default hypothesis scored 0 where the fixed version scored 1. Keywords sometimes sufficed on their own (URL detection was perfect because every Twitter link starts with the same prefix). Three discovery algorithms (Alpha, Inductive, Heuristics Miner) recovered the same coherent process model from the resulting log, so the logs are usable, though the authors concede the mined model covers only the topics they thought to define and analysis stopped at each conversation’s first message.

For a personal knowledge backend over conversational history, the transferable finding is the labeling economics: roughly 300 hand-labeled examples plus threshold tuning were enough to structure 200,000-plus messages, because the pre-trained model carries the semantics and only the decision boundary needs calibration. That pattern (define a schema of event types, verify on a tiny sample, sweep the archive) maps directly onto turning a conversation archive into a queryable timeline of typed events, and adding a new event type never degrades existing ones since each label is inferred independently. The caveats transfer too: results come from short English tweets in three consumer-support domains, hypothesis phrasing is brittle and opaque, and the schema must be known in advance, which is exactly the part a memory system over open-ended conversations cannot assume.

Read from: full text

FABLES: Evaluating faithfulness and content selection in book-length summarization

Kim et al., COLM 2024, arXiv:2404.01261

The paper reports the first large-scale human evaluation of faithfulness in summaries of book-length documents (mean 121K tokens). The authors sidestep two chronic problems at once: data contamination, by using 26 fictional books published in 2023 or 2024, and annotation cost, by hiring 14 Upwork annotators who had already read the books for pleasure. Five models (GPT-3.5-Turbo, GPT-4, GPT-4-Turbo, Mixtral, Claude-3-Opus) each summarize every book via hierarchical merging; GPT-4 decomposes each summary into atomic, decontextualized claims; annotators label each claim faithful, unfaithful, partially supported, or unverifiable, with quoted evidence. The result is 3,158 annotated claims across 130 summaries, collected for \$5.2K, about \$40 per summary.

Three findings stand out. First, Claude-3-Opus was the most faithful summarizer at 90.9% faithful claims, ahead of GPT-4 and GPT-4-Turbo (about 78%), GPT-3.5-Turbo (72%), and Mixtral (69%); on this corpus of 26 recent novels, only five of 130 summaries were entirely accurate. Second, unfaithful claims cluster around character or relationship states (38.6%) and events (31.5%), and half require indirect, multi-hop reasoning over the narrative to refute, which distinguishes the task from entity-centric fact checking. Third, automatic verification largely fails: on a seven-book subset (723 claims), the best auto-rater, Claude-3-Opus given the entire book as context, reached only 58.2 F1 on detecting unfaithful claims, and it beat both BM25 retrieval (24.9 F1) and human-annotated evidence spans. Beyond faithfulness, coding of the summary-level comments yielded a taxonomy of omission errors: 33% to 65% of summaries omit key events, every model makes chronology errors, and the long-context models systematically over-weight book endings (at least 20% of Claude-3-Opus summaries).

For a personal knowledge backend built over conversational history, the negative results matter most. Retrieval-then-verify, the backbone of most memory and RAG validation schemes, collapses when the query is a claim about states, relationships, or arcs distributed across a long record rather than a localizable fact; even full-context reading by the strongest model of early 2024 could not reliably catch fabrications. A memory system that summarizes long histories into durable records will inherit exactly these failure modes: silent omission of key events, recency bias toward the end of the record, and confident state claims that no single passage can confirm or deny. The paper argues, credibly, that claim-level verification over long documents is itself a benchmark for long-context understanding, and its human-annotation protocol (readers who already know the source) is a workable template for auditing any summarization layer in a memory pipeline. Scope caveats: fiction only, English, 26 books, February 2024 model checkpoints.

Read from: pages 1-14 of 41 (full main text and conclusion; appendices skipped)

The NarrativeQA Reading Comprehension Challenge

Tomáš Kočiský, Jonathan Schwarz, Phil Blunsom, Chris Dyer, Karl Moritz Hermann, Gábor Melis, Edward Grefenstette; TACL 2018 (arXiv Dec 2017); arXiv:1712.07040

NarrativeQA is a reading comprehension benchmark built from 1,567 full stories, split evenly between Project Gutenberg books and scraped movie scripts, each matched to a human-verified Wikipedia plot summary. The central design move is annotation from the summary alone: Mechanical Turk workers wrote about 30 question and answer pairs per story while never seeing the full text, with copying blocked by instruction and by JavaScript limits on the annotation site. Because the questions were written against an abstractive retelling, answering them from the full story cannot reduce to span lookup; the reader must locate and integrate narrative information about entities, events, and their relations, sometimes across passages separated by hundreds of pages. The paper opens with a survey showing why earlier datasets (CNN/Daily Mail, CBT, SQuAD, NewsQA, MS MARCO, SearchQA) reward shallow salience or local span extraction instead.

The dataset holds 46,765 question and answer pairs; questions average 9.8 tokens and answers 4.73. Stories average around 62,000 tokens with a maximum near 430,000, against summaries averaging about 650. Only 29.57 percent of answers appear verbatim in the stories, versus 44.05 percent in the summaries. The experiments make the gap concrete: on the summaries task a simplified BiDAF span predictor reaches 36.30 test Rouge-L, within reach of the 57.02 human score, but on full stories the same model collapses to 6.22 while a retrieve-then-read pipeline (200-token chunks selected by question similarity, then AS Reader) tops out around 14, barely above a no-context Seq2Seq baseline. Retrieval itself is the bottleneck: the authors found chunk selection from a short question hard even for humans unfamiliar with the story. Scope limits apply in the same breath: the story count is modest, all evaluation is n-gram overlap or MRR rather than semantic judgment, human scores were measured on summaries only, and the 2017 baselines predate transformers, so the absolute numbers are floors, not ceilings.

For the Python memory harness this is the charter document for the question set and the injection experiment. It supplies the cited method for the certified-unknown novel: write questions from a summary-level view so answers demand narrative memory, then run the bare-model spot check that must fail without support, exactly the failure NarrativeQA demonstrates at book scale. The corpus-to-tree stage with entities and dated facts is the harness answer to the retrieval bottleneck the paper names, the three injection arms replay its no-context, retrieved-chunk, and summary-context contrast under schema-validated LLM calls across swappable model tiers, and the LLM-judge scoring replaces the overlap metrics the paper itself strains against with short abstractive answers.

Read from: all 12 pages of the PDF, full text extracted with pdfplumber.

Text Segmentation as a Supervised Learning Task

Omri Koshorek, Adir Cohen, Noam Mor, Michael Rotman, Jonathan Berant; NAACL 2018; arXiv:1803.09337

This short paper reframes text segmentation, the job of cutting a document into topically coherent spans, as supervised learning rather than the unsupervised clustering and graph methods that dominated prior work. The blocker had always been data: existing benchmarks were tiny (Choi’s 920 synthetic concatenations, five Manifesto documents, a hundred-odd Wikipedia pages), and the synthetic ones do not resemble natural prose. The authors mine English Wikipedia tables of contents to build Wiki-727K, a corpus of 727,746 documents with hierarchical segment boundaries, split 80/10/10, then train a two-level model: a bidirectional LSTM builds each sentence embedding by max-pooling over word2vec inputs, and a second bidirectional LSTM reads the sentence sequence and emits a per-sentence probability that it ends a segment. Inference is greedy, cutting wherever the probability exceeds a threshold tuned on the validation set.

The results establish that natural data separates real methods from artifact-fitters. Under the Pk metric (Beeferman et al.’s sliding-window boundary error, with k set to half the mean segment length), the model scores 22.13 on the full Wiki-727K test set and 18.24 on the Wiki-50 subsample, against a random baseline near 53 and human performance of 14.97 on Wiki-50. GraphSeg, which beats everyone on the synthetic Choi data (5.6 to 7.2), collapses to 63.56 on Wiki-50, worse than random; word-similarity cues that distinguish concatenated unrelated documents fail on topic shifts inside one document. Scope limits ride along with the wins: the model is trained on top-level Wikipedia sections only, loses to Chen et al. on the chemistry Elements set (33.82 vs 20.1 word-level Pk, blamed on GoogleNews embeddings), and the human ceiling itself is soft since annotators saw few documents. Runtime is linear, 1.6 seconds per document on CPU against GraphSeg’s 23.6.

For the memory harness this paper governs the corpus-to-tree stage, the very first cut that decides what a leaf node even is before entity extraction, dated-fact capture, or any schema-validated LLM call sees the text. Its contribution to that stage is less the LSTM than the evaluation contract: Pk against ground-truth boundaries gives the chunking step a real metric instead of eyeballed chunk quality, so an LLM-based segmenter routed through the swappable model tiers can be scored the same way as the cheap TextTiling arm, and segmentation error can be held apart from downstream injection-arm and LLM-judge effects. The GraphSeg collapse is also the standing warning that any segmenter must be validated on natural transcripts, not synthetic splices.

Read from: all 5 pages of the PDF (full paper, pypdf text extraction).

PDX: A Data Layout for Vector Similarity Search

Kuffo, Krippner, Boncz 2025, SIGMOD 2025, arXiv:2503.04422

The paper proposes PDX (Partition Dimensions Across), a storage layout for embedding collections that groups vectors into blocks and, within each block, stores each dimension's values contiguously, the way Parquet rowgroups hold columns. On top of it sits PDXsearch, a search framework that computes distances dimension-by-dimension across 64 vectors at a time in tight scalar loops that compilers auto-vectorize. The design targets a specific weakness of the standard vector-by-vector layout: pruning algorithms such as ADSampling and BSA, which stop reading a vector's dimensions once a partial distance shows it cannot reach the top-k, interleave bounds checks with distance math and lose their advantage against plain SIMD linear scans. PDXsearch scans an IVF bucket in three phases (a full scan of the first block to set a threshold, a warmup that fetches dimensions in exponentially growing steps, and a prune phase that skips discarded vectors once fewer than about 20 percent remain), keeping the code branchless.

On four CPUs (Intel Sapphire Rapids, AMD Zen4 and Zen3, Graviton4) and ten datasets from 16 to 1536 dimensions, the auto-vectorized PDX distance kernels averaged 2.0x over hand-written SIMD kernels on the horizontal layout, and 5.5x when 8 or fewer dimensions are read, the case pruning creates. In IVF index search on Zen4 at the highest recall, ADSampling on PDX ran 4.6x faster than its own SIMD horizontal version and 2.2x faster than FAISS IVF_FLAT, reaching 7.2x over FAISS on the OpenAI/1536 dataset on Intel. The paper also shows that horizontally stored ADSampling can lose to FAISS outright, so the pruning idea only pays off given the right layout. A companion algorithm, PDX-BOND, orders dimensions by how far the query sits from each dimension's block mean; it needs no preprocessing or transformation of the vectors and is exact, and it beat FAISS, USearch, and Milvus on exact search by 2.5x, 3.7x, and 4.0x on average. All results are single-threaded, in-memory, float32, on collections of 0.3 to 10 million vectors; graph indexes like HNSW are discussed but untested.

For a personal knowledge backend over conversational history, the relevant findings are the exact-search numbers and PDX-BOND's no-preprocessing property. A memory store that grows by appending new conversation embeddings daily cannot afford PCA retraining on every ingest, and at personal scale (well under a million vectors) exact search stays feasible, which sidesteps recall tuning entirely. At organizational scale the IVF results apply, and the layout composes with existing indexes rather than replacing them. The broader lesson for persistent AI memory work is that retrieval cost is partly a systems problem: the same algorithm swung from slower than a linear scan to 4.6x faster on the strength of data layout alone.

Read from: full text

End-to-end Neural Coreference Resolution

Lee, He, Lewis, Zettlemoyer. EMNLP 2017. arXiv:1707.07045

Lee and colleagues present the first coreference resolver trained end to end from gold clusters alone, with no syntactic parser and no hand-engineered mention detector, resources every prior competitive system required. The model treats every text span up to ten words as a candidate mention and embeds each one from bidirectional LSTM boundary states, an attention-weighted head-word vector, and a size feature. A factored score, two unary mention scores plus a pairwise antecedent score, decides for each span which earlier span, if any, it corefers with; a dummy antecedent fixed at zero lets the model abstain. The full pairwise space is quartic in document length, so the unary scores prune candidates to 0.4 spans per word with at most 250 antecedents each, and training maximizes the marginal likelihood of gold antecedents, which means mention detection emerges as a byproduct of the clustering objective.

On the English CoNLL-2012 OntoNotes test set the single model reaches 67.2 average F1 across MUC, B-cubed, and CEAR, and a five-model ensemble reaches 68.8, beating the prior state of the art by 1.5 and 3.1 points. The gains are mostly recall: the width cap discards under 2 percent of training mentions where prior pipelines lost over 9 percent, and pruning still recovers 92.7 percent of gold mentions. Dev ablations credit distance and width features with 3.8 F1, GloVe embeddings 2.4, speaker and genre metadata 1.4, and head-finding attention 1.3. Swapping in oracle gold mentions adds 17.5 F1 (to 85.2), so mention detection and anaphoricity, not pair scoring, carry most of the remaining error. Scope is narrow: one benchmark, English only, metadata assumed available, and training documents truncated to 50 sentences.

The analysis states the task’s difficulty in terms a consumer can act on. Learned attention matches syntactic heads 68 to 93 percent of the time with no head supervision, but the characteristic failure is conflating relatedness with identity: the model links “The flight attendants” to “The pilots,” and “Prince Charles and his new wife Camilla” to “Charles and Diana,” because the embeddings sit close together. It also lacks world knowledge, tying a pronoun to “some ships” on plurality cues alone.

For a memory harness that must resolve pronouns and aliases to entities before fact extraction, this paper defines what is being bought when that step is a model call. Even the strongest 2017 system leaves roughly a third of the averaged metric unclaimed, and its residual errors are exactly alias-resolution errors: a relatedness-for-identity confusion files facts about Diana under Camilla and silently corrupts the store. The factored scoring is the useful interface, since mention confidence and link confidence are separable signals a harness can threshold independently rather than trusting cluster output wholesale, and the oracle gap warns that most of the risk sits in deciding what counts as a mention at all.

Read from: full text

EvoTaxo: Building and Evolving Taxonomy from Social Media Streams

Yiyang Li et al., arXiv:2603.19711, 2026-03

Abstract (verbatim)

Constructing taxonomies from social media corpora is challenging because posts are short, noisy, semantically entangled, and temporally dynamic. Existing taxonomy induction methods are largely designed for static corpora and often struggle to balance robustness, scalability, and sensitivity to evolving discourse. We propose EvoTaxo, a LLM-based framework for building and evolving taxonomies from temporally ordered social media streams. Rather than clustering raw posts directly, EvoTaxo converts each post into a structured draft action over the current taxonomy, accumulates structural evidence over time windows, and consolidates candidate edits through dual-view clustering that combines semantic similarity with temporal locality. A refinement-and-arbitration procedure then selects reliable edits before execution, while each node maintains a concept memory bank to preserve semantic boundaries over time. Experiments on two Reddit corpora show that EvoTaxo produces more balanced taxonomies than baselines, with clearer post-to-leaf assignment, better corpus coverage at comparable taxonomy size, and stronger structural quality. A case study on the Reddit community /r/ICE_Raids further shows that EvoTaxo captures meaningful temporal shifts in discourse. Our codebase is available [here](#).

Bearing on this project

Builds and EVOLVES taxonomies from temporally ordered streams, accumulating structural evidence over time windows with per-node concept memory to preserve semantic boundaries. That is the living-tree-that-ingests-over-time idea, applied to taxonomy, with results on two Reddit corpora.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

BoostTaxo: Zero-Shot Taxonomy Induction via Boosting-Style Agentic Reasoning and Constraint-Aware Calibration

Yancheng Ling et al., arXiv:2605.12520, 2026-04

Abstract (verbatim)

Taxonomy induction is crucial for organizing concepts into explicit and interpretable semantic hierarchies. While existing methods have achieved promising results, their generalization, structural reliability, and efficiency remain limited, hindering their performance in zero-shot and large-scale scenarios. To overcome these limitations, we introduce BoostTaxo, a boosting-style LLM framework for zero-shot taxonomy induction. It takes a set of domain terms as inputs and performs parent identification in a coarse-to-fine manner, employing retrieval-augmented definition refinement, hybrid parent candidate selection, candidate rating, and structure-aware score calibration to improve taxonomy construction. Specifically, a lightweight LLM is used to efficiently filter candidate parents, while a large-scale LLM is employed to rank and score candidate parents for fine-grained parent selection. Structural features are further incorporated to calibrate candidate edge weights and enhance the reliability of the induced taxonomy. The unified BoostTaxo is evaluated on three public benchmark datasets, namely WordNet, DBLP, and SemEval-Sci, and achieves superior or comparable performance to state-of-the-art methods in zero-shot taxonomy induction. The ablation study validates the contribution of the hybrid parent candidate selection and the structure-aware score calibration to the overall performance. Further analysis investigates the impact of candidate selection size on taxonomy quality and presents representative case and failure studies, providing deeper insights into the effectiveness and limitations of the proposed framework.

Bearing on this project

Zero-shot taxonomy induction via boosting-style agentic reasoning. Further evidence the induction space is populated.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Lost in the Middle: How Language Models Use Long Contexts

Nelson F. Liu et al., TACL 2024 (arXiv 2023), arXiv:2307.03172

This paper asks how well language models actually use the long contexts they can now accept, and it answers with controlled position experiments rather than aggregate benchmarks. The authors build two tasks where exactly one item in the context matters: multi-document question answering over NaturalQuestions-Open (2,655 queries, one answer-bearing Wikipedia passage among k minus 1 Contriever-retrieved distractors) and a synthetic key-value retrieval task over JSON objects of random UUID pairs. In both, they vary two knobs independently, the total context length (10, 20, or 30 documents; 75, 140, or 300 key-value pairs) and the position of the relevant item, then measure accuracy as a function of position across MPT-30B-Instruct, LongChat-13B (16K), GPT-3.5-Turbo, and Claude-1.3, each in base and extended-context variants.

The central finding is a U-shaped performance curve: accuracy is highest when the relevant item sits at the very beginning or very end of the context and drops sharply in the middle. The drop is not small; GPT-3.5-Turbo loses more than 20 points on multi-document QA, and in the 20- and 30-document settings its worst-case mid-context accuracy falls below its 56.1 percent closed-book score, meaning the documents actively hurt. Extended-context variants track their base models almost exactly when the input fits both windows, so a bigger window does not mean better use of it. Mechanistic probes sharpen the picture: encoder-decoder models stay flat only within their training-time sequence length, query-aware contextualization (repeating the query before the data) fixes the synthetic task but barely moves multi-document QA, and even non-instruction-tuned MPT-30B shows the same U shape. A retriever-reader case study finds reader accuracy saturating far before retriever recall; going from 20 to 50 retrieved documents buys roughly 1.5 percent. The scope is 2023-era models at modest lengths (roughly 2K to 16K tokens), with an extractive accuracy metric, so absolute numbers will not transfer to current models even if the protocol does.

For the memory harness, this paper governs the injection arm, the stage that packs scored memory rows (tree summaries, entities, dated facts) into a budgeted context before the schema-validated answer call. It dictates the ordering policy: the highest-scored rows go to the edges of the assembled prompt, never buried mid-block, and it supplies the baseline the ordering experiment must beat or replicate. It also cautions against the reflex of injecting more rows just because the model tier's window allows it, since marginal rows past saturation add cost and can subtract accuracy. The position-sweep protocol itself is reusable as a harness diagnostic under LLM-judge scoring.

Read from: full extracted text of pages 1 through 12 of the 18-page PDF (all main sections, conclusion, and references; appendices skimmed only via in-text mentions).

User Feedback in Human-LLM Dialogues: A Lens to Understand Users But Noisy as a Learning Signal

Liu, Zhang, Choi 2025, User Feedback in Human-LLM Dialogues, EMNLP 2025, arXiv:2507.23158

The paper asks whether implicit user feedback in chat logs (a user saying “could you label the y-axis?” rather than clicking a thumbs-down) can be detected reliably and then used to improve a model. Working with two real interaction corpora, LMSYS-chat-1M and WildChat, the authors densely annotate 109 conversations (75 from LMSYS, 34 from WildChat), labeling every user turn after the first with a six-way ontology inherited from Don-Yehiya et al. 2024: positive feedback, four negative types (rephrasing, make aware with or without correction, ask for clarification), and no feedback. Inter-annotator agreement is Cohen’s kappa 0.70 for binary and 0.60 for fine-grained labels. They then build a GPT-4o-mini classifier with in-context examples and, in the second half, test whether using the content of negative feedback, not just its polarity, produces better training data: a stronger model regenerates the unsatisfying response conditioned on the user’s complaint, and the result is distilled into 7B models (Vicuna and Mistral) by supervised fine-tuning on 20K conversations.

The descriptive findings are the cleaner half. Feedback is common and skews negative: in later turns it appears in more than half of user utterances, while praise is rare. Refusals explain almost none of it (under 3 percent of sampled responses in either corpus). More troubling for naive harvesting, prompts that elicit positive feedback are slightly more toxic and lower quality than random ones; manual inspection traces this to users praising successful jailbreaks. Their detection prompt roughly doubles prior accuracy on the dense setting (41.6 versus 31.5 percent binary). The training results are mixed and scope-bound: feedback-conditioned regeneration beats regeneration from scratch on MTBench (80 short human-written questions; Vicuna-7b reaches 6.86 versus 6.38), but on WildBench (500 longer, more complex real-user questions) it loses (23.38 versus 27.97 for the same setup), and plain distillation from the stronger model on random conversations does best of all.

For anyone building a personal knowledge backend over conversational history, the annotation study is the useful part: it establishes that later turns are dense with correction signal marking unmet intent, exactly the turns a memory system should weight when inferring what a user actually wants, while the training results warn that this signal cannot be consumed raw. Positive reactions encode bad behavior as well as good, and inferring unspoken intent can erode adherence to well-specified requests. Both lessons rest on two 2023-era public corpora and 7B models, one of which (LMSYS) reflects evaluation play rather than real work, so organizational logs may behave differently.

Read from: full text

Evaluating Very Long-Term Conversational Memory of LLM Agents

Maharana et al. 2024, LoCoMo: Evaluating Very Long-Term Conversational Memory, ACL 2024, arXiv:2402.17753

The paper builds LoCoMo, a dataset of 50 very long two-person conversations, each averaging about 300 turns and 9,200 tokens across up to 35 sessions, roughly nine times the length of MSC, the previous standard. Rather than collect real conversations over months, the authors generate them: two gpt-3.5-turbo agents, each seeded with an expanded persona from MSC and a temporal event graph of up to 25 causally linked life events spanning 6 to 12 months, converse using a memory architecture borrowed from Park et al.’s generative agents (per-session summaries as short-term memory, per-turn observations as long-term memory). Agents can also share web-searched images and react to them. Human annotators then edited about 15 percent of turns and replaced or removed about 19 percent of images to fix long-range inconsistencies. On top of the dataset sit three tasks: question answering across five reasoning types (single-hop, multi-hop, temporal, open-domain, adversarial), event summarization scored with an adapted FactScore, and multimodal dialogue generation.

The headline finding is that every tested configuration sits far below the human QA benchmark of 87.9 F1. GPT-4-turbo with a 4K window is the best base model at 32.1 overall. Extending context helps ordinary recall (gpt-3.5-turbo-16k climbs from 24.1 to 37.8 as its window grows from 4K to 16K) but collapses on adversarial questions, falling to 2.1 F1 against 70.2 for GPT-4-turbo at 4K: given more context, the model hallucinates answers to unanswerable questions and misattributes events to the wrong speaker. Temporal reasoning stays weak everywhere. On event summarization the long-context model actually underperforms the 4K base model (39.9 versus 45.9 FactScore F1), and the common failure modes are missed causal links, hallucinated details, and wrong speaker attributions. All of this is on 50 English, LLM-generated-then-edited conversations, so the numbers describe that synthetic corpus, not real chat logs.

The most transferable result concerns memory representation. RAG over raw dialogue logs helps modestly, but RAG over “observations,” per-turn assertions distilled about each speaker’s life, is the strongest configuration: top-5 observations reach 41.4 overall F1, and performance degrades as more are retrieved, which the authors read as a signal-to-noise problem. Session summaries retrieve well but answer poorly, suggesting summarization loses the facts QA needs. For a personal knowledge backend over conversational history, that ordering (distilled assertions beat both raw transcripts and summaries, and retrieval precision beats retrieval volume) is a directly usable design finding, and the adversarial collapse is a warning that scale of context is not the same as memory.

Read from: pages 1-10 of 19 (body complete; remainder is references and appendix)

Efficient and robust approximate nearest neighbor search using Hierarchical Navigable Small World graphs

Malkov & Yashunin, Hierarchical Navigable Small World graphs, TPAMI, arXiv:1603.09320

The paper introduces HNSW, a graph index for approximate K-nearest-neighbor search that needs no auxiliary coarse-search structure. Elements are inserted one at a time; each is assigned a maximum layer drawn from an exponentially decaying distribution, and a proximity graph is built at every layer it occupies. Search starts at the sparse top layer, greedily descends to a local minimum, then drops a layer and repeats, so the structure behaves like a probabilistic skip list with proximity graphs in place of linked lists. The second contribution is a neighbor-selection heuristic: rather than linking a new element to its M closest candidates, it keeps only candidates closer to the new element than to any already-selected neighbor, which recovers a relative neighborhood subgraph and preserves connectivity across cluster boundaries.

The authors argue $O(\log N)$ search and $O(N \log N)$ construction, with the caveat that the analysis assumes exact Delaunay graphs and constant node degree; empirical support comes from low-dimensional data ($d=4$ and $d=8$ random vectors up to 10M elements), and they state that verifying the resilience of the approximation at $d=128$ would need infeasibly large datasets. Practical settings fall out of the experiments: M between 5 and 48, ground-layer cap $M_{\max} = 2M$, level normalizer $m_L = 1/\ln(M)$, efConstruction near 100, and index overhead of roughly 60 to 450 bytes per object. On the ann-benchmarks datasets (1M SIFT, 1.2M GloVe, 2M CoPhIR, 1M DEEP, 60k MNIST) HNSW beats Annoy, FLANN, FALCONN, and VP-tree by wide margins, and on the wiki-8 dataset under Jensen-Shannon divergence it improves on plain NSW by almost three orders of magnitude in distance computations. Against Faiss product quantization on a 200M SIFT subset, HNSW reaches higher recall at faster query times and builds in 42 minutes versus 11 to 12 hours, but needs 64 Gb of RAM against Faiss's 23.5 to 30, and its query-time scaling there visibly deviates from pure logarithm.

For a conversational-memory backend, HNSW is the algorithm inside most current vector stores, so its tradeoffs are the field's tradeoffs. Insertion is truly incremental with no reshuffling, which fits an append-only history; it handles arbitrary metrics, including edit distance on a 1M DNA dataset. Two limits matter at organizational scale: the paper defers element update and deletion to future work, and distributed search is awkward because every query funnels through the top layer, with skip-graph techniques proposed but not implemented. RAM residency of the full graph is the cost of its speed.

Read from: full text

A Survey on Open Information Extraction

Christina Niklaus, Matthias Cetto, André Freitas, Siegfried Handschuh, COLING 2018, arXiv:1806.05599

This survey maps the first decade of Open Information Extraction, the task of pulling relational tuples of the form (arg1; rel; arg2) from text without a predefined relation inventory. It opens with the three founding requirements from Banko et al. 2007 (unsupervised extraction in a single corpus pass, shallow features that survive heterogeneous text, efficiency at Web scale), then organizes the systems that followed into four families: learning-based (TextRunner, WOE, OLLIE, ReNoun), rule-based (ReVerb, KrakeN, Exemplar, PropS, PredPatt), clause-based restructuriers (ClausIE, Stanford Open IE), and systems that capture inter-proposition relationships (OLLIE’s modifiers, CSD-IE, NestIE, MinIE, Graphene). The through-line is drift from the founding constraints: later systems lean on dependency parsers, trading the original speed and domain-independence claims for precision.

The mechanisms it establishes are concrete. ReVerb’s POS-based regular expression covers about 85 percent of verb-mediated relational phrases in English while filtering incoherent and uninformative extractions. OLLIE bootstraps pattern templates from ReVerb seed tuples and adds attribution and clausal-modifier fields so hypothetical or reported claims are not stored as asserted fact. MinIE annotates polarity, modality, and attribution outside the tuple to keep extractions minimal. The evaluation critique is the survey’s sharpest contribution: most systems were hand-scored on a few hundred sentences of news, Wikipedia, or Web text, with precision-style metrics and recall proxies like yield, so cross-system numbers rarely compare. Stanovsky and Dagan’s benchmark (over 10,000 extractions across 3,200 sentences, built on assertedness, minimal propositions, and completeness principles) and RelVis with its six error classes are the proposed fixes, and at writing time almost no one had adopted them. The scope limits are stated in the same section: nearly all systems are English-only, and canonicalization and coreference are named as open problems, not solved ones.

For the memory harness, this paper governs the fact-row generation stage, where segmented conversation text becomes dated entity-and-fact rows before tree placement. It supplies the design vocabulary (assertedness, so beliefs and hypotheticals never enter the tree as facts; minimality, so each row carries one claim; attribution fields for who said what) that the schema-validated LLM extraction calls should enforce regardless of which model tier runs them. It equally warns the evaluation design: yield-style counts and hand-scored spot checks are the documented traps, so the LLM-judge scoring of the three injection arms needs a fixed gold slice and explicit error classes, RelVis-style, rather than raw extraction counts.

Read from: the full text of all 13 pages, extracted with PyMuPDF (sections 1 through 5 plus references).

NoLiMa: Long-Context Evaluation Beyond Literal Matching

Modarressi, Deilamsalehy, Dernoncourt, Bui, Rossi, Yoon, Schütze; ICML 2025 (PMLR 267); arXiv:2502.05167

The paper builds a needle-in-a-haystack benchmark in which the question and the hidden fact share almost no words, so a model cannot find the needle by pattern matching and must instead follow a latent association. A needle like “Actually, Yuki lives next to the Semper Opera House” answers the question “Which character has been to Dresden?” only if the model knows the opera house is in Dresden; two-hop variants stretch the chain further (Dresden to Saxony). The authors quantify the gap they are closing with ROUGE precision between question and relevant context: vanilla NIAH scores 0.905 R-1, NoLiMa scores 0.069. The benchmark uses 58 question-needle pairs built from 5 needle groups, haystacks assembled from snippets of 10 open-license books and filtered (with Llama 3.3 70B plus manual review) to remove accidental distractors and false answers, and 26 needle placements per length, giving 7,540 tests per context length.

Across 13 models claiming at least 128K context, short-context performance is strong (most base scores above 84 percent) but decays fast. At 32K tokens, 11 of 13 fall below half their base score; GPT-4o, the best performer, drops from 99.3 to 69.7. Defining effective length as the longest context holding 85 percent of the base score, most models sit at 2K or below; GPT-4o reaches 8K against a claimed 128K. The same Llama 3.1 70B that reaches 32K effective length on RULER manages 2K here, which isolates literal matching as the crutch. Two-hop questions and inverted fact order (keyword before character name) degrade further, and an aligned-depth analysis of the last 2K tokens shows the drop tracks total context length rather than needle position, implicating attention capacity rather than position encoding. CoT prompting helps (two-hop at 16K goes from 42.7 to 56.7 for Llama 3.3 70B) but does not fix it; on a hard subset, GPT-o1 and DeepSeek-R1 still fall below 50 percent of base at 32K. Restoring literal matches via multiple choice lifts two-hop 32K from 25.9 to 87.2, while a single distractor sentence sharing words with the question cuts GPT-4o’s effective length to 1K.

For a conversational-history backend, the message is that dumping months of transcripts into context and asking associative questions (“when did I decide X”) will fail well before the advertised window, especially when the query is phrased differently than the original conversation, which is the normal case. Retrieval must do semantic bridging before the model sees anything, keep contexts short, and avoid surfacing lexically similar but irrelevant passages, since those actively mislead. The scope caveat: results come from synthetic single-fact needles in book-text haystacks up to 32K (128K for two models), not from real dialogue data.

Read from: full text

Industry-scale Knowledge Graphs: Lessons and Challenges

Noy, Gao, Jain, Narayanan, Patterson, Taylor; ACM Queue 17(2) / CACM 62(8), 2019; DOI 10.1145/3331166

This is an experience paper, not an empirical study: practitioners from Google, Microsoft, Facebook, eBay, and IBM (expanding on a 2018 ISWC panel) each describe their production knowledge graph, then pool the challenges that recur across all five. It proceeds company by company through design decisions, then through shared open problems: entity disambiguation, type membership, changing knowledge, extraction from unstructured sources, and operating at scale.

The scale numbers frame everything. Google holds roughly 1 billion entities and 70 billion assertions; Microsoft's Bing graph about 2 billion entities and 55 billion facts; Facebook about 50 million entities and 500 million assertions; eBay expects 100 million products and over a billion triples; IBM's framework is proven past 100 million documents and 5 billion relationships. All five use strongly typed entities and relations under a schema or ontology; correctness at this size depends on machinery rather than curation. Bing's single entry for the actor Will Smith is assembled from 108,000 facts across 41 websites, against 200 Will Smiths in Wikipedia alone, which is why the authors name disambiguation the top industry challenge despite years of research. The paper offers mechanisms rather than solutions, and mostly self-reported ones with no benchmarks. Google bootstraps its schema from a small set of low-level structures plus metatypes (types as instances of types) to validate invariants at scale, such as painters predating their paintings. Facebook keeps provenance on every assertion, defines correctness as always being able to explain why an assertion was made, and rebuilds the whole graph from sources daily through an idempotent build. eBay funnels all writes through a replicated log feeding a low-latency document store and a separate analytical graph store. IBM uses polymorphic stores, treats source evidence as a primitive, and defers entity resolution to query time so that context can settle ambiguous names.

For a personal knowledge backend built over conversational history, the transferable material is architectural. Provenance-first assertions, identity managed separately from surface forms, late-binding disambiguation, and the log-plus-specialized-stores pattern all apply directly to conflicting, evolving conversational facts, and IBM's layering of general, domain, and private customer knowledge is essentially the personal-versus-shared split an organizational rollout would need. The authors also flag personalized on-device graphs, small enough to ship to a phone, as an open direction they invite research on. As a representation reference the paper shows the dominant industrial pattern, typed property graphs with schema-level reasoning, but it describes no query languages, storage internals, or evaluation, so it orients rather than specifies.

Read from: full text

RouteLLM: Learning to Route LLMs with Preference Data

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M Waleed Kadous, Ion Stoica; arXiv (cs.LG), 2024 (v4 Feb 2025); ICLR 2025; arXiv:2406.18665

RouteLLM is a training framework for learning routers that decide, per query and before any generation happens, whether a strong expensive model or a weak cheap one should answer. The router is a win prediction model estimating the probability that the strong model beats the weak one on a given query, paired with a threshold that converts that probability into a routing decision; sweeping the threshold traces a cost-quality curve. The authors train four router architectures (similarity-weighted Bradley-Terry ranking, matrix factorization over model and query embeddings, a BERT classifier, and a Llama 3 8B causal classifier) on 65k pairwise human preference labels from Chatbot Arena, with models clustered into ten tiers to fight label sparsity, and evaluate on decontaminated MMLU, MT Bench, and GSM8K against a random-routing baseline.

The headline mechanism is data augmentation. Trained on Arena data alone, routers do well only on MT Bench (matrix factorization needs half the GPT-4 calls of random to recover 50 percent of the quality gap) and sit at chance on MMLU and GSM8K, which the authors trace to distribution mismatch via a benchmark-dataset similarity score. Adding roughly 1500 golden-labeled MMLU validation questions, under 2 percent of the training data, lifts every router about 20 percent past random on MMLU; adding 120K GPT-4-judge labels (about \$700 to collect) raises MT Bench APGR improvement past 50 percent. Best routers cut cost 3.66x on MT Bench at 95 percent of GPT-4 quality, though only 1.4-1.5x on MMLU and GSM8K, and gains depend on augmentation data resembling the target distribution. The strongest result for reuse: the same routers, untouched, route Claude 3 Opus/Sonnet and Llama 3.1 70B/8B pairs at comparable quality, so the learned signal is query difficulty rather than model identity. Scope is binary strong-versus-weak routing only, and router rankings shift across benchmarks without a clean explanation.

For the memory harness, this governs the model-tier assignment stage: the layer where every schema-validated LLM call (segmentation of the corpus into the tree, entity extraction, dated-fact extraction) is bound to a model tier behind a swappable interface. The harness's planned sensitivity experiment fixes one tier per stage and measures degradation; RouteLLM supplies the natural extension, a per-call router in front of the same interface, and its model-pair transfer result argues the router would survive a tier swap without retraining. The catch carries over too: routing quality tracks how well training data matches the harness's actual call distribution, so per-stage judge-labeled samples would be the price of entry. The three injection arms and LLM-judge scoring sit downstream and are untouched by routing.

Read from: full extracted text of all 16 pages, including appendices.

ReasoningBank: Scaling Agent Self-Evolving with Reasoning Memory

Ouyang et al. 2026, ICLR 2026, arXiv:2509.25140

ReasoningBank is a memory framework for LLM agents that face a stream of tasks with no ground-truth feedback. Instead of storing raw trajectories (Synapse) or success-only workflows (AWM), it distills each finished episode into structured memory items, each with a title, a one-sentence description, and content holding the abstracted strategy or pitfall. An LLM-as-a-judge labels each trajectory success or failure without reference answers, and both outcomes feed extraction: successes yield validated strategies, failures yield preventative lessons. At each new task the agent retrieves the top-k items by embedding similarity and injects them into its system instruction; new items are consolidated back by simple addition, closing the loop. The authors also propose memory-aware test-time scaling (MaTTS), which spends extra compute per task (k parallel rollouts with self-contrast, or k sequential self-refinements) and aggregates across the redundant attempts to curate memory, rather than converting each rollout independently.

On WebArena (684 tasks, Map excluded) with Gemini-2.5-flash, ReasoningBank lifts overall success rate from 40.5 (no memory) to 48.8, and MaTTS at k=5 reaches 51.8; gains of +7.2 and +4.6 hold for Gemini-2.5-pro and Claude-3.7-sonnet. On SWE-Bench-Verified, resolve rate rises from 34.2 to 38.8 (flash) and 54.0 to 57.4 (pro), and Mind2Web improves across cross-task, cross-website, and cross-domain splits. Steps per task drop as well, up to 1.4 versus no memory, with the largest cuts on successful runs (2.1 fewer steps on Shopping, a 26.9 percent reduction), so memory shortens good paths rather than truncating bad ones. Ablations on WebArena-Shopping with flash show failures matter: success-only extraction gives 46.5, adding failures 49.7, while AWM degrades from 44.4 to 42.2 when fed failures. The judge itself is only 72.7 percent accurate there, yet simulated accuracy from 70 to 90 percent barely moves the results. A case study shows memory items maturing from procedural rules into compositional checking strategies over the task stream.

For a personal knowledge backend over conversational history, the transferable finding is representational: distilled, self-contained strategy units with title, description, and content beat both raw transcripts and success-only procedures, and failure records carry real signal if the schema is built to hold lessons rather than replays. The robustness to a noisy self-judge suggests curation can run without human labels, which is what organizational scale requires. Scope limits apply: benchmarks are web browsing and repository-level software fixes, streams are hundreds of tasks, consolidation is pure addition with no pruning or deduplication, so behavior over years of accumulation is untested.

Read from: pages 1-12 of 30

Unifying Large Language Models and Knowledge Graphs: A Roadmap

Pan et al. 2024, IEEE TKDE, arXiv:2306.08302

This is a survey and roadmap, not an experimental paper. The authors organize research on combining LLMs and knowledge graphs into three frameworks: KG-enhanced LLMs (graphs injected during pre-training, at inference, or used to probe and interpret models), LLM-augmented KGs (language models applied to KG embedding, completion, construction, graph-to-text generation, and question answering), and Synergized LLMs + KGs, where the two are trained or operated jointly for representation and reasoning. Under each framework they build a fine-grained taxonomy and tabulate representative methods, roughly a hundred systems spanning 2018 to 2023, from ERNIE and K-BERT through RAG, KEPLER, DRAGON, StructGPT, and Think-on-Graph, each tagged by technique and by the kind of graph or model involved.

Because it synthesizes rather than measures, its findings are the trade-offs it documents. Pre-training injection yields the strongest downstream performance on knowledge-intensive tasks but never permits knowledge updates without retraining; retrieval at inference handles frequently changing open-domain facts but the underlying model was never trained to exploit them, so results can be sub-optimal. For KG completion, encoder-style methods (KG-BERT and successors) must score every candidate entity and cannot handle unseen entities, while generator-style methods generalize and skip ranking but may produce entities that do not exist in the graph. The probing work it collects is sobering: LAMA-style cloze tests show LLMs recall popular facts, but a Wikidata study of low-frequency entities found scaling fails to appreciably improve memorization of tail facts, and a causal analysis (Shaobo et al. 2022) found models complete facts from positionally close words rather than knowledge-bearing ones.

For a personal knowledge backend over conversational history, two threads matter most. First, the LLM-augmented KG construction pipeline (entity discovery, coreference resolution, relation extraction, plus end-to-end systems like Grapher and PiVE, which uses a small T5 to iteratively correct a larger model's extracted triples) is exactly the machinery for distilling structured, queryable memory from raw transcripts, and the survey's own trade-off analysis argues for inference-time retrieval over parameter injection whenever the store must stay editable. Second, the agent-style reasoning line (StructGPT, Think-on-Graph) shows LLMs traversing a graph step by step without retraining, giving interpretable answers that cite paths, which is the pattern an organizational memory would need for auditability. The future-directions section flags the open problems such a system inherits: hallucination detection against the graph, the ripple effects of editing one fact, and injecting knowledge into black-box models where prompts are bounded by context length.

Read from: pages 1-24 of 28

QuALITY: Question Answering with Long Input Texts, Yes!

Richard Yuanzhe Pang et al., NAACL 2022, arXiv:2112.08608

QuALITY is a multiple-choice question answering dataset over English passages averaging about 5,000 tokens, mostly Project Gutenberg science fiction plus Slate and other nonfiction, all CC-BY or more permissive. Its central move is procedural. Twenty-two hired writers each read a full passage and write ten four-option questions; every question then passes through two validation gates. Untimed validation, with three to five annotators who read the whole text, certifies that a single answer is correct and unambiguous (6,737 of 7,620 questions survive, 88.4 percent). Speed validation gives five annotators only 45 seconds with the passage; questions that a majority of these skimmers get wrong, while untimed readers get them right, are labeled HARD. That subset is 3,360 questions, 49.9 percent of the dataset, and it is the paper’s evidence that the questions cannot be answered by search or skimming.

The numbers earn the design. Human majority-vote accuracy is 93.5 percent overall and 89.1 percent on HARD, while the best baseline (DeBERTaV3-large with DPR sentence extraction, after intermediate training on RACE) reaches only 55.4 percent, 46.7 on HARD. Question-only baselines hit 43.3 percent, and even oracle extraction using the correct answer as the retrieval query tops out at 78.3 on dev, so relevant-excerpt retrieval alone does not solve the task. Predicting the option with highest lexical overlap scores 26.6 percent, near chance. Scope limits sit next to these results: the corpus is mainstream US English written by a relatively privileged population, passages cap near 8k tokens, the baselines predate modern long-context models, and the pipeline cost \$9.10 per question, so the method is expensive to replicate at scale.

For the memory harness, QuALITY governs the question-set construction stage, the point where the certified-unknown corpus is turned into probes before any injection arm runs. Its two-gate protocol translates directly: an unambiguity check (schema-validated LLM calls can play the untimed annotators, with the judge tier swappable) and a shallow-access check, where a bare model or a snippet-only pass stands in for the 45-second skimmers. A question enters the counterfactual-miss or corrected-facts band only if it fails under shallow access and succeeds under full-narrative access, which is exactly what makes the three injection arms comparable and keeps LLM-judge scoring honest, since four discriminating options give the judge a closed decision rather than a free-text similarity call. The 26.6 percent lexical-overlap floor is the target any generated distractor set should match.

Read from: full extracted text of pages 1 to 8 (abstract through conclusion, ethics, and contributions) of the 23-page arXiv PDF; remaining pages are references and appendices, not read.

LLM Evaluators Recognize and Favor Their Own Generations

Arjun Panickssery, Samuel R. Bowman, Shi Feng; NeurIPS 2024 (oral), Advances in Neural Information Processing Systems 37; arXiv:2404.13076

The paper asks whether self-preference, the tendency of an LLM evaluator to score its own outputs above equal-quality alternatives, is driven by self-recognition: does the model prefer a text because it can tell the text is its own? Working on news summarization with 1,000 articles each from XSUM and CNN/DailyMail, the authors have GPT-4, GPT-3.5, and Llama-2-7b-chat each play generator, evaluator, and authorship identifier. They measure both properties in pairwise form (choose the better summary, or pick which of two is yours) and individual form (rate one summary, or answer yes/no on authorship), converting output token probabilities into confidence scores and averaging over both option orderings to cancel position bias, which is severe: GPT-4, GPT-3.5, and Llama reverse pairwise preferences under reordering 25, 58, and 89 percent of the time.

Out of the box, all three models show self-preference, and all three exceed chance at self-recognition; GPT-4 reaches 73.5 percent accuracy distinguishing its summaries from two other LLMs and humans, while weaker models mostly cannot separate themselves from stronger ones. Fine-tuning on as few as 500 pairwise examples pushes GPT-3.5 and Llama 2 above 90 percent self-recognition, and across all fine-tuned variants, including control models trained on length, vowel count, readability, and degenerate tasks, self-preference strength falls on a single linear trend against self-recognition score, transferring across datasets. The authors invalidate the reverse causal story (evaluators slightly disfavor fine-tuned models' generations, average preference 0.46), yet they are explicit that this is evidence toward causation, not proof: results cover one task family, three models, one prompt design, and lack ground-truth quality controls, so disproportionate preference is shown only indirectly, through paired self-preference scores summing above one.

For the memory harness, this paper governs the LLM-judge scoring stage that grades answers from the three injection arms on the query set. The stage-wise model-sensitivity experiment swaps the model tier behind extraction and answering, from a local 7B through frontier; if the judge were drawn from the same family as the arm under test, this result predicts inflated scores for that arm, corrupting the comparison. The design consequence is to pin one fixed strong judge across every arm and never let judge identity covary with the system under test, to strip authorship cues from judged transcripts where feasible, and to aggregate judge calls over both answer orderings, since the ordering-bias numbers here are as large as the effect being measured. The tree-building and extraction stages are untouched by this result; it constrains scoring only.

Read from: full paper text, pages 1 to 10 of the PDF (all body sections and conclusion) plus appendices A and B, extracted via pypdf; citation metadata from gap_survey_entries.json.

Generative Agents: Interactive Simulacra of Human Behavior

Park, O'Brien, Cai, Morris, Liang, and Bernstein, UIST 2023, arXiv:2304.03442

The paper builds twenty-five LLM-driven agents, each seeded with one paragraph of natural language identity, and runs them in Smallville, a Sims-like sandbox town, to test whether an architecture wrapped around a language model (gpt-3.5-turbo; GPT-4 was invitation-only at the time) can produce believable long-horizon behavior. The architecture has three parts. A memory stream records every experience as a timestamped natural language object; retrieval scores each record as an equally weighted sum of recency (exponential decay, factor 0.995 per game hour since last access), importance (the LLM rates poignancy 1 to 10 at creation), and relevance (cosine similarity of embeddings against a query), min-max normalized. Reflection fires when summed importance of recent events passes a threshold of 150, roughly two or three times per agent day: the model proposes three salient questions from the 100 most recent records, retrieves against each, and writes cited higher-level inferences back into the stream, forming trees whose leaves are raw observations. Planning drafts a day in five to eight chunks, then recursively decomposes to hour and then 5 to 15 minute actions; plans and reflections re-enter the stream so retrieval mixes all three.

In a controlled evaluation, 100 Prolific raters ranked interview responses from the full architecture, three ablations, and a crowdworker-authored baseline. The full system scored highest (TrueSkill μ 29.89) and each removed component cost believability, down to μ 21.21 with no memory at all; the fully ablated condition, which stands in for prior first-order prompting work, trailed the full system by a standardized effect of $d = 8.16$, and it was statistically indistinguishable from the human crowdworkers. In a two-game-day end-to-end run, knowledge of a party spread from one agent to thirteen (4% to 52%) and a mayoral candidacy from one to eight, with zero hallucinated claims of knowing; network density rose from 0.167 to 0.74 with 1.3% hallucinated acquaintance claims; five of twelve invitees attended the party. All of this holds only at this scale: 25 agents, two simulated days, one hand-built town, at a cost of thousands of dollars in tokens.

For a personal knowledge backend, the transferable finding is that a flat log plus scored retrieval is not enough; the ablations show believable use of history required the synthesized layer too, and the failure catalog (missed retrievals, embellished recall, one agent conflating a neighbor named Adam Smith with the economist) is a checklist of what such a backend must guard against. This paper fixed the memory-stream, reflection, retrieval-triple vocabulary that most subsequent persistent-memory work inherits. The authors also flag memory hacking, planted false events, as an open robustness problem, which becomes an access-control question at organizational scale.

Read from: full text

HERCULES: Hierarchical Embedding-based Recursive Clustering Using LLMs for Efficient Summarization

Gabor Petnehazi et al., arXiv:2506.19992, 2025-06

Abstract (verbatim)

The explosive growth of complex datasets across various modalities necessitates advanced analytical tools that not only group data effectively but also provide human-understandable insights into the discovered structures. We introduce HERCULES (Hierarchical Embedding-based Recursive Clustering Using LLMs for Efficient Summarization), a novel algorithm and Python package designed for hierarchical k-means clustering of diverse data types, including text, images, and numeric data (processed one modality per run). HERCULES constructs a cluster hierarchy by recursively applying k-means clustering, starting from individual data points at level 0. A key innovation is its deep integration of Large Language Models (LLMs) to generate semantically rich titles and descriptions for clusters at each level of the hierarchy, significantly enhancing interpretability. The algorithm supports two main representation modes: `direct` mode, which clusters based on original data embeddings or scaled numeric features, and `description` mode, which clusters based on embeddings derived from LLM-generated summaries. Users can provide a `topic_seed` to guide LLM-generated summaries towards specific themes. An interactive visualization tool facilitates thorough analysis and understanding of the clustering results. We demonstrate HERCULES's capabilities and discuss its potential for extracting meaningful, hierarchical knowledge from complex datasets.

Bearing on this project

Hierarchical embedding-based recursive clustering with LLM summarization. Uses recursive k-means, i.e. HARD assignment - one item, one cluster. Part of the evidence that the hierarchical-summary line partitions by default, which is what makes the covering question worth asking.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Beyond the Context Window: A Cost-Performance Analysis of Fact-Based Memory vs. Long-Context LLMs for Persistent Agents

Pollertlam & Kornsuwannawit (Bricks Technology), arXiv preprint, 2026, arXiv:2603.04814

The paper asks a deployment question: for an assistant that must remember a user across sessions, is it better to resend the whole conversation history to a long-context model every turn, or to extract facts once into a memory store and retrieve only what each query needs? The authors build the memory side on the open-source Memo pipeline (GPT-5-nano extracts atomic flat-typed facts, text-embedding-3-small embeds them into pgvector, top-20 retrieval feeds a GPT-5-mini reader) and compare it against long-context inference with GPT-5-mini and GPT-OSS-120B on three benchmarks: LongMemEval (500 conversations, about 101,600 tokens each), LoCoMo (10 dialogues, 1,986 questions), and PersonaMem v2 (222 conversations, 589 questions). Accuracy is scored by GPT-5-mini as judge with a three-vote majority; cost is modeled from actual API spend with a 90 percent prompt-caching discount applied to the long-context side from turn two onward.

Long-context GPT-5-mini wins decisively on factual recall: 92.85 percent versus 57.68 on LoCoMo and 82.40 versus 49.00 on LongMemEval, gaps of 35.2 and 33.4 points. On PersonaMem v2 the gap shrinks to 7.3 points, and the memory system edges out long-context GPT-OSS-120B (62.48 versus 60.50), because persona attributes are exactly the stable facts flat extraction captures well. The cost picture inverts with use. Extraction compresses a 101,600-token history to about 2,909 retrieved tokens (roughly 35:1), a one-time write of \$0.0430 per user, after which each query costs about \$0.0013; the cached long-context turn still scales with context length. At 100k tokens the memory system becomes cheaper after about ten turns and saves 26 percent by turn twenty; the break-even falls from 13 turns at 30k to 9 at 500k. An error analysis names what compression loses: precise temporal markers, implicit coreferences, and updates that never propagate (a “twice a week” fact surviving a later revision to three). All accuracy results are scoped to Memo’s flat extraction; the authors caution they do not generalize to structured memory architectures like Zep or temporal semantic memory.

For a personal knowledge backend over conversational history, this paper supplies the economic argument and its price: at organizational scale, per-query cost stays near-flat only if you accept extraction loss, and the three failure modes are a concrete checklist of what a fact schema must handle (timestamps, role resolution, update propagation). The finding that answer-model capability, not context length, limited GPT-OSS-120B is a useful caution against attributing retrieval failures to the memory layer. The static-snapshot protocol, with no incremental updates during querying, marks the gap between these benchmarks and a live system.

Read from: full text

Just What You Desire: Constrained Timeline Summarization with Self-Reflection for Enhanced Relevance

Muhammad Reza Qorib et al., arXiv:2412.17408, 2024-12

Abstract (verbatim)

Given news articles about an entity, such as a public figure or organization, timeline summarization (TLS) involves generating a timeline that summarizes the key events about the entity. However, the TLS task is too underspecified, since what is of interest to each reader may vary, and hence there is not a single ideal or optimal timeline. In this paper, we introduce a novel task, called Constrained Timeline Summarization (CTLs), where a timeline is generated in which all events in the timeline meet some constraint. An example of a constrained timeline concerns the legal battles of Tiger Woods, where only events related to his legal problems are selected to appear in the timeline. We collected a new human-verified dataset of constrained timelines involving 47 entities and 5 constraints per entity. We propose an approach that employs a large language model (LLM) to summarize news articles according to a specified constraint and cluster them to identify key events to include in a constrained timeline. In addition, we propose a novel self-reflection method during summary generation, demonstrating that this approach successfully leads to improved performance.

Bearing on this project

Constrained timeline summarization with self-reflection. The paper that surfaced timeline summarization as a field during the o8-27 search.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Zep: A Temporal Knowledge Graph Architecture for Agent Memory

Rasmussen, Paliychuk, Beauvais, Ryan, Chalef (Zep AI), arXiv preprint, 2025, arXiv:2501.13956

The paper describes Zep, a commercial memory service for LLM agents, and its core engine Graphiti, a temporally aware knowledge graph built continuously from conversation. Incoming messages become episode nodes, a non-lossy raw store; an LLM pipeline (gpt-4o-mini, with a reflection step to curb hallucination) extracts entities and facts from each message plus the last four for context, embeds them in 1024 dimensions, and resolves duplicates against existing nodes through hybrid search and an LLM judgment. Above the entity layer, communities of strongly connected entities carry map-reduce style summaries, maintained incrementally by label propagation rather than Leiden so the graph can extend without full recomputation. The distinctive piece is the bi-temporal model: every fact edge carries four timestamps, two for when the system learned and expired it and two for when it held true in the world, and new edges that contradict old ones invalidate them rather than overwrite them. Retrieval runs cosine similarity, BM25, and graph breadth-first search in parallel, reranks (RRF, MMR, mention frequency, node distance, or cross-encoder), and formats roughly 1.6k tokens of facts with validity ranges plus entity summaries.

On MemGPT's Deep Memory Retrieval benchmark (500 conversations of only about 60 messages each), Zep scores 94.8% with gpt-4-turbo against MemGPT's 93.4%, and 98.2% with gpt-4o-mini; the authors themselves note that simply stuffing the full conversation into context scores 94.4% and 98.0%, so DMR barely discriminates. The stronger evidence is LongMemEval, with conversations averaging 115k tokens: Zep beats the full-context baseline by 15.2% with gpt-4o-mini (63.8% vs 55.4%) and 18.5% with gpt-4o (71.2% vs 60.2%), while cutting latency about 90% (2.58s vs 28.9s with gpt-4o). Gains concentrate in temporal reasoning, multi-session synthesis, and preference questions; single-session-assistant questions actually drop 17.7% under gpt-4o, an admitted failure. MemGPT could not be run on LongMemEval at all.

For a personal memory backend over conversational history, this is close to a reference design: episodic raw storage under a derived semantic layer, entity resolution as the hard ongoing problem, and explicit fact validity intervals as the answer to knowledge that changes, which flat retrieval handles badly. The retrieval numbers also argue that a graph layer earns its cost mainly at scale; below context-window size it buys little accuracy. The paper doubles as a critique of the field's evaluation habits: DMR is too small and too fact-retrieval shaped, and no benchmark yet tests fusing conversation with structured business data, which is exactly the organizational-scale case. All results depend on chat-style corpora and OpenAI models.

Read from: full text

Collaborative Memory: Multi-User Memory Sharing in LLM Agents with Dynamic Access Control

Rezazadeh, Li, Lou, Zhao, Wei, Bao (Accenture Center for Advanced AI), arXiv preprint (under review), 2025, arXiv:2505.18279

The paper asks how a system serving many users with many LLM agents can share memory across user boundaries without leaking anything a given user is not entitled to see. Its answer is a framework, not a learned model: permissions are encoded as two time-varying bipartite graphs (user-to-agent and agent-to-resource), and memory is split into a private tier scoped to the originating user and a shared tier open to cross-user retrieval. Every memory fragment carries immutable provenance: creation time, contributing user, contributing agents, and resources touched. A fragment is readable in a given user-agent context only if its contributing agents fall within the user's current agent permissions and its source resources fall within the agent's current resource permissions, so revoking an edge in the graph retroactively hides fragments that depended on it. Read and write policies, configurable globally, per agent, or per user, filter and transform fragments on the way in and out; the implementation runs everything on gpt-4o with text-embedding-3-large retrieval, a coordinator that routes subqueries to specialist agents, and an aggregator that composes the final answer.

Three scenarios carry the evaluation, all at small scale. On MultiHop-RAG (609 news articles across six domains, 2,556 multi-hop questions, five users, six domain agents), fully shared memory holds accuracy above 0.90 while cutting external resource calls by up to 61 percent at 50 percent query overlap and 59 percent at 75 percent overlap, relative to per-user isolated memory. A second scenario uses 200 synthetic business queries under asymmetric roles (only a Strategy Director sees all four agents); with no ground truth available it reports only that partial sharing reduces resource calls for every role. The third uses SciQAG science QA (five categories, five users, 100 queries) with permissions granted through step t4 and revoked through t8: accuracy rises and falls with the access graph, resource calls decline over time as memory substitutes for retrieval, and logged access matrices show no fragment crossing a permission boundary. The authors concede the settings are simulated, the user counts modest, and that LLM-based policy enforcement can still occasionally breach policy.

For a personal knowledge backend built over one person's conversational history, the relevant machinery is the provenance-stamped fragment plus retrospective permission check: it is exactly the primitive needed if a single-user memory tree ever grows organizational tenants, since access can change after fragments exist without rewriting the store. The efficiency result also matters, as it quantifies memory's role as a cache that displaces repeated retrieval. Within the field, the paper marks the point where LLM memory work imported access-control formalism (RBAC and ABAC lineage) rather than assuming one global store, though it establishes a formulation and feasibility evidence, not a benchmark others can yet be measured against.

Read from: full text

From Isolated Conversations to Hierarchical Schemas: Dynamic Tree Memory Representation for LLMs

Rezazadeh, Li, Wei, Bao (Accenture Center for Advanced AI), ICLR 2025, arXiv:2410.14052

MemTree is an online memory system that stores an LLM agent's accumulating history as a dynamically grown tree rather than a flat embedding table. Each node holds aggregated text, an embedding of that text, and links to parent and children; depth corresponds to abstraction, with general summaries near the root and specifics at the leaves. New information is inserted by traversing from the root: at each node, cosine similarity against the children is compared to a depth-adaptive threshold (θ grows exponentially with depth, so deeper, more specific nodes demand closer matches). A sufficiently similar child is descended into; otherwise a new leaf is attached and traversal stops. Every ancestor on the path then has its summary rewritten by an LLM call to fold in the new content, and its embedding refreshed. Insertion is $O(\log N)$, and the authors connect the scheme to online top-down hierarchical clustering, proving a $\beta/3$ approximation of the optimal Moseley-Wang revenue under a well-separatedness assumption. Retrieval collapses the tree: the query embedding is compared against every node at once, filtered by a threshold, and the top-k nodes are returned.

The evaluation, run with GPT-4o plus text-embedding-3-large and separately Llama-3.1-70B, covers four benchmarks. On the 15-round Multi-Session Chat sessions, simply feeding GPT-4o the full history wins (95.6 percent accuracy) since each dialogue is under a thousand tokens; among memory systems MemTree edges MemoryStream (84.8 vs 84.4) and beats MemGPT (70.4). The interesting result comes from MSC-E, the authors' 70-session extension to 200 dialogue rounds: there MemTree (82.5 overall) beats both full-history prompting (78.0) and MemoryStream (80.7), and holds accuracy steady across evidence positions where the naive baseline degrades. On QuALITY single-document QA with Llama-3.1-70B, MemTree (59.8) roughly matches the offline RAPTOR (59.0) despite building its index incrementally. On MultiHop RAG (609 news articles, 2,556 questions), MemTree reaches 80.5 overall, within 0.5 points of RAPTOR and above GraphRAG (78.3), and on temporal queries (68.4) it beats every baseline including human-annotated evidence. The learned tree over those 609 documents has 3,164 nodes, depth 13, branching factor 2.1. Per-item insertion averages 10 seconds; cumulative cost runs about 1.4x the offline methods.

For a personal knowledge backend over conversational history, the central lesson is that the online-versus-offline distinction matters more than the specific structure: RAPTOR and GraphRAG require full corpus rebuilds to absorb new material, while MemTree pays a modest ongoing tax to stay current, which is the right trade for a system ingesting conversations continuously. Two caveats travel with the results: parent summaries are LLM-rewritten on every insertion, so ancestor content drifts with the insertion order and inherently favors recent material (the authors show accuracy rising from 82.1 to 84.2 for late-positioned evidence), and the corpora tested top out at hundreds of documents, well short of organizational scale.

Read from: pages 1-12 of 34 (full main text; remainder is references and appendices)

Knowledge Graph Embedding for Link Prediction: A Comparative Analysis

Rossi, Firmani, Matinata, Merialdo, and Barbosa, ACM TKDD, 2021, arXiv:2002.00819

The paper retrains and re-evaluates 16 embedding-based link prediction models from scratch, spanning three families (tensor decomposition, geometric, and deep learning), against AnyBURL, a rule-mining baseline, on the five standard benchmarks: FB15k, WN18, FB15k-237, WN18RR, and YAGO3-10. Beyond the usual global metrics (Hits@K, mean rank, mean reciprocal rank), the authors define per-fact structural features of the training graph and correlate each with prediction accuracy: the number of “peers” (alternative valid answers seen in training), a TF-IDF style Relational Path Support score for paths up to three hops, relation properties such as symmetry and transitivity, and the degree of the original reified node for facts derived from FreeBase hyperedges.

Three findings carry the paper. First, the rule-based baseline holds its own: AnyBURL ranks at or near the top on almost every dataset and metric while learning its rules in 100 to 1,000 seconds, against roughly 1 to 300 hours of training for the embedding models; among those, only ComplEx with N3 regularization is consistently comparable. Second, graph structure largely determines difficulty: more source peers help (they are worked examples of the answer), more target peers hurt, and higher path support helps nearly every model. On FB15k and WN18 the leakage from inverse relations is so severe that the path-support score alone, used as a scoring function with two-hop paths on a 512-fact sample, reaches 93.55 percent H@1 on WN18. Third, evaluation plumbing matters: the policy for breaking score ties, rarely even reported, produces wildly incomparable numbers. CapsE scores 73.01 H@1 on FB15k under the min policy and 1.93 under average, and a perfect 100 versus 0 on YAGO3-10; the authors trace the ties to saturating activations (ReLU, and sigmoid in its near-flat regions) and show a leaky-ReLU variant mostly closes the gap.

For a personal knowledge backend over conversational history, the transferable lessons are less about any single model than about method. A graph distilled from one person’s conversations will be sparse and skewed, with few peers and thin path support per fact, which is precisely the regime where every tested model performs worst; the cheap, interpretable rule-based approach is the sensible starting point before any embedding machinery. And the tie-policy result is a standing caution for persistent-memory evaluation generally: retrieval benchmarks can flatter or bury a system through unstated protocol details. Scope limit throughout: all results come from graphs of 14k to 123k entities sampled from FreeBase, WordNet, and YAGO, well below organizational scale.

Read from: full text

Bitemporal Property Graphs to Organize Evolving Systems

Rost, Fritzsche, Schons, Zimmer, Gawlick, Rahm; arXiv preprint, 2021; arXiv:2111.13499; a summary report on a one-year Oracle and University of Leipzig cooperation.

This paper reports the outcomes of a joint project between the University of Leipzig and Oracle on organizing relationships within multi-dimensional time-series data, with IoT sensor networks (an aircraft’s instrumented components is the running example) as the motivating case. Rather than a single deep result, it presents four connected artifacts: TPGM+, a bitemporal property graph model; T-PGQL, a temporal extension of Oracle’s PGQL query language; CGN, a continuous event-detection mechanism; and BiTeGra, a prototype database that stores the graphs in a relational system. The paper walks through each in turn, with the query language treated in the most detail.

The core of TPGM+ is bitemporality carried down to the property level. Every vertex, edge, and property value gets two close-open time intervals, valid-time (when the fact held in the world) and transaction-time (when the database learned it), plus five integrity constraints covering uniqueness, referential integrity of edges and of properties, constant edge endpoints, and constant labels. The property-level intervals are the delta over the authors’ earlier TPGM, which had to copy an entire element whenever one property changed. T-PGQL exposes the intervals through dot-notation accessors (`VAL_FROM`, `TX_TIME`, and so on), SQL:2011 period predicates (`CONTAINS`, `OVERLAPS`, `PRECEDES`), a `FOR TX_TIME` clause for snapshot and range time travel with `AS OF`, `FROM TO`, `BETWEEN AND`, and `ALL` variants, and `FIRST/LAST` aggregates. CGN adapts Oracle’s Continuous Query Notification to graphs, distinguishing notifications on any change to matched elements from notifications only when the query result changes. BiTeGra maps graphs to bitemporal relational tables and weighs schema choices explicitly: one big vertex-and-edge table pair versus a table per label, and properties as columns versus a key-value property table, with hybrid schemas (HyVE, HyPe) that record the per-label choice in metadata. The paper presents no benchmarks or evaluation; the language covers retrieval only, not data manipulation, and the system is a prototype.

For a personal knowledge backend over conversational history, the transferable idea is the two-clock design: separating when something was true from when the system recorded it is exactly what distinguishes “he changed jobs in May” from “I learned this in July,” and doing it per property rather than per entity avoids duplicating a whole record when one fact changes. The `FOR TX_TIME AS OF` construction is a ready-made vocabulary for asking what the system believed at a past moment, useful for auditing an agent’s evolving beliefs. The schema-mapping section is also a compact survey of how property graphs sit on relational storage, though the absence of any performance data means the design tradeoffs are argued, not measured.

Read from: full text

LLM-based Multi-Agent Blackboard System for Information Discovery in Data Science

Salemi et al. 2026, LLM-Based Multi-Agent Blackboard System, arXiv:2510.01285

The paper revives the classic blackboard architecture (Hearsay-II lineage, Erman et al. 1980) as a communication paradigm for LLM multi-agent systems, targeting data discovery: answering data science questions when the relevant files sit unidentified inside a large, heterogeneous data lake. Instead of a master agent assigning subtasks to workers it must already understand, the main agent posts a request describing what it needs to a shared blackboard. Helper agents, each responsible for one cluster of the data lake plus one web search agent, monitor the board and volunteer responses only when they judge themselves capable. The data lake is partitioned offline by giving file names to Gemini 2.5 Pro for clustering; each file agent then studies its cluster's structure in advance. The main agent runs a ReAct loop with five actions (plan, reason, execute code, request help, answer) capped at 10 steps, and its final output is a Python program that loads the right files and computes the answer.

Across KramaBench plus discovery-augmented versions of DSBench and DA-Code, the blackboard beats DS-GRU (all files in context), RAG over the lake (E5-large, top 5 files), and a master-slave variant that is otherwise identical, with 13 to 57 percent relative end-to-end gains over the best baseline. The pattern holds across four backbones (Gemini 2.5 Pro and Flash, Claude 4 Opus, Qwen3-Coder 30B): with Gemini 2.5 Pro, macro average rises from 24.0 (master-slave) to 28.5. File discovery F1 reaches 0.561 versus 0.513 for master-slave and 0.247 for RAG, again Gemini 2.5 Pro only. The advantage grows with lake size: the two largest lakes, DA-Code (145 files) and KramaBench Astronomy (1,556 files), show 70 and 211 percent relative gains over master-slave. Costs are real: roughly 2.3 times RAG's per-question spend, though runtime stays comparable (132 to 145 seconds across all three systems on a 50-question sample). Content-embedding clustering outperforms filename clustering, and more clusters generally help.

For a personal knowledge backend over conversational history, the transferable finding is architectural: routing by self-selection beats routing by a central index of who knows what, precisely because a coordinator's model of each partition goes stale. Partition the corpus, let each partition's agent pre-digest its holdings offline, broadcast queries, and let holders volunteer. That pattern sidesteps the retrieval weakness the authors document (embedding retrieval collapses to 0.247 F1 on this heterogeneous data) and scales in the direction an organization's archive grows. The caveats travel with it: results cover data lakes of at most 1,556 files, tabular and scientific domains, and a fixed two-level hierarchy.

Read from: pages 1-14 of 44 (full main text; remainder is references and appendices)

RAPTOR: Recursive Abstractive Processing for Tree-Organized Retrieval

Sarathi, Abdullah, Tuli, Khanna, Goldie, Manning; ICLR 2024; arXiv:2401.18059

RAPTOR replaces the flat chunk store of standard retrieval augmentation with a summary tree built from the bottom up. The corpus is split into 100-token chunks (whole sentences preserved), embedded with SBERT, and then soft-clustered with Gaussian Mixture Models over UMAP-reduced embeddings; cluster count is chosen by BIC, and a chunk can join more than one cluster. Each cluster is summarized by gpt-3.5-turbo, the summaries are re-embedded, and the cycle repeats until clustering is no longer feasible, yielding a tree whose leaves are raw text and whose upper nodes are progressively broader abstractions. Because clustering runs on embeddings rather than document position, distant but related passages can share a parent, which is what separates RAPTOR from adjacency-based hierarchies like LlamaIndex or the recursive summarization of Wu et al. (2021). Construction cost scales linearly in tokens. At query time, the authors' preferred strategy simply collapses the tree: every node at every layer competes in one cosine-similarity search, and nodes are added until a token budget is reached (2,000 tokens, roughly the top 20 nodes, chosen after collapsed tree beat layer-by-layer traversal on 20 QASPER stories).

In controlled comparisons with UnifiedQA-3B as reader, adding the tree improved every retriever tested (SBERT, BM25, DPR) on all three benchmarks: NarrativeQA (1,572 book and movie transcripts), QASPER (5,049 questions over 1,585 NLP papers), and QuALITY (multiple-choice over passages averaging 5,000 tokens). With stronger readers the gains held: on QASPER, RAPTOR reached F-1 of 53.1 with GPT-3 and 55.7 with GPT-4, beating DPR by 1.8 to 4.5 points and edging past long-context models like CoLT5 XL (53.9). The headline result is QuALITY with GPT-4: 82.6 percent accuracy against a prior best of 62.3, and a 21.5-point lead over CoLISA on the hard subset. All corpora are single long documents or paper collections, not open-domain web scale, and an annotation study found about 4 percent of generated summaries contained minor hallucinations, though these did not propagate upward or measurably hurt QA. A layer ablation on one QuALITY story showed full-tree search (73.68) beating any single layer (about 58), evidence that the abstraction hierarchy itself, not just better chunking, does the work.

For a memory backend over conversational history, the transferable idea is that recall should span abstraction levels: some queries need the raw utterance, others need the season-of-work summary, and a collapsed search over both lets the query pick its own granularity. RAPTOR is also a batch-built, static index; conversations accrete, so an incremental variant of its cluster-summarize loop is the open engineering question, and its 4 percent summary hallucination rate is a caution for any system whose upper layers become the record people trust.

Read from: pages 1-12 of 23 (full main text; references and appendices skipped)

TraceMem: Weaving Narrative Memory Schemata from User Conversational Traces

Yiming Shu et al., arXiv:2602.09712, 2026-02

Abstract (verbatim)

Sustaining long-term interactions remains a bottleneck for Large Language Models (LLMs), as their limited context windows struggle to manage dialogue histories that extend over time. Existing memory systems often treat interactions as disjointed snippets, failing to capture the underlying narrative coherence of the dialogue stream. We propose TraceMem, a cognitively- inspired framework that weaves structured, narrative memory schemata from user conversational traces through a three-stage pipeline: (1) Short-term Memory Processing, which employs a deductive topic segmentation approach to demarcate episode boundaries and extract semantic representation; (2) Synaptic Memory Consolidation, a process that summarizes episodes into episodic memories before distilling them alongside semantics into user-specific traces; and (3) Systems Memory Consolidation, which utilizes two-stage hierarchical clustering to organize these traces into coherent, time-evolving narrative threads under unifying themes. These threads are encapsulated into structured user memory cards, forming narrative memory schemata. For memory utilization, we provide an agentic search mechanism to enhance reasoning process. Evaluation on the LoCoMo benchmark shows that TraceMem achieves state-of-the-art performance with a brain- inspired architecture. Analysis shows that by constructing coherent narratives, it surpasses baselines in multi-hop and temporal reasoning, underscoring its essential role in deep narrative comprehension. Additionally, we provide an open discussion on memory systems, offering our perspectives and future outlook on the field. Our code implementation is available at: <https://github.com/YimingShu-teay/TraceMem>

Bearing on this project

This is the closest published work to the whole design. Segmentation into episodes, episodic summaries, distillation into traces, hierarchical clustering into time-evolving narrative threads, memory cards, agentic search. It occupies the thread idea end to end and is SOTA on LoCoMo. The one visible difference is that TraceMem CLUSTERS traces into threads, which is a partition (a trace belongs to one thread), where this project places one unit in N threads at full weight. That distinction is thin, unverified, and the only thing standing between this design and being a reimplementation.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Developing Time-Oriented Database Applications in SQL

Richard T. Snodgrass, Morgan Kaufmann (Series in Data Management Systems), 1999, book, 528 pp. in this PDF

This is an orientation to a book, not a summary of a paper, and it draws on the parts that state the argument: the preface, chapter 1, and the opening of the bitemporal chapter. Snodgrass translates roughly 1600 research papers on temporal databases into guidance for people writing ordinary SQL. The preface opens with a development story that reads like a confession. A team adds one DATE column, finds the primary key no longer holds, adds a second DATE column and inherits a crop of off by one bugs where a comparison should have been less than or equal, or where a query needs a “+ 1 DAY” correction, then learns that two reports run days apart cannot be reconciled at all. The claim is that this is a category error, not sloppiness. The teaching runs through case studies from real installations: oil field records, cattle location data, a Danish mortgage bank.

The organizing scheme is three orthogonal triples. The temporal data types are instant, interval, and period. The kinds of time are user defined time (an uninterpreted value), valid time (when a fact was true in the modeled world), and transaction time (when a fact was stored). The kinds of statement are current, sequenced, and nonsequenced, and that split applies to queries, modifications, views, and constraints. Snodgrass is blunt about tool support as of 1999. SQL-92 offers DATE, TIME, TIMESTAMP, and INTERVAL but nothing for valid or transaction time; sequenced statements, which he calls the most useful case, get no support whatsoever; SQL3's Part 7, SQL/Temporal, which adds a PERIOD constructor, was not expected to reach ballot before the next millennium. The rest falls to the developer, with code fragments tested against DB2 UDB, Informix, Access, SQL Server, Sybase, Oracle8, and UniSQL, none fully compliant.

The bitemporal chapter is what bears on memory systems. Nykredit, a Danish mortgage bank carrying nine million loans against eight million customers and seven million properties, tracks both times on its property ownership table, giving each row six columns: two foreign keys, VT_Begin and VT_End, TT_Start and TT_Stop. The payoff is a repair procedure. When someone reports bad data, staff find when the error was stored, roll the table back to that transaction time, read the valid-time history as it then stood, and fix it; because transaction time is append only, the fix is itself logged. That is the operation a conversational memory backend needs when a stored belief turns out to be wrong. The costs are equally plain: a sequenced primary key cannot be expressed as a PRIMARY KEY constraint and needs an assertion, modifications multiply into nine forms, and the discipline lives in the application, not the engine.

Read from: pages 17-33 and 301-308 of 528

Cognitive Architectures for Language Agents

Sumers, Yao, Narasimhan, Griffiths 2024, Cognitive Architectures for Language Agents, TMLR, arXiv:2309.02427

This is a conceptual paper, not an experimental one. The authors argue that LLM-based agents recapitulate a sixty-year lineage running from Post’s production systems through Markov algorithms to cognitive architectures like Soar, with an LLM playing the role of a probabilistic production system: where a production rewrites string XYZ to XWZ, an LLM defines a distribution over completions of a prompt. On that analogy they build CoALA, a framework that describes any language agent along three dimensions: memory (working memory plus long-term episodic, semantic, and procedural stores), an action space split into external grounding actions and internal retrieval, reasoning, and learning actions, and a decision cycle that plans (propose, evaluate, select) and then executes one grounding or learning action per loop.

The paper’s product is the organizing power of that scheme rather than any benchmark number. Table 2 casts prominent agents into the framework and exposes real structural differences: SayCan is evaluation-only over a fixed set of 551 robot skills with no internal actions; ReAct adds reasoning but has no long-term memory or learning; Voyager exercises all four action types, learning by writing code skills into procedural memory; Generative Agents write LLM reflections over episodic memory into semantic memory; Tree of Thoughts has elaborate propose-evaluate-select decision-making but no persistence at all. From the gaps the survey exposes, the authors derive directions: learned retrieval procedures are essentially unstudied, deletion and modification of memory (unlearning) is neglected, and almost no agent treats learning as a deliberately chosen action rather than a fixed schedule. All of this is a reading of the 2022-2023 agent literature, so its claims about what is understudied are claims about that snapshot, not measurements.

For a personal knowledge backend built over conversational history, the useful contribution is the memory typology and its access rules. The paper’s worked example is close to that exact problem: a retail assistant with read-only semantic memory (the inventory), read-write episodic memory (past interactions), and episodes summarized by the LLM before storage rather than logged raw. Reflection, in the Generative Agents sense, is the mechanism that turns accumulated episodes into reusable semantic facts, which is the step a conversation archive needs to scale beyond retrieval. The paper also gives the field a shared vocabulary; most persistent-memory systems since can be read as instantiations of its episodic-semantic-procedural split. Its caveat is worth keeping too: writing to procedural memory (an agent’s own code or prompts) is flagged as the riskiest form of learning, and longer context windows may shift how much long-term memory matters at all.

Read from: full text

When Good OCR Is Not Enough: Benchmarking OCR Robustness for Retrieval-Augmented Generation

Lin Sun et al., arXiv:2605.00911, 2026-04

Abstract (verbatim)

Industrial Retrieval-Augmented Generation (RAG) systems depend on optical character recognition (OCR) to transform visual documents into text. Existing OCR benchmarks rely on character-level metrics, which inadequately measure downstream RAG effectiveness under real-world conditions. We introduce an OCR benchmark for industrial RAG systems covering 11 challenging document types, including extreme layouts, high-resolution pages, complex or watermarked backgrounds, historical documents with non-standard reading orders, visually decorated text, and documents containing tables and mathematical formulas. Evaluating recent SOTA OCR models under a controlled OCR-first RAG pipeline shows clear performance degradation on realistic industrial documents despite strong conventional benchmark scores. We find that high OCR accuracy does not necessarily translate into strong downstream RAG performance: structural and semantic errors can cause substantial retrieval failures even when WER/CER remains low. Further analysis shows that this mismatch is category-dependent, arises through both retrieval-side and downstream generation-side failures, and remains stable across representative OCR-first pipeline choices. The benchmark is publicly available at <https://github.com/Qihoo360/InduOCRBench>.

Bearing on this project

Second paper on OCR robustness for RAG, 2026. Confirms the line is populated rather than a one-off.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Toward a Theory of Hierarchical Memory for Language Agents

Talebirad, Parsaee, Szepesvari, Nadiri, and Zaiane; ICLR 2026; arXiv:2603.21564

This is a theory paper, not an empirical one. The authors observe that a wide range of long-context and agentic memory systems, from document-hierarchy retrievers like RAPTOR and GraphRAG to conversational memories like xMemory and H-MEM to agent execution-trace managers like MemoBrain and StackPlanner, all implement the same underlying pipeline, and they formalize it as three operators. Extraction (α) maps raw data to a graph of atomic information units; coarsening (C), a pair of a grouping function π and a representative function ρ , partitions units and assigns each group a compressed representative, iterated to build a multi-level hierarchy; traversal (τ) takes a query and a token budget and returns a subset of atoms. They prove simple information-theoretic facts about this structure: each deterministic, lossy coarsening level strictly decreases entropy, and by the data processing inequality mutual information with any query is monotonically nonincreasing up the hierarchy, so the atom layer sets the information ceiling for everything above it.

The central contribution is the self-sufficiency spectrum and the coarsening-traversal coupling. Self-sufficiency, $SS = I(G; \rho(G)) / H(G)$, measures how much of a group's information its representative preserves, from LLM summaries near 1 to bare category labels near 0. The coupling claim is that this value dictates viable retrieval: self-sufficient representatives support collapsed search over all levels at once (RAPTOR's pattern), while referential ones require top-down refinement (H-MEM, xMemory); mismatching the two wastes budget. The authors state plainly that this is consistent with existing evidence but has not been tested in a controlled comparison. Appendices add a computable LLM-based proxy SS_{θ} , a Fano bound capping branching factor at $2^{((B+1)/(1-\epsilon))}$ given B bits of routing information, an optimal branching factor of e under uniform cost, a depth bound $L \geq \lceil \log_r(N/C) \rceil$ tying corpus size N and context window C to required levels, and a coherence condition on π (within-group affinity must exceed between-group affinity) that justifies top-down pruning. The whole framework assumes a static hierarchy; insertion, restructuring, and retrieval-driven decay break the Markov argument, which the authors name as the most pressing open problem.

For a personal knowledge backend built over conversational history, the paper supplies a design vocabulary rather than results: it says to choose summary style and search strategy jointly, that finer-grained extraction raises the retrieval ceiling, and that the depth bound predicts when one summary layer stops sufficing as a corpus grows toward organizational scale. Table 1's mapping of eleven systems onto (α , C , τ) is also a compact survey of the persistent-memory field as of early 2026.

Read from: full text

Let Me Speak Freely? A Study on the Impact of Format Restrictions on Performance of Large Language Models

Zhi Rui Tam, Cheng-Kuang Wu, Yi-Lin Tsai, Chieh-Yen Lin, Hung-yi Lee, Yun-Nung Chen. EMNLP 2024 Industry Track (arXiv:2408.02442, cs.CL 2024).

This paper asks whether forcing a language model to answer in a structured format degrades the quality of what it says, not just how it says it. The authors compare three progressively looser regimes on the same tasks: constrained decoding (JSON-mode), format-restricting instructions for JSON, XML, and YAML schemas, and a two-step NL-to-Format conversion where the model answers freely and then reformats. They test three reasoning datasets (GSM8K, Last Letter, Shuffled Objects) and four classification datasets across GPT-3.5-Turbo, Claude-3-Haiku, Gemini-1.5-Flash, LLaMA-3-8B, and Gemma-2-9B, averaging over nine prompt variants and using an LLM as a near-perfect answer parser validated against human annotation.

The central finding is that stricter constraints hurt reasoning more. On GSM8K, adding a schema to the JSON instruction drops Claude-3-Haiku from 86.99 to 23.44 and GPT-3.5-Turbo from 74.70 to 49.25 while inflating variance. The damage is not a parsing artifact: LLaMA-3-8B shows a 38.15 point Last Letter gap in JSON at a parsing error rate of only 0.148 percent. Mechanisms matter too; on Last Letter, JSON-mode placed the answer key before the reasoning key in 100 percent of GPT-3.5-Turbo responses, silently converting chain-of-thought into direct answering. Classification is the counterpoint: JSON-mode often matches or beats free text there, apparently by restricting the answer space. NL-to-Format roughly preserves natural-language accuracy, and a cheap corrective reformatting call repairs most parsing failures. All of this comes from small, low-cost models on simple two-field schemas; the authors state they could not test GPT-4o or 70B-class models, and gpt-4o-mini with OpenAI's stricter JSON-Schema mode narrows but does not close the gap (94.57 free text versus 91.71 on GSM8K).

For the memory harness, this governs the schema-validated LLM call layer that every extraction stage shares: distilling the corpus into tree nodes, building entity records, and emitting dated fact rows all demand structured output from swappable model tiers, and this paper says the schema itself can be the failure cause. A stage that collapses under a cheap constrained model may be reacting to the constraint, not the model, so the stage-wise sensitivity results need a constrained-versus-prompted ablation before blaming the tier, and the paper's regimes supply the arms. Its practical guidance maps cleanly: reasoning-heavy distillation favors loose formats or NL-to-Format with a repair pass, while classification-shaped steps such as entity linking can safely keep strict schemas. The caveat travels with the result: evidence covers small models and two-field schemas only, close to but not identical with the harness's richer record shapes.

Read from: full extracted text of the 19-page PDF (main body pages 1-8 plus the limitations section and appendices B through E).

Balancing Novelty and Saliency: Adaptive Learning to Rank Entities for Timeline Summarization of High-impact Events

Tuan Tran et al., arXiv:1701.03947, 2017-01

Abstract (verbatim)

Long-running, high-impact events such as the Boston Marathon bombing often develop through many stages and involve a large number of entities in their unfolding. Timeline summarization of an event by key sentences eases story digestion, but does not distinguish between what a user remembers and what she might want to re-check. In this work, we present a novel approach for timeline summarization of high-impact events, which uses entities instead of sentences for summarizing the event at each individual point in time. Such entity summaries can serve as both (1) important memory cues in a retrospective event consideration and (2) pointers for personalized event exploration. In order to automatically create such summaries, it is crucial to identify the “right” entities for inclusion. We propose to learn a ranking function for entities, with a dynamically adapted trade-off between the in-document saliency of entities and the informativeness of entities across documents, i.e., the level of new information associated with an entity for a time point under consideration. Furthermore, for capturing collective attention for an entity we use an innovative soft labeling approach based on Wikipedia. Our experiments on a real large news datasets confirm the effectiveness of the proposed methods.

Bearing on this project

2017, and it matters because it does entity SALIENCE ranking for timeline summarization - which is this project's promotion threshold, asked and answered nine years earlier in a different domain.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Wikidata: A Free Collaborative Knowledgebase

Vrandečić and Krotzsch, Communications of the ACM 57(10), 2014, DOI 10.1145/2629489

This is a systems and design article by Wikidata’s project director and its data model lead, describing the collaboratively edited knowledgebase Wikimedia launched in October 2012 to hold Wikipedia’s factual data in one place. The motivating problem is duplication: the same fact (the population of Rome, say) lived in hundreds of independently edited language articles and disagreed across them, while the underlying data, spread over 30 million articles in 287 languages, was nearly impossible to extract. The article walks through the design decisions (open editing, community-controlled schema, plurality of conflicting claims, secondary rather than primary data, multilingual by construction), the data model, and early growth statistics, and situates the project against Semantic MediaWiki, Freebase, OpenCyc, DBpedia, and YAGO.

The core representational contribution is a layered statement model rather than plain triples. Items carry property-value pairs, with a small set of datatypes (quantity, item, string, time, coordinates, URL); properties are themselves first-class pages with labels and declared datatypes, and the community, not a fixed ontology, decides which properties exist. Statements take subordinate property-value pairs called qualifiers, so “population of Rome = 2,651,040” can carry “as of 2010” and “method: estimation” attached to the assertion rather than the item. Every claim can carry references, themselves lists of property-value pairs, and contributors can mark statements preferred or deprecated to manage contradictory sources; the model also distinguishes “value unknown” and “no value” from mere incompleteness. By August 2014 this held 15.8 million items, 43.2 million statements (23.2 million with references), and 1,176 properties, with property use heavily skewed toward “instance of.” The numbers document adoption, not data quality; the authors note the data is far from complete and complex query support did not yet exist.

For a personal knowledge backend built over conversational history, the transferable pieces are the qualifier-and-reference structure (every stored fact carries its temporal context and its source, which is exactly the provenance a memory system needs when conversations contradict each other), the preferred/deprecated mechanism for coexisting conflicting claims instead of forced resolution, and stable never-reused IDs decoupled from labels. The bootstrapping story also matters: Wikidata gained its initial item set by importing an existing artifact (language links), and about 90 percent of its half million daily edits came from bots, with human curation layered on top, a plausible division of labor for automated extraction plus owner review. The scope limit is that Wikidata’s schema governance depends on a large volunteer community; a single-user or single-organization system inherits the representation but must supply curation by other means.

Read from: full text

Large Language Models are not Fair Evaluators

Peiyi Wang et al., ACL 2024 (long paper), 2024, arXiv:2305.17926 / aclanthology.org/2024.acl-long.511

This paper demonstrates that the common practice of asking an LLM to compare two candidate responses is fragile in a specific, exploitable way: the ranking depends on which answer appears first in the prompt. The authors take the standard Vicuna benchmark template, in which GPT-4 or ChatGPT scores two assistant responses on a 1 to 10 scale, and simply swap the order of the responses. The verdict frequently flips even though the template explicitly instructs the judge to ignore presentation order. They quantify this with a conflict rate, the fraction of examples on which the two orderings yield contradictory verdicts, then propose three calibrations and validate them against 80 questions hand-annotated win/tie/lose by three of the authors.

The numbers are stark. With ChatGPT as judge comparing Vicuna-13B against ChatGPT, Vicuna's win rate is 2.5 percent in the first slot and 82.5 percent in the second, a conflict rate of 82.5 percent; GPT-4 shows the opposite lean, favoring the first position, at a 46.3 percent conflict rate. Bias concentrates where responses are close in quality: score gaps of 1 or less flip often, gaps of 3 or more rarely do. The fixes are Multiple Evidence Calibration (generate the explanation before the score, sample k times; k of 3 is the sweet spot), Balanced Position Calibration (run both orderings and average the 2k scores), and Human-in-the-Loop Calibration (an entropy score over the win/tie/lose outcomes flags unstable items for human review). MEC plus BPC lifts agreement with the human majority from 52.7 to 62.5 percent for GPT-4 and from 44.4 to 58.8 percent for ChatGPT; escalating the top 20 percent highest-entropy items to humans reaches roughly human-average accuracy at 39 percent less annotation cost. The scope is narrow, though: one 80-question benchmark, two judge models via the OpenAI API, pairwise templates only, and no account of why the bias exists.

For the memory harness this governs the LLM-judge scoring stage, the step that grades answers produced by the three injection arms after the corpus has been distilled into the tree of entities and dated facts. Any pairwise comparison across arms must run both orderings through the schema-validated judge call and average, whatever model tier is doing the judging, and single-order verdicts on close pairs should be treated as noise. The entropy-based escalation maps directly onto the small hand-labeled calibration set: route only the low-agreement items to Justin instead of labeling uniformly.

Read from: full text of the paper, all 11 pages including the conclusion and limitations, extracted with pdfplumber.

NovelQA: Benchmarking Question Answering on Documents Exceeding 200K Tokens

Cunxiang Wang et al., ICLR 2025 (arXiv 2024), arXiv:2403.12766

NovelQA is a long-context question-answering benchmark built from 89 English novels (65 public domain, 24 copyrighted) averaging over 200,000 tokens, well past the 60K-token ceiling of prior long-range datasets. Expert annotators, mostly English literature students who had already read the books they chose, wrote 2,305 question tuples, each carrying a question, a golden answer, four multichoice options, and located textual evidence from the novel. Half the questions came from ten-plus templates targeting known LLM weaknesses; the rest were free-form. Questions divide into multi-hop (35.0 percent), single-hop (42.8 percent), and detail (22.2 percent) complexity bands and seven aspect types (times, meaning, span, setting, relation, character, plot). Quality control kept only 79.4 percent of submissions and an inter-annotator test on multichoice answers hit 94.6 percent Cohen's kappa.

The evaluations establish that even 128K-200K context models fall well short of readers. Claude-3.5-Sonnet, the best model tested, reached 62.30 percent generative accuracy against a 90 percent human baseline (97 percent multichoice); Llama-3.1-70B managed 51.50 percent. Accuracy collapses when evidence sits past the 100K-token mark: the weighted generative average drops from 46.84 to 25.36 percent, a decline at the end of very long texts rather than the expected lost-in-the-middle dip. Generative scoring uses GPT-4 as judge, validated at 89.25 percent kappa against two human evaluators on 800 outputs, a defensible arrangement here because the questions are objective and evidence-grounded. Two caveats travel with these numbers: close-book runs show models already know the famous novels (GPT-4 scores 60.94 percent multichoice with no text at all), so open-book scores partly measure memorization, and the benchmark is English-only, novels-only, with truncation from the front when inputs overflow.

For the harness, NovelQA governs the question-design and evidence-annotation stage of the three injection arms. Its per-question evidence location is the template to copy when the corpus-to-tree ingest builds entity and dated-fact records: every probe question in the counterfactual and corrected-facts bands should carry a golden answer plus a pointer to the exact source span, so that a miss can be traced to retrieval versus synthesis. Its complexity-by-aspect taxonomy gives the schema for the schema-validated question-generation calls across model tiers, and its judge validation protocol (objective questions, kappa against human raters on a sample) is the procedure the LLM-judge scoring arm should reproduce before its scores are trusted. The close-book control is equally worth copying: without it, contamination masquerades as memory.

Read from: extracted text of pages 1-15 of the 23-page PDF, covering the full main paper (abstract through conclusion and references) plus the limitations, ethics, and opening appendix sections.

Searching for Best Practices in Retrieval-Augmented Generation

Wang et al., EMNLP 2024, arXiv:2407.01219

The paper treats a RAG pipeline as a sequence of swappable modules (query classification, chunking, embedding, vector database, retrieval, reranking, repacking, summarization, generator fine-tuning) and searches for the best method at each position. Since testing every combination is infeasible, the authors run a greedy three-step procedure: compare candidate methods per module in isolation, then optimize one module at a time within a full pipeline while holding the others fixed, then assemble recommended configurations. The end-to-end evaluation uses a Llama2-7B-Chat generator over a Milvus index of 10 million English Wikipedia passages plus 4 million medical texts, scored across commonsense reasoning, fact checking, open-domain QA, multihop QA, and medical QA, with RAGAs-style RAG metrics, subsampled to at most 500 examples per dataset.

The per-module findings come with specific conditions. Sentence-level chunking at 512 tokens scored best for faithfulness (97.59) on a corpus of just sixty pages of a Lyft 10-K, so the chunk-size result is thin evidence. LLM-Embedder matched BAAI/bge-large-en on MS MARCO retrieval at a third the size. On TREC DL19/20, HyDE plus hybrid search (BM25 weighted at alpha 0.3 against dense scores) gave the best retrieval quality, while query rewriting and decomposition did not help. For reranking on MS MARCO, monoT5 balanced quality and cost; TILDEv2 was 225 times faster (0.02 versus 4.5 seconds per query) at some accuracy loss. In the full pipeline, query classification (a BERT classifier deciding whether to retrieve at all, 0.95 accuracy) raised the average score from 0.428 to 0.443 and cut latency from 16.41 to 11.58 seconds per query. Reverse repacking, putting the most relevant document nearest the query, beat forward and sides. Fine-tuning the generator on a mix of one relevant and one random document made it most resistant to retrieval noise. The two recipes: best performance uses Hybrid with HyDE, monoT5, reverse, Recomp (average 0.483, about 11.7 s/query); the balanced recipe swaps in plain hybrid retrieval and TILDEv2, cutting latency to roughly 1.5 seconds at similar quality.

For a conversational-history backend, the transferable lessons are the deciding-when-to-retrieve gate (the single cheapest win in both accuracy and latency), hybrid sparse-plus-dense retrieval with a light reranker, and ordering context so the best evidence sits nearest the query. The caveat matters for memory work generally: these numbers come from public QA corpora and a 7B generator, and the paper itself concedes chunking was only lightly explored, so the recipes are starting priors rather than settled law for personal or organizational memory stores.

Read from: pages 1-16 of 22

CogniFold: Always-On Proactive Memory via Cognitive Folding

Suli Wang et al., arXiv:2605.13438, 2026-05

Abstract (verbatim)

Existing agent memory remains predominantly reactive and retrieval-based, lacking the capacity to autonomously organize experience into persistent cognitive structure. Toward genuinely autonomous agents, we introduce CogniFold, a brain-inspired “always-on” agent memory designed for the next generation of proactive assistants. CogniFold continuously folds fragmented event streams into self-emerging cognitive structures, bootstrapping progressively higher-level cognition from incoming events and accumulated knowledge. We ground this by extending Complementary Learning Systems (CLS) theory from two layers (hippocampus, neocortex) to three, adding a prefrontal intent layer. Emulating the prefrontal cortex as the locus of intentional control and decision-making, CogniFold achieves this through graph-topology self-organization: cognitive structures proactively assemble under the stream, merge when semantically similar, decay when stale, relink through associative recall, and surface intents when concept-cluster density crosses a threshold. We evaluate structural formation using CogEval-Bench, demonstrating that CogniFold uniquely produces memory structures that match cognitive expectations and concept emergence. Furthermore, across eight downstream benchmarks – two probing long-term conversational memory (LoCoMo, LongMemEval) and six spanning other cognitive domains – we validate that CogniFold simultaneously performs robustly on conventional memory tasks. Our code is available at <https://github.com/OpenNorve/CogniFold>.

Bearing on this project

The architecture scoop. An always-on proactive agent memory that continuously FOLDS event streams into self-emerging structures, grounded by extending Complementary Learning Systems theory - the same theoretical spine this project identified on 08-19 - with similarity merging, staleness decay, and intent surfacing when concept-cluster density crosses a threshold. Ships code, evaluates on LoCoMo and LongMemEval. The only visible difference from this design is the trigger: concept-cluster density versus a session-boundary hook.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Same Author or Just Same Topic? Towards Content-Independent Style Representations

Wegmann, Schraagen, Nguyen 2022, Same Author or Just Same Topic?, RepL4NLP 2022, arXiv:2204.04907

The paper attacks a confound in learning writing-style embeddings: models trained on authorship verification (AV, deciding whether two texts share an author) can score well by memorizing what an author writes about rather than how. The authors propose two changes to the training task. First, a contrastive AV setup (CAV): instead of a binary pair, the model sees an anchor utterance and must pick which of two candidates shares its author, trained with a triplet loss. Second, content control (CC): the different-author distractor is sampled from the same conversation as the anchor, from the same domain (subreddit), or at random, so that at the strictest level topic can no longer separate authors. They fine-tune siamese BERT, cased BERT, and RoBERTa models (RoBERTa consistently won on dev) on 210k training tasks built from a 2018 Reddit sample of 100 English-language subreddits, 600 conversations each, with a 70/15/15 author-disjoint split.

Conversation-controlled tasks are the hardest form of AV: every model scores lowest on that test set (around .69 AUC for models trained on it) and highest on the random one (up to .79 for the no-control model), consistent with the distractors carrying fewer content cues. The decisive evidence comes from a new STEL-Or-Content task, which pits a same-style paraphrase against a same-content, different-style sentence. Base RoBERTa picks style only 5 percent of the time; the CAV model with conversation control reaches .42 accuracy on the formal/informal dimension, against .24 or less for domain or no control. The authors argue this is not mere punishment of lexical overlap, since the same model holds even AV performance across all three test sets and normal STEL performance. Agglomerative clustering of its embeddings (7 clusters, 14,756 utterances) surfaces concrete stylistic signatures: terminal-punctuation habits, lowercase i, contraction spelling, apostrophe versus right single quotation mark, line breaks. All of this is one model family on one platform in one language; the conversation label also does not transfer directly to non-conversational corpora, a limit the paper states itself.

For a personal knowledge backend built over conversational history, the operative lesson is that “same source” signals are contaminated: any retrieval or user-modeling layer that embeds utterances will silently entangle who is speaking with what they discuss, and disentangling requires deliberately constructed hard negatives, here obtained for free from conversation structure. That trick, mining supervision from the thread a message sits in, scales to organizational chat archives without labeling. The paper also models an evaluation discipline that persistent-memory work mostly lacks: do not trust the training metric; build a probe where the shortcut and the target answer disagree.

Read from: full text

Machine Knowledge: Creation and Curation of Comprehensive Knowledge Bases

Weikum, Dong, Razniewski, Suchanek; Foundations and Trends in Databases, 2021; arXiv:2009.11564

This is a 289-page survey of how large knowledge bases (KBs) are built and maintained automatically from web and text sources. The authors organize the field into a pipeline of four phases: discovery (source selection, entity detection, taxonomy construction), canonicalization (entity linking and matching), augmentation (extracting attributes and relations, then growing an open schema when the fixed one runs out), and cleaning (constraint reasoning, rule mining, embeddings, provenance and temporal scoping). Each phase gets a methods chapter, ranging from hand-written regex patterns and distant supervision through CRFs, graph algorithms, and transformer-based extractors, followed by case studies of YAGO, DBpedia, NELL, Wikidata, and industrial graphs at Google and Amazon.

What the survey establishes is less a single result than a set of engineering conclusions drawn across two decades of systems. The quality bar they set is error rates below 5 percent, ideally under 1 percent, and they show the two goals of correctness and coverage pull in opposite directions: conservative extraction from premium sources (Wikipedia, IMDB, MayoClinic) gets the low-hanging fruit, then that initial KB seeds distant supervision for riskier open extraction. Public KBs hold millions of entities and up to billions of statements; proprietary ones run one to two orders of magnitude larger. Entities and types must come first, because attributes and relations are only as good as the canonicalized backbone under them. And no KB construction is a one-shot task: it is a life-cycle with humans (architects and curators) permanently in the loop, since no single method wins across the precision, recall, and cost trade-offs.

Two threads bear directly on a personal memory backend over conversational history. First, the wrap-up names the personal KB as an open research problem: an “augmented memory” on the user’s own device, built from a longitudinal digital footprint, with unresolved questions about what to capture, how to contextualize statements for interpretability, and how users manage the life-cycle. That is nearly a spec for a conversation-derived memory tree, and it inherits the survey’s machinery (canonicalization so one entity has one node, temporal scoping so facts age, completeness assessment so the system knows what it does not know). Second, the language-model discussion is a useful caution: they show a 2021-era model completes “Paris is located on the banks of the [?]” correctly but fails on Dylan’s protest songs, results they call “sort of correct, but completely useless,” which is their argument that latent model knowledge cannot replace an explicit, curated store. Scoped to pre-GPT-3-era models, but the structural point about verifiable symbolic memory versus parametric recall still frames the field.

Read from: pages 1-20 and 229-236 of 289

Efficient Guided Generation for Large Language Models

Brandon T. Willard, Rémi Louf; arXiv (cs.CL), 2023; implemented as the Outlines library; arXiv:2307.09702

Willard and Louf reformulate constrained text generation as transitions between the states of a finite-state machine. The standard way to force an LLM's output to match a regular expression or context-free grammar is to scan the entire vocabulary at every decoding step, zeroing the logits of tokens that would break the constraint; that costs $O(N)$ per token, with N (vocabulary size) commonly above 50,000. Their alternative precomputes an index. Before generation starts, every vocabulary string is walked through the constraint's FSM from every viable start state (their Algorithm 3), producing a map from FSM states to the exact set of tokens acceptable in each state (Algorithm 4). At decode time the sampler tracks the current FSM state and looks up the valid-token mask in a hash map, so masking costs $O(1)$ on average. The paper then extends the scheme past regular expressions to context-free grammars via pushdown automata and augmented LALR(1) parser operations, with a trie indexing parser stack configurations, which reaches formats like JSON, Python, and SQL.

The empirical claims are modest but pointed. Against the Guidance library on GPT-2 ($N = 50,257$), Outlines holds near-flat runtime as generated length grows while Guidance climbs past 100 seconds by 60 tokens; the comparison is a single timed sample per setting, and the authors themselves hedge about configuration oversights, so read it as an asymptotic argument rather than a benchmark. The memory trade is quantified once: naive un-reduced indices for an augmented Python grammar run about 50 MB. Worked examples on GPT2-medium show the character of the guarantee, and its limit: a regex-guided model always emits a well-formed year or IP address, but the values (1952 for Chomsky's birth year, 2.2.6.1 for Google DNS) are wrong. Structure is guaranteed; truth is not.

For the memory harness this is the contract mechanism at the extraction stage, the step that turns corpus text into tree nodes, entity records, and dated fact rows through schema-validated LLM calls. Because the FSM index is built from the schema and the vocabulary alone, the method is model agnostic, which is exactly what the swappable model tier design needs: when the local 7B is served through Ollama with Outlines-style enforcement, schema-invalid output becomes a property switched off at decode time rather than an error rate to count, and retry loops disappear from that stage. It says nothing about the downstream arms: injection strategy and LLM-judge scoring still measure semantic quality, since guided decoding guarantees only that every record parses, never that its contents are right.

Read from: the full 18-page paper (all sections, pages 1-18).

Recursively Summarizing Books with Human Feedback

Wu et al. (OpenAI Alignment team), arXiv preprint, 2021, arXiv:2109.10862

The paper trains GPT-3 models (6B and 175B parameters, 2048-token context) to summarize entire fiction novels by fixed recursive task decomposition: a chunking algorithm splits the book, the model summarizes each chunk, then summarizes concatenations of those summaries, recursing to depth 3 for books of hundreds of thousands of words. Each task is conditioned on prior summaries at the same level (“previous context”) so summaries flow in order. A single model handles every level, trained first by behavioral cloning on human demonstrations, then by reinforcement learning against a reward model fit to human comparisons, following Stiennon et al. (2020). The framing is scalable oversight: labelers judge each subtask against only its local input, so no one has to read the whole book, which the authors estimate makes each data point over 50x cheaper than end-to-end collection.

On 40 popular books published in 2020 (unseen in pretraining, primarily narrative fiction averaging over 100K words), the best 175B RL model produced summaries rated 6 of 7 about 5 percent of the time and 5 of 7 over 15 percent, but the average stayed well below the human-written summaries; labeler agreement on relative model quality was near 80 percent. RL clearly beat behavioral cloning at 175B, less so at 6B, and comparisons became more label-efficient than demonstrations past 10-20K labels while being 3x faster to collect. Training only on the first subtree generalized as well as training on the full tree; the final full-tree model actually got worse, which the authors tie to compounding lower-level errors and auto-induced distributional shift. The 175B models set state of the art on BookSum (beating non-oracle baselines by 3-4 ROUGE points and every baseline on BERTScore), and their depth 1 summaries let a zero-shot 3B UnifiedQA model score competitively on NarrativeQA.

For a personal knowledge backend built over conversational history, this is the canonical evidence that hierarchical summarization can compress arbitrarily long corpora while keeping facts traceable back to source text, and that summary trees support downstream question answering without retrieval over raw text. Its failure modes are equally instructive: local context produced a *Pride and Prejudice* summary that turned a request for a dance into a marriage proposal, book-level output read as event lists rather than coherent narrative, and themes spread thinly across many chunks vanished. Any memory tree that distills leaves independently inherits exactly these risks (within this study’s scope: fiction, GPT-3-era models, fixed decomposition), so cross-leaf context and error auditing at composition points are design requirements, not refinements.

Read from: pages 1-20 of 37

MemoryAgentBench: Evaluating Memory in LLM Agents via Incremental Multi-Turn Interactions

Yuanzhe Hu, Yu Wang, Julian McAuley (UC San Diego), ICLR 2026, arXiv:2507.05257

The paper builds a benchmark for what it calls memory agents: LLM systems that accumulate information across many turns rather than reading one long document. Drawing on memory science, the authors name four competencies a memory system needs: accurate retrieval, test-time learning, long-range understanding, and selective forgetting. Their core methodological move is to feed context incrementally, chunk by chunk, wrapped in user messages with explicit memorize-this instructions, so that the agent's storage mechanism actually fires; questions come only after ingestion is complete. They restructure existing long-context datasets (document QA, LongMemEval, intent classification, novel summarization) and add two of their own, EventQA for temporal recall over novels and FactConsolidation for overwriting facts with later contradictory ones. The result is 2,071 questions over contexts of 103K to 1.44M tokens, run against long-context agents, three flavors of RAG (BM25, embedding models, structure-augmented systems like GraphRAG and HippoRAG-v2), and agentic memory products including MemGPT, MIRIX, Memo, Cognee, and Zep, most with GPT-4o-mini as backbone.

No method masters all four competencies. RAG agents beat their raw backbone on accurate retrieval (HippoRAG-v2 averages 65.1 against GPT-4o-mini's 49.2), but long-context models win test-time learning and long-range understanding, where retrieving top-k passages cannot deliver a global view; GraphRAG scores 0.4 on novel summarization. Selective forgetting breaks everyone: no system exceeds 28 percent on multi-hop fact updates, and even o4-mini, which reaches 80 on a 6K-token version, collapses to 14 at 32K, so the failure is length-driven, not task-driven. The commercial memory products fare worst overall (Memo 21.1, Cognee 20.6, Zep 24.0, against 60.6 for plain long-context GPT-5-mini). Ablations show smaller chunks help retrieval but hurt holistic tasks, and a stronger backbone lifts MIRIX by 9.7 points while barely moving simple RAG. Scope caveats: one backbone family dominates the main table, and budget limited which agents were tested.

For anyone building a personal knowledge backend over conversational history, the incremental-ingestion protocol is exactly the deployment shape, and the findings are cautionary in a specific way: retrieval-first architectures, which is what a summarized archive with search amounts to, handle lookup well but fail at superseding stale facts and at questions requiring the whole corpus. That argues for explicit update-and-consolidation machinery (newest-wins provenance, periodic re-summarization) rather than append-and-retrieve alone. For the field, the paper supplies a shared vocabulary of four competencies and hard evidence that current commercial memory layers lag a plain long context, which reframes where the open problems actually are.

Read from: pages 1-10 of 33

Remember When It Matters: Proactive Memory Agent for Long-Horizon Agents

Yifan Wu et al., arXiv:2607.08716, 2026-07

Abstract (verbatim)

In long-horizon tasks, decision-relevant state is often scattered across an expanding trajectory, while the action agent must surface it and act. As trajectories grow, task requirements, environment facts, prior attempts, diagnoses, and open subgoals can be buried in the context window or pushed beyond it, failing to influence decisions when needed. We call this failure mode “behavioral state decay”. We study memory as an active intervention mechanism rather than passive retrieval. A separate memory agent runs alongside an unmodified action agent, updating a structured memory bank from the recent trajectory and deciding whether to inject a memory-grounded reminder or remain silent. The module is plug-and-play with frontier action agents and existing agent harnesses. Across Terminal-Bench 2.0 and τ^2 -Bench, it improves pass@1 for both weaker and stronger action agents, with gains of +8.3 pp on Terminal-Bench and +6.8 pp on τ^2 -Bench. Ablations show that selective intervention outperforms passive bank exposure, always-on injection, advisor-only guidance, and general retrieval. As an early step toward open-weight memory policies, we train Qwen3.5-27B on SETA using SFT and GRPO, improving validation reward and achieving partial transfer to Terminal-Bench.

Bearing on this project

Already known from the 08-19 dossier, where it was recorded as naming the adjacent failure (behavioral state decay) but only intra-task. Its horizon is within a task; this project’s is cross-session over years.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

PASK: Toward Intent-Aware Proactive Agents with Long-Term Memory

Zhifei Xie et al., arXiv:2604.08000, 2026-04

Abstract (verbatim)

Proactivity is a core expectation for AGI. Prior work remains largely confined to laboratory settings, leaving a clear gap in real-world proactive agent: depth, complexity, ambiguity, precision and real-time constraints. We study this setting, where useful intervention requires inferring latent needs from ongoing context and grounding actions in evolving user memory under latency and long-horizon constraints. We first propose DD-MM-PAS (Demand Detection, Memory Modeling, Proactive Agent System) as a general paradigm for streaming proactive AI agent. We instantiate this paradigm in Pask, with streaming IntentFlow model for DD, a hybrid memory (workspace, user, global) for long-term MM, PAS infra framework and introduce how these components form a closed loop. We also introduce LatentNeeds-Bench, a real-world benchmark built from user-consented data and refined through thousands of rounds of human editing. Experiments show that IntentFlow matches leading Gemini3-Flash models under latency constraints, while identifying deeper user intent.

Bearing on this project

Intent-aware proactive agents with long-term memory. Part of the proactive-memory cluster that formed through 2026.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Improving Unsupervised Dialogue Topic Segmentation with Utterance-Pair Coherence Scoring

Linzi Xing, Giuseppe Carenini; SIGDIAL 2021; arXiv:2106.06719

Xing and Carenini attack dialogue topic segmentation (DTS), the task of cutting a conversation into topically coherent runs of utterances, in the unsupervised setting that the field is stuck with because labeled dialogue boundaries are scarce; the labeled sets that exist are small or synthetic and get spent on evaluation. Their move is to keep the classic TextTiling inference (coherence scores over adjacent utterance pairs, converted to depth scores, boundary wherever depth exceeds a mean-minus-half-sigma threshold) but replace its surface-feature coherence measure with a fine-tuned Next Sentence Prediction BERT. Since no gold coherence labels exist, they manufacture training data from open-domain corpora (DailyDialog in English, NaturalConv in Chinese): adjacent utterance pairs that follow basic dialogue-act flows such as Questions-Inform become positives, and negatives come from swapping in a non-adjacent turn from the same dialogue or a turn from a different-topic dialogue, trained with a marginal ranking loss. This yields 91,581 English and 599,148 Chinese paired samples with no human annotation.

On DialSeg 711 the full model reaches Pk 26.80, WindowDiff 28.24, and F1 0.776 against 40.44, 44.63, and 0.608 for plain TextTiling, and it beats every baseline on Doc2Dial (Pk 45.23) and the Chinese ZYS set (Pk 40.99), though the Chinese gains are smaller and unexplained. The ablation is the sharpest finding: dropping the dialogue-flow constraint on positive sampling costs far more than dropping the topic constraint on negatives (Pk 32.60 versus 26.95 on DialSeg 711), so knowing that a question and its answer belong together matters more than knowing two dialogues discuss different topics. The paper also shows why document-trained segmenters underperform here: TeT variants powered by GloVe or vanilla BERT, which win on monologue text, fail to consistently beat bare TextTiling on conversational data, because signals learned from monologue prose mislead on fragmented, informal turns. Scope limits ride along with these numbers: all three test sets are goal-oriented or synthetic dialogues, inference is tied to TextTiling, and the method needs a large open-domain corpus in the target language.

For the memory harness, this governs the corpus-to-tree ingestion stage, the cut from raw conversation logs into segments before any entity extraction, dated-fact rowing, or tree placement happens. It is the published warning that a document-tuned chunker on chat transcripts is the wrong default, and its Pk and WindowDiff protocol supplies the ground-truth metric frame for scoring whatever segmenter, deterministic TextTiling arm or schema-validated LLM call through the swappable model tiers, feeds the three injection arms downstream; segmentation quality then becomes measurable input to the LLM-judge scores rather than an eyeballed assumption.

Read from: full paper, all 11 pages including conclusions and references, via pypdf text extraction.

Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering

Xu, Cruz, Guevara, Wang, Deshpande, Wang, and Li (LinkedIn), SIGIR 2024 industry track, arXiv:2404.17723

This is a five-page industry paper describing a question-answering system deployed inside LinkedIn's customer service organization. The problem is retrieval over historical issue-tracking tickets (Jira), where conventional RAG treats tickets as plain text, chunks them to fit embedding-model context windows, and thereby loses two things: the internal structure of each ticket and the explicit links between tickets. The authors replace the flat text corpus with a knowledge graph built in two phases. Intra-ticket parsing turns each ticket into a tree whose nodes are sections (summary, description, steps to reproduce, priority, root cause, comments), using rule-based extraction for predefined fields and an LLM guided by a YAML template for the rest. Inter-ticket connection then links trees by explicit references recorded in Jira (clone, related-to) and by implicit edges added when the cosine similarity of ticket-title embeddings crosses a threshold. Node text is embedded with pretrained models (E5, BERT) into a vector database (Qdrant); the graph lives in Neo4j.

At query time an LLM parses the user's question into named entities keyed to template sections plus an intent (for example, steps to reproduce). Entity values are matched against section-specific node embeddings, node scores are summed to the ticket level to rank the top K tickets, and the LLM then rewrites the query with the winning ticket ID and translates it into a Cypher query that walks the tree to the intent node. A fallback reverts to plain text retrieval if query execution fails. On a curated golden dataset, with GPT-4 and E5 held fixed in both arms, the method lifts MRR from 0.522 to 0.927 (a 77.6 percent gain), Recall@3 from 0.640 to 1.000, and BLEU from 0.057 to 0.377. In a randomized split of the support team over roughly six months, median per-issue resolution time fell 28.6 percent (P50 from 7 to 5 hours, mean from 40 to 15). The golden dataset is small and internal, the graph template is hand-built, and the benchmark numbers reflect one domain, so the figures establish a deployment result rather than a general benchmark.

For a personal knowledge backend over conversational history, the transferable idea is the dual-level representation: parse each unit (here a ticket, there a session or leaf) into a section tree, then connect units by explicit references and thresholded embedding similarity, so retrieval can return a coherent subgraph instead of arbitrary chunks. The section-aware scoring, which sums similarity only over nodes of the queried type, is a concrete mechanism for intent-scoped retrieval, and the paper's own future-work list (automatic template extraction, dynamic graph updates from queries) marks exactly the parts that were manual here and would need to be automatic at organizational scale.

Read from: full text

Asymmetric Capacity Allocation in Self-Refinement Pipelines

Zhuoyi Yang et al., arXiv:2608.21345, 2026-08

Abstract (verbatim)

Self-refinement, typically structured as generation, critique, and revision, is a widely adopted paradigm for improving LLM generation and serves as a core mechanism in many LLM agents. While the three stages involve different cognitive demands, most existing approaches conveniently treat the model size as an implementation detail rather than a subject of study, which may lead to a waste of resources. Little work has systematically examined how model size affects each stage or whether effective self-refinement requires equally capable models for generation, critique, and revision. We present the first stage-wise model size study of the self-refinement pipeline on 5 benchmarks from different domains using 6 model sizes of Qwen3 and 4 model sizes of Gemma 3. We conclude that larger generators and refiners generally improve the pipeline, whereas an undersized refiner can even harm performance. Second, performance is highly insensitive to the size of the critic, although including even a small critic consistently outperforms omitting critique altogether. Our findings demonstrate that model capacity should not be allocated uniformly across self-refinement pipelines. Instead, different stages exhibit distinct size scaling characteristics, providing practical guidance for designing more computationally efficient multi-stage language model systems.

Bearing on this project

Kills the per-stage tier angle. Allocating model capacity asymmetrically across pipeline stages is the question this project's sensitivity table was going to answer. Published this month.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Beyond Static Summarization: Proactive Memory Extraction for LLM Agents

Chengyuan Yang et al., arXiv:2601.04463, 2026-01

Abstract (verbatim)

Memory management is vital for LLM agents to handle long-term interaction and personalization. Most research focuses on how to organize and use memory summary, but often overlooks the initial memory extraction stage. In this paper, we argue that existing summary-based methods have two major limitations based on the recurrent processing theory. First, summarization is “ahead-of- time”, acting as a blind “feed-forward” process that misses important details because it doesn’t know future tasks. Second, extraction is usually “one-off”, lacking a feedback loop to verify facts, which leads to the accumulation of information loss. To address these issues, we propose proactive memory extraction (namely ProMem). Unlike static summarization, ProMem treats extraction as an iterative cognitive process. We introduce a recurrent feedback loop where the agent uses self-questioning to actively probe the dialogue history. This mechanism allows the agent to recover missing information and correct errors. Our ProMem significantly improves the completeness of the extracted memory and QA accuracy. It also achieves a superior trade-off between extraction quality and token cost.

Bearing on this project

Proactive memory EXTRACTION rather than recall - an iterative self-questioning loop that probes dialogue history to recover missed detail. Different axis from proactive retrieval; relevant to the distiller stage.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

DTELS: Towards Dynamic Granularity of Timeline Summarization

Chenlong Zhang et al., arXiv:2411.09297, 2024-11

Abstract (verbatim)

The rapid proliferation of online news has posed significant challenges in tracking the continuous development of news topics. Traditional timeline summarization constructs a chronological summary of the events but often lacks the flexibility to meet the diverse granularity needs. To overcome this limitation, we introduce a new paradigm, Dynamic-granularity TimeLine Summarization, (DTELS), which aims to construct adaptive timelines based on user instructions or requirements. This paper establishes a comprehensive benchmark for DTELS that includes: (1) an evaluation framework grounded in journalistic standards to assess the timeline quality across four dimensions: Informativeness, Granular Consistency, Factuality, and Coherence; (2) a large-scale, multi-source dataset with multiple granularity timeline annotations based on a consensus process to facilitate authority; (3) extensive experiments and analysis with two proposed solutions based on Large Language Models (LLMs) and existing state-of-the-art TLS methods. The experimental results demonstrate the effectiveness of LLM-based solutions. However, even the most advanced LLMs struggle to consistently generate timelines that are both informative and granularly consistent, highlighting the challenges of the DTELS task.

Bearing on this project

Timeline summarization with dynamic granularity - i.e. how coarse or fine the timeline should be. Directly relevant to the arc-grouping question (how N chapter cells become 5 sections).

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Extract, Define, Canonicalize: An LLM-based Framework for Knowledge Graph Construction

Bowen Zhang, Harold Soh; EMNLP 2024 (main conference); arXiv:2404.03868

EDC is a three-phase framework for building knowledge graphs from text with prompted LLMs and no fine-tuning of the base models. Prior LLM-based KG construction stuffed the whole schema into the prompt, which breaks once the schema outgrows the context window; EDC inverts the order. Phase one runs open information extraction, pulling free-form [subject, relation, object] triplets with no schema in sight. Phase two asks the LLM to write a natural language definition for each relation the extraction induced. Phase three canonicalizes: definitions are embedded with a sentence transformer, vector search proposes the nearest schema candidates, and the LLM verifies whether each substitution is legitimate in context, which curbs the over-generalization that pure embedding clustering suffers. The framework runs in two modes, aligning to a given target schema (rejecting triplets with no equivalent) or growing a canonical schema from empty. An optional refinement pass, EDC+R, feeds prior triplets plus relations fetched by a trained Schema Retriever (a fine-tuned E5-mistral-7b embedder) back into extraction as a hint.

On WebNLG, REBEL, and Wiki-NRE (schemas of 159, 200, and 45 relations, sizes that break prompt-embedded baselines), EDC matches or beats the trained specialists REGEN and GenIE, with GPT-4, GPT-3.5-turbo, and a local Mistral-7b all viable for extraction. Refinement helps everywhere: with GPT-3.5 on WebNLG, Partial F1 climbs from 0.746 to 0.794, and ablating the Schema Retriever gives most of that gain back. In self-canonicalization, human-judged precision drops only from 0.982 to 0.956 on WebNLG while relation count falls from 529 to 200, far ahead of CESI at 0.724. The caveats are real: only relations get canonicalized (entity de-duplication is future work), only the extraction phase was tested across model tiers while everything else ran GPT-3.5, and the pipeline is call-heavy at roughly \$0.009 per example even on cheap models.

For the harness, EDC is the closest published shape to the stage where raw dated fact rows and entity mentions must merge into a canonical entity knowledge base. Its define-then-verify canonicalization is directly portable Python: schema-validated extraction calls through swappable model tiers map onto EDC's finding that a local 7B holds up at extraction while verification may need a stronger tier, and its target-alignment rejection of non-matching triplets gives the measured-rejection band published company. Its human-scored precision protocol (agreement 0.94) is a sanity anchor for the LLM-judge scoring that grades what the three injection arms retrieve, though EDC says nothing about segmentation, tree ingestion, or injection ordering; it governs canonicalization alone.

Read from: full text of all 17 pages (body pages 1 to 9 covering abstract, method, experiments, conclusion, and limitations, plus references), extracted via pypdf.

OCR Hinders RAG: Evaluating the Cascading Impact of OCR on Retrieval-Augmented Generation

Junyuan Zhang et al., arXiv:2412.02592, 2024-12

Abstract (verbatim)

Retrieval-augmented Generation (RAG) enhances Large Language Models (LLMs) by integrating external knowledge to reduce hallucinations and incorporate up-to-date information without retraining. As an essential part of RAG, external knowledge bases are commonly built by extracting structured data from unstructured PDF documents using Optical Character Recognition (OCR). However, given the imperfect prediction of OCR and the inherent non-uniform representation of structured data, knowledge bases inevitably contain various OCR noises. In this paper, we introduce OHRBench, the first benchmark for understanding the cascading impact of OCR on RAG systems. OHRBench includes 8,561 carefully selected unstructured document images from seven real-world RAG application domains, along with 8,498 Q&A pairs derived from multimodal elements in documents, challenging existing OCR solutions used for RAG. To better understand OCR's impact on RAG systems, we identify two primary types of OCR noise: Semantic Noise and Formatting Noise and apply perturbation to generate a set of structured data with varying degrees of each OCR noise. Using OHRBench, we first conduct a comprehensive evaluation of current OCR solutions and reveal that none is competent for constructing high-quality knowledge bases for RAG systems. We then systematically evaluate the impact of these two noise types and demonstrate the trend relationship between the degree of OCR noise and RAG performance. Our OHRBench, including PDF documents, Q&As, and the ground truth structured data are released at: <https://github.com/opendatalab/OHR-Bench>

Bearing on this project

Kills the input-degradation angle. This is exactly the question of how OCR quality cascades into a retrieval-augmented pipeline, with a benchmark. The OCR tax control in this project's Chinese corpus is therefore an established method applied to a new setting, not a new instrument - weaker as a claim, stronger as science, and now citable.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

TimelineReasoner: Advancing Timeline Summarization with Large Reasoning Models

Liancheng Zhang et al., arXiv:2605.12518, 2026-04

Abstract (verbatim)

The proliferation of online news poses a challenge to extracting structured timelines from unstructured content. While recent studies have shown that Large Language Models (LLMs) can assist Timeline Summarization (TLS), these approaches primarily treat models as passive generators. The emergence of Large Reasoning Models (LRMs) presents an opportunity to reason over events actively, enabling iterative evidence acquisition, the detection of missing events, and the validation of temporal consistency. To systematically leverage the reasoning capabilities of LRMs, we propose TimelineReasoner, a novel framework that shifts TLS from static generation to an active, reasoning-driven process. Unlike prior work, TimelineReasoner adopts a two-stage framework: Global Cognition, which tracks events at a macroscopic level and continuously updates a global event memory, and Detail Exploration, which identifies informational gaps and refines the timeline via targeted document retrieval. To support this, TimelineReasoner incorporates several specialized mechanisms, including an Event Scraper for retrieving temporal event descriptions, a Timeline Updater for refining the timeline, and a Supervisor for detecting gaps in the timeline and guiding retrieval. Experimental results on open-domain TLS datasets demonstrate that TimelineReasoner significantly outperforms existing LLM-based TLS methods in terms of timeline accuracy, coverage, and coherence. On closed-domain TLS datasets, our method performs on par with or exceeds state-of-the-art approaches. This work not only pushes the boundaries of TLS but also highlights the broader potential of LRM-based reasoning frameworks for timeline summarization.

Bearing on this project

Timeline summarization with large reasoning models, 2026. The LLM-era version of the entity axis.

Read from: **abstract only**. Added 2026-08-27 during the novelty reassessment. Full-text read outstanding; do not cite specifics beyond the abstract until that is done.

Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena

Lianmin Zheng et al.; NeurIPS 2023 Datasets and Benchmarks Track; 2023; arXiv:2306.05685

This is the systematic study that made “LLM-as-a-judge” a defensible evaluation method rather than a convenience. The authors build two human-preference benchmarks: MT-bench, 80 hand-written multi-turn questions across eight categories (writing, roleplay, extraction, reasoning, math, coding, and two knowledge areas), and Chatbot Arena, a crowdsourced platform where users vote in anonymous head-to-head battles. Against roughly 3K controlled expert votes and 3K crowd votes, they test whether a strong LLM judge reproduces human preference. They define three judging modes (pairwise comparison, single-answer grading, and reference-guided grading) and audit each for failure.

The headline result is that GPT-4 agrees with human experts at 85 percent on non-tie votes, above the 81 percent humans reach with each other, with the caveat that agreement is easiest when models differ sharply in quality; on close pairs it falls toward 70 percent. The bias audit supplies the numbers a practitioner needs. Position bias is severe: Claude-v1 was order-consistent in only 23.8 percent of cases and only GPT-4 exceeded 60 percent, so the paper’s conservative fix is to judge both orderings and count disagreement as a tie. A “repetitive list” verbosity attack fooled Claude-v1 and GPT-3.5 on 91.3 percent of padded answers versus 8.7 percent for GPT-4. Self-enhancement bias is observed (GPT-4 favors itself by about 10 percent win rate) but the authors state their data cannot confirm it as causal. Judges also fail on math they can solve in isolation; a reference-guided prompt, where the judge answers first and grades against its own answer, cut the failure rate from 14 of 20 to 3 of 20. Single-answer grading tracked pairwise results closely, which licenses the cheaper mode. The study measures helpfulness only, folding accuracy, relevance, and creativity into one score.

For the memory harness, this paper governs the final scoring stage: after the corpus is segmented into a tree of entities and dated facts, and after answers are produced under the three injection arms through schema-validated calls at each model tier, an LLM judge scores those answers against gold. The design consequences are direct. The judge grades against a fixed reference answer (the setting where reference-guided grading nearly eliminated math failures), positions get swapped or a stable single-answer rubric is used, verbose answers are treated as suspect, and the judge model stays pinned rather than varying with the arm. The 80 to 85 percent agreement figure is the calibration target the hand-labeled set should match, while the near-random consistency of weaker judges warns against scoring with the same small local models being tested.

Read from: full main body, pages 1 to 10 of the arXiv v4 PDF (abstract through conclusion; appendices and references not read), via pypdf text extraction.

A Comprehensive Survey on Automatic Knowledge Graph Construction

Zhong, Wu, Li, Peng, and Wu; ACM Computing Surveys 56(4), 2024; arXiv:2302.05019

This is a survey of more than 300 methods for building knowledge graphs automatically from raw data. The authors organize construction into three stages: knowledge acquisition (entity discovery, coreference resolution, relation extraction), knowledge refinement (graph completion and fusion with other graphs), and knowledge evolution (conditional and temporal knowledge, dynamics over time). They define a knowledge graph as $G = \{E, R, T, F_k\}$, where T is a set of triples (head, tail, relation) and F_k is background knowledge, a rule set or schema that constrains which facts count as knowledge rather than mere data; a graph without this component, they argue, is just a data graph. The survey frames its coverage with the HACE view of big data (heterogeneous, autonomous, complex, evolving) and compares itself against nine prior surveys, claiming broader coverage of conditions, coreference, and semi-structured sources.

The construction-from-text sections lay out a stable pipeline vocabulary. Entity discovery splits into named entity recognition, fine-grained entity typing, and entity linking (disambiguation against existing nodes, creating a node when none exists). Coreference resolution runs as mention detection followed by antecedent scoring, with the end-to-end span-ranking model of Lee et al. as the standard deep architecture. Relation extraction divides by configuration: open extraction without a schema, schema-bound classification, distant supervision with its noise problem and the at-least-one multi-instance assumption, plus document-level, few-shot, and joint variants. For each task the survey tracks the same progression, from rules and wrappers through statistical models (CRFs, SVMs) to CNN, LSTM, attention, GCN, and pre-trained-language-model designs. It also catalogs practical resources: about two dozen KG projects with sizes (Wikidata at 96M+ items, YAGO at 2B+ facts, ConceptNet at 34M+ items) and tools from spaCy and OpenNRE to the end-to-end gBuilder.

For a personal knowledge backend built over conversational history, the useful contributions are the pipeline decomposition itself (a memory system doing entity linking and coreference over chat logs is running exactly these tasks), the insistence that a schema or background knowledge separates a knowledge graph from a data graph, and the storage discussion (relational, key/value, and native graph stores such as Neo4j and RDF systems). The open problems the authors flag map directly onto that setting: long intricate contexts requiring coreference, logic, and commonsense reasoning; knowledge dynamics as facts change; and federated learning for privacy when scaling to an organization, where encrypted entity alignment remains unsolved. The survey predates LLM-based extraction, so its method coverage stops at BERT-era models.

Read from: pages 1-20 (introduction, definitions, resources, preprocessing, entity discovery, coreference, relation extraction through distant supervision) and pages 33-37 (knowledge dynamics, storage, discussion, conclusion); skipped pages 21-32 (few-shot, document-level, and joint extraction; knowledge refinement; most of knowledge evolution) and the references.

UniversalNER: Targeted Distillation from Large Language Models for Open Named Entity Recognition

Wenxuan Zhou, Sheng Zhang, Yu Gu, Muhao Chen, Hoifung Poon; ICLR 2024; arXiv:2308.03279

The paper asks whether a small model can match a frontier model on one broad task class rather than on everything, and answers with what it calls targeted distillation through mission-focused instruction tuning. Instead of diversifying instructions the way Alpaca and Vicuna do, the authors diversify inputs: they sample 50,000 passages of at most 256 tokens from the Pile's 22 subcorpora, have ChatGPT (gpt-3.5-turbo-0301, temperature 0) tag every entity and its type with no type list supplied, and after filtering obtain 45,889 passage-annotation pairs covering 240,725 entities and 13,020 distinct entity types. LLaMA 7B and 13B are then tuned on a conversation-style template that asks one entity type per query and expects a JSON list back, with negative entity types sampled in proportion to corpus frequency so the model learns to return empty lists.

To evaluate, they assemble the largest open NER benchmark to date: 43 datasets across 9 domains (biomedicine, clinical, programming, social media, law, finance, and others), scored by strict entity-level micro F1. Zero-shot, UniNER-7B reaches 41.7 average F1 and UniNER-13B 43.4, against 34.9 for ChatGPT itself and roughly 14 for Vicuna-7B; generalist distillates like Alpaca sit near zero, over 30 points behind. The ablations carry as much information as the headline: frequency-based negative sampling beats no sampling by 21.9 F1, asking all types in one query costs 3.3 F1 versus one type per query, and definition-style type labels cost 11.8 on standard benchmarks while improving tolerance to type paraphrasing. With supervised finetuning added, UniNER-7B hits 84.78 average F1 on 20 datasets, above InstructUIE-11B at 81.16. The scope is real but narrow in the same breath: English only, NER only, a 2023-era teacher, and even the winning zero-shot average of 41.7 says open-type extraction remains far from solved.

For the memory harness, this paper governs the entity extraction stage of the corpus-to-tree pipeline, where schema-validated LLM calls pull entities and dated facts into the KB before the three injection arms and LLM-judge scoring ever run. It calibrates what the swappable cheap tier can be expected to do there: a 7B-class model can plausibly own extraction, but only a specialized one; a generalist small model is the named failure mode, trailing by about 30 F1. Two of its mechanisms transfer directly without any distillation of our own: one-entity-type-per-call prompting over all-in-one requests, and deliberate negative queries so the extractor learns that an empty list is a valid answer.

Read from: extracted text of pages 1 through 13 of the 19-page PDF, covering the full main paper (introduction through conclusion on page 9) plus part of the references; the remaining pages are references and appendix tables.